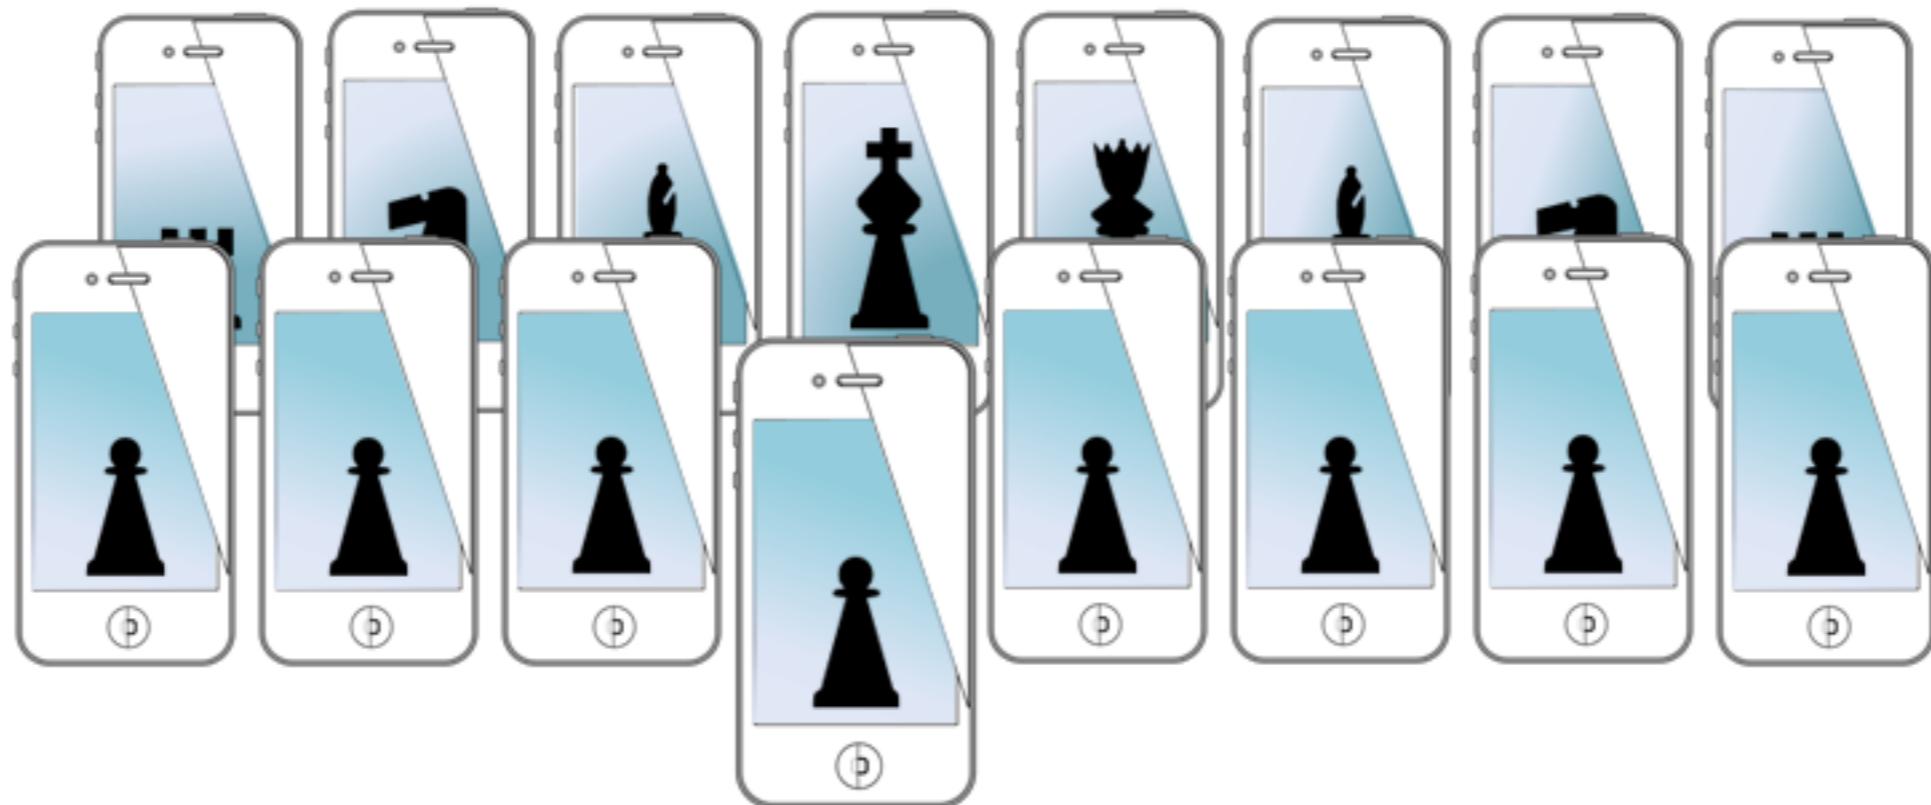


MOBILE SENSING LEARNING



CS5323 & 7323
Mobile Sensing and Learning

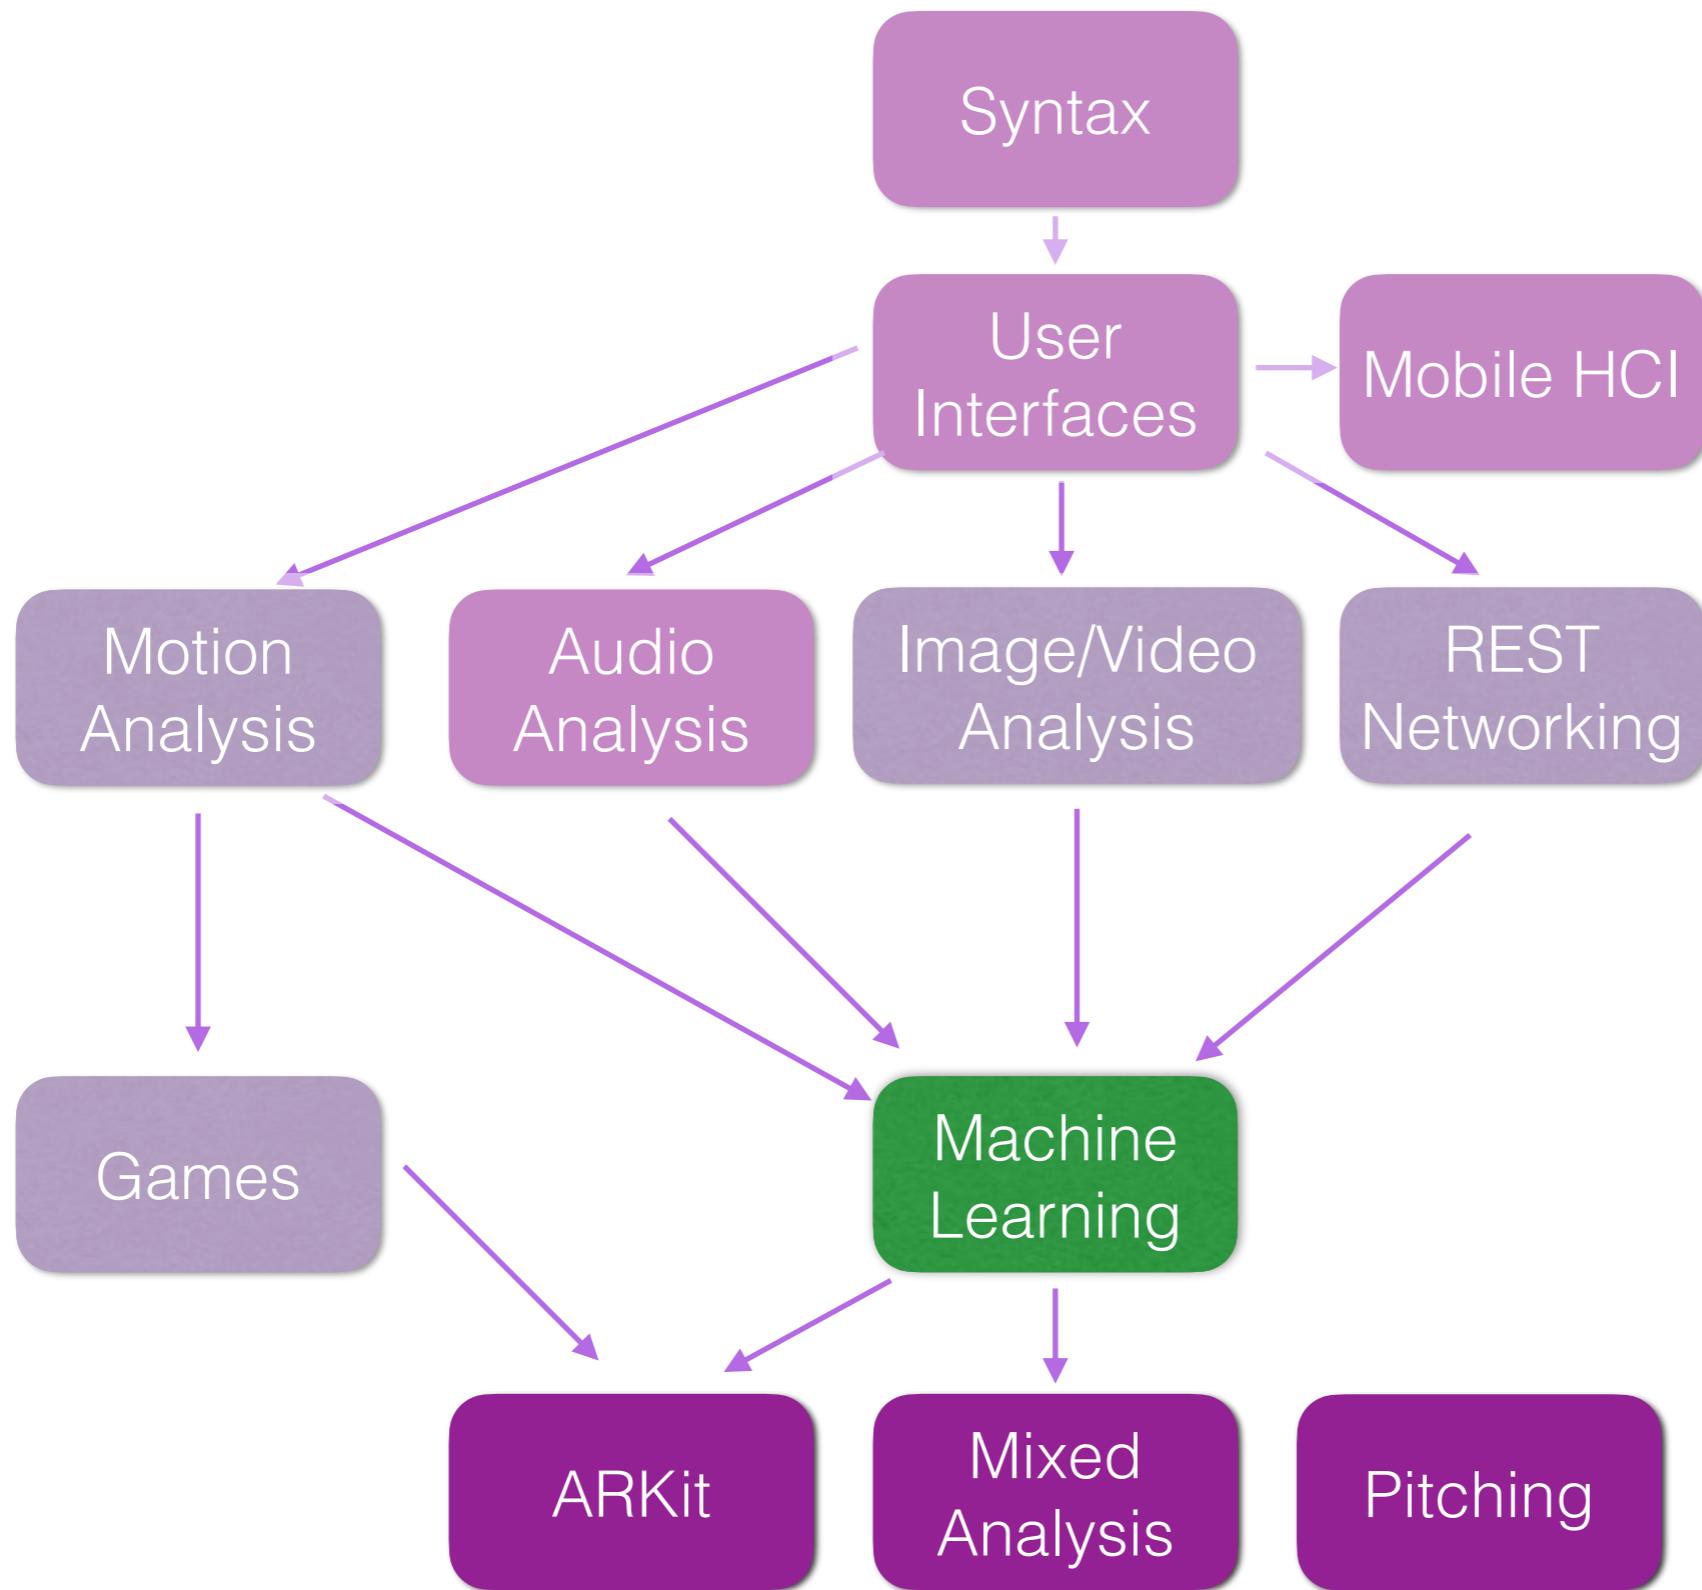
machine learning with Apple

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

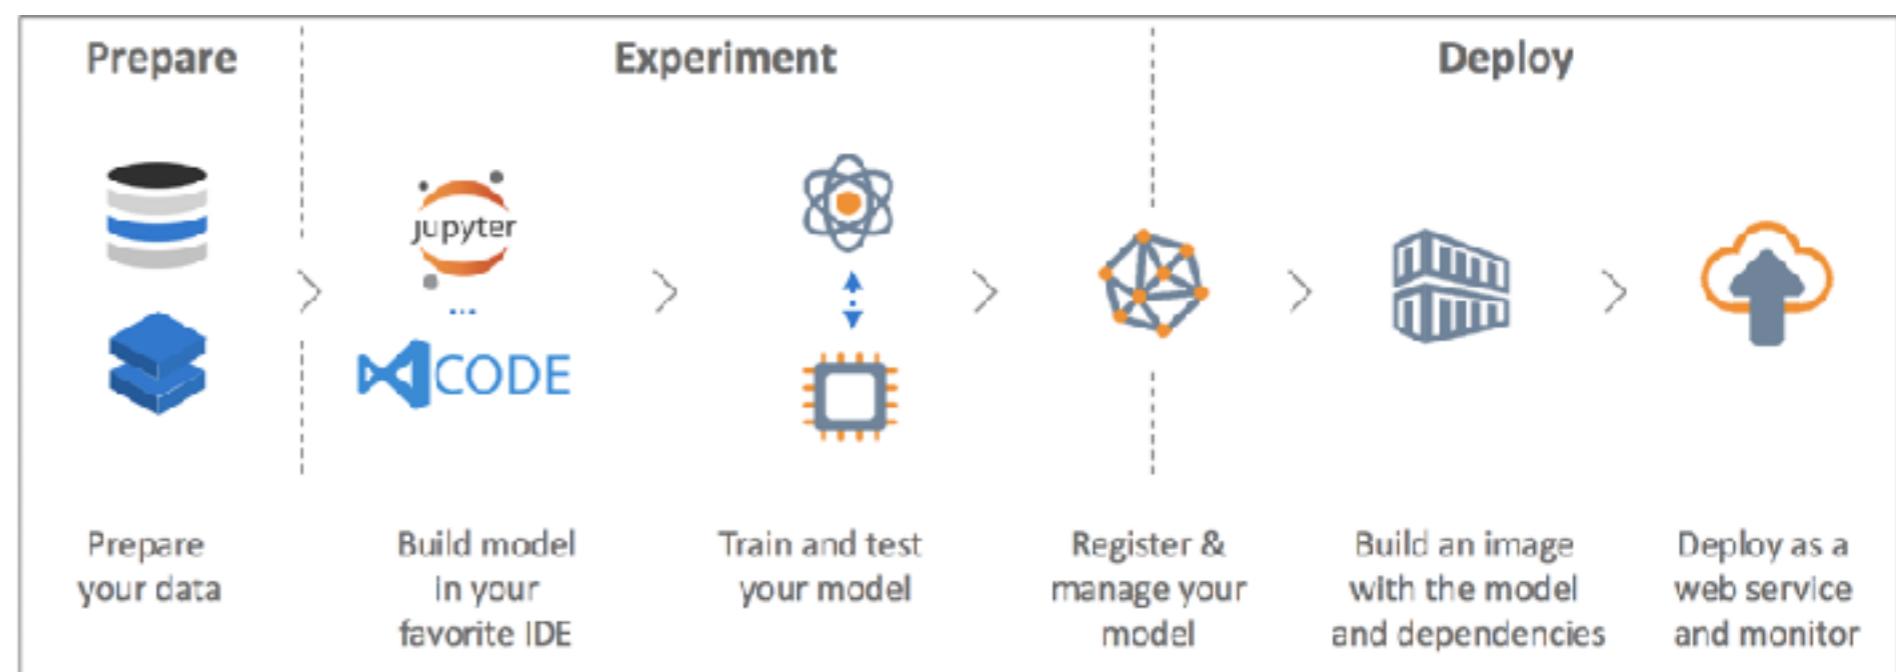
course logistics and agenda

- logistics:
 - grading update
 - final flipped module (next time!!!)
 - **A5** is due soon-ish, try to make it a first draft of your final project!
- Agenda
 - CoreML/CreateML
 - **final project proposal** is also due at same time
 - today: talk about old projects!

class overview



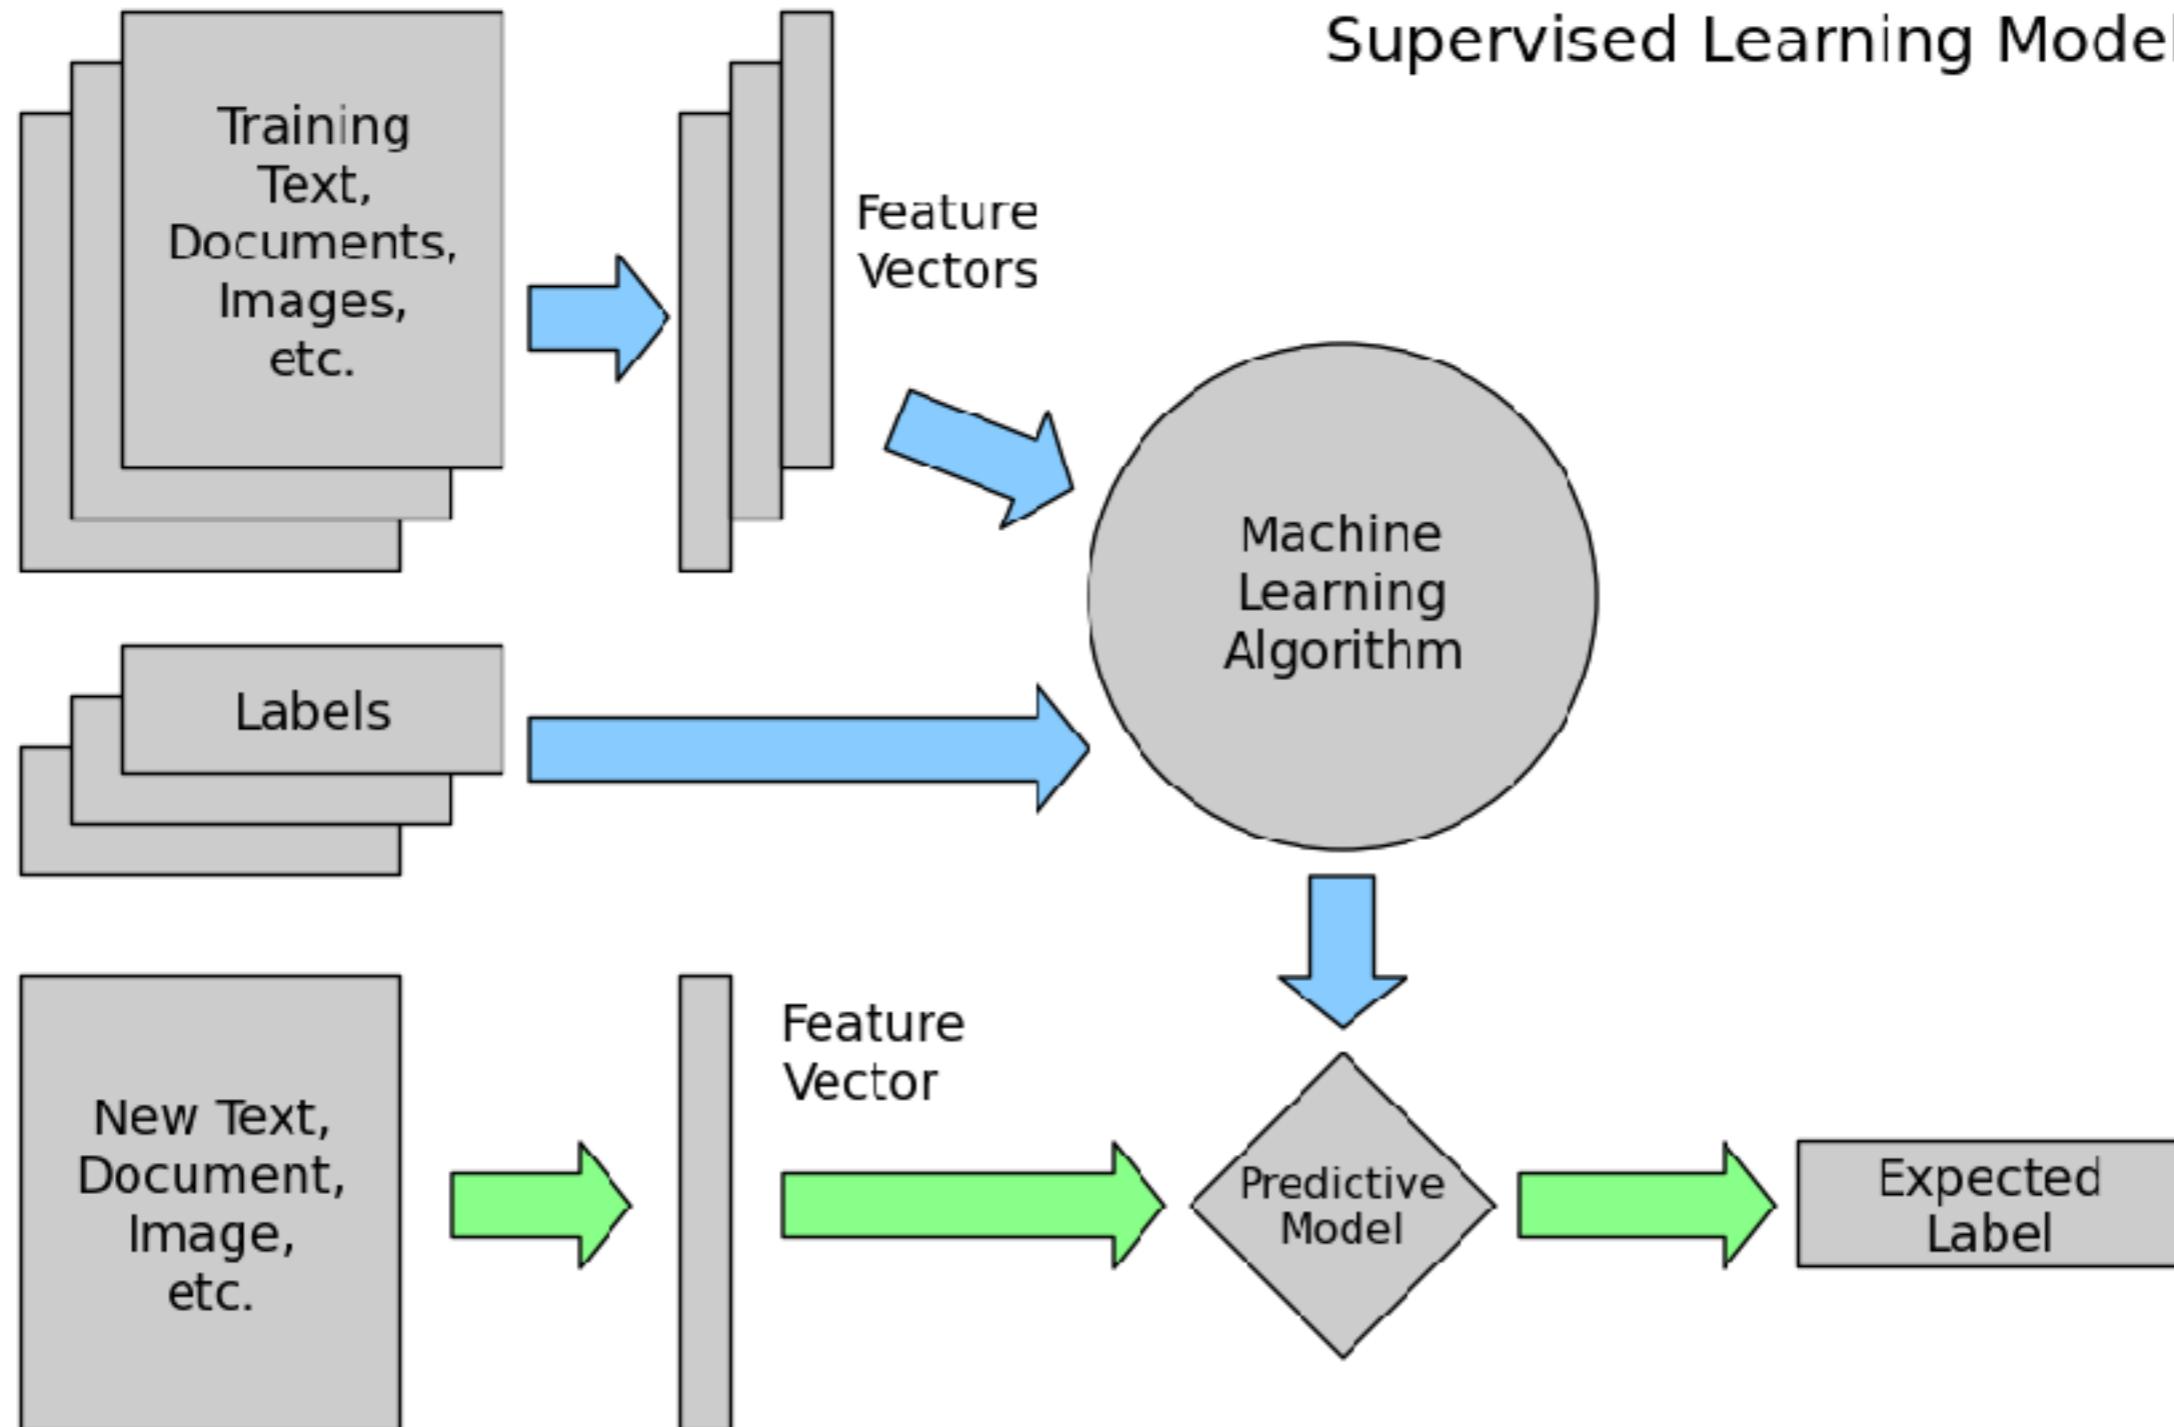
lab five town hall



machine learning

The image shows the Google Cloud Platform homepage. At the top, there is a navigation bar with links for "Why Google", "Products", "Solutions", "Customers", "Developers", "Support", and "Partners". On the right side of the navigation bar are links for "Contact sales" and "Try it now". Below the navigation bar, there is a search bar with the placeholder "Search this site" and a magnifying glass icon. The main content area features a large image of a server rack with many cables. Overlaid on this image is the text "Prediction API" next to a blue hexagonal icon containing a white line graph. Below this, there is a description: "Use Google's machine learning algorithms to analyze data and predict future outcomes using a familiar RESTful interface." At the bottom left of the main content area is a blue button with the text "Try it now". Overlaid on the entire main content area is a large, semi-transparent white text box containing the message "This is a Coach Course in ML Right Now!!".

machine learning models



types of data

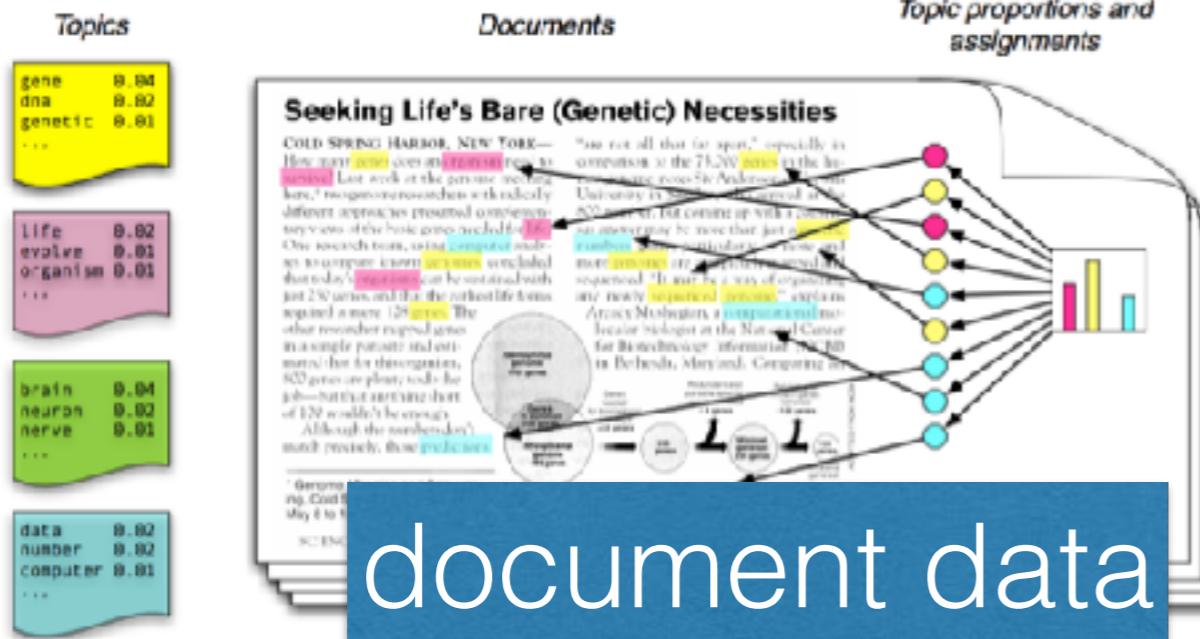


Figure source: Blei, D. M. (2012). Probabilistic topic models. *Communications of the ACM*, 55(4), 77-84.

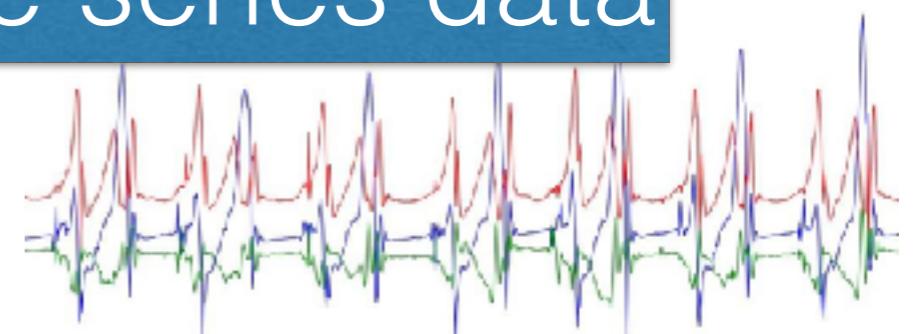


table data

Attributes, columns, variables, fields, characteristics, Features
 Objects, records, rows, points, samples, cases, entities, instances

TID	Pregnant	BMI	Age	Diabetes
1	Y	33.6	41-50	positive
2	N	26.6	31-40	negative
3	Y	23.3	31-40	positive
4	N	28.1	21-30	negative
5	N	43.1	31-40	positive
6	Y	25.6	21-30	negative
7	Y	31.0	21-30	positive
8	Y	35.3	21-30	negative
9	N	30.5	51-60	positive
10	Y	37.6	51-60	positive

time series data



features and labels

- actually, feature vectors
- **classic example:** the iris dataset—table data



setosa

versicolor

virginica

- 4 features
 - sepal length in cm [5.1, 3.5, 1.4, 0.2] setosa
 - sepal width in cm [5.7, 2.8, 4.5, 1.3] versicolor
 - petal length in cm [7.6, 3.0, 6.6, 2.1] virginica
 - petal width in cm

50 examples each

features

- most common is numeric and categorical
 - vector quantization (numeric to categories)
 - text:
 - bag of words
 - term frequency inverse document frequency
 - text embeddings from neural nets
 - graphs
 - used to quantize
- take data mining!
or python machine learning!

common ML algorithms

nonparametric

- nearest neighbor
- k-nearest neighbor (KNN)
- kernel density estimator

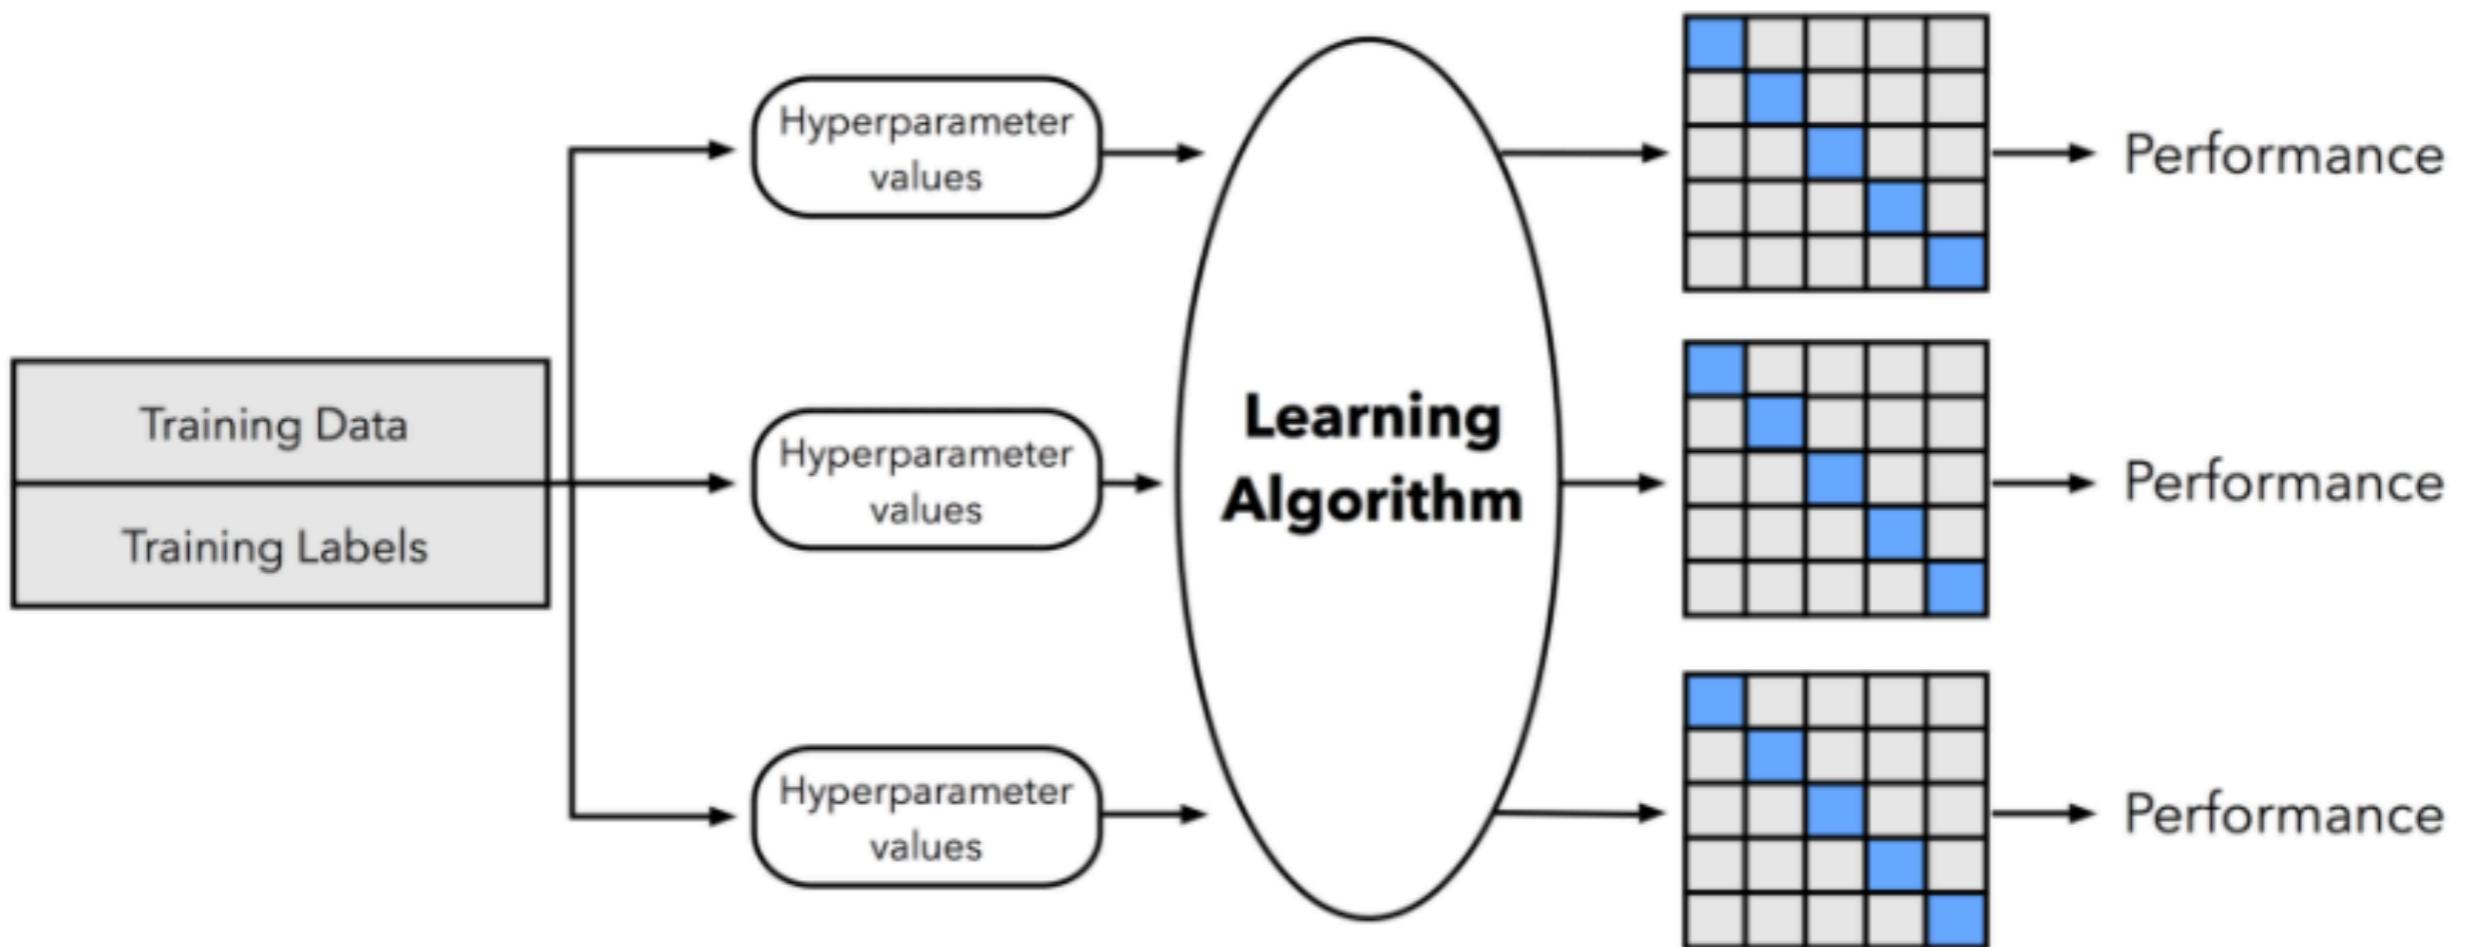
parametric

- decision tree
- random forest/boosted trees
- logistic regression
- neural networks
- gaussian mixtures

- support vector machines
- and many many more...

finding the best ML model

- try a bunch of stuff until it works well enough



http://ethen8181.github.io/machine-learning/model_selection/model_selection.html

turi create demo, with SFrame



Carlos Guestrin · 2nd

Senior Director of AI and Machine Learning at Apple &
Amazon Professor of Machine Learning at University of
Washington



python_short_examples >
TuriExample.ipynb

turicreate.SFrame

`class turicreate.SFrame (data=None, format='auto', _proxy=None)`

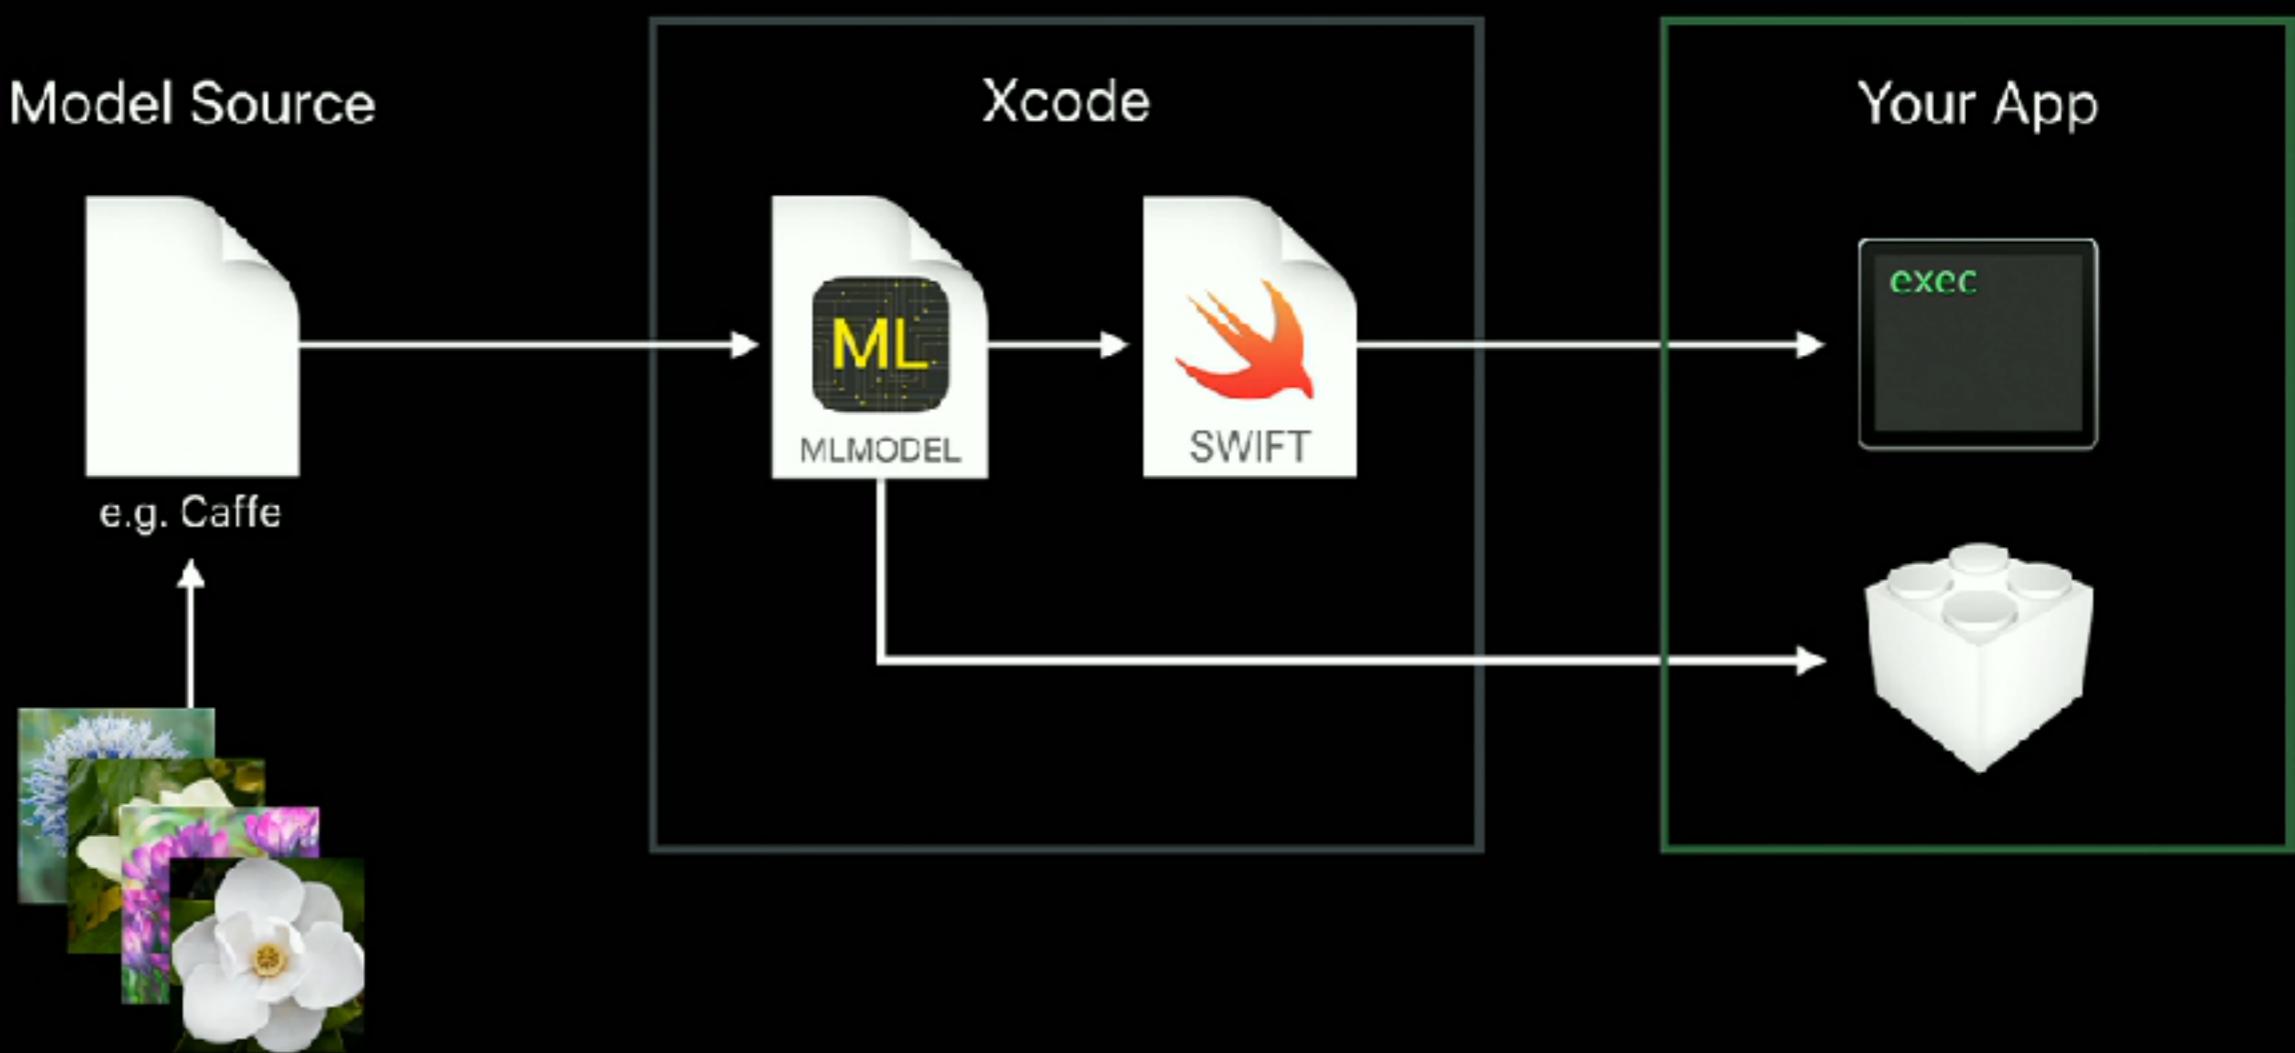
SFrame means scalable data frame. A tabular, column-mutable dataframe object that can scale to big data. The data in SFrame is stored column-wise, and is stored on persistent storage (e.g. disk) to avoid being constrained by memory size. Each column in an SFrame is a size-immutable `SArray`, but SFrames are mutable in that columns can be added and subtracted with ease. An SFrame essentially acts as an ordered dict of SArrays.

Currently, we support constructing an SFrame from the following data formats:

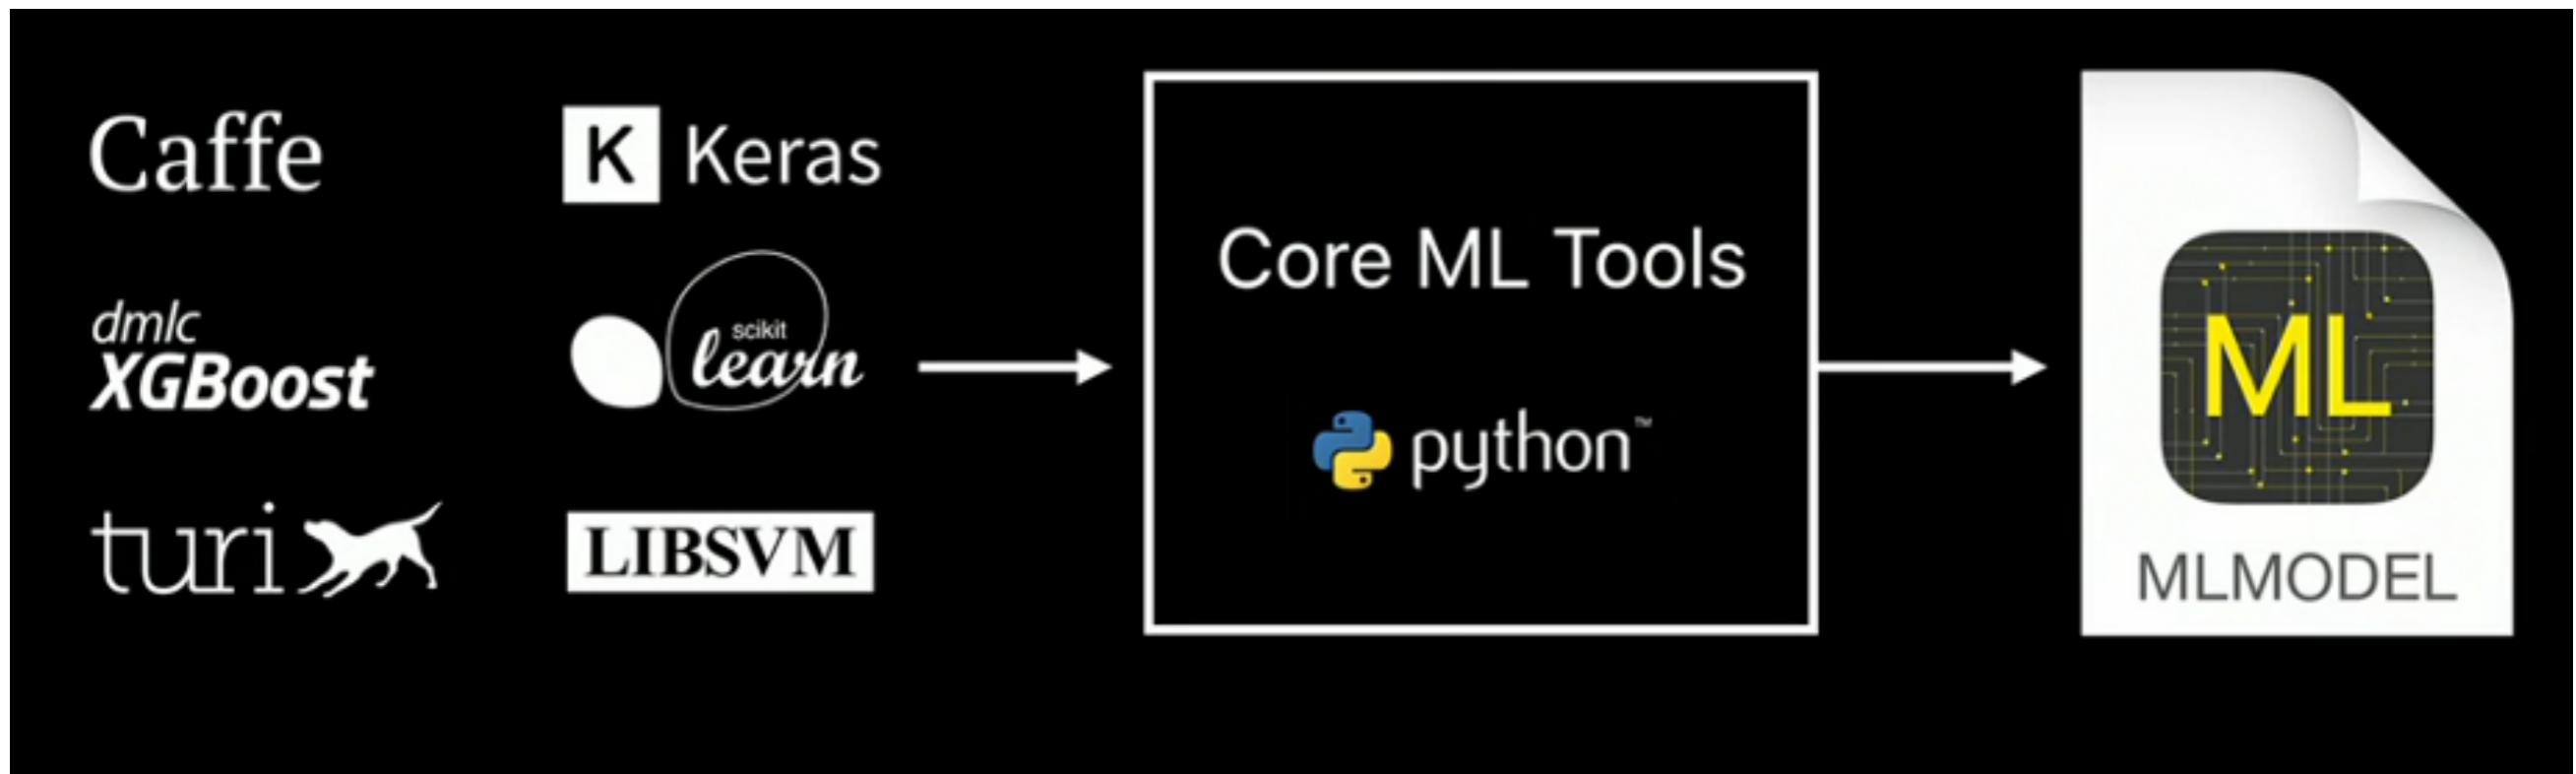
- csv file (comma separated value)
- sframe directory archive (A directory where an sframe was saved previously)
- general text file (with csv parsing options, See `read_csv()`)
- a Python dictionary
- pandas.DataFrame
- JSON

CoreML

Conversion Workflow



CoreML



but... we want more than
pre-trained models

installing CoreMLTools

- built into **TuriCreate**

```
model.export_coreml("MyModel.mlmodel")
```

- also available for other libraries, like **sklearn**: so create a conda environment and install coremltools

- conda install sklearn numpy (+others) ...
- pip install coremltools

```
clf = RandomForestClassifier(n_estimators=50)
print("Training Model", clf)

clf.fit(X,y)

print("Exporting to CoreML")

coreml_model = coremltools.converters.sklearn.convert(
    clf,
    ["accelX"]*50+["accelY"]*50+["accelZ"]*50, # feature names (optional)
    "Direction" # label name (optional)

# save out as a file
coreml_model.save('rf.mlmodel')
```

using CoreML

- drag into project

▼ Machine Learning Model

Name RandomForestAccel

Type Tree Ensemble Classifier

Size 28 KB

Author unknown

Description description not included

License unknown

▼ Model Class

C RandomForestAccel ⓘ
Automatically generated Swift model class

▼ Model Evaluation Parameters

Name	Type	Des
▼ inputs		
input	MultiArray (Double 150)	
▼ outputs		
classLabel	String	
classProbability	Dictionary (String → Double)	

```
/// Class for model loading and prediction
@available(macOS 10.13, iOS 11.0, tvOS 11.0, watchOS 4.0, *)
class RandomForestAccel {
    var model: MLModel

    /**
     Construct a model with explicit path to mlmodel file
     - parameters:
     - url: the file url of the model
     - throws: an NSError object that describes the problem
    */
    init(contentsOf url: URL) throws {
        self.model = try MLModel(contentsOf: url)
    }

    /// Construct a model that automatically loads the model from the bundle
    convenience init() {
        let bundle = Bundle(for: RandomForestAccel.self)
        let assetPath = bundle.url(forResource: "RandomForestAccel",
            try! self.init(contentsOf: assetPath!)
    }

    /**
     Make a prediction using the structured interface
     - parameters:
     - input: the input to the prediction as RandomForestAccelInput
     - throws: an NSError object that describes the problem
     - returns: the result of the prediction as RandomForestAccelOutput
    */
    func prediction(input: RandomForestAccelInput) throws -> RandomForestAccelOutput
}
```

CoreML

- assumes some understanding of **final flipped module**
- combining our web server, mongo, iOS, CoreML



and now its time for a demo

Create ML

making machine learning easy to use for developers
that are not data scientists



Image

Image classification
Object detection
Style transfer NEW



Video

Action classification NEW
Style transfer NEW



Motion

Activity classification



Sound

Sound classification



Text

Text classification
Word tagging



Tabular

Tabular classification
Tabular regression

these are also built into TuriCreate!

could be good for final projects!!!

Task focused toolkits

Recommender Systems

Image Classification

Drawing Classification

Sound Classification

How it works

Advanced Usage

Deployment to Core ML

Image Similarity

Object Detection

One-Shot Object Detection

Style Transfer

Activity Classification

Text Classifier

How Does This Work?

Training and making predictions for a sound classifier model is a three stage process:

1. Signal preprocessing
2. A pretrained neural network is used to extract deep features
3. A custom neural network is used to make the predictions

Details below about each stage.

At a high level, the preprocessing pipeline does the following:

- The raw pulse code modulation data from the wav file is converted to floats on a [-1.0, +1.0] scale.
- If there are two channels, the elements are averaged to produce one channel.
- The data is resampled to only 16,000 samples per second.
- The data is broken up into several overlapping windows.
- A [Hamming Window](#) is applied to each windows.
- The [Power Spectrum](#) is calculated, using a [Fast Fourier Transformation](#).
- Frequencies above and below certain thresholds are dropped.
- [Mel Frequency Filter Banks](#) are applied.
- Finally the natural logarithm is taken of all values.

The preprocessing pipeline takes 975ms worth of audio as input (exact input length depends on sample rate) and produces an array of shape (96, 64).

https://apple.github.io/turicreate/docs/userguide/sound_classifier/

could be good for final projects!!!

Task focused toolkits

Recommender Systems

Image Classification

Drawing Classification

Sound Classification

Image Similarity

Object Detection

One-Shot Object Detection

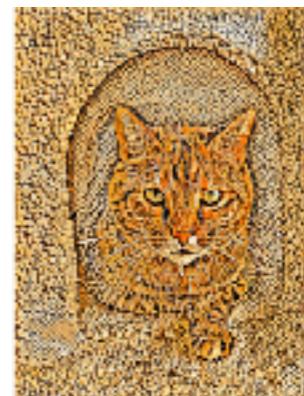
[Style Transfer](#)

How it works

Deployment to Core ML

Activity Classification

Text Classifier



Style transfer model

The technique used in Turi Create is based on "["A Learned Representation For Artistic Style"](#)". The model is compact and fast and hence can run on mobile devices like an iPhone. The model consists of 3 convolutional layers, 5 residual layers (2 convolutional layers in each) and 3 upsampling layers each followed by a convolutional layer. There are a total of 16 convolutional layers.

There are three aspects about this technique that are worth noting:

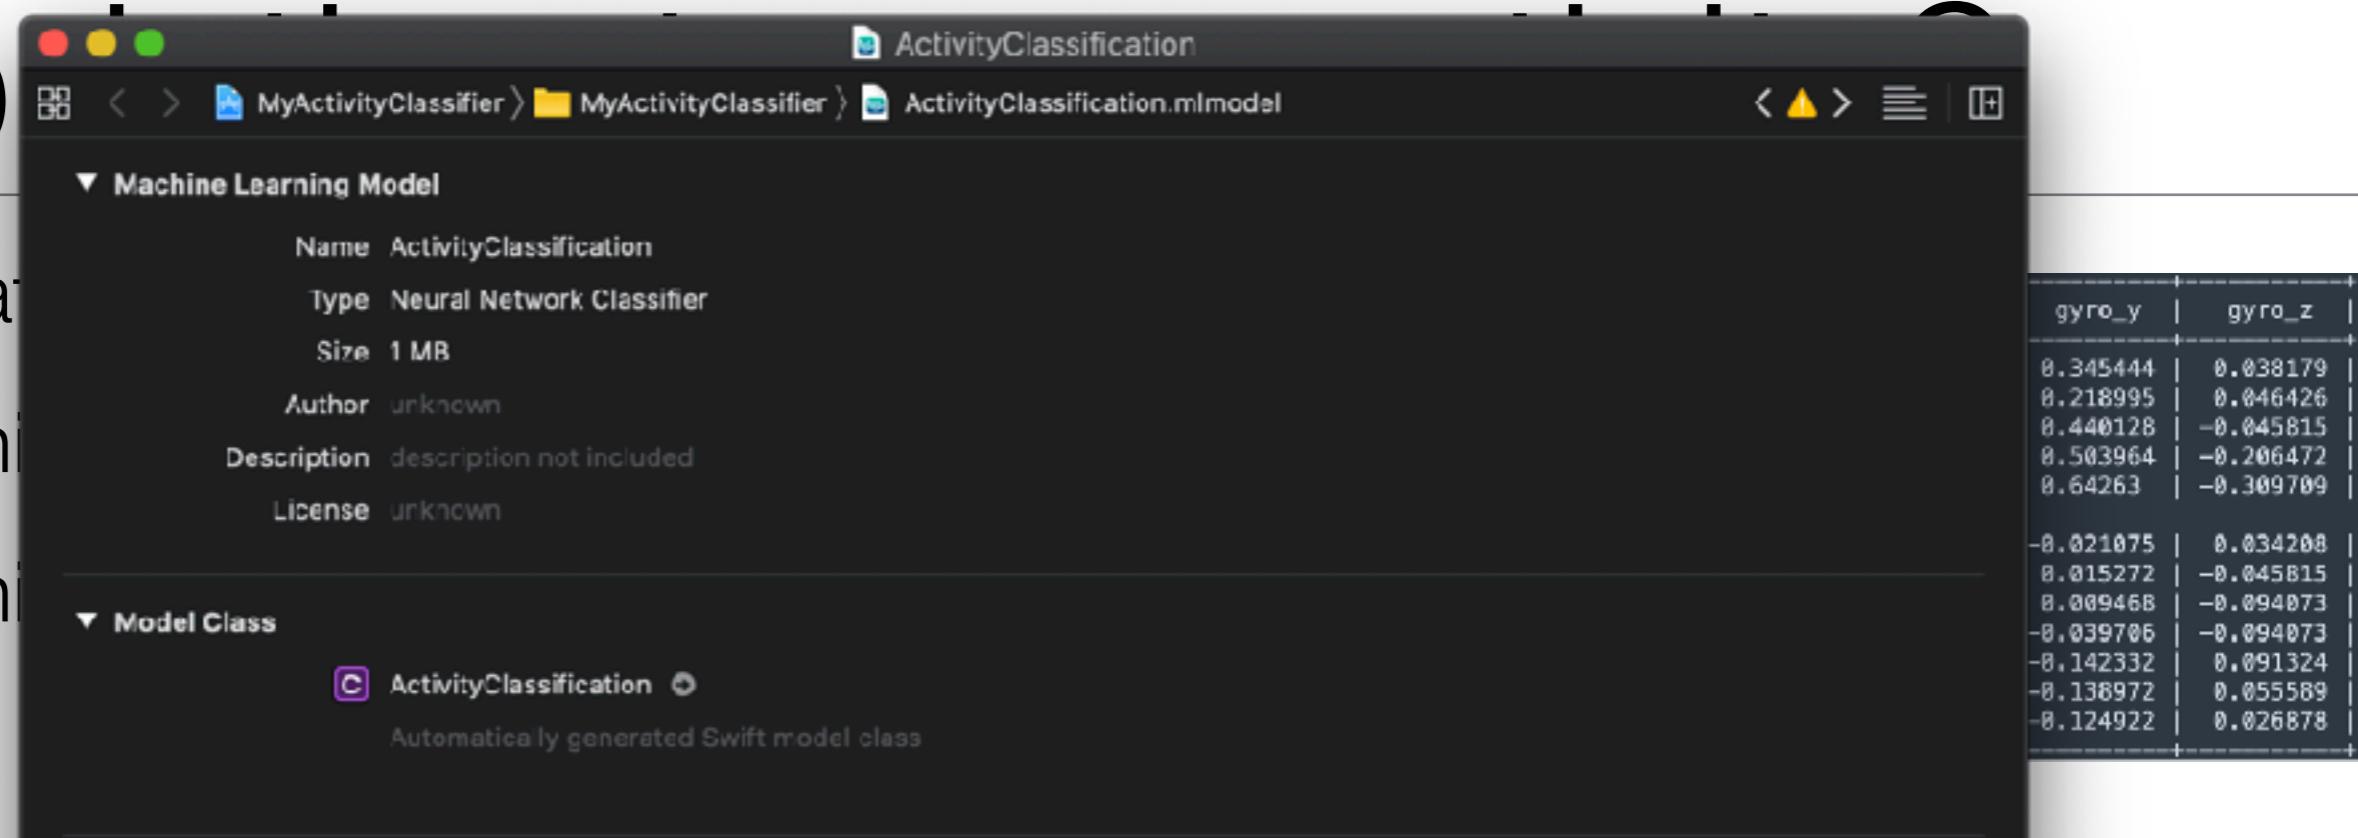
- It is designed to be incredibly fast at stylizing images, allowing deployment on device. As a trade off, the model creation takes longer.
- A single model can incorporate a large number of styles without any significant increase in the size of the model.
- The model can take input of any size and output a stylized image of the same size.

During training, we employ [Transfer Learning](#). The model uses the visual semantics of an already trained VGG-16 network to understand and mimic stylistic elements.

https://apple.github.io/turicreate/docs/userguide/sound_classifier/

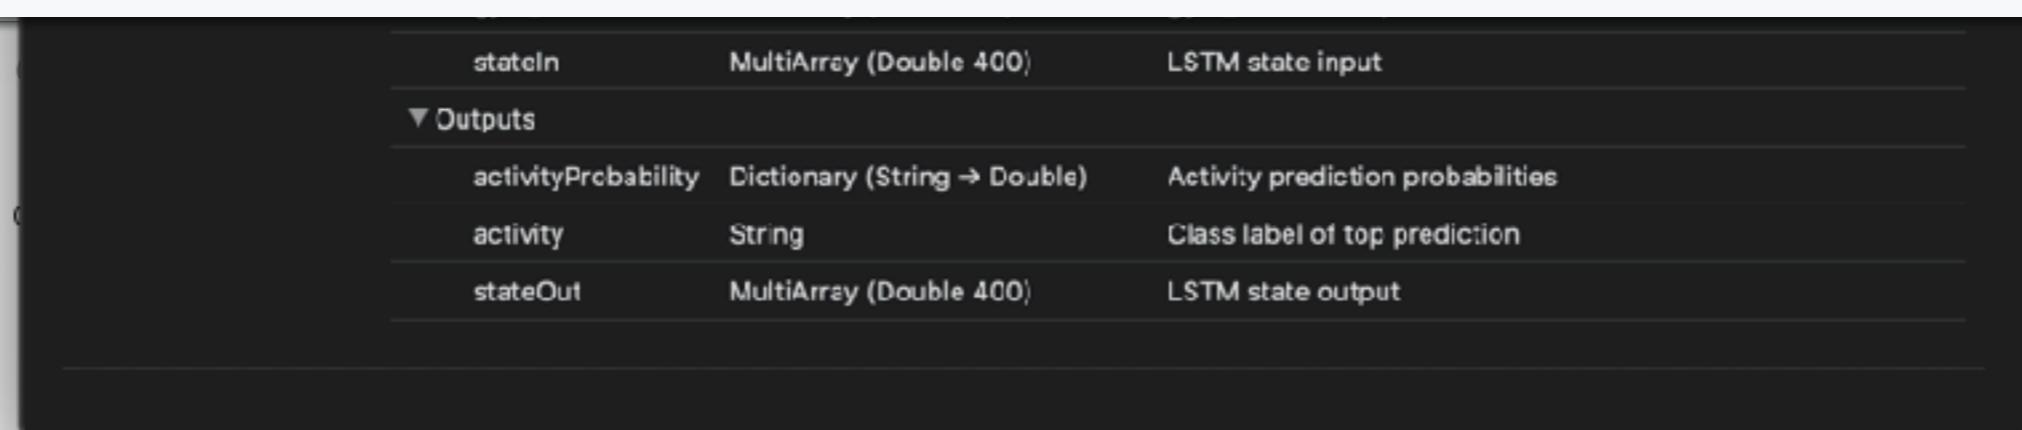
up

- format
- organize
- organize
- train



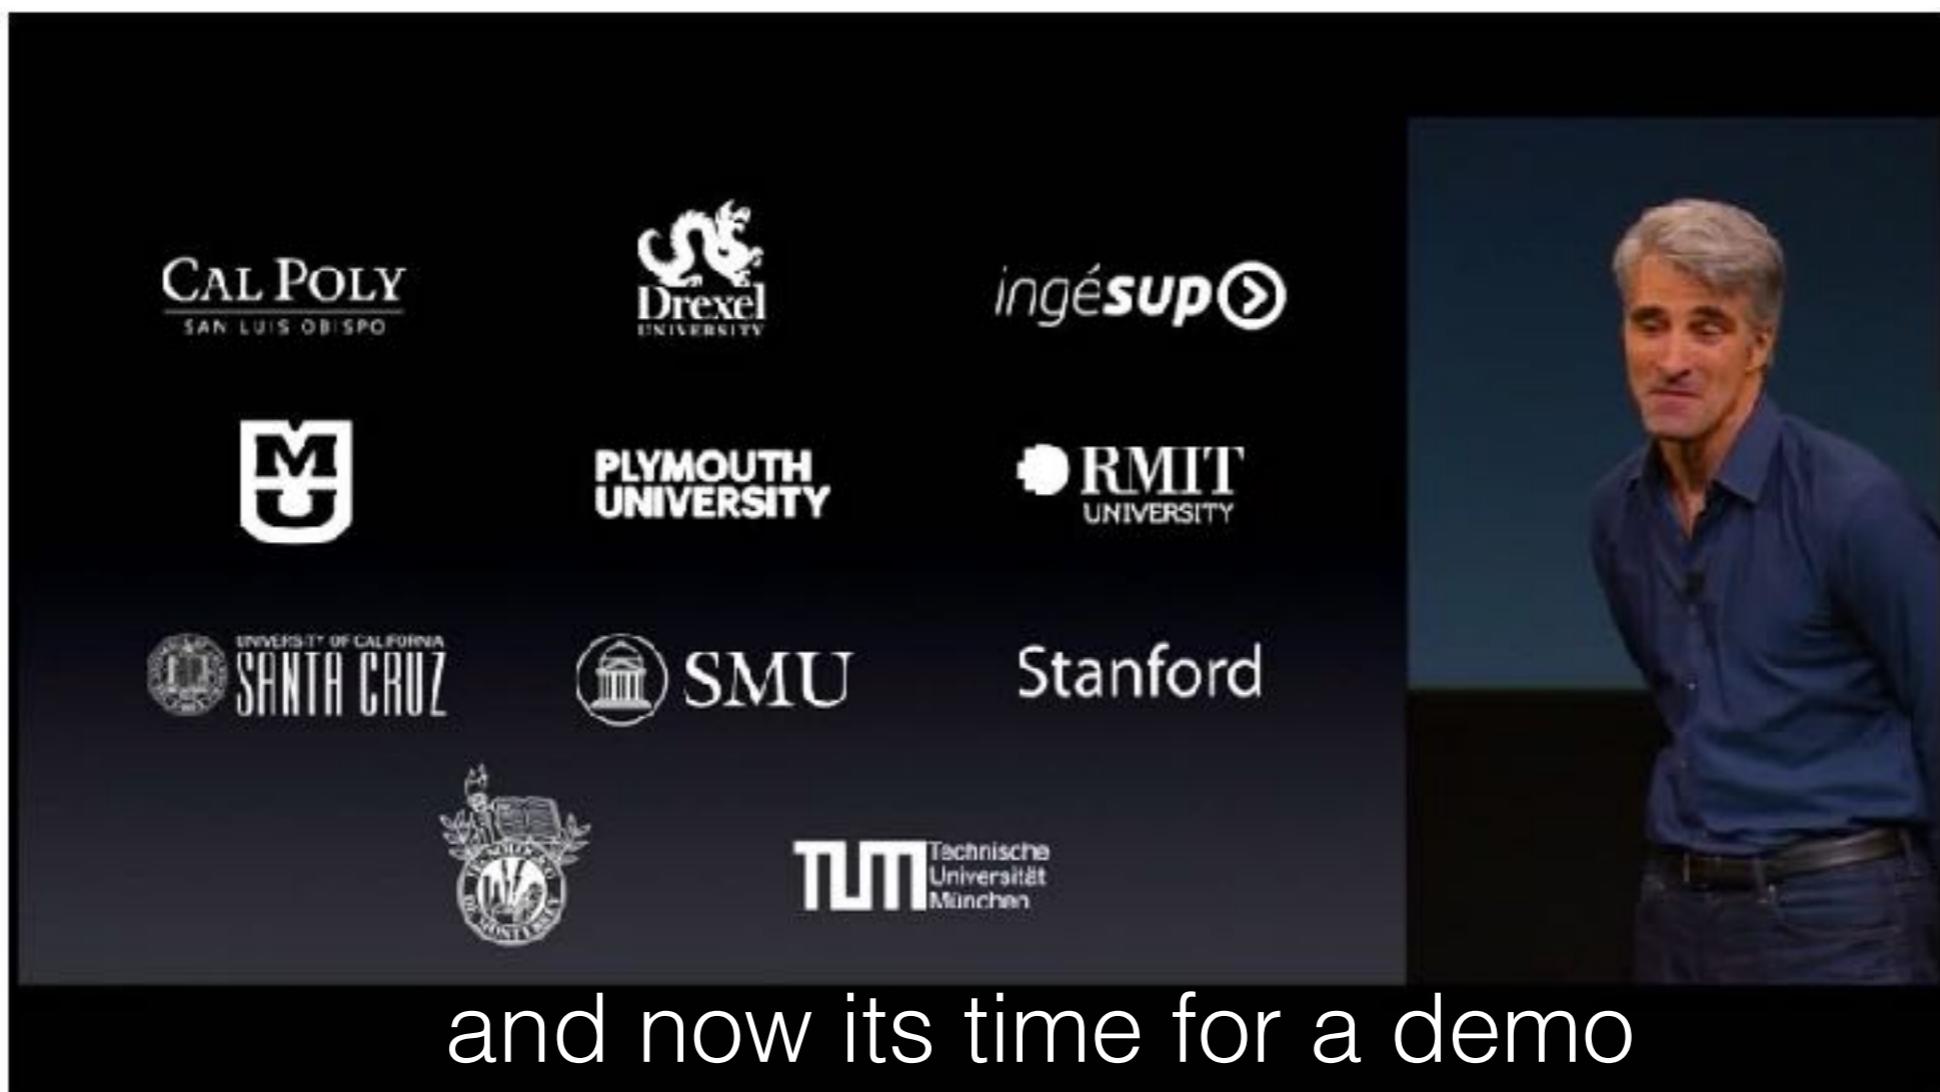
```
// Perform model prediction
let modelPrediction = try!
    activityClassificationModel.prediction(acc_x: accelDataX, acc_y: accelDataY, acc_z: accelDataZ,
                                             gyro_x: gyroDataX, gyro_y: gyroDataY, gyro_z: gyroDataZ, stateIn: stateOutput)

// Update the state vector
stateOutput = modelPrediction.stateOut
```



CreateML

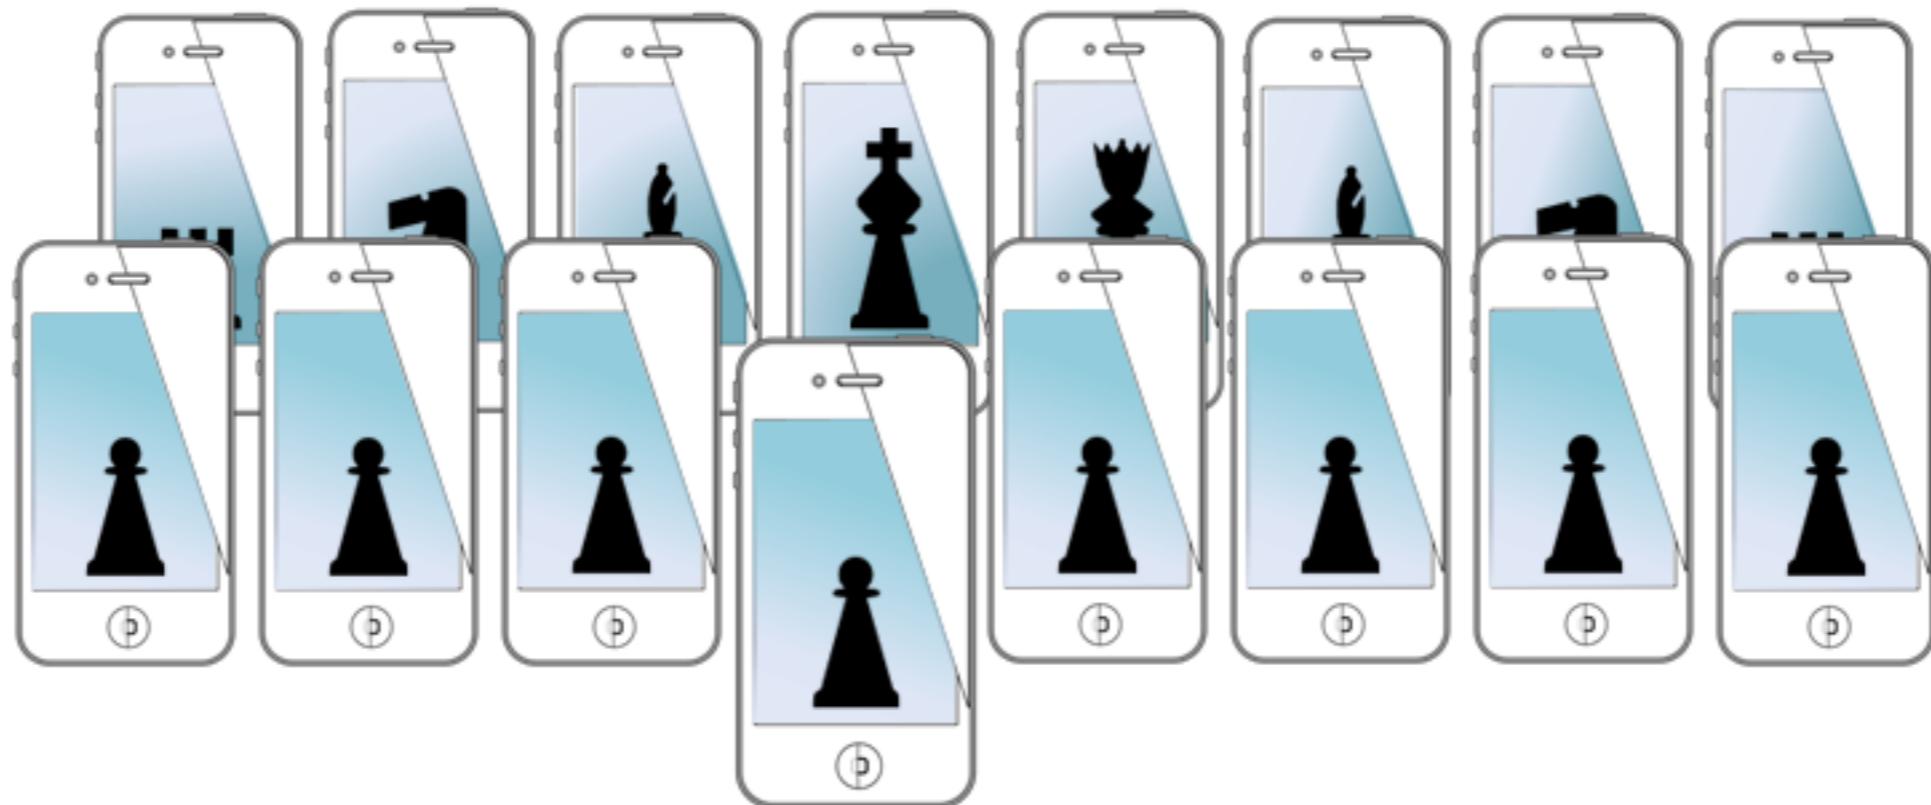
- extended demo from beginning to end
- combining our tornado, mongo, iOS, CoreML



And now...

- final project discussion
- and proposal!!

MOBILE SENSING LEARNING



CS5323 & 7323
Mobile Sensing and Learning

machine learning crash course

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University