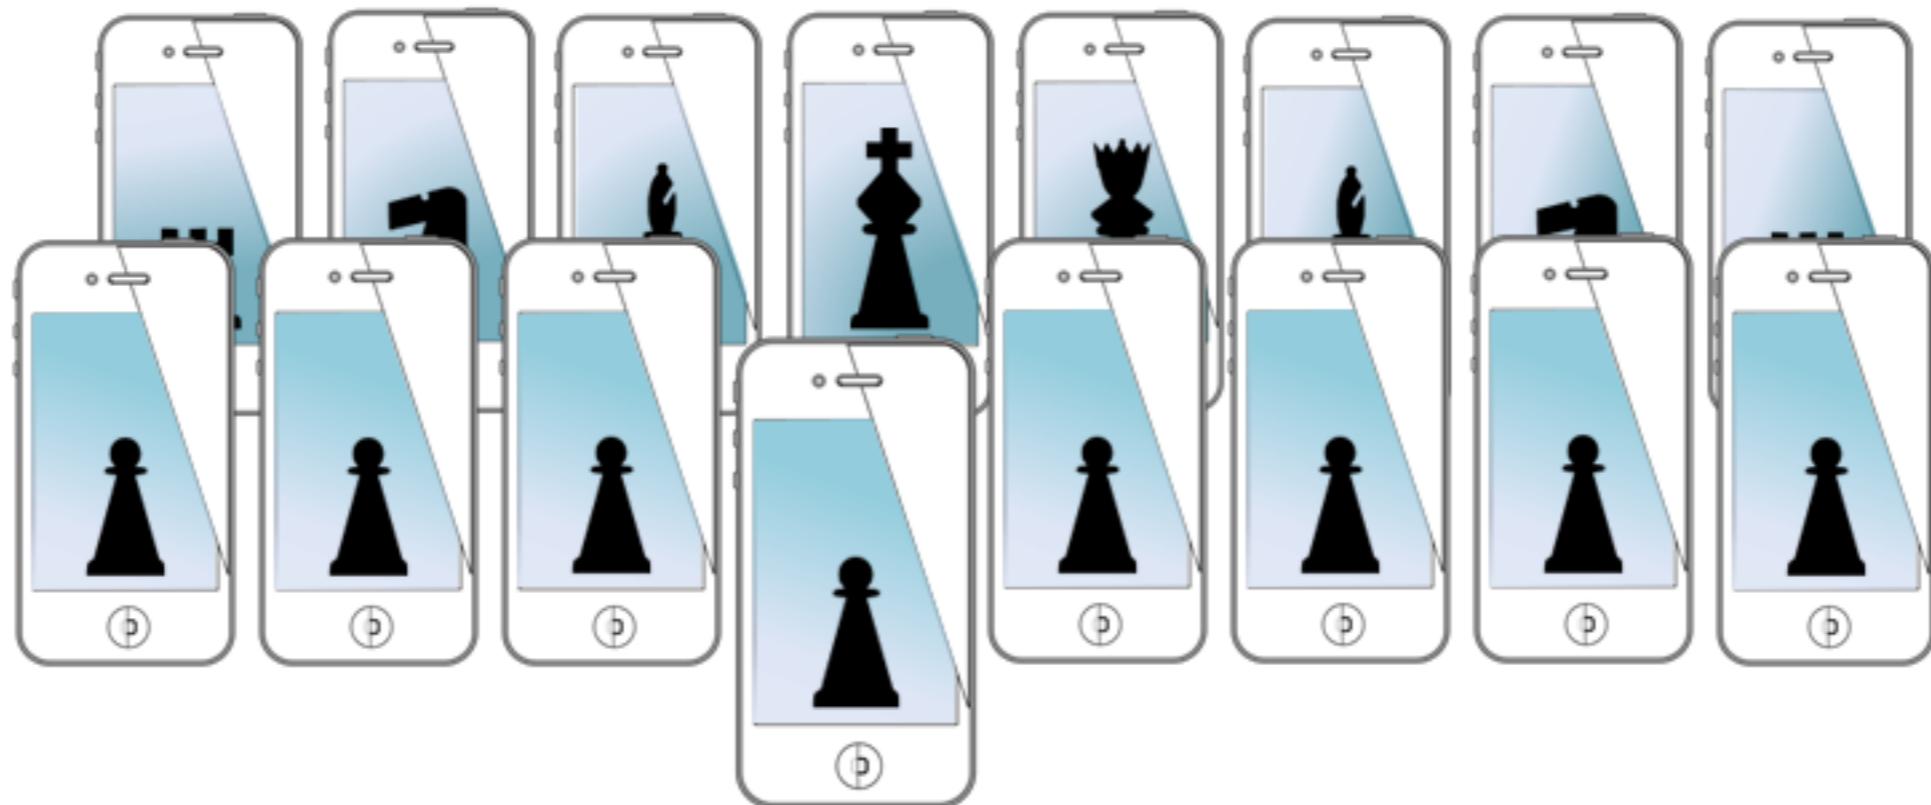


# MOBILE SENSING LEARNING



**CS5323 & 7323**  
Mobile Sensing and Learning

CoreML and ARKit

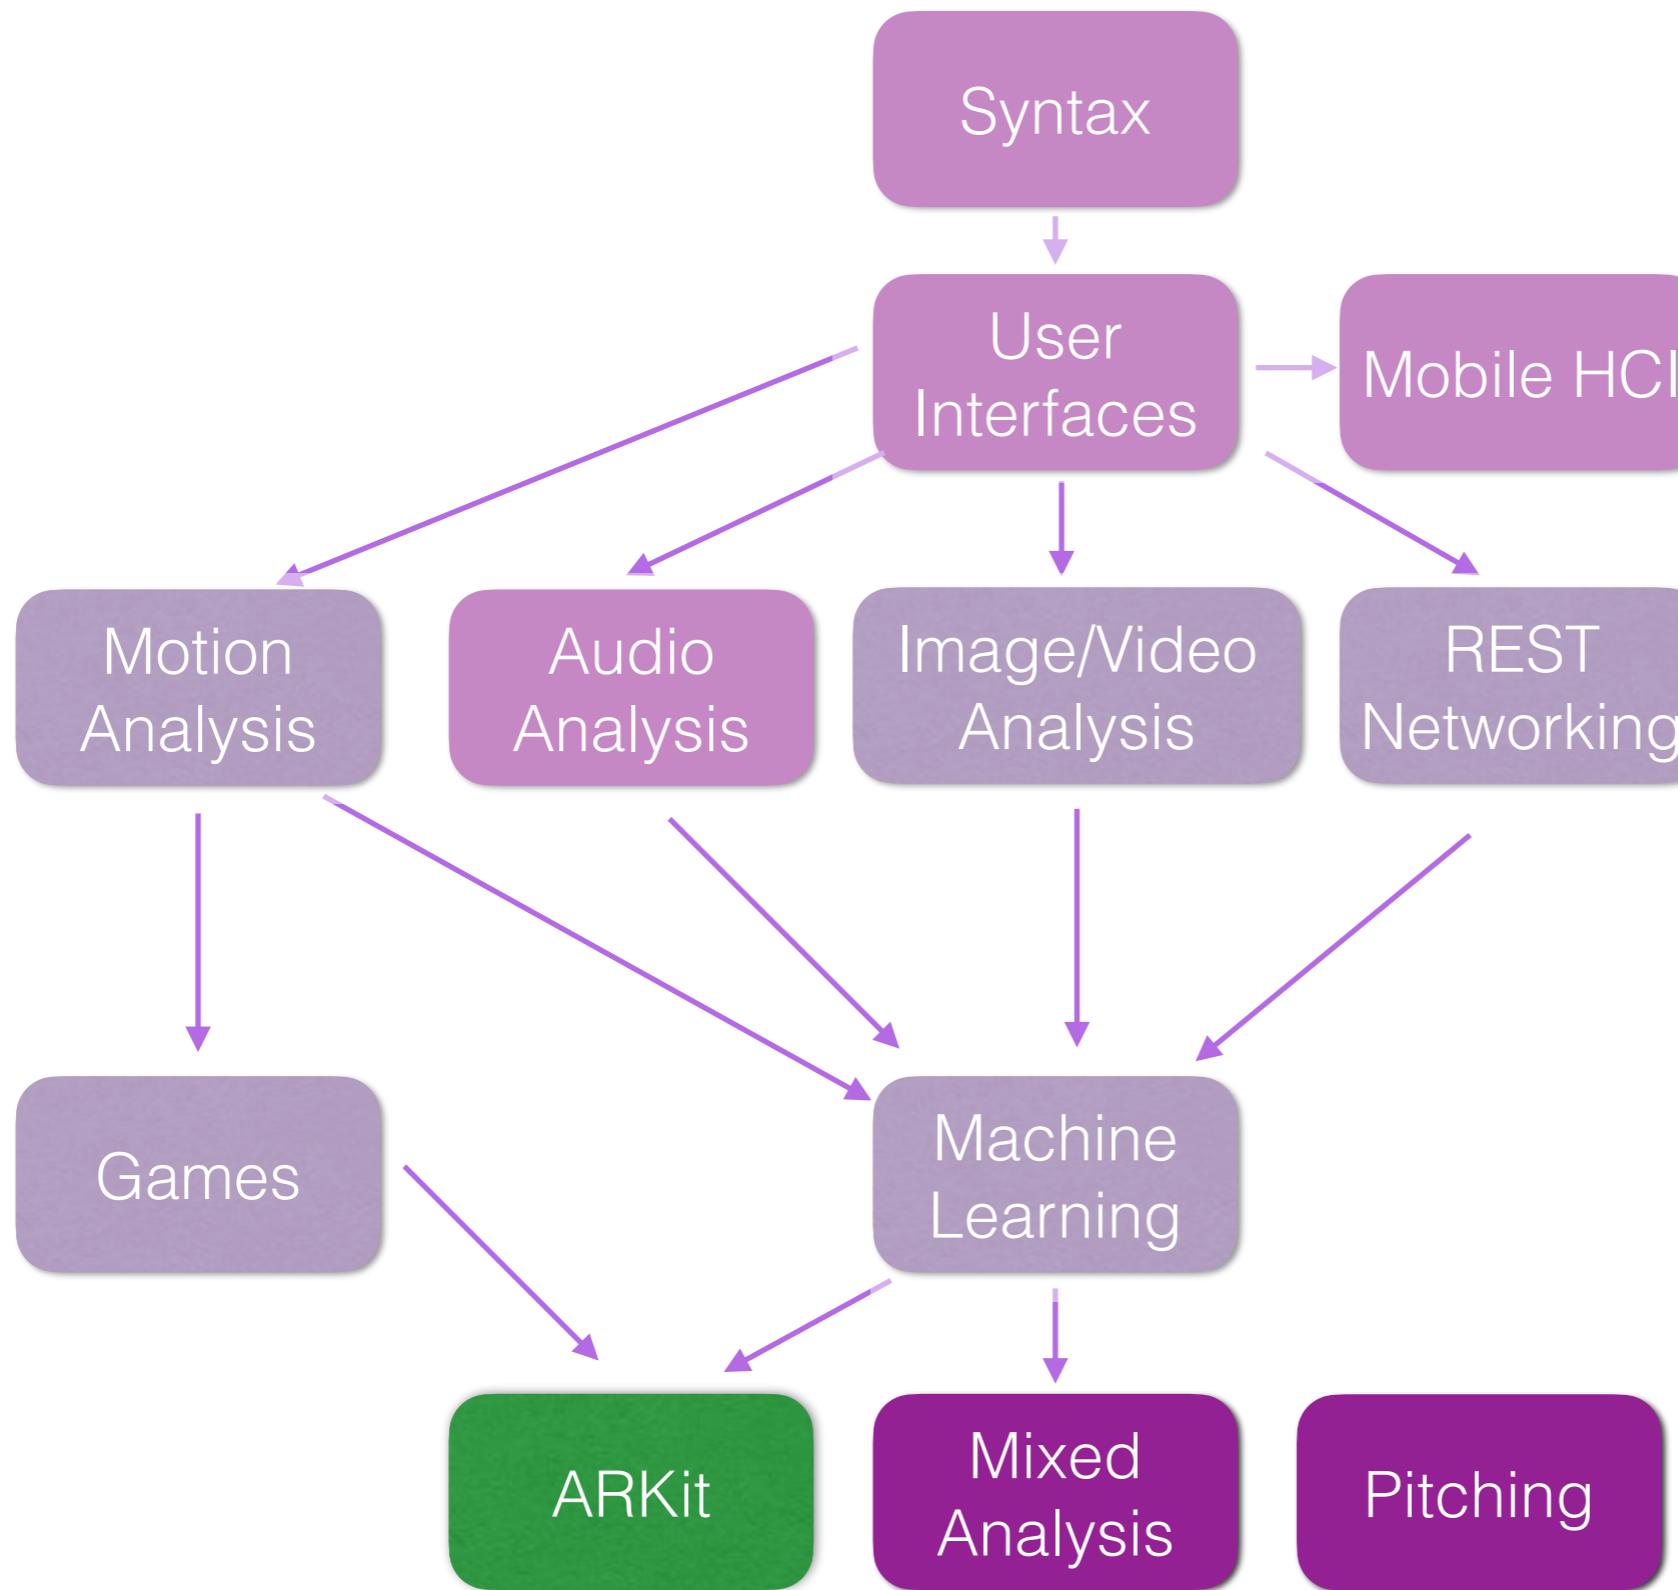
Eric C. Larson, Lyle School of Engineering,  
Computer Science, Southern Methodist University

# course logistics and agenda

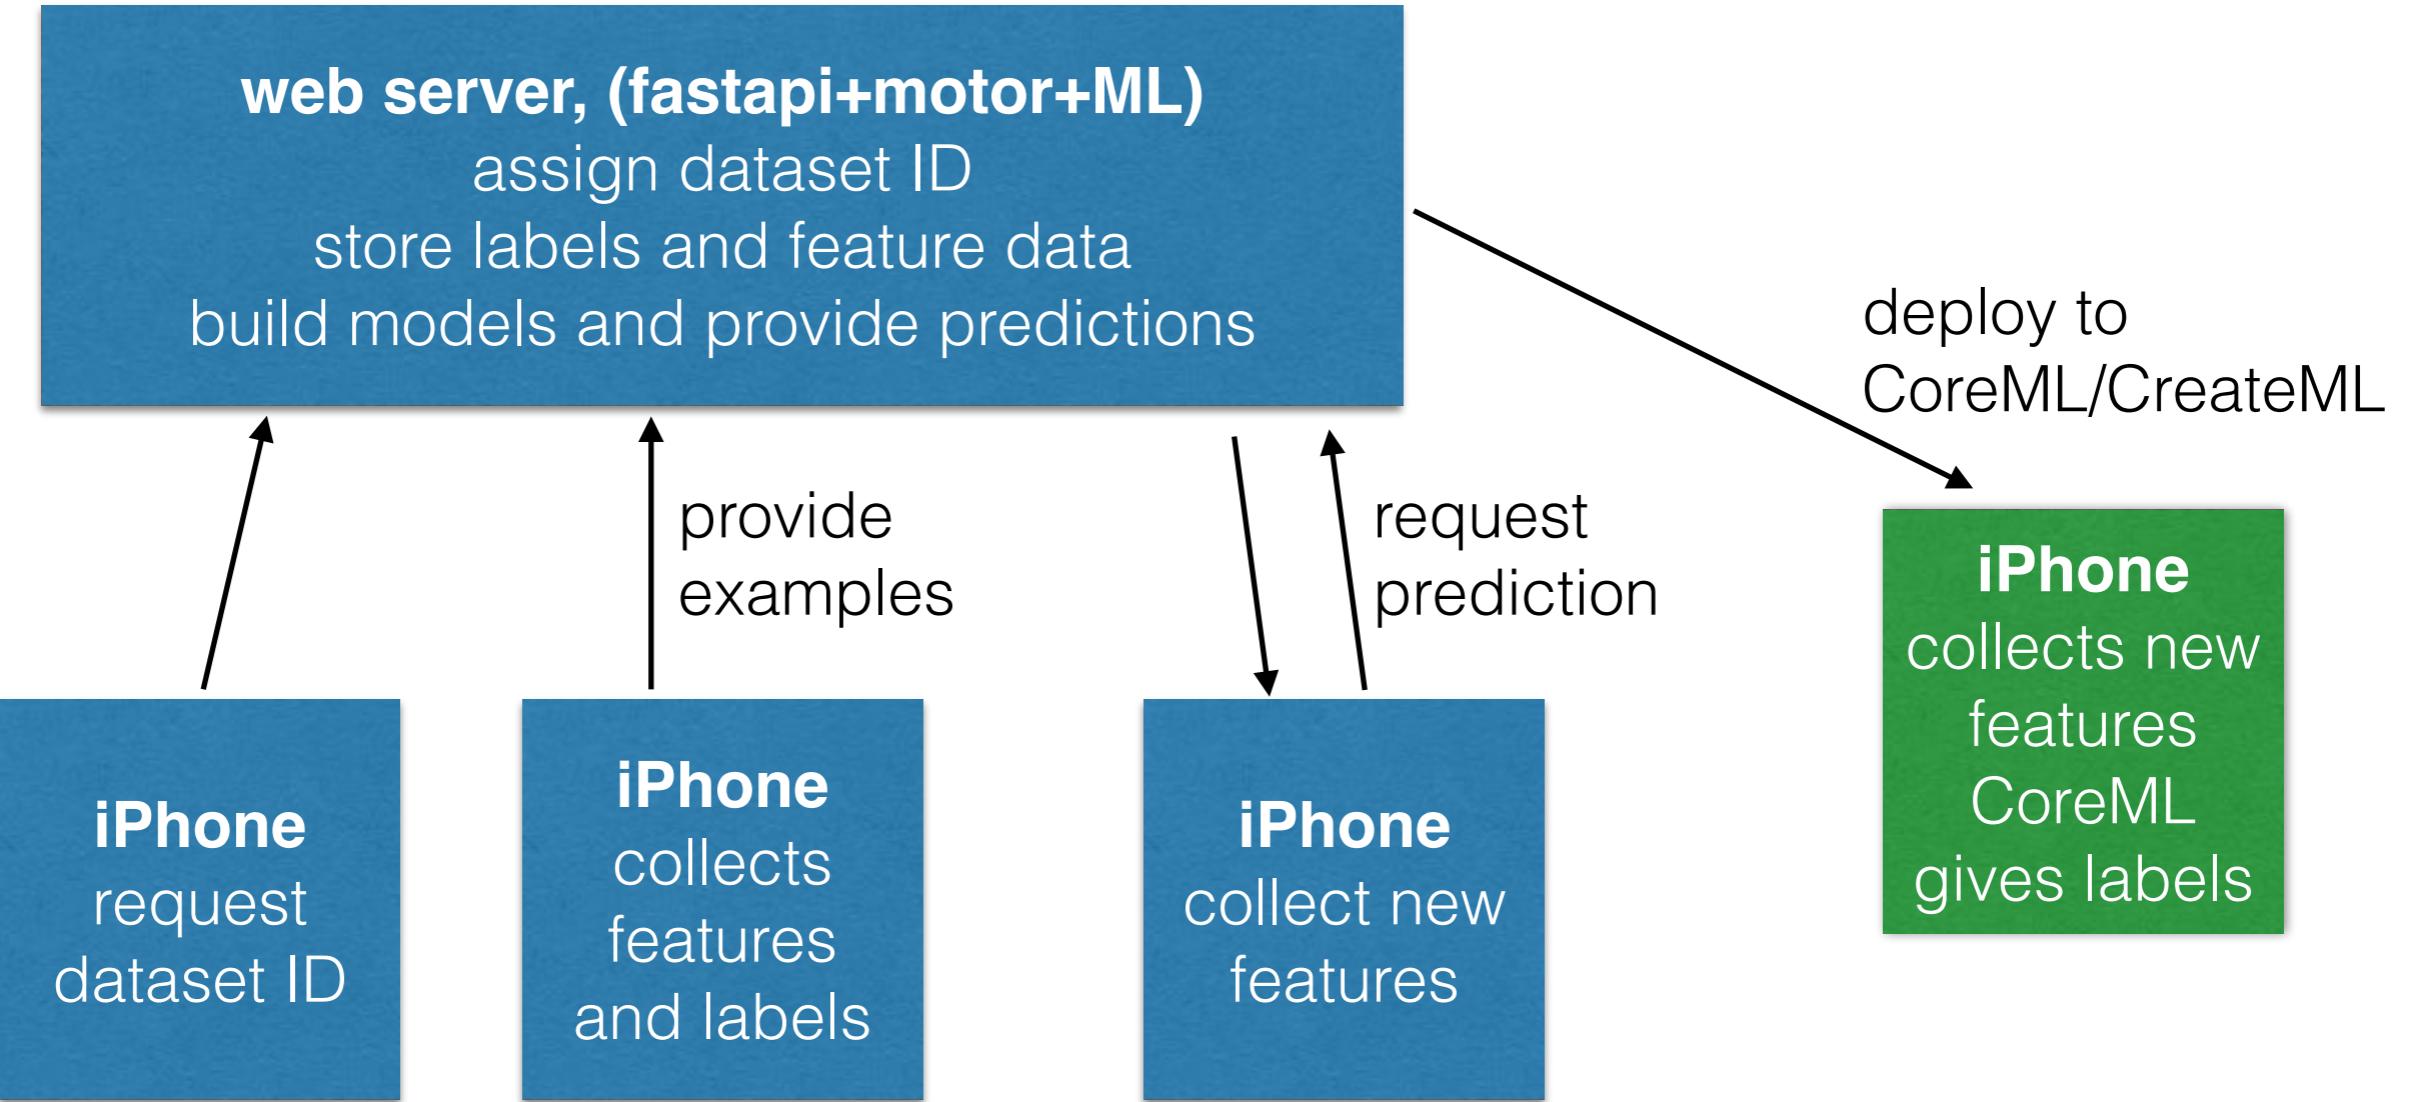
- agenda:
  - ARKit
  - ARKit demo
  - ARKit Object Recognition
  - If time: Using YOLO with ARKit integration

<https://developer.apple.com/videos/play/wwdc2017/602/>

# class overview



# assignment 5 demo



all uploads handled via JSON formatted HTTP requests  
most of code is already functional

# Lab Five and Proposals

- Questions on lab five
- Questions on proposals

## Members:

**Design:** No, NOT opting into the MOD

**Description:** I want to develop a guitar chord recognition program that allows users to identify various guitar chords through their mobile phones. The reason for the design is that guitar players often cannot find the right chord in the process of learning the guitar. When a user finds good music, they can use chord analysis software to analyze how to play accompaniments. The app will analyze a guitar chord in real time, send the data to the server for online analysis, and display the name and fingering of the current guitar chord at the bottom of the UI.

There is an app on the app store that does this, called Chord AI. It is unclear if the app is very reliable.

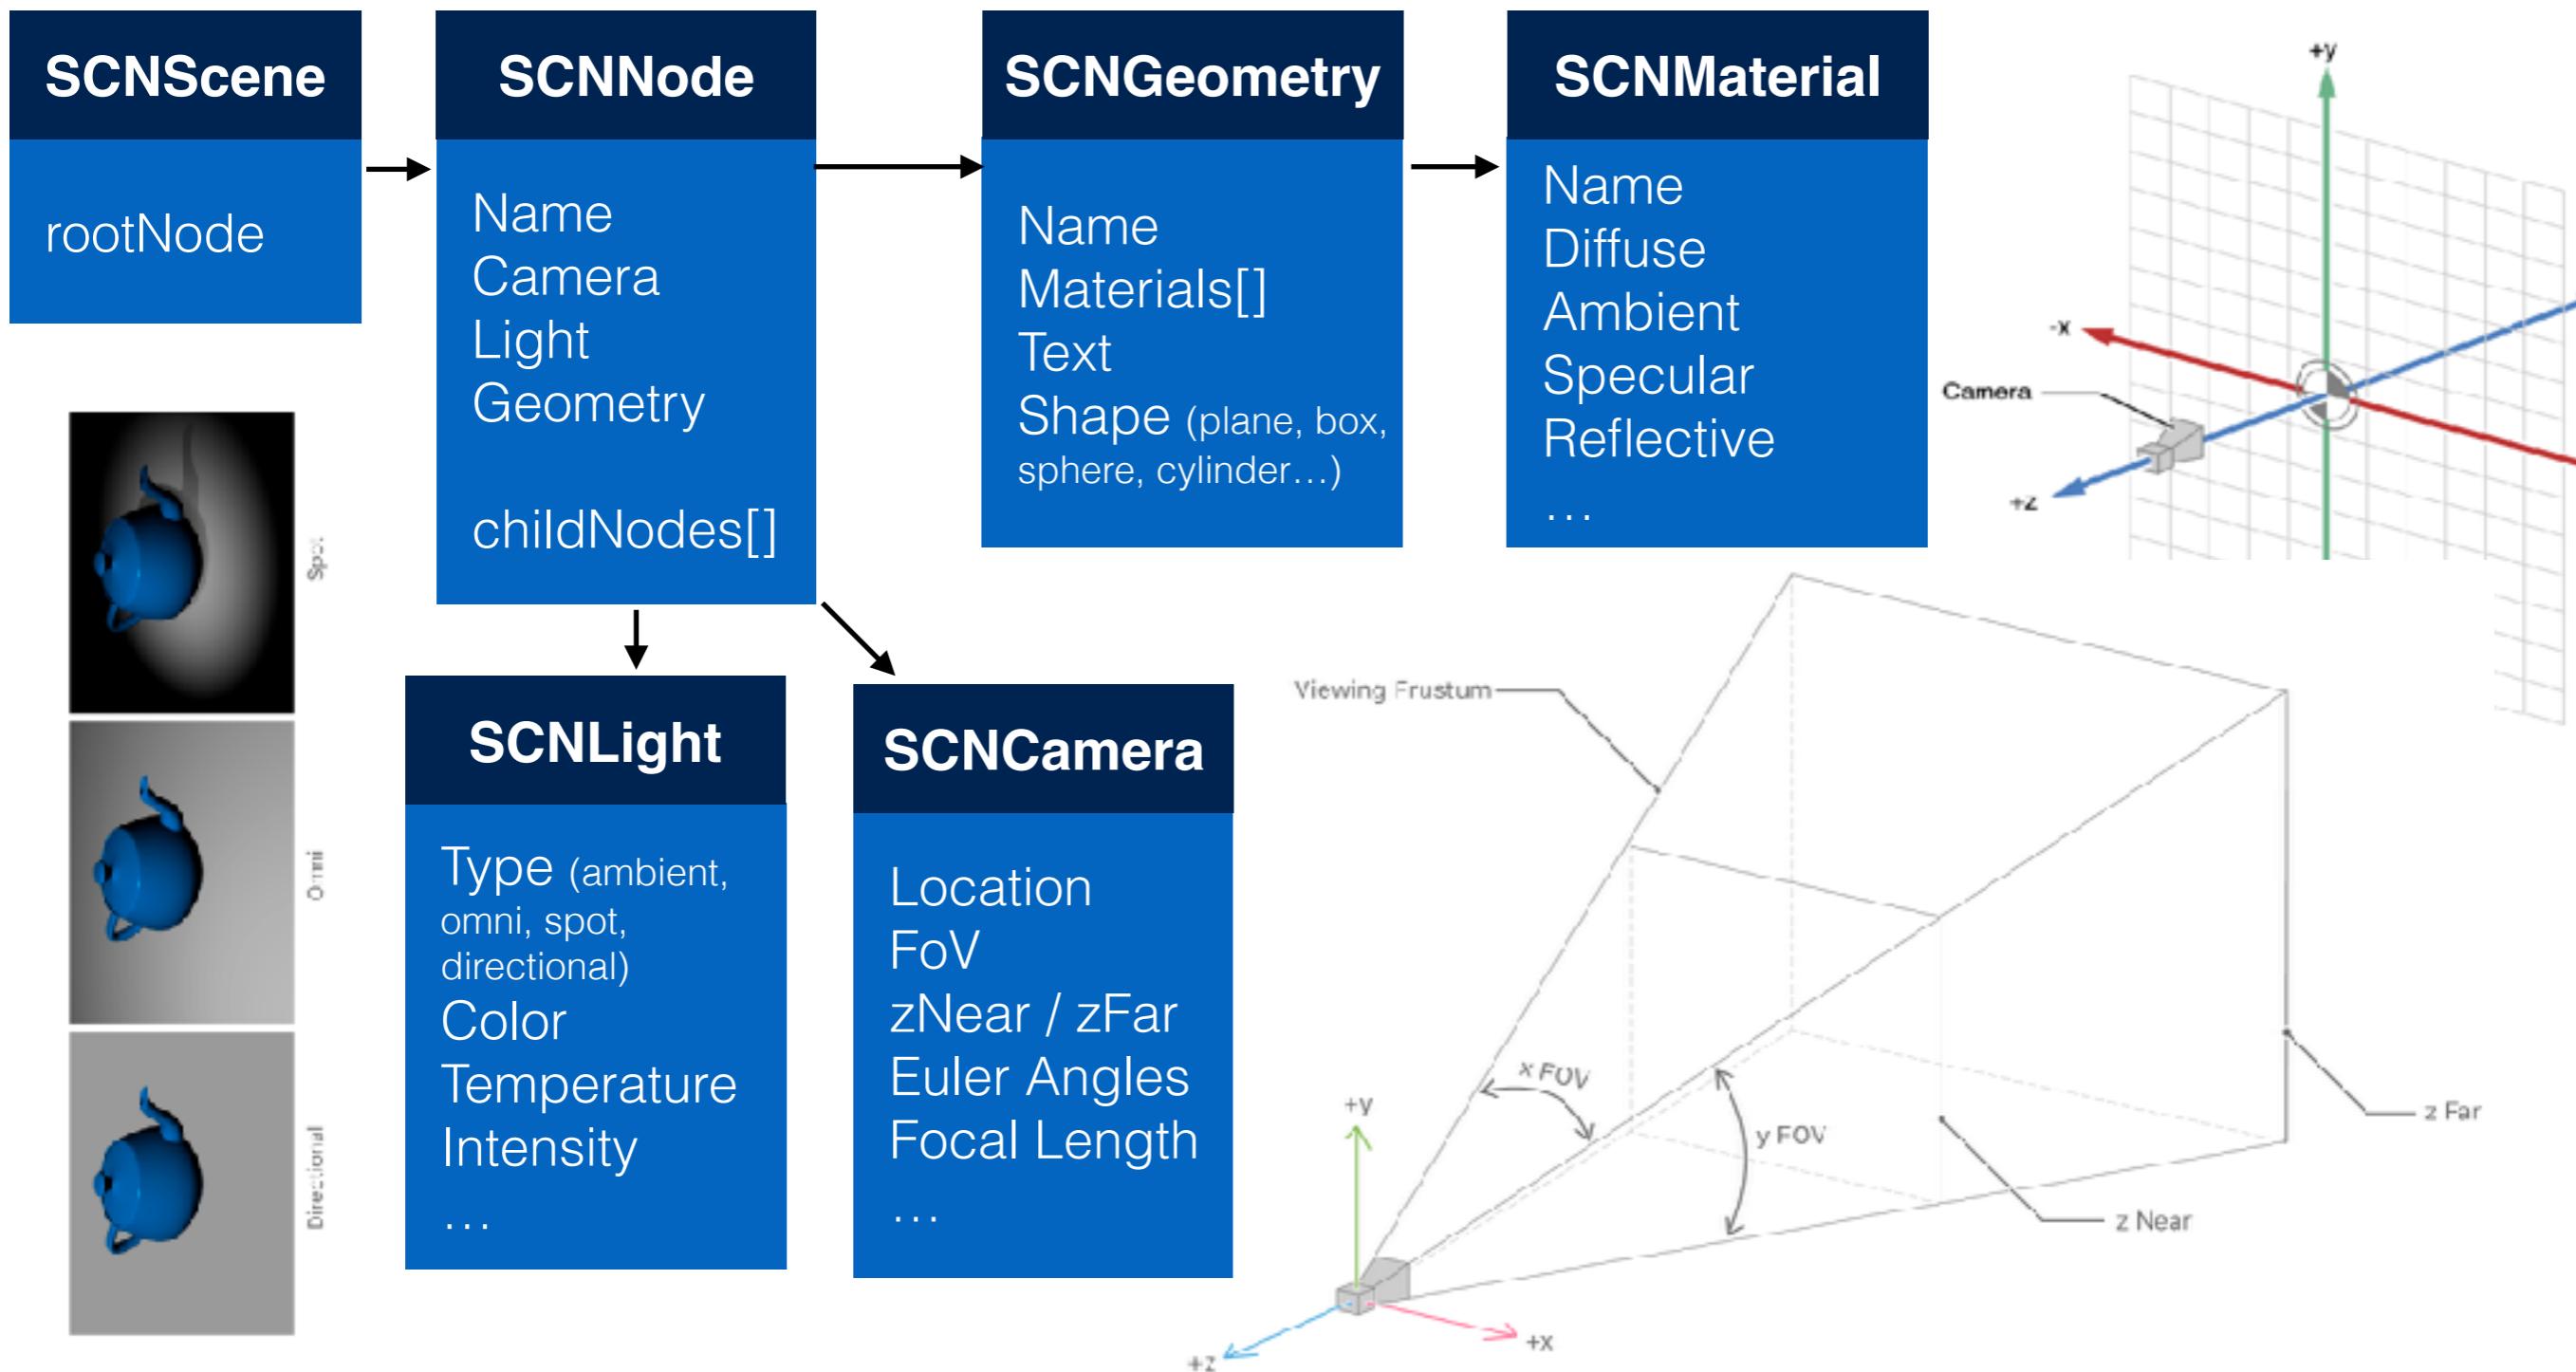
## Constraints:

1. The program will be able to detect and analyze the guitar chords collected by the microphone. Since there are nearly 1000 guitar chords, the accuracy of the analyzed chords will mainly be the main chords that are common and used by most people. These include major and minor variants of different base notes including the sharps and flats. The chord will always have the main note in the base (i.e., no inversions).
2. The only interface of this program is the logic of real-time analysis, so the Prediction function will be automatically opened and analyzed on the main page, always running.
3. The real-time analysis audio interface of the program will also display the guitar fingering diagram corresponding to the actual chord.
4. The program will use ML as a Service to run as a server and will send data HTTP requests to ML on the server to request prediction results.

## Elements of Labs:

1. UI
2. Audio Processing
3. Machine Learning as a Service

# work flow in 3D scenes



SCNNode is the base for nearly everything in simulation env.

but... we want the SCNScene  
**to be the real world**  
ARKit anchors the SCNScene to  
objects in the immediate environment

```
// Setup view
let view = self.view as SCNView
view.scene = scene
```

```
// Setup view
let view = self.view as ARSCNView
view.scene = scene
```

# ARKit basics



Tracking

World tracking  
Visual inertial odometry  
No external setup



Rendering

Easy integration  
AR views  
Custom rendering



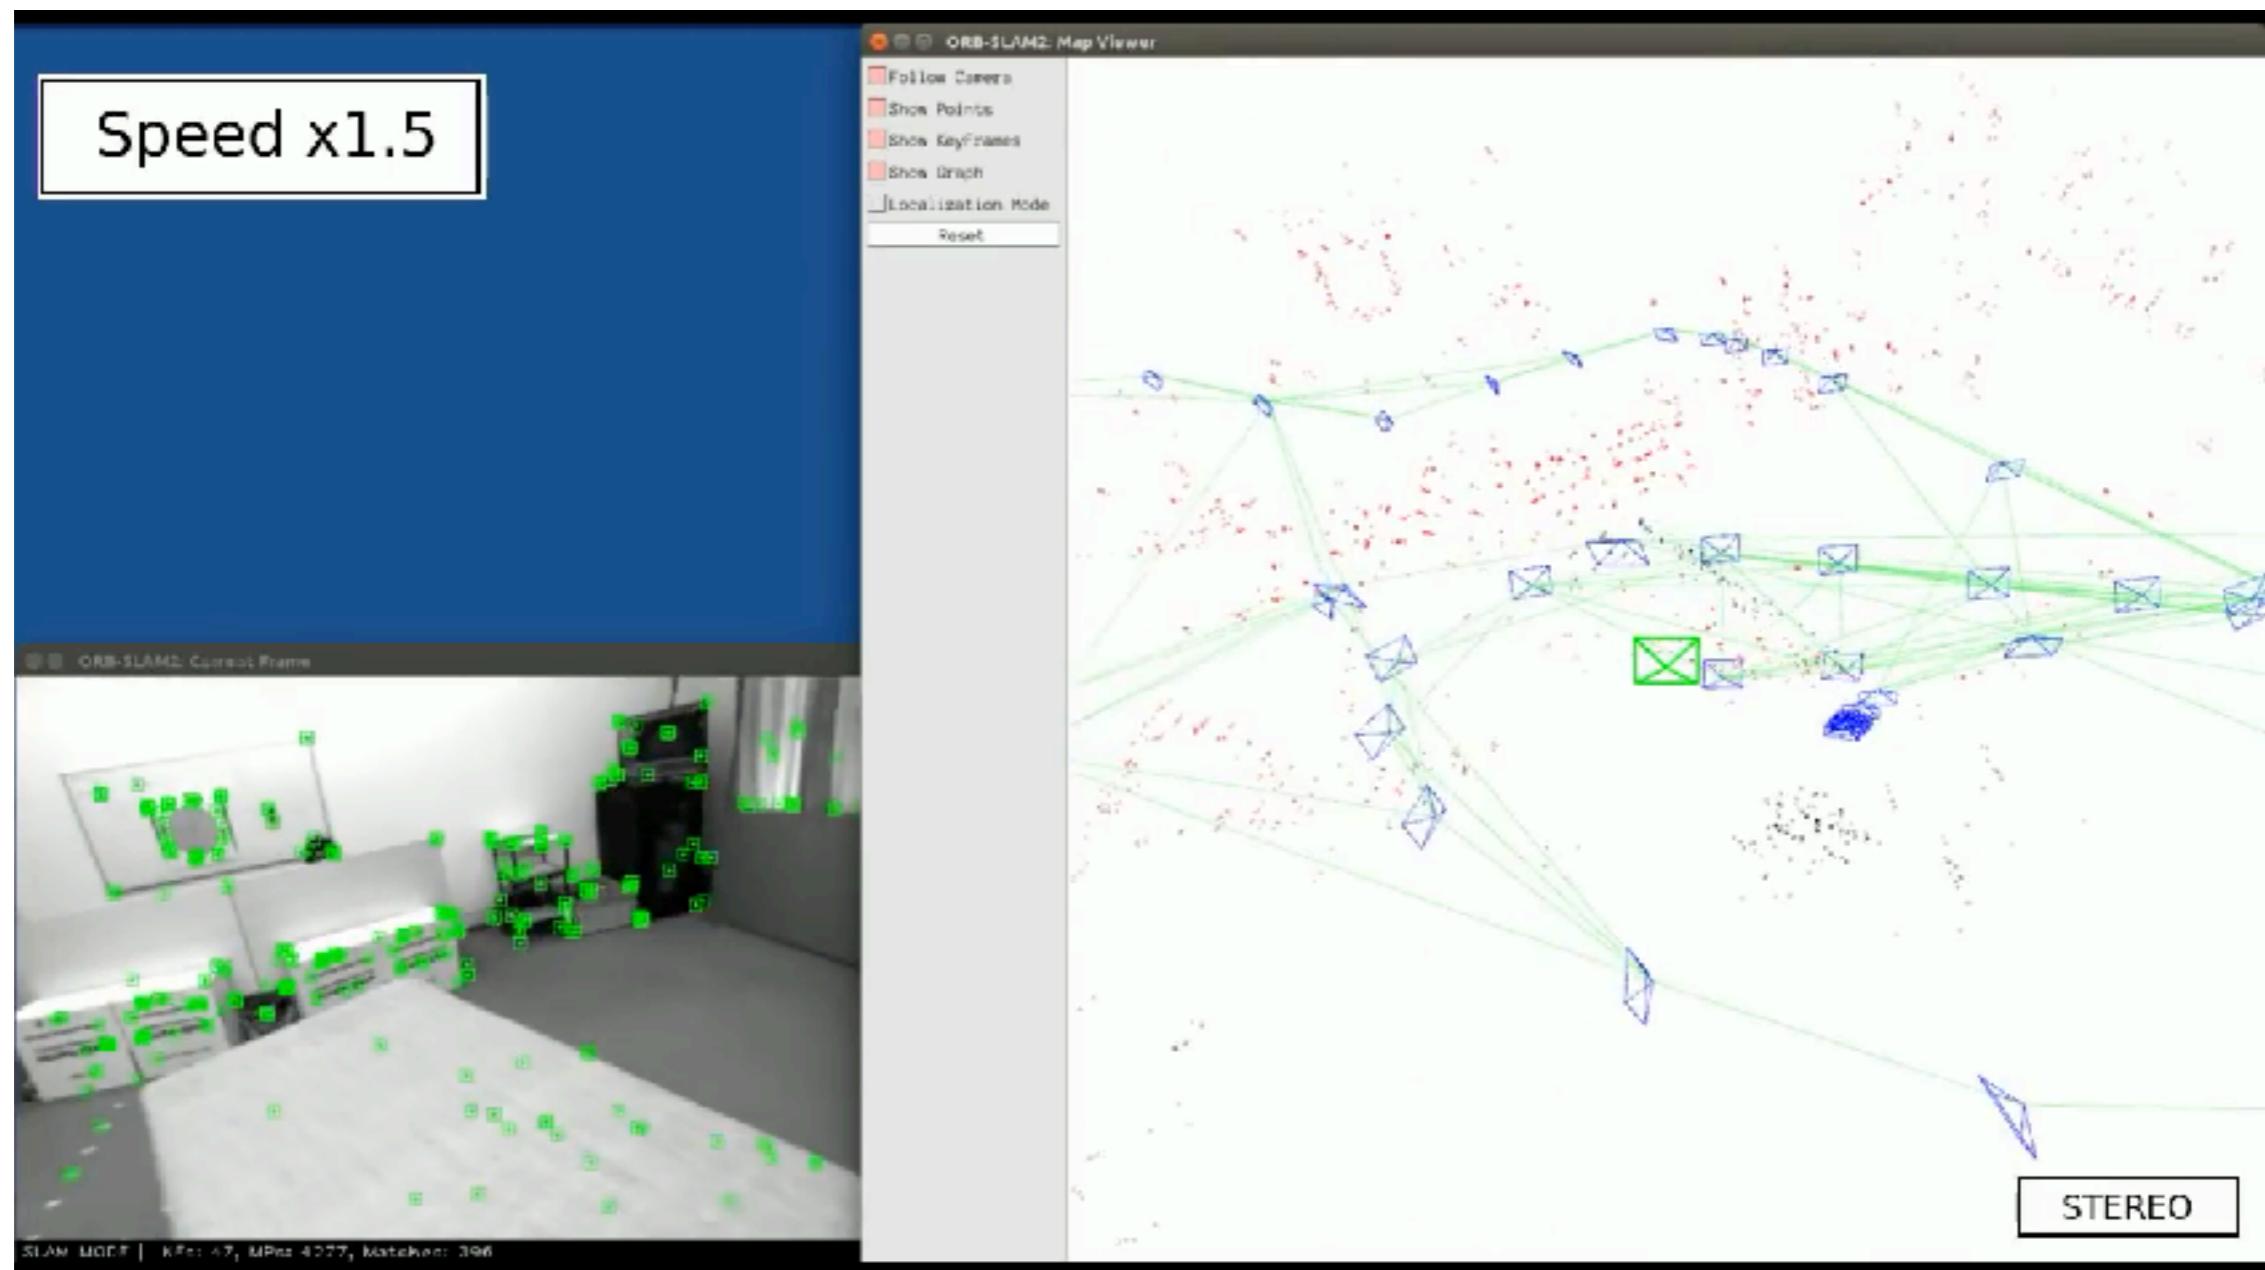
Scene Understanding

Plane detection  
Hit-testing  
Light estimation

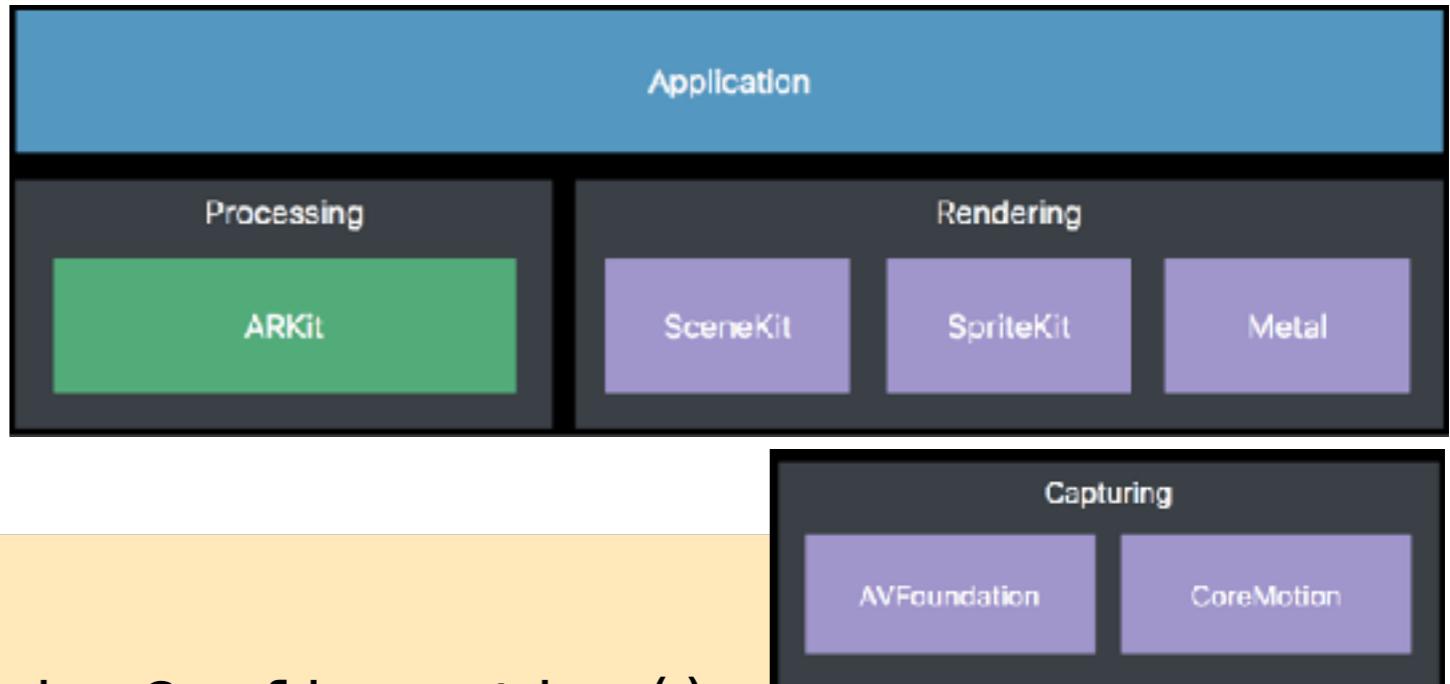
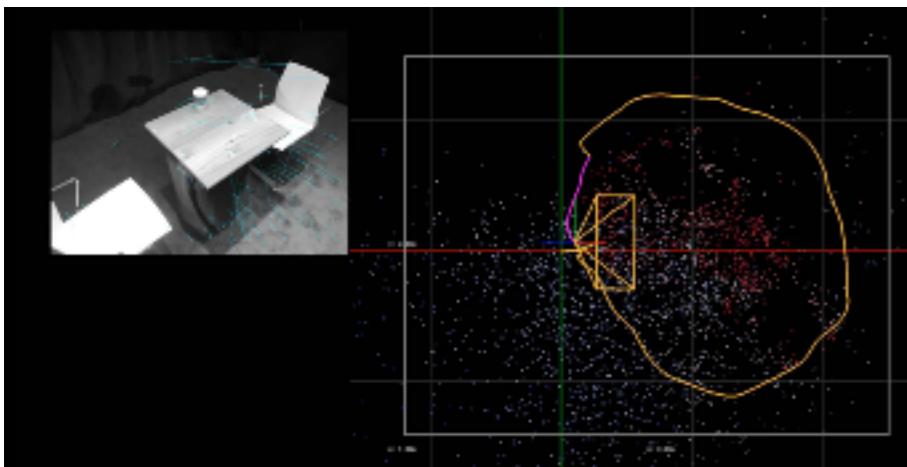
[https://  
developer.apple.  
com/videos/play/  
wwdc2017/602/](https://developer.apple.com/videos/play/wwdc2017/602/)

# SLAM

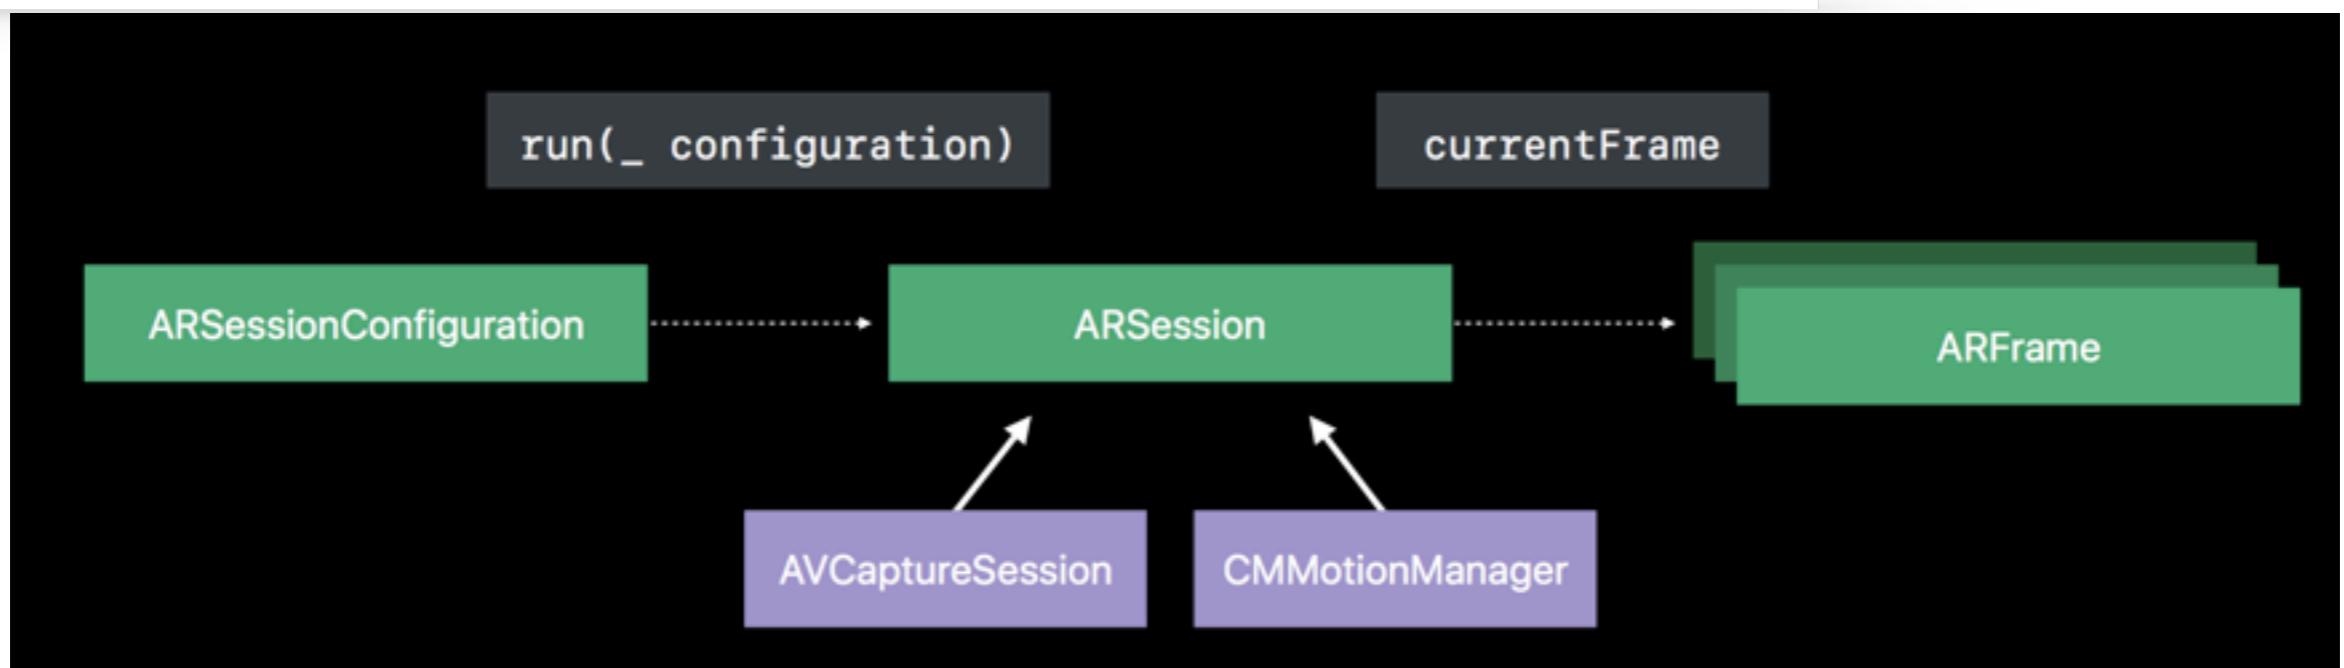
- ORB-SLAM: Uses Key points in Image
- Simultaneous Localization and Mapping



# ARKit system integration



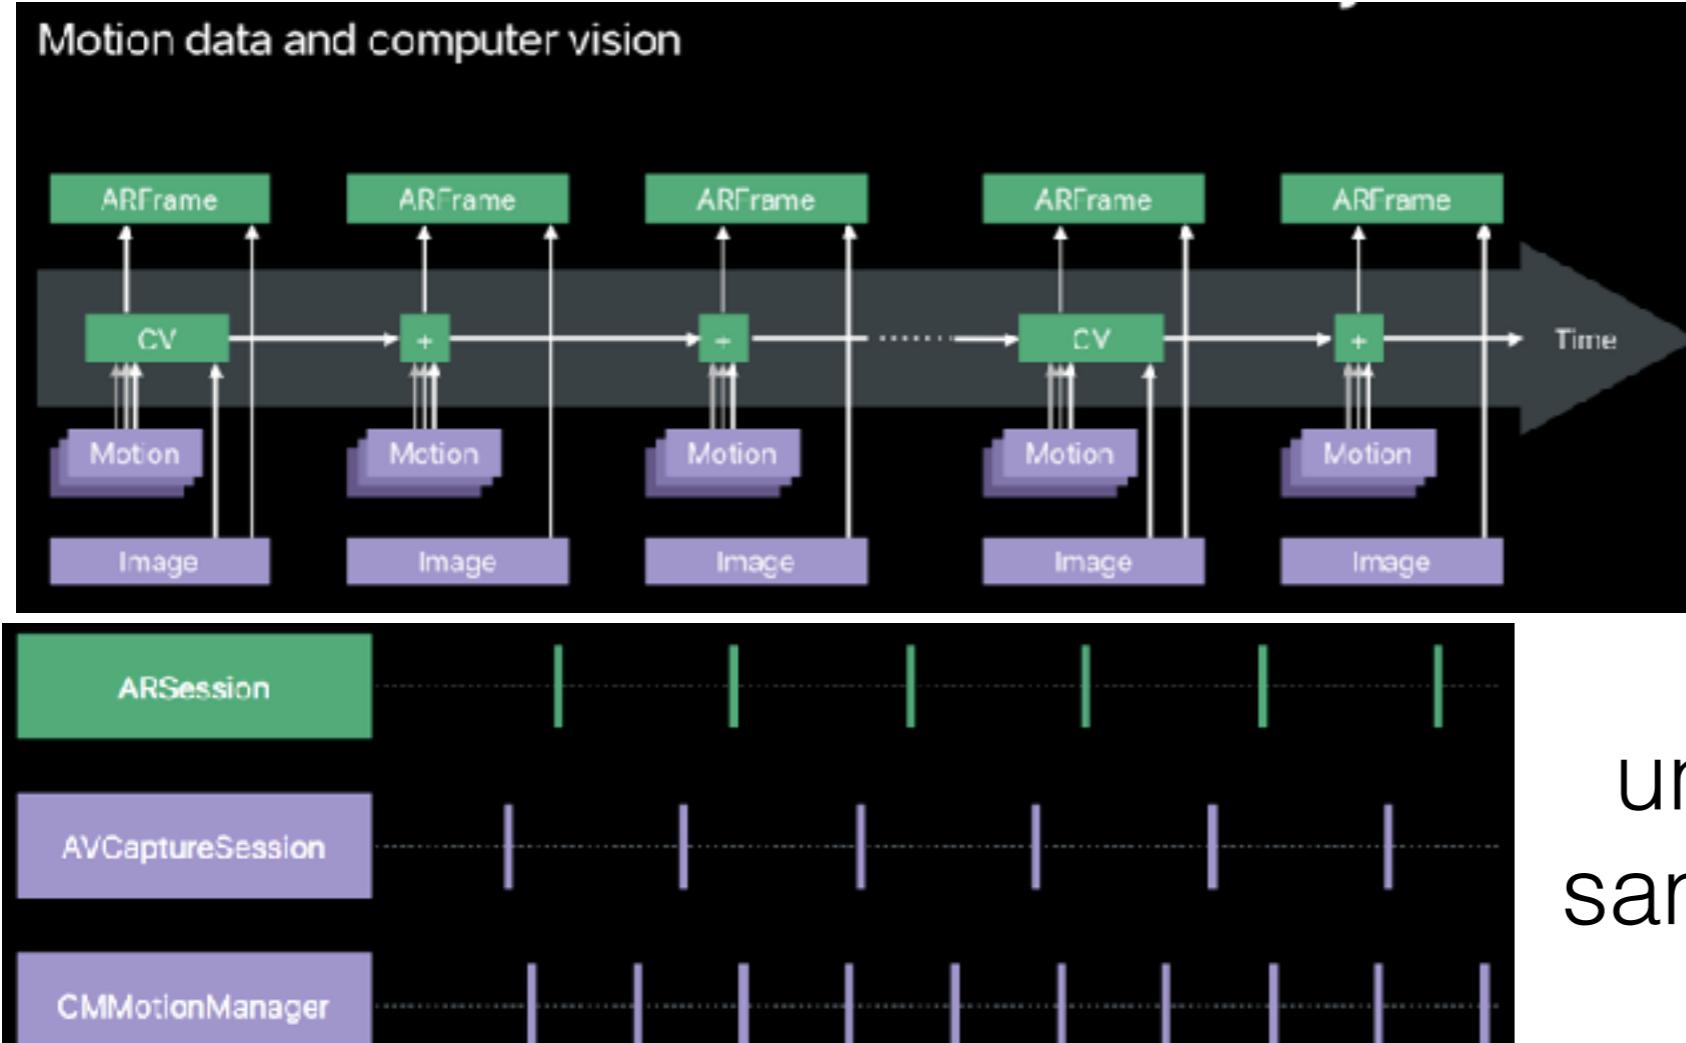
```
let session = ARSession()  
let conf = ARWorldTrackingSessionConfiguration()  
session.run(conf)
```



<https://developer.apple.com/videos/play/wwdc2017/602/>

# ARKit Over Time

Motion data and computer vision



unified  
sampling

```
// Access the latest frame
func session(_: ARSession, didUpdate: ARFrame)
// Handle session errors
func session(_: ARSession, didFailWithError: Error)
```

delegation

<https://developer.apple.com/videos/play/wwdc2010/101/>

# the AR world

poll ARSession, get ARFrame

```
// grab the current AR session frame from the scene, if possible
guard let currentFrame = sceneView.session.currentFrame else {
    return
}

// setup some geometry for a simple plane
let imagePlane = SCNPlane(width:sceneView.bounds.width/6000,
                           height:sceneView.bounds.height/6000)

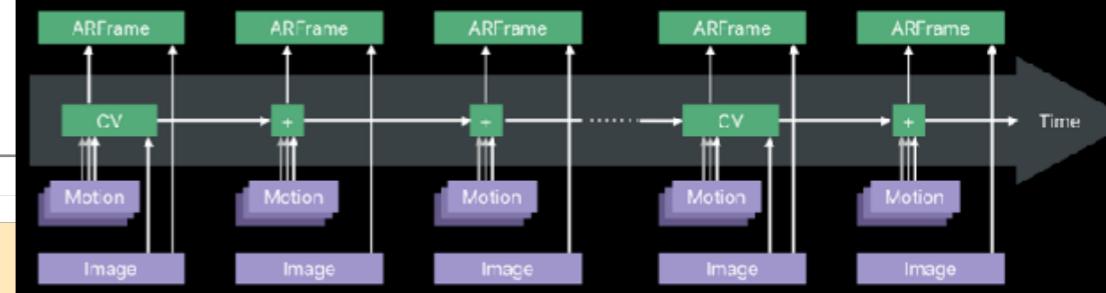
// add the node to the scene
let planeNode = SCNNNode(geometry:imagePlane)
sceneView.scene.rootNode.addChildNode(planeNode)

// update the node to be a bit in front of the camera inside the AR session

// step one create a translation transform
var translation = matrix_identity_float4x4
translation.columns.3.z = -0.1

// step two, apply translation relative to camera for the node
planeNodesimdTransform = matrix_multiply(currentFrame.camera.transform, translation )
```

Motion data and computer vision



create a plane

add to world

translate

# operations

Figure 1-8 Matrix configurations for common transformations

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Identity

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ tx & ty & tz & 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} sx & 0 & 0 & 0 \\ 0 & sy & 0 & 0 \\ 0 & 0 & sz & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \theta & \sin \theta & 0 \\ 0 & -\sin \theta & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Rotate around X axis

$$\begin{bmatrix} \cos \theta & 0 & -\sin \theta & 0 \\ 0 & 1 & 0 & 0 \\ \sin \theta & 0 & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Rotate around Y axis

$$\begin{bmatrix} \cos \theta & \sin \theta & 0 & 0 \\ -\sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Rotate around Z axis

```
// step one create a translation transform
var translation = matrix_identity_float4x4
translation.columns.3.z = -0.1

// step two, apply translation relative to camera for the plane
planeNodesimdTransform = matrix_multiply(currentFrame.c
```

## SIMD: single instruction, multiple data

### Creating Transform Matrices

func `SCNMatrix4MakeTranslation`(Float, Float, Float)

Returns a matrix describing a translation transformation.

func `SCNMatrix4MakeRotation`(Float, Float, Float, Float)

Returns a matrix describing a rotation transformation.

func `SCNMatrix4MakeScale`(Float, Float, Float)

Returns a matrix describing a scale transformation.

# ARKit and SceneKit



- extended demo
- branches:**
- (1) bare implementation,
  - (2) image extension,
  - (3) stylization (with SCNActions)

A dark grey rectangular background featuring logos of nine partner universities: CAL POLY SAN LUIS OBISPO, Drexel UNIVERSITY, ingéSUP®, RMIT UNIVERSITY, PLYMOUTH UNIVERSITY, Stanford, UNIVERSITY OF CALIFORNIA SANTA CRUZ, SMU, and TUM Technische Universität München. To the right of the logos is a photograph of Steve Jobs speaking on stage.

and now its time for a demo

# object detection in AR

---

- need to perform detection of images and 3D objects
  - to anchor AR in the world
  - for shared AR experiences
- keep in mind
  - automatic object detection is a computational process
  - ...but is very easy to get started

# image recognition in ARKit

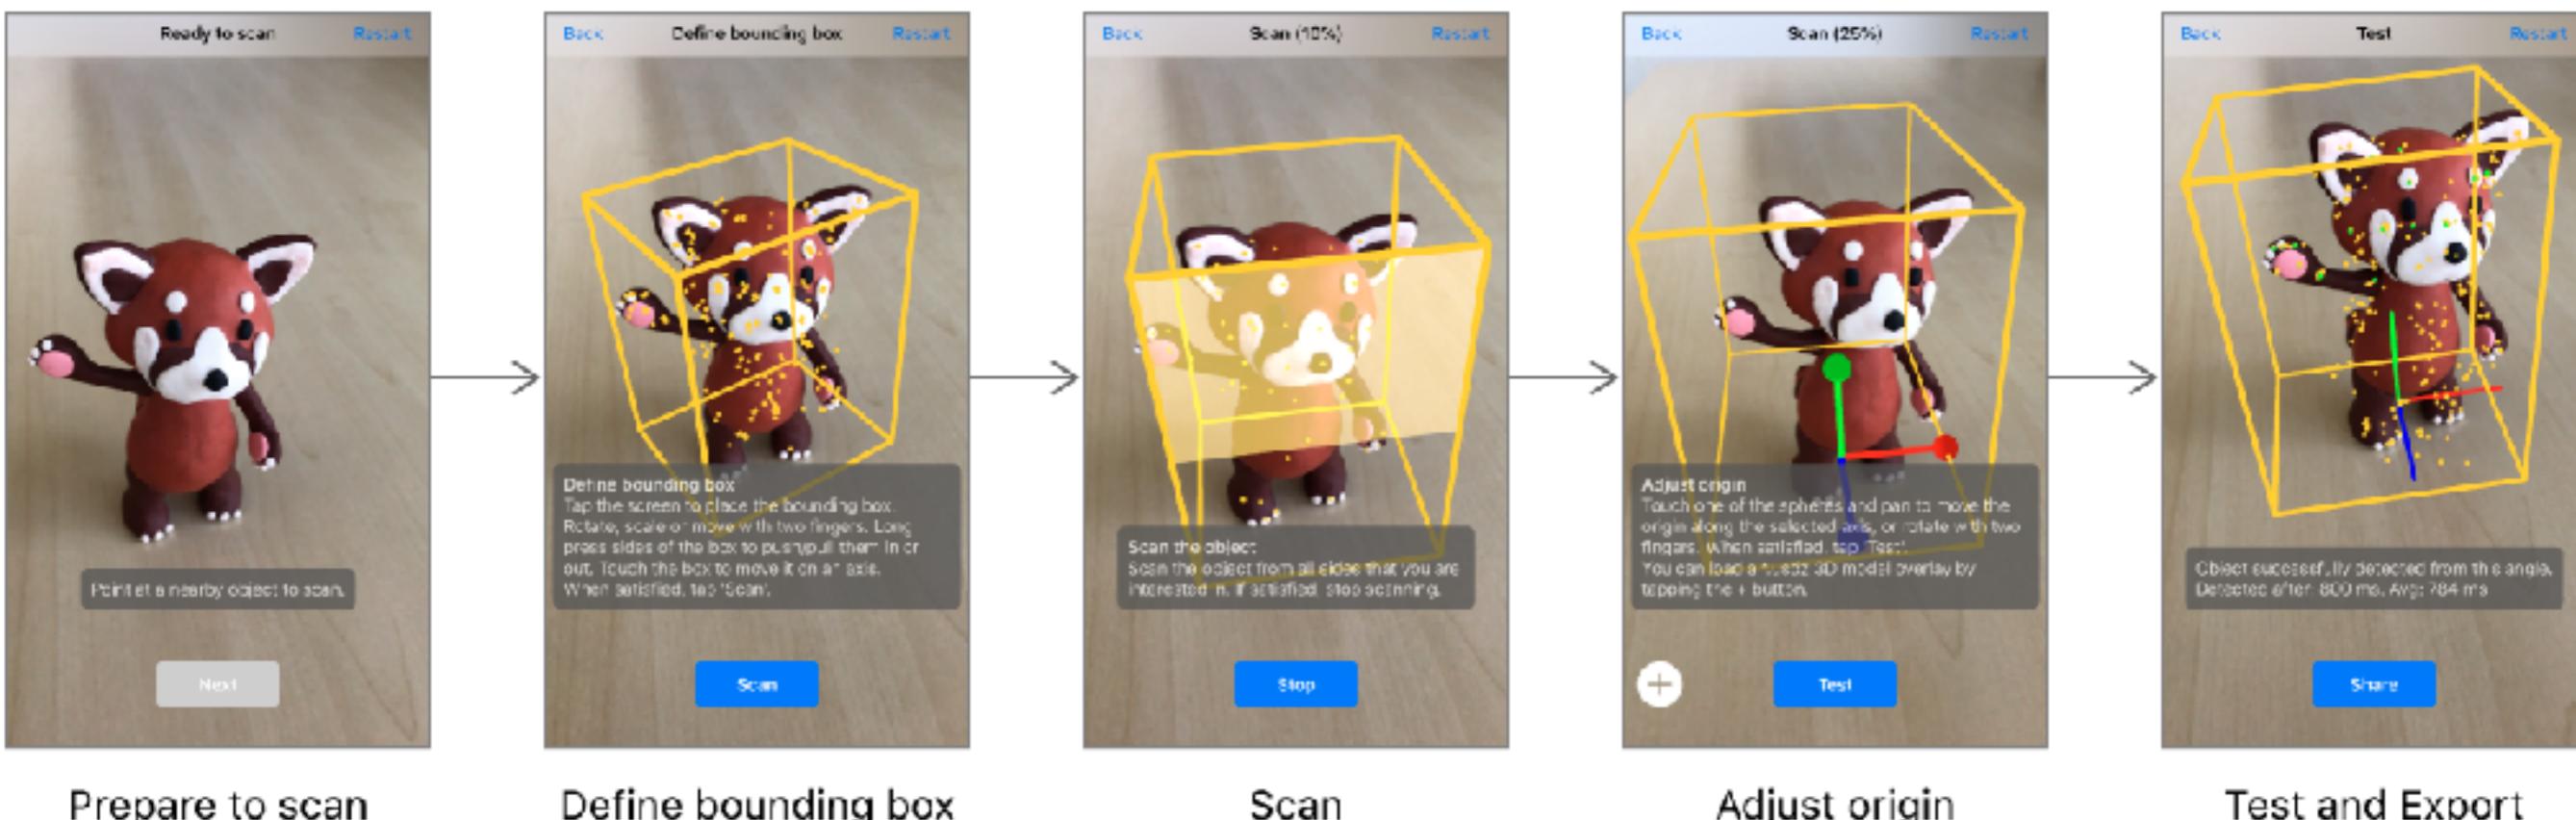
- detection based on scans of keypoints in example image, projected onto detected planes
- drag and drop image into Assets file as an “ARImage”



<https://developer.apple.com/videos/play/wwdc2018/610/>

# object recognition in ARKit

- detection based on scans of an object
- scanning app available here:
  - [https://developer.apple.com/documentation/arkit/scanning\\_and\\_detecting\\_3d\\_objects](https://developer.apple.com/documentation/arkit/scanning_and_detecting_3d_objects)



Prepare to scan

Define bounding box

Scan

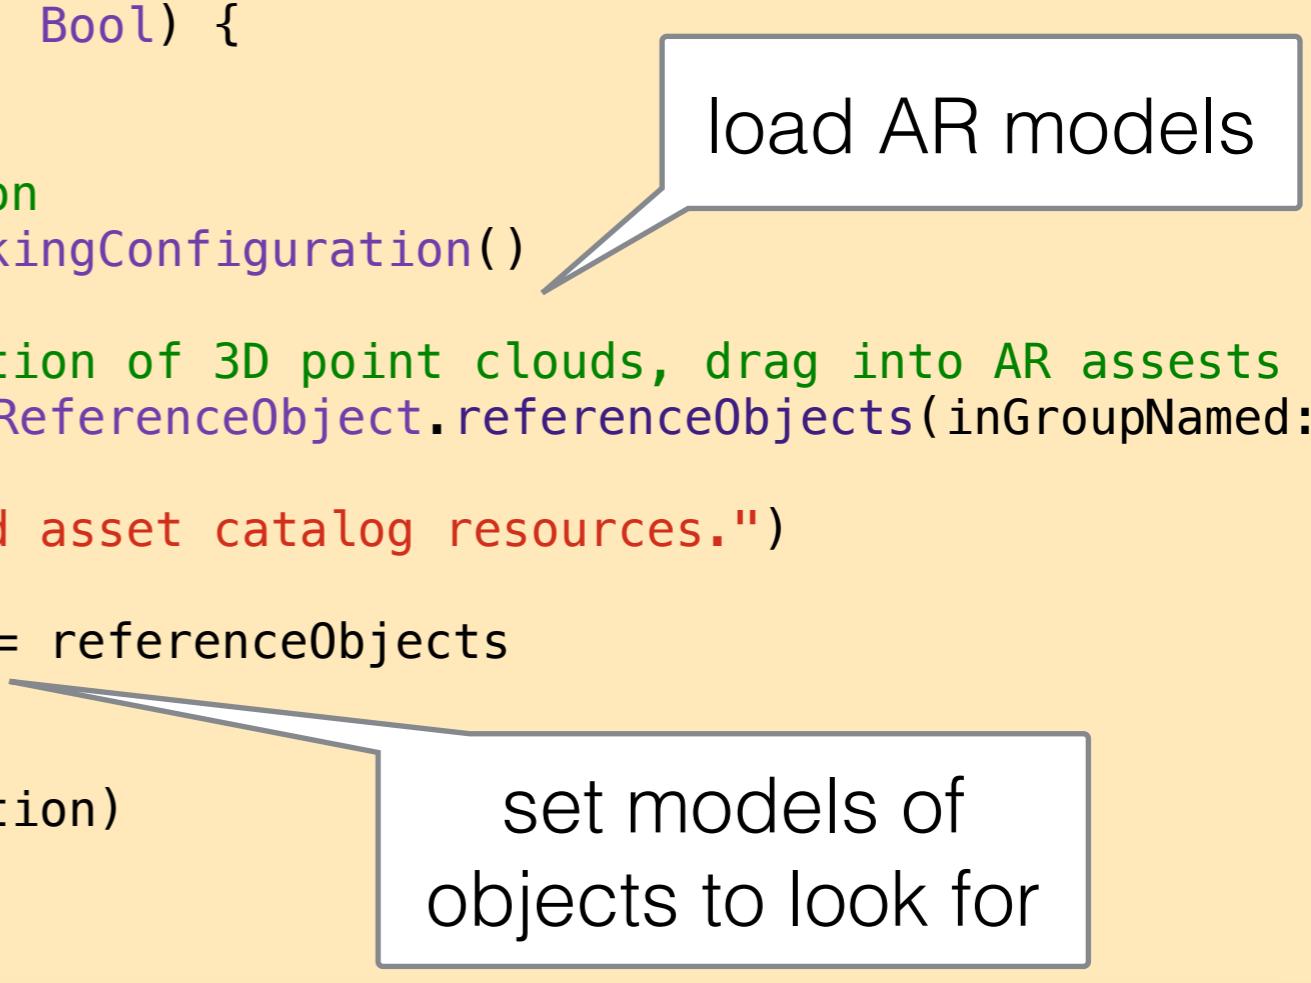
Adjust origin

Test and Export

# object recognition in ARKit

- drag **aobject** into assets repository
- tell ARKit to begin looking for object

```
override func viewWillAppears(animated: Bool) {  
    super.viewWillAppear(animated)  
  
    // Create a session configuration  
    let configuration = ARWorldTrackingConfiguration()  
  
    // here is where we setup detection of 3D point clouds, drag into AR assets  
    guard let referenceObjects = ARReferenceObject.referenceObjects(inGroupNamed:  
        "gallery", bundle: nil) else {  
        fatalError("Missing expected asset catalog resources.")  
    }  
    configuration.detectionObjects = referenceObjects  
  
    // Run the view's session  
    sceneView.session.run(configuration)  
}
```



load AR models

set models of objects to look for

# object recognition in ARKit

- use delegation to respond when object is found
- add new nodes anchored to the ARAnchor

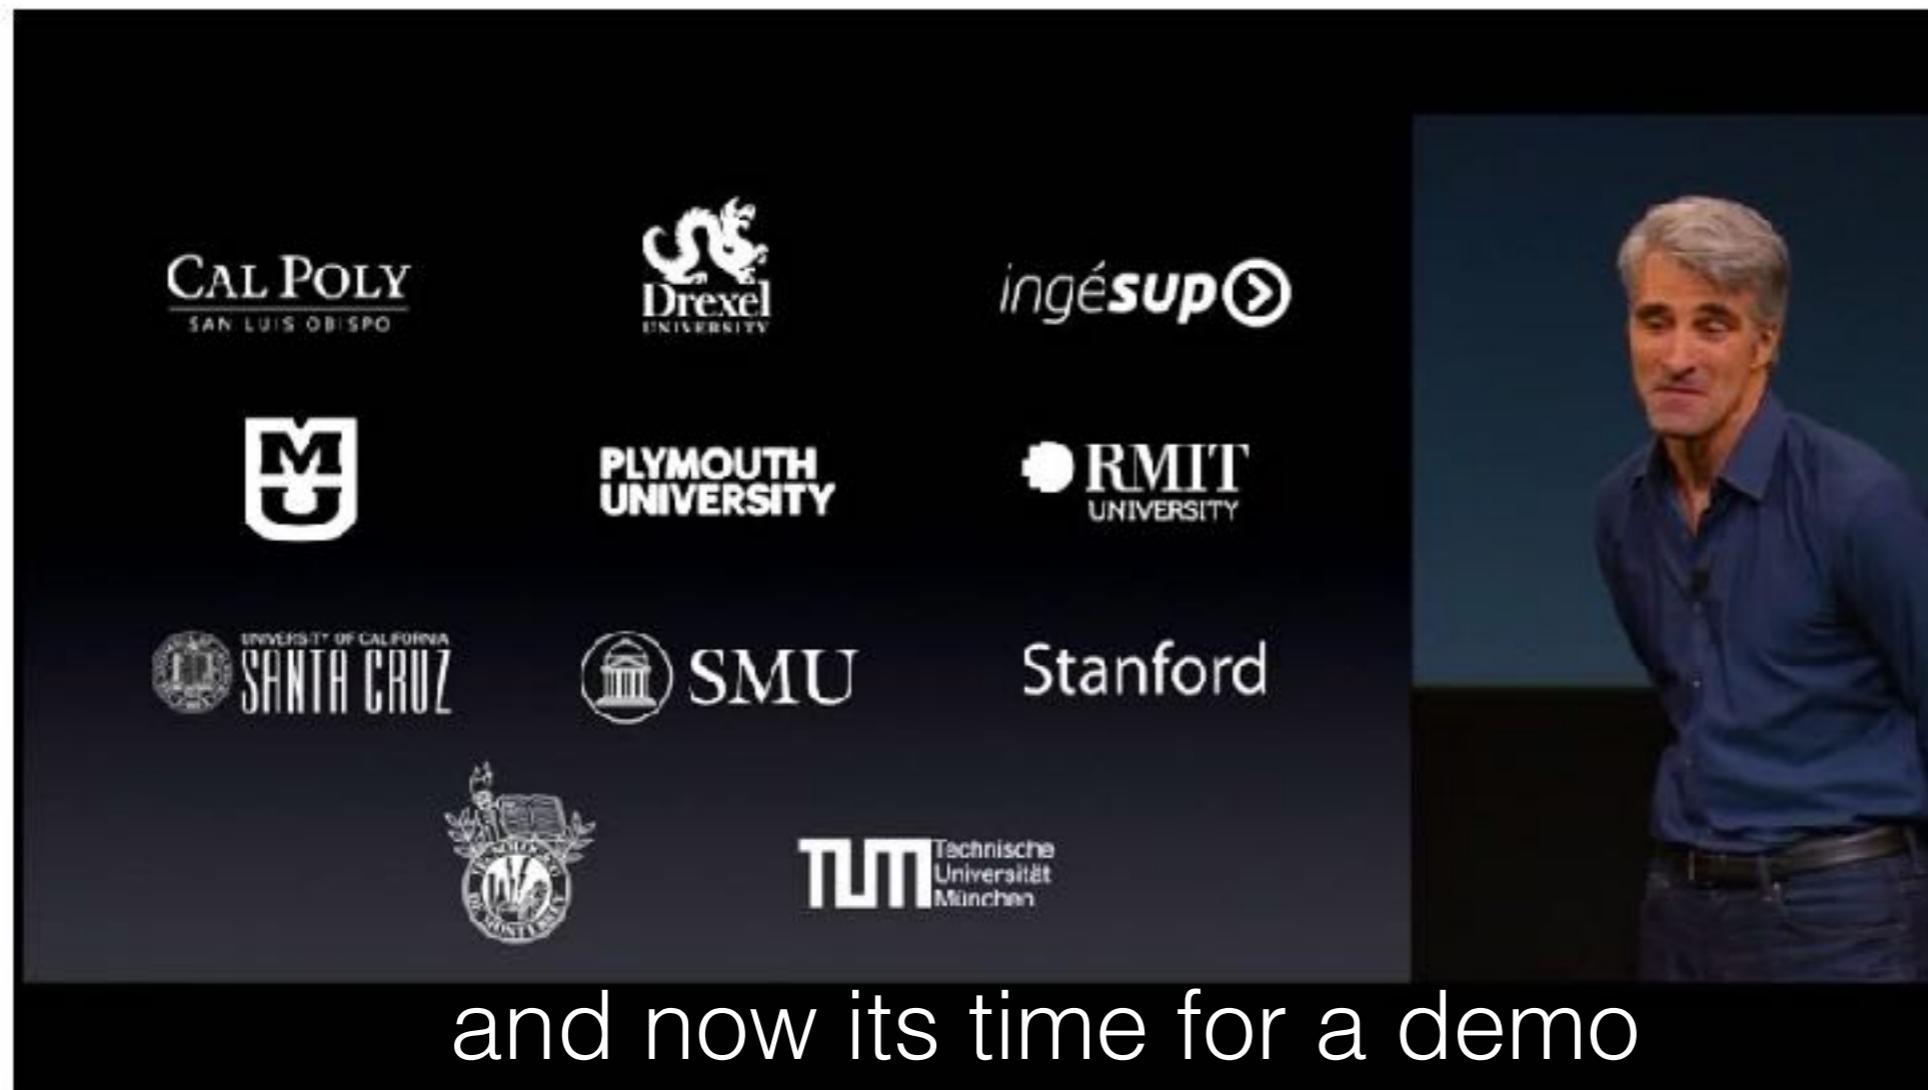
```
func renderer(_ renderer: SCNSceneRenderer,  
             didAdd node: SCNNode,  
             for anchor: ARAnchor) {  
  
    if let objectAnchor = anchor as? ARObjectAnchor {  
  
        newNode = ... what you want it to be ...  
        print(objectAnchor.name! + " found")  
        node.addChildNode(newNode)  
    }  
}
```

delegate function called  
when new node is added to  
ARSCNScene

true if recognized as  
ARObject from gallery

# ARKit with ObjectRecognition

- find the object!
- place images around the object!



A collage of university logos and a photograph of Steve Jobs. The logos include CAL POLY, Drexel UNIVERSITY, ingéSUP, PLYMOUTH UNIVERSITY, RMIT UNIVERSITY, UNIVERSITY OF CALIFORNIA SANTA CRUZ, SMU, Stanford, and TUM Technische Universität München. To the right of the logos is a photograph of Steve Jobs speaking on stage.

and now its time for a demo

# additional UI elements

if time

- planes
- simple nodes
- actions
- UI Text



# Planes as Content Holders

- planes are placeholders for flat content like images and videos
- planes are the geometry for a node in SceneKit



```
// add in a video near the detected object
let planarNode = SCNNode()
let contentPlane = SCNPlane(width: CGFloat(0.1), height: CGFloat(0.05))
let avMaterial = SCNMaterial()
```

```
avMaterial.diffuse.contents = ...image or video...
contentPlane.materials = [avMaterial]
```

```
planarNode.geometry = contentPlane
node.addChildNode(planarNode)
```

setup image or video

set material of the plane

plane is geometry of node



# Simple Nodes in SceneKit

- define geometry with pre-made shape, like box
- make new node with given geometry



```
let boxNode:SCNNode? = SCNNode()  
  
let box = SCNBox(width: CGFloat(0.1), height: CGFloat(0.1),  
                  length: CGFloat(0.1), chamferRadius: 0.01)  
box.firstMaterial?.diffuse.contents = UIColor(white: 1.0, alpha: 0.8)  
box.firstMaterial?.isDoubleSided = true
```

```
boxNode!.geometry = box // make this node a box!
```

setup initial color

```
node.addChildNode(box!)
```

```
let alphaAction = SCNAction.fadeOpacity(to: 0.1, duration: 5)  
box?.runAction(alphaAction)
```

add animation, change alpha!



# Actions in SceneKit

- any movement scaling, or change of added nodes should be done via actions
- actions can be reused on different nodes!

```
// add some animation  
let prevScale = uiObject.scale  
uiObject.scale = SCNVector3(CGFloat(prevScale.x)/3, CGFloat(prevScale.y)/3,  
                           CGFloat(prevScale.z)/3)
```

**not yet on screen:** will start smaller,  
then scale back to original size

```
let scaleAction = SCNAction.scale(to: CGFloat(prevScale.x), duration: 2)  
scaleAction.timingMode = .easeIn
```

```
node.addChildNode(uiObject)
```

make visible

define scaling action and params

```
uiObject.runAction(scaleAction, forKey: "scaleAction")
```

run action immediately

# UI Text Elements in SceneKit



```
func createTextNode(textString: String) -> SCNNode? {  
    let textNode: SCNNode? = SCNNode()  
    textNode!.geometry = setupTextParameters(textString: textString)  
    textNode!.scale = SCNVector3Make(0.001, 0.001, 0.001)  
    textNode!.position = SCNVector3Make(-0.1, 0.1, 0.0)  
    textNode?.castsShadow = true  
}
```

lower left corner

```
return textNode  
}
```

if plane detection is enabled

```
func setupTextParameters(textString: String) -> SCNTex{  
    let text = SCNTex(string: textString, extrusionDepth: 1)  
  
    let material = SCNMaterial()  
    material.diffuse.contents = UIColor.white  
  
    text.flatness = 0  
    text.isWrapped = true  
    text.materials = [material]  
    return text  
}
```

extrusion defines depth of the text (flat versus extended)

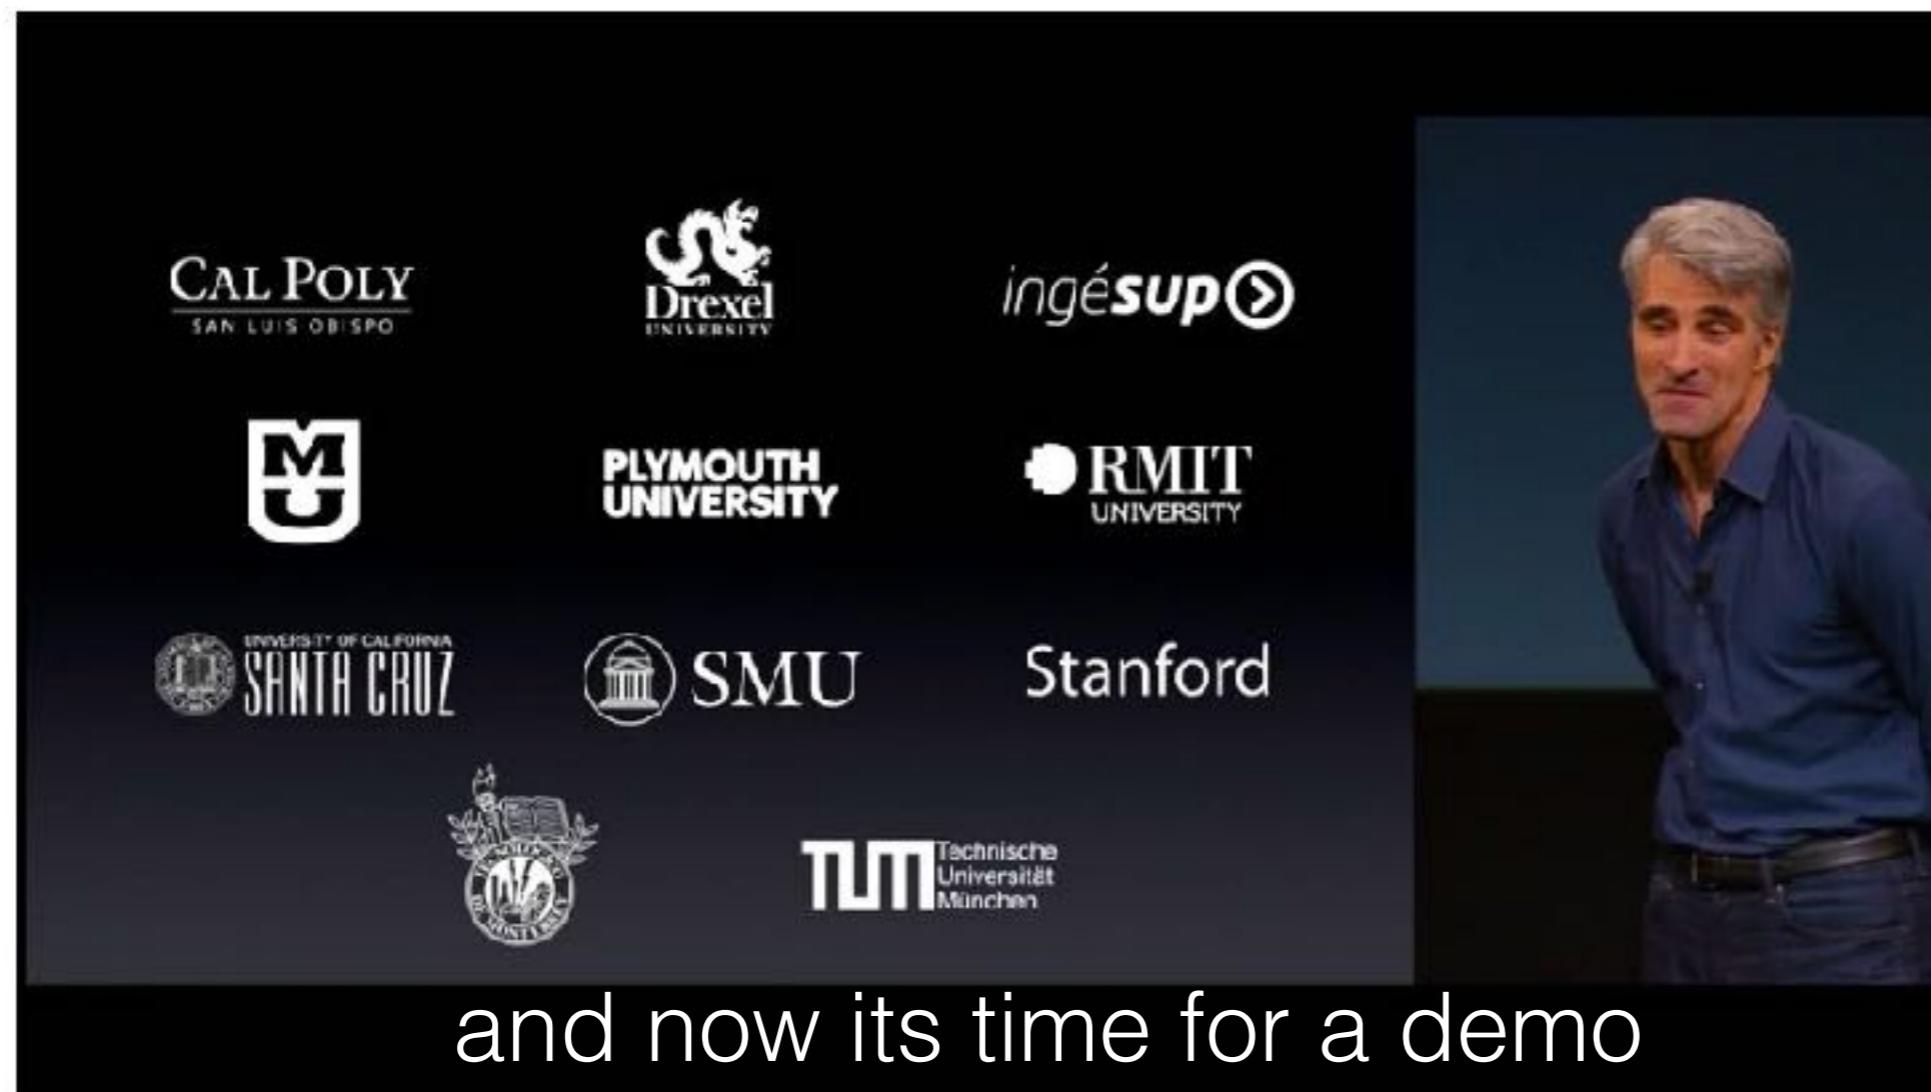
color and other properties are set in the material of the node



<https://developer.apple.com/documentation/scenekit/scntext>

# ARKit with ObjectRecognition

- add to demo some more interesting UI elements



and now its time for a demo

# for next time...

---

- Speech Recognition
- Pitching (Required)
- ~Fin~

# Back Up Slides

---

# Detecting Objects in Images

---

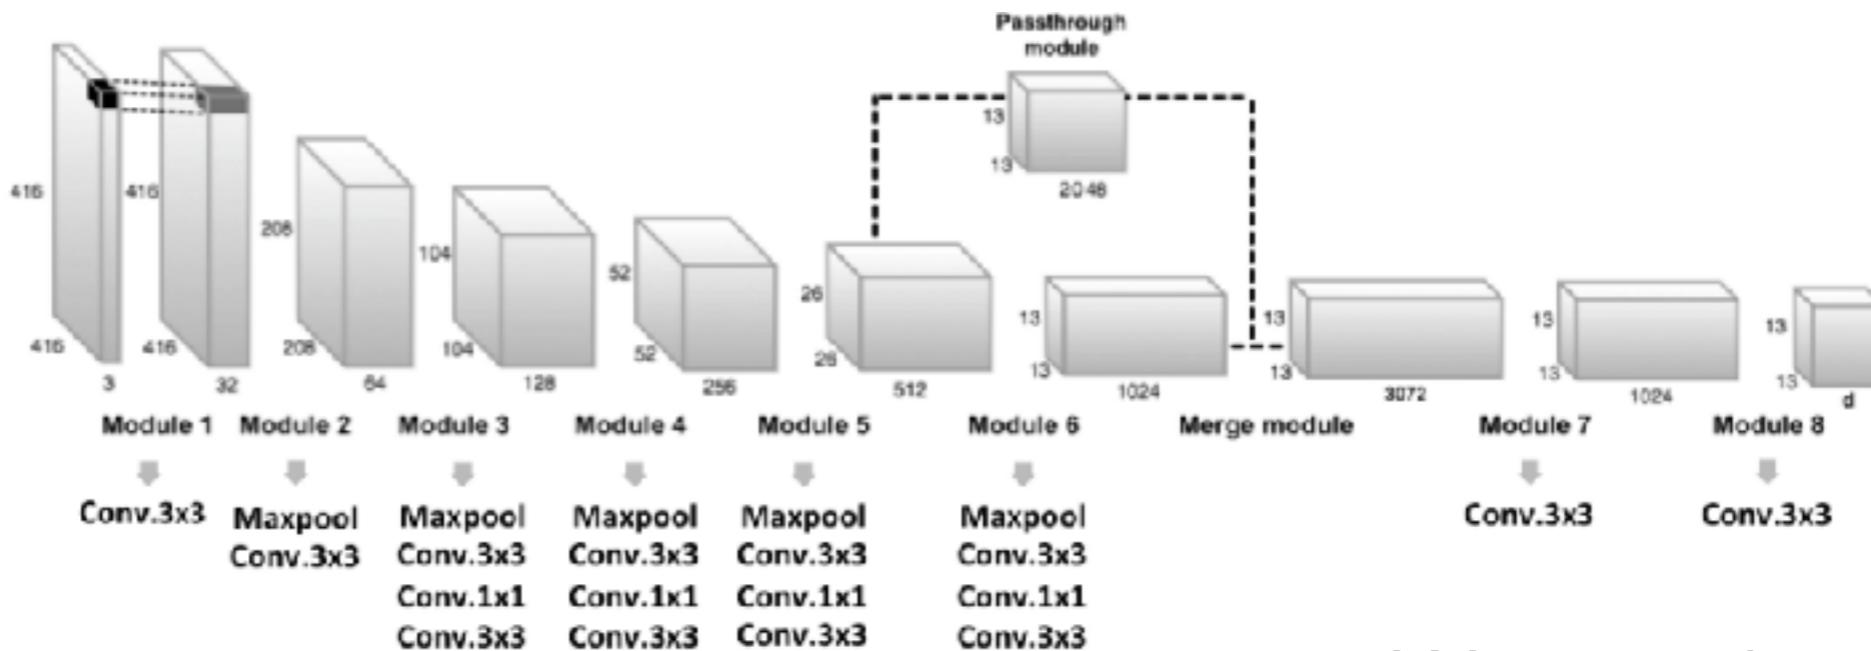
- Using Turi Create
  - SFrames
  - Simple Machine Learning
- Object Detection with YOLO
  - Creation
  - Exporting to CoreML
  - Incorporating into ARKit

# didn't we already detect objects?

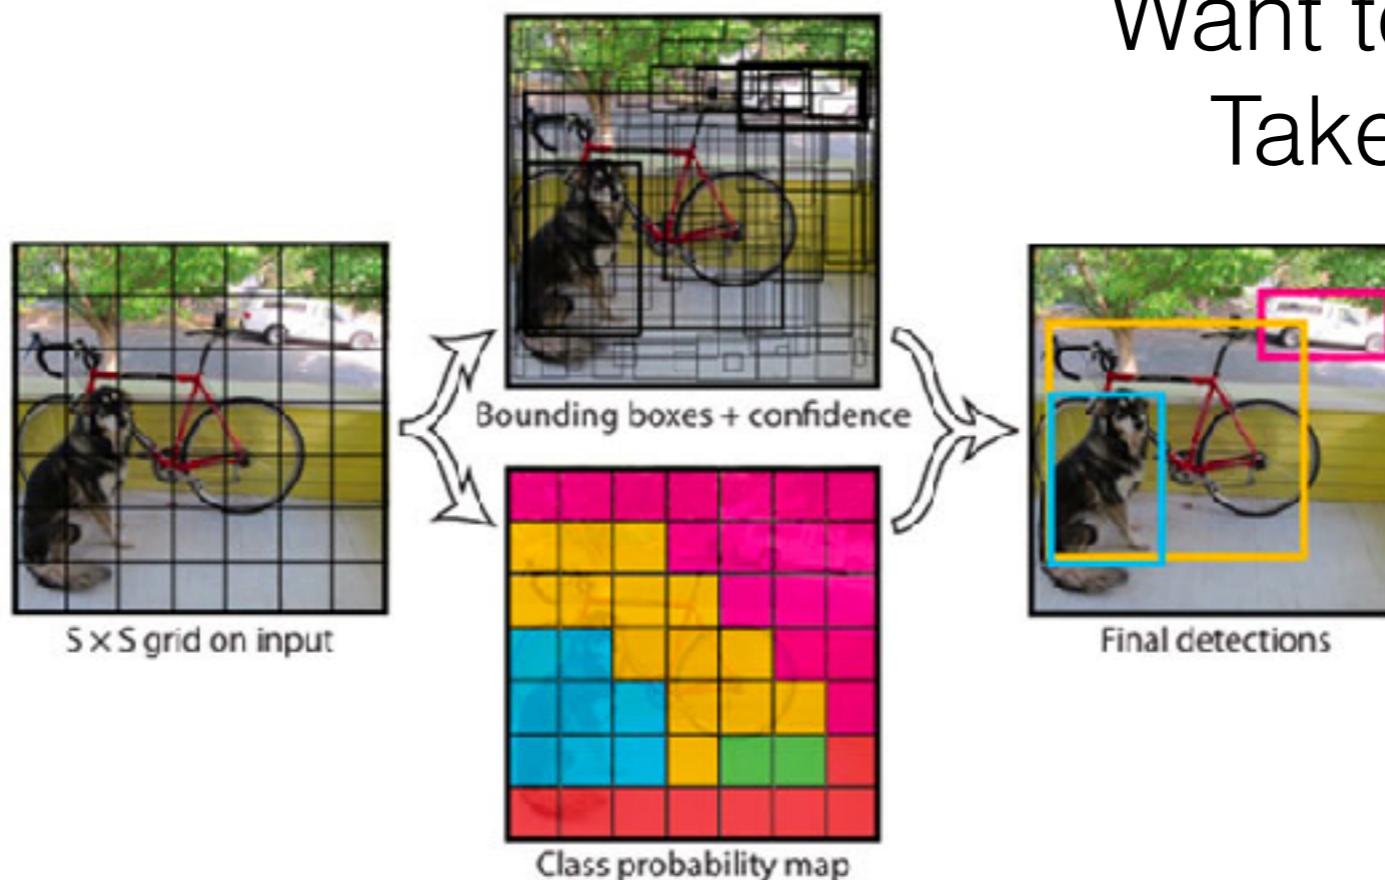
---

- with ARKit, yes
- but what if we want to detect all sorts of objects
- and we need to do it very fast (multiple times per second)
- and we only have examples of the types of objects we want, not a scan of the exact object
- enter... Turi Create

# YOLO



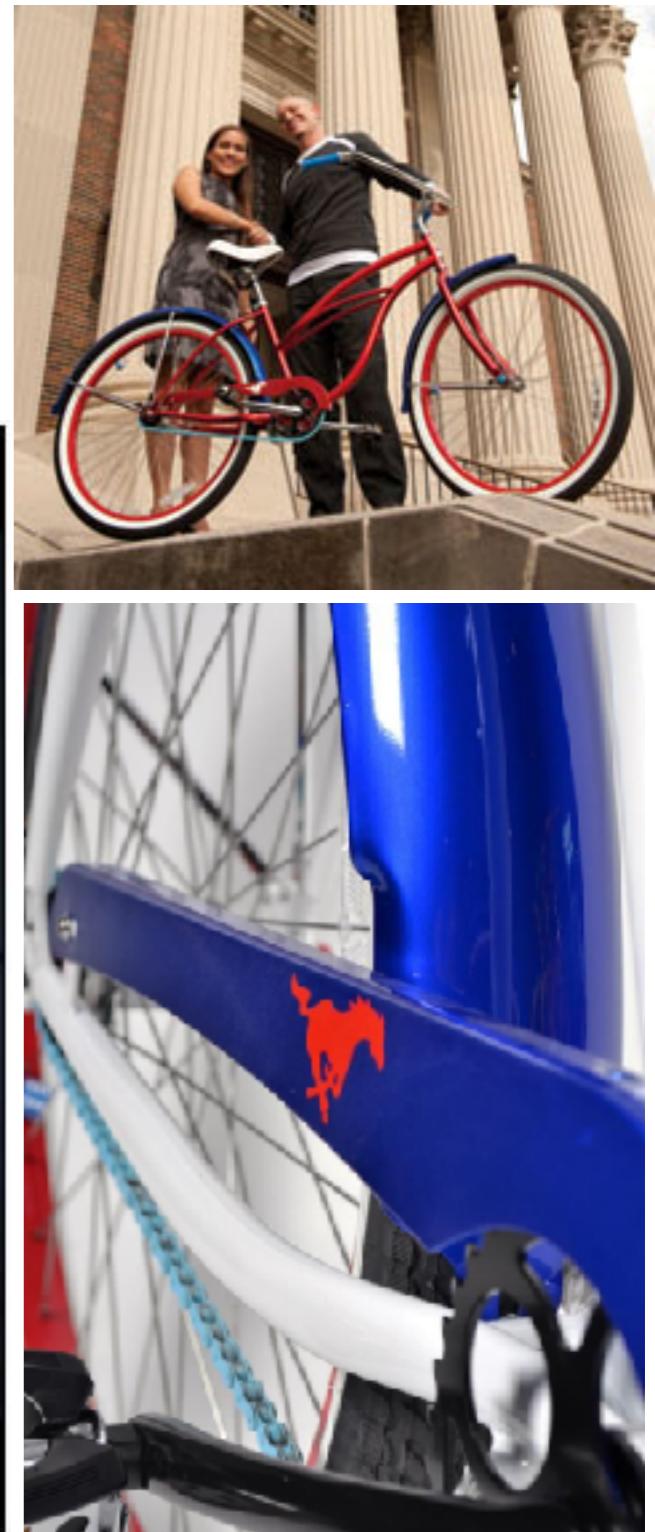
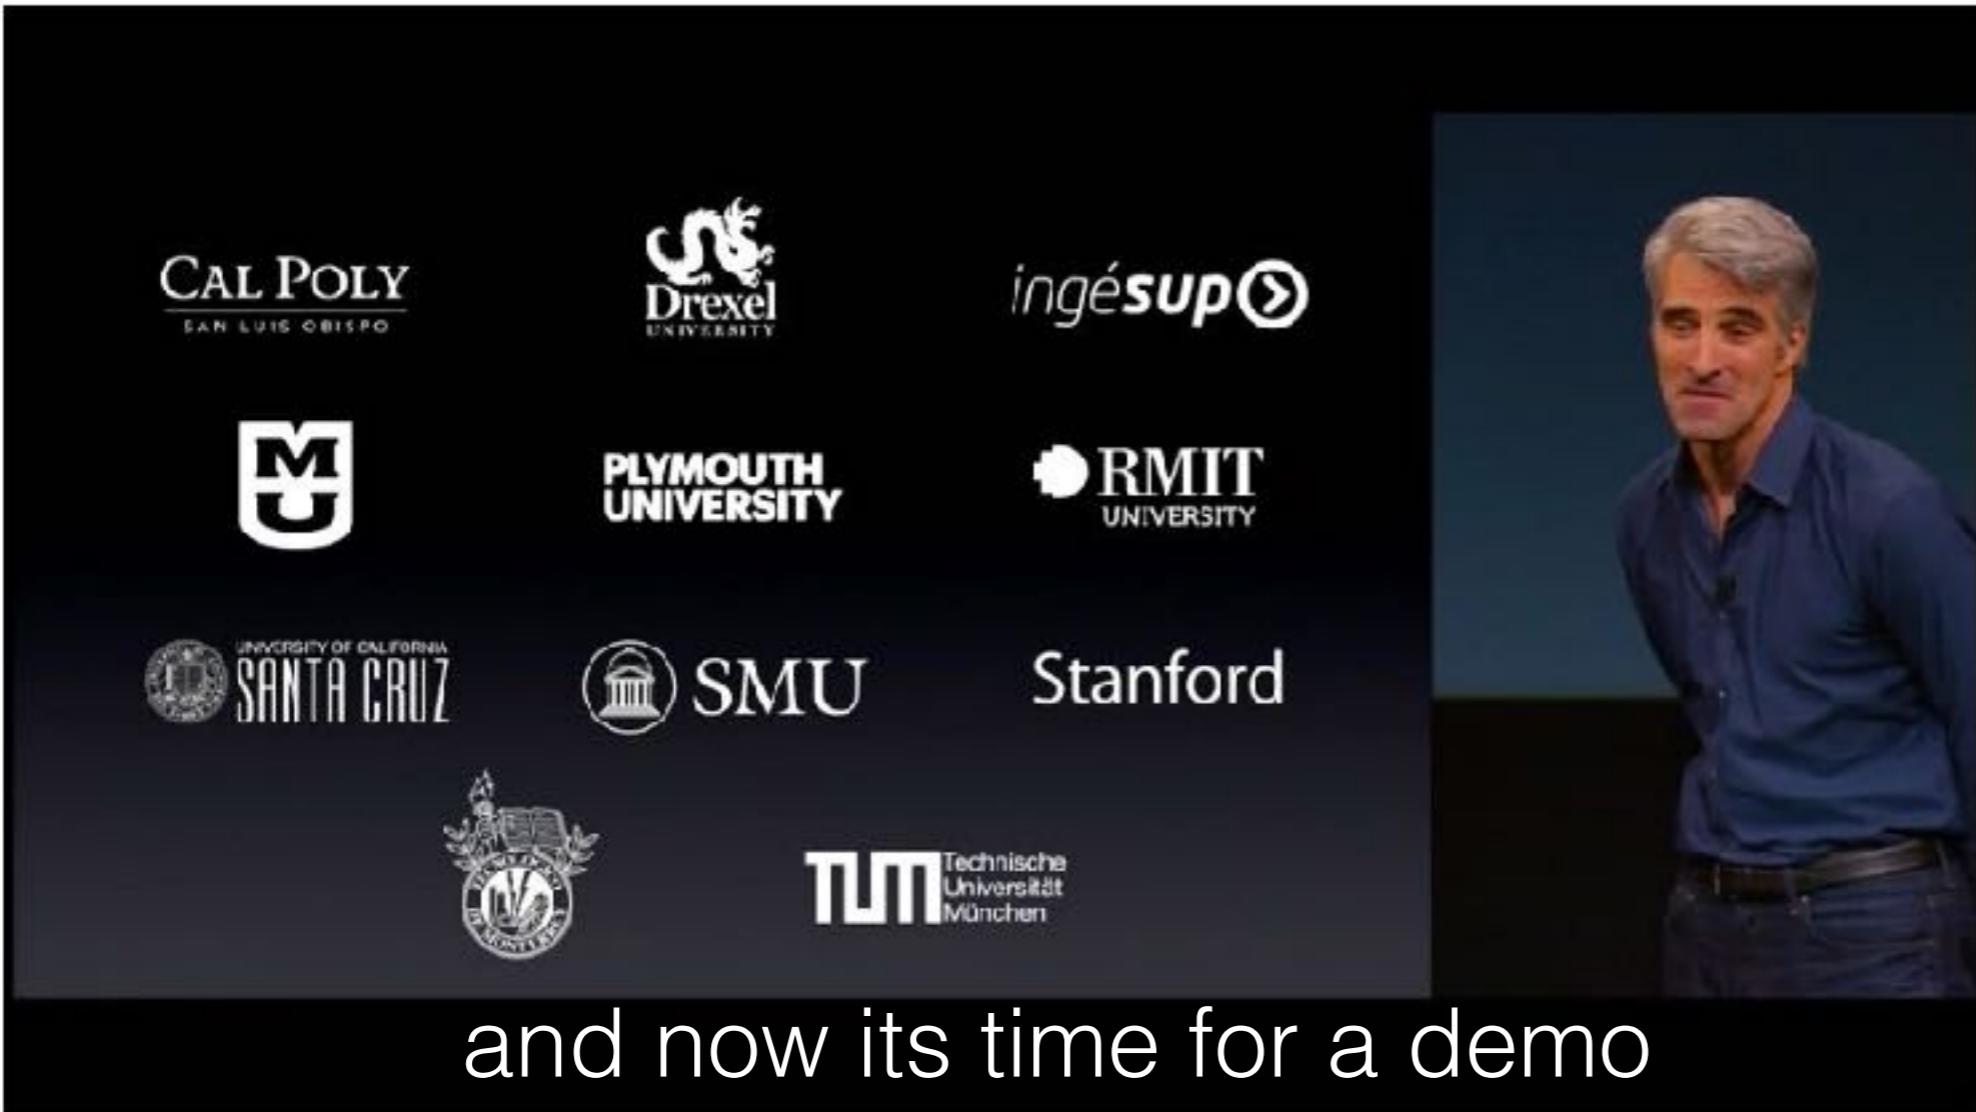
Want to know more?  
Take CS8321...



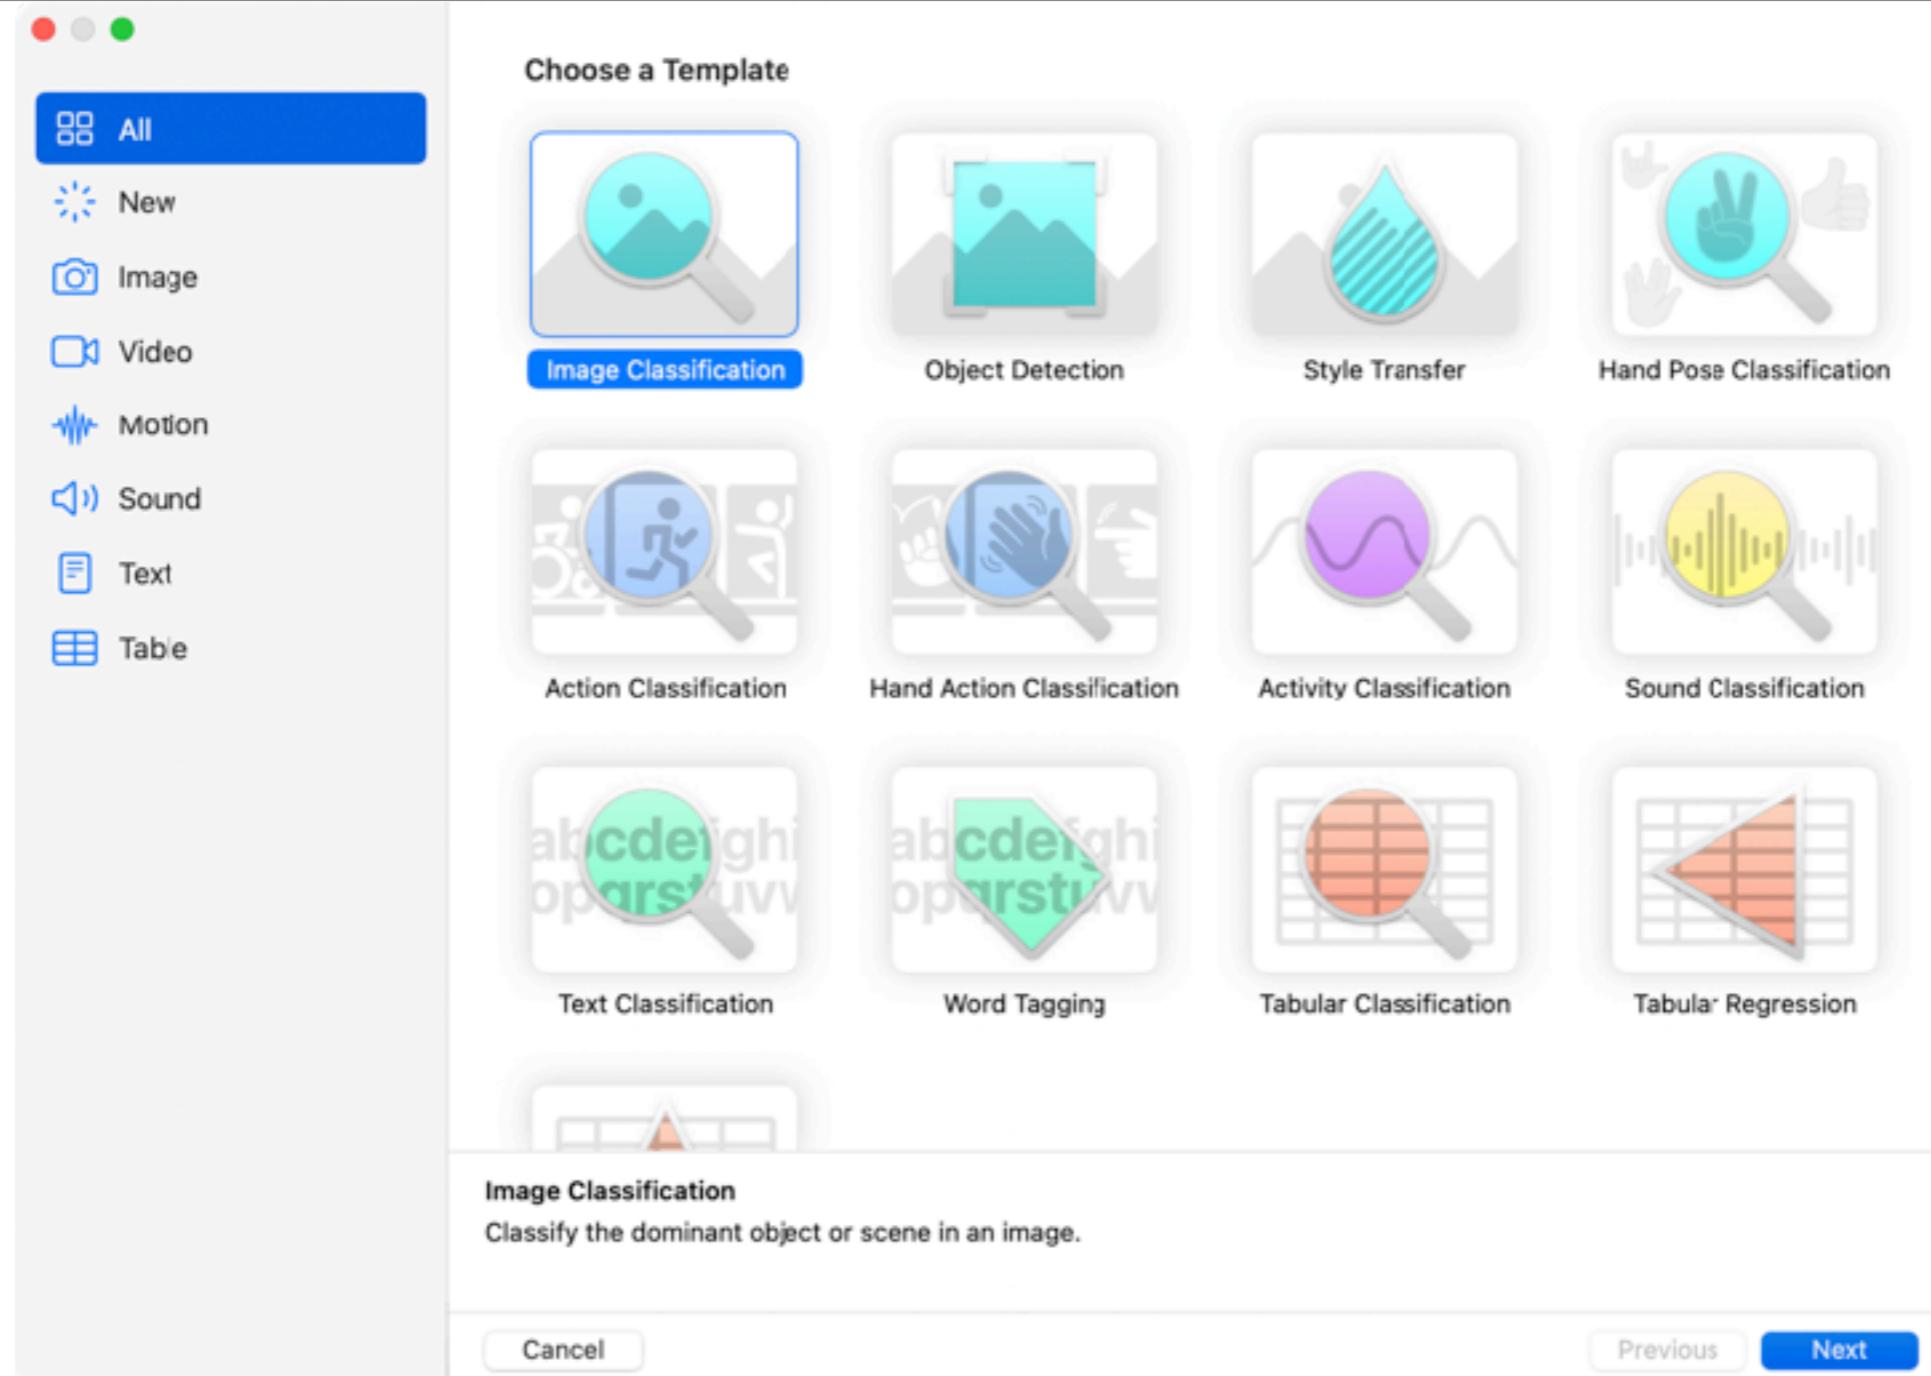
<https://datascience.stackexchange.com/questions/42509/yolo-algorithm-understanding-training-data>

# create an object detector with Turi

- Turi Create, high level overview
  - SFrames and Object Detectors



# Bonus: Create ML (iOS 15.0+)



<https://developer.apple.com/documentation/creatempl/>

# Create ML Stylizer

The screenshot shows a user interface for a machine learning project titled "ExampleStyleTransfer".

**Project:** ExampleStyleTransfer

**Model Sources:** ExampleStyleTransfer 1 (selected), Iteration 500

**Training Tab:** Active. Shows "Stylized Validation at Iteration 700" (a stylized dog face) and "Style" (The Great Wave off Kanagawa). Metrics: Style Loss 22.510 at Iteration 700, Content Loss 12.915 at Iteration 700.

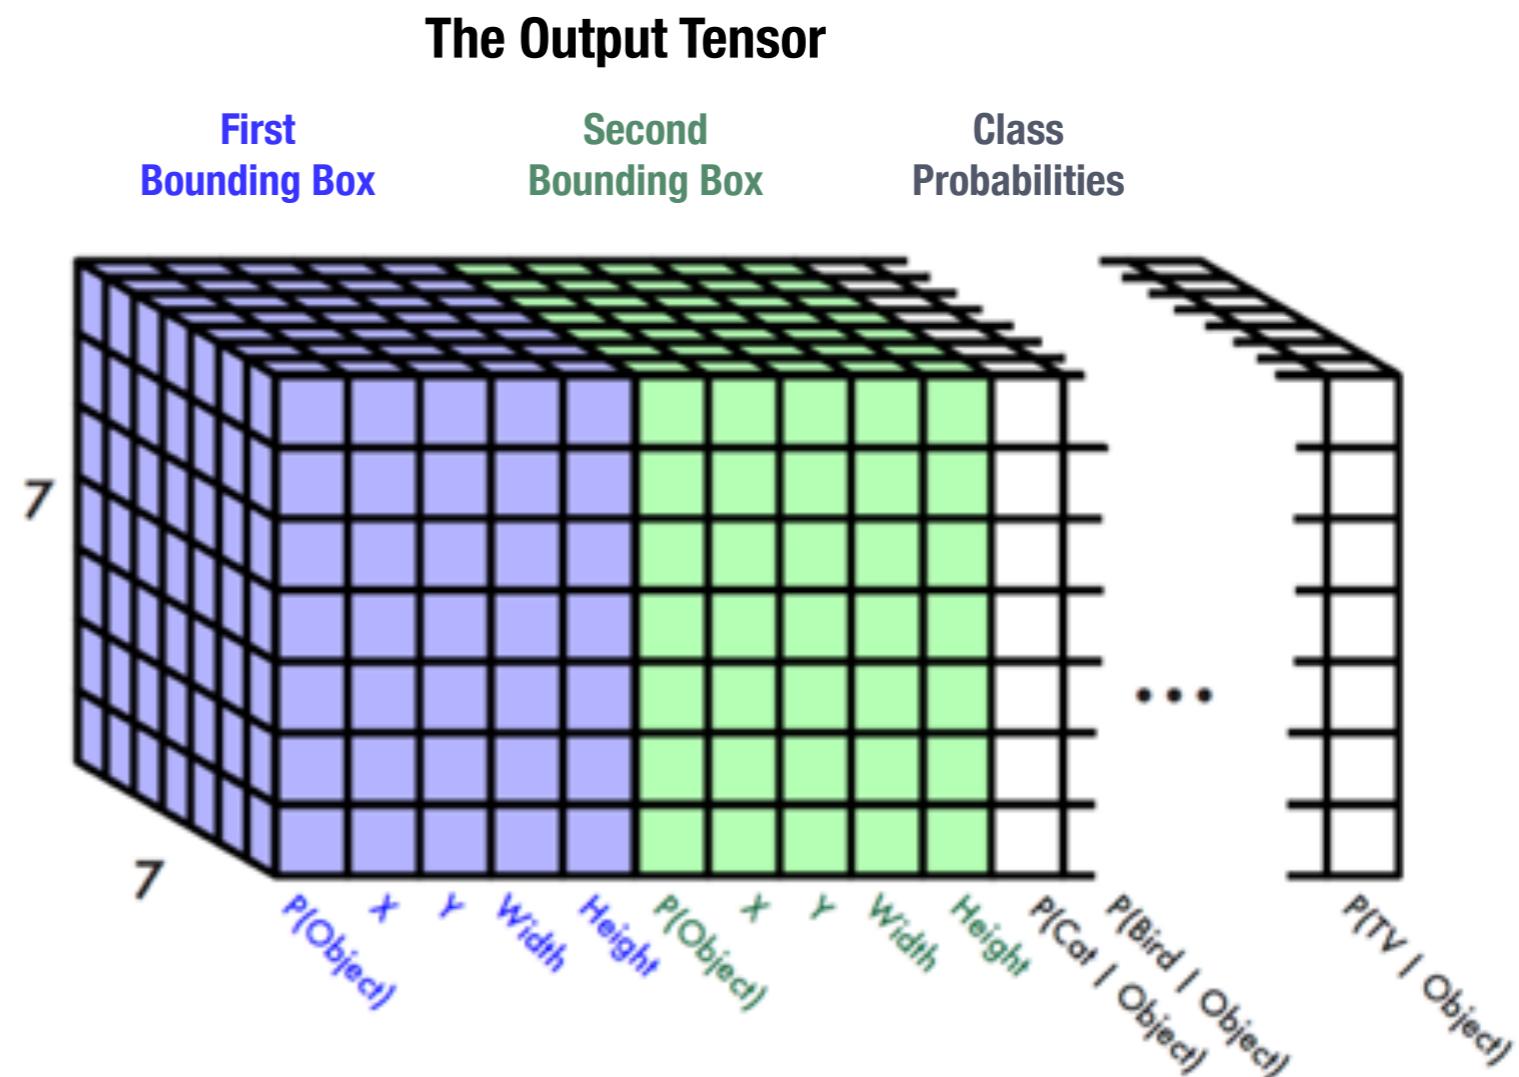
**Activity Log:** A list of events from November 10, 2021, including Training Completed at 6:35 PM, Training Extended at 6:26 PM, and various file operations like Snapshot Saved and Snapshot created at Iteration 500.

**Graphs:** Two line graphs showing loss over iterations. The left graph shows a sharp initial drop in style loss, followed by a noisy plateau around 20. The right graph shows a more stable, fluctuating content loss over time.

**Status:** Completed 700 iterations.

# YOLO output tensor

- is there an object?
- where is the object?
- how large is the Object?
- what is the object?
- if competing, which
- object is most likely?



# using the vision API

---

- Overview:
  - streaming video will require constant images from the AR session, **use delegation**
  - inside delegate, need to **grab** ARFrame, **convert** to pixel buffer, and **process** on a **background** thread
  - **handle vision** results in a new function
  - **update** the UI on the **main thread**



# setting up vision

```
let model = PersonBike()  
private var requests = [VNRequest]()  
  
@discardableResult  
func setupVision(useCPUOnly:Bool) {  
  
    //MARK: One, Setup Vision  
    let visionModel = try VNCoreMLModel(for: model.model)  
  
    // use this request to setup the object recognition with Vision  
    let objectRecognition = VNCoreMLRequest(model: visionModel,  
                                              completionHandler:  
                                                self.handleObjectRecognitionResult)  
  
    objectRecognition.imageCropAndScaleOption = .scaleFill  
    objectRecognition.usesCPUOnly = [True / False]  
    self.requests = [objectRecognition]  
}  
  
load model from Turi  
setup vision wrapper  
save this request for later
```



# using vision with AR

ARDelegate function,  
called at 60 FPS

```
func renderer(_ renderer: SCNSceneRenderer, updateAtTime time: TimeInterval) {  
    //MARK: Two, Get Image Frame from AR  
    guard let frame = sceneView.session.currentFrame else { return }  
  
    guard self.currentBuffer == nil, else { return } // limit FPS of detection  
    self.currentBuffer = frame.capturedImage // the pixels to process  
  
    // run in the background so that AR doesn't suffer performance  
    DispatchQueue.global(qos: .background).async {  
        // setup input image for the request  
        let imageRequestHandler = VNImageRequestHandler(  
            cvPixelBuffer: self.currentBuffer,  
            orientation: ORIENTATION,  
            options: [:])  
  
        imageRequestHandler.perform(self.requests)  
    }  
}
```

start processing request  
that we setup previously

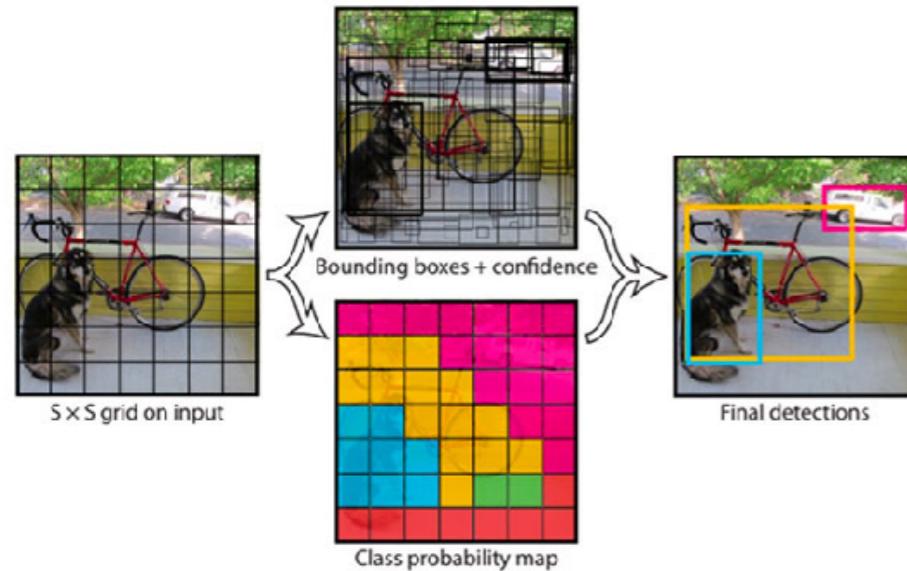
limit FPS of detection

setup image buffer

tell request what pixels to  
process



# handling output request



YOLO Output:

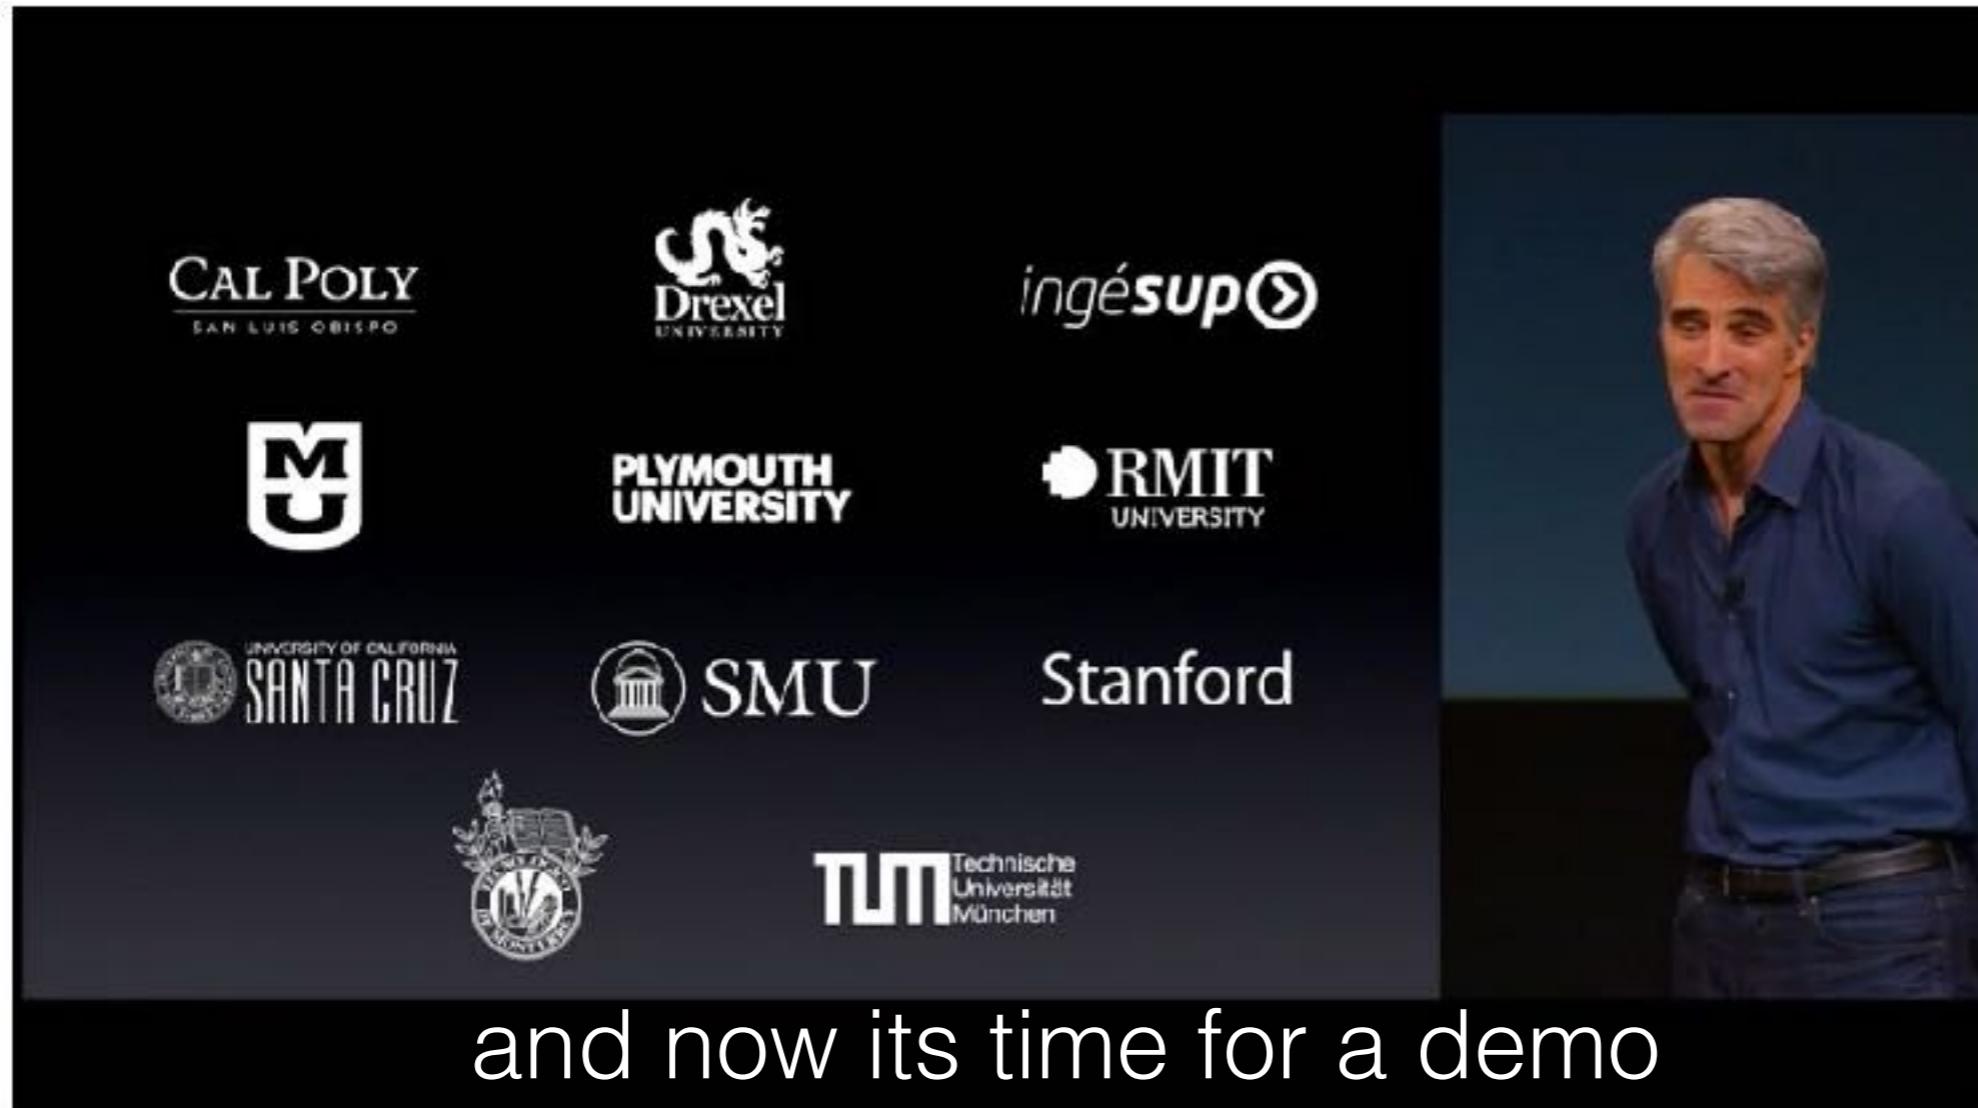
- list of bounding boxes for every square in grid
- list of all possible classes (bike, person)

can have multiple objects in the scene

```
func handleObjectRecognitionResult(_ request: VNRequest, error: Error?) {
    // perform all the UI updates on the main queue
    if let results = request.results { // if we have valid results, else its nil
        DispatchQueue.main.async(execute: {
            self.drawVisionRequestResults(results)
            self.updateOverlay() —————— go over, if time!
            // set as nil so we can process next ARFrame Image
            self.currentBuffer = nil —————— process the next
        })                                captured frame now
    }
}
```

# object detection

- YOLO demo
- If time, go over creating overlays with CATransactions



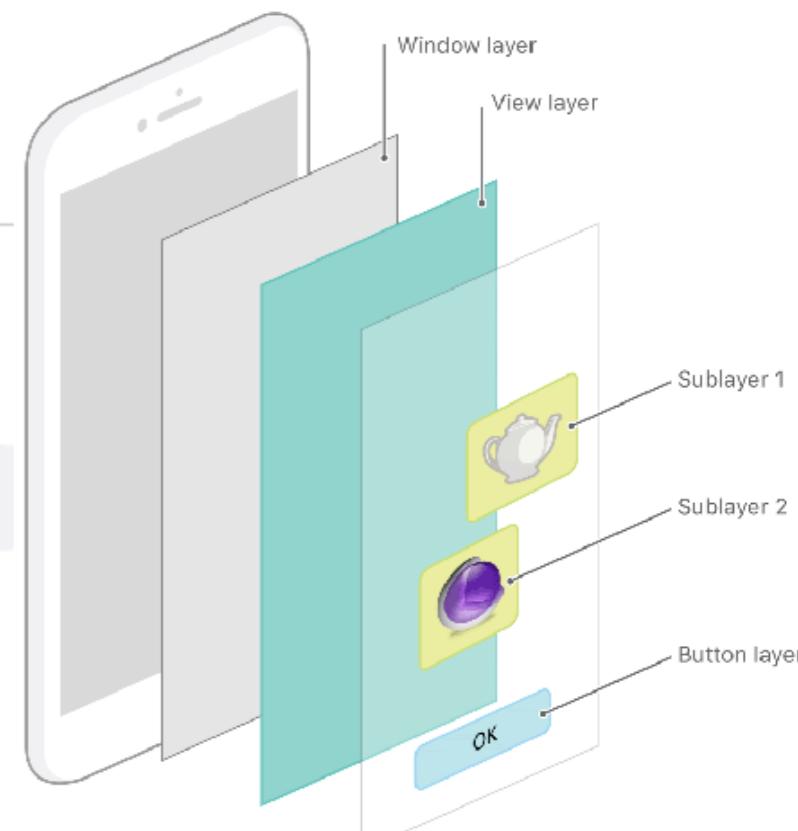
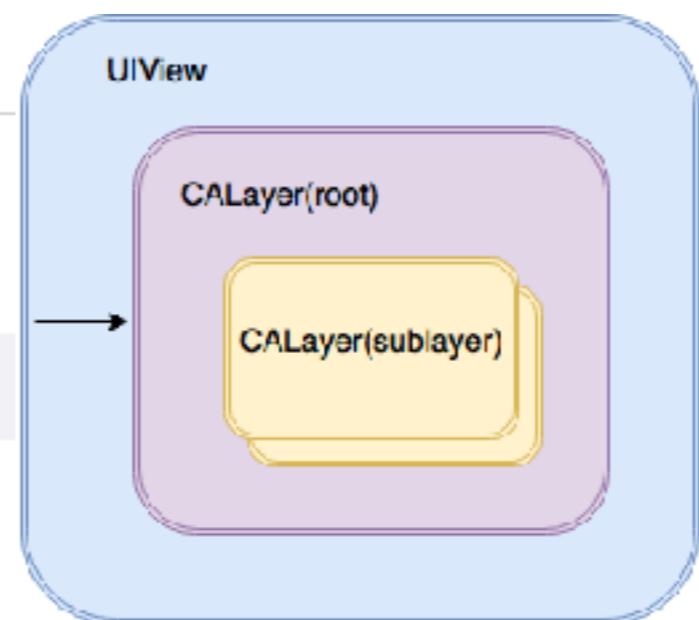
and now its time for a demo

# CATransactions

A mechanism for grouping multiple layer-tree operations into atomic updates to the render tree.

## Declaration

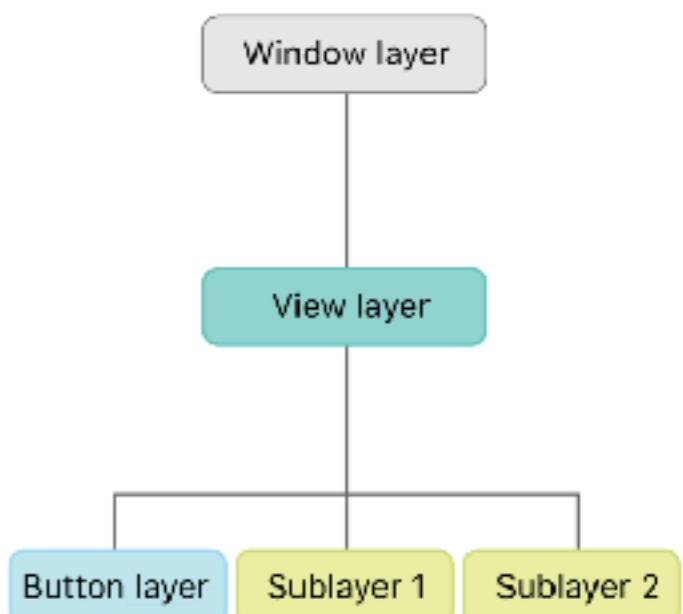
```
class CATransaction : NSObject
```



## Overview

**CATransaction** is the Core Animation mechanism for batching multiple layer-tree operations into atomic updates to the render tree. Every modification to a layer tree must be part of a transaction. Nested transactions are supported.

Core Animation supports two types of transactions: *implicit* transactions and *explicit* transactions. Implicit transactions are created automatically when the layer tree is modified by a thread without an active transaction and are committed automatically when the thread's runloop next iterates. Explicit transactions occur when the application sends the **CATransaction** class a **begin()** message before modifying the layer tree, and a **commit()** message afterwards.



# Overlays from YOLO

```
detectionOverlay = CALayer() // container layer that has all the renderings
detectionOverlay.name = "DetectionOverlay"

// set the initial bounds, will transform when we know more about the image
detectionOverlay.bounds = CGRect(x: 0.0, y: 0.0,
                                  width: self.view.bounds.width,
                                  height: self.view.bounds.height )

self.sceneView.layer.addSublayer(detectionOverlay)

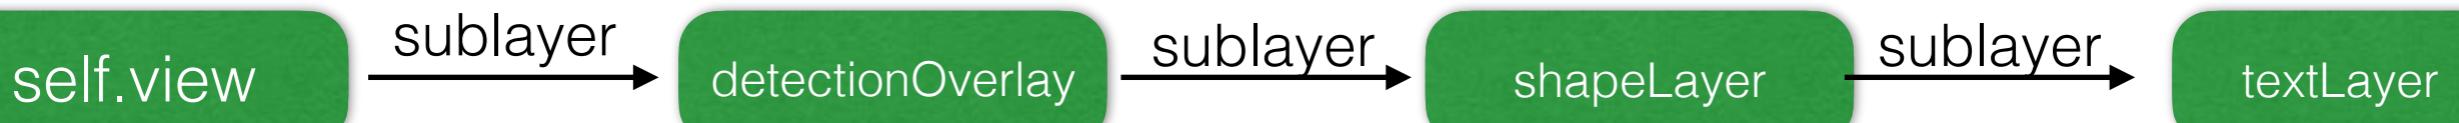
func updateOverlay() {

    let bounds = self.view.bounds

    // ... Some magic transforms to the view ...
    // this tries to get the best mapping we can from the cropping that
    // Core Vision used, but It may not be 100% perfect
    // Scales, and adds magic numbers to the X and Y positions

    detectionOverlay.setNeedsDisplay() // sets display for all subviews
}
```

# Overlays from YOLO



```
// now draw everything
CATransaction.begin()
detectionOverlay.sublayers = nil // remove all the old recognized objects

// for each result, create an overlay layer to display on it
for observation in results where observation is VNRecognizedObjectObservation {

    // Select the label with the highest confidence and get bounds
    let topLabelObservation = objectObservation.labels[0]
    let objectBounds = VNImageRectForNormalizedRect(objectObservation.boundingBox,
                                                    Int(captureImageSize.width), Int(captureImageSize.height))

    // get UI box for object
    let shapeLayer = self.createRoundedRectLayerWithBounds(objectBounds, ...)

    // show the label and confidence in the UI Box
    let textLayer = self.createTextSubLayerInBounds(objectBounds, ...)

    shapeLayer.addSublayer(textLayer)// add text to box
    detectionOverlay.addSublayer(shapeLayer) // add box to the UI
}

CATransaction.commit() // now commit everything so that it displays
```