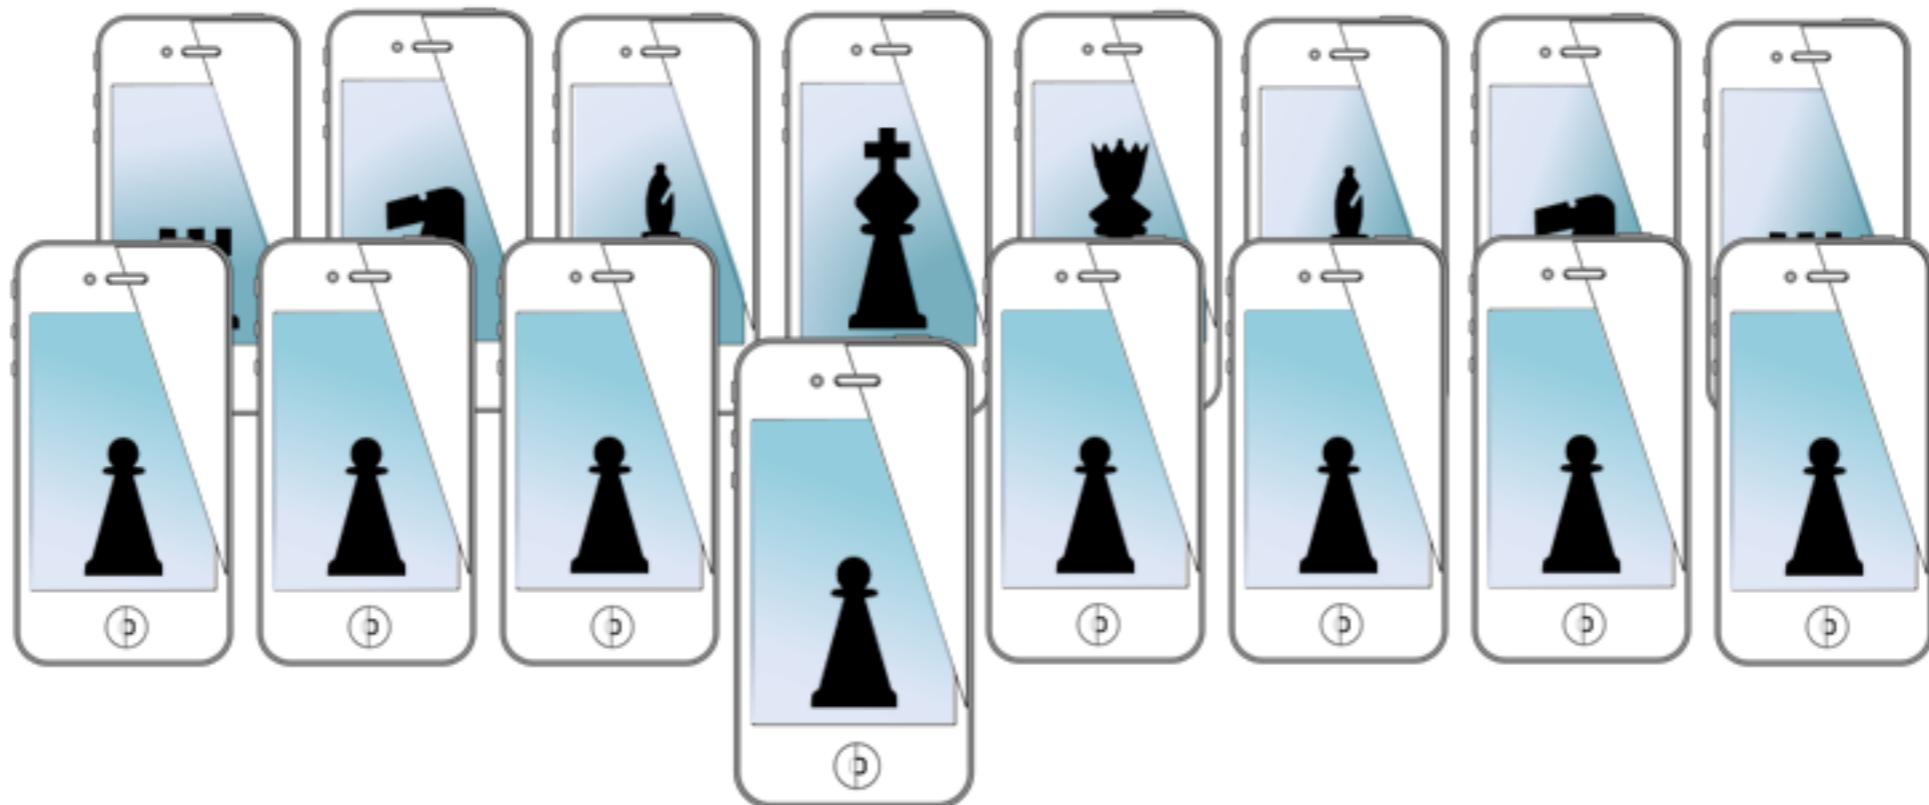


MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

audio graphing, sampled data, & accelerate

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

course logistics

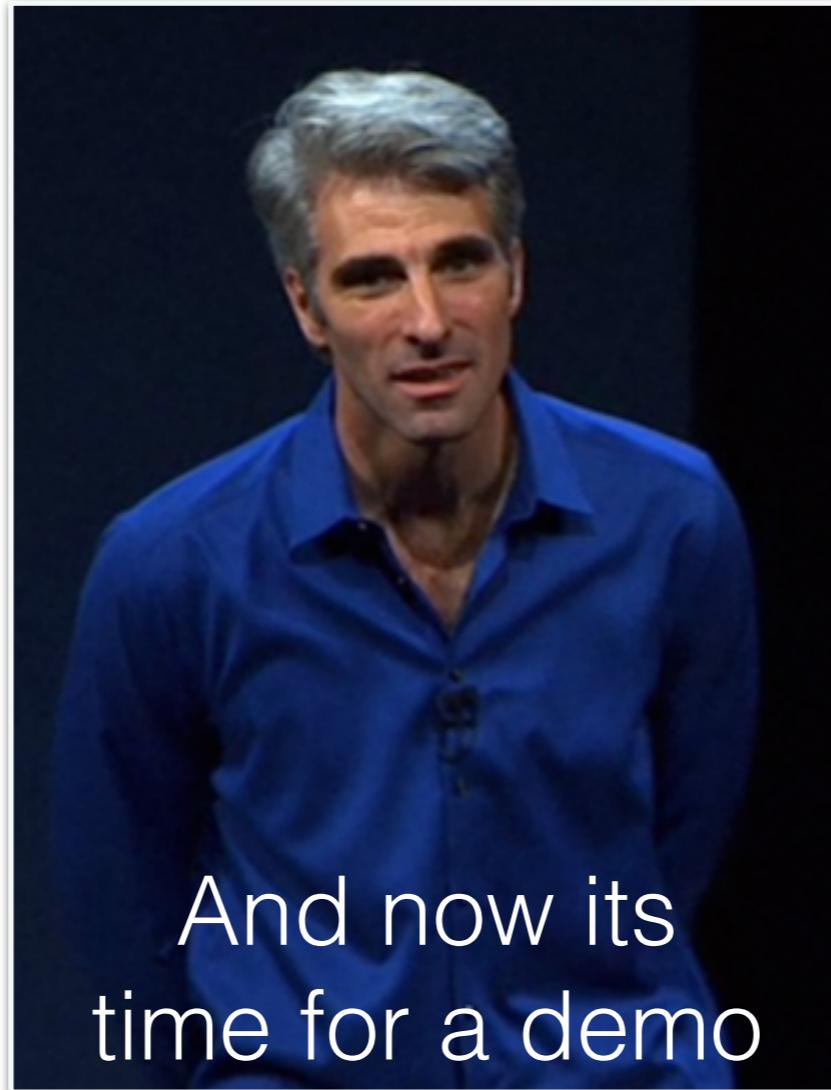
- flipped module on audio next time!

agenda

- dealing with sampled data
- the accelerate framework
 - massive digital signal processing library
- graphing audio fast (well, graphing anything)
 - must use lowest level graphing, Metal

sample from the mic

- recall: data from the microphone on novocaine



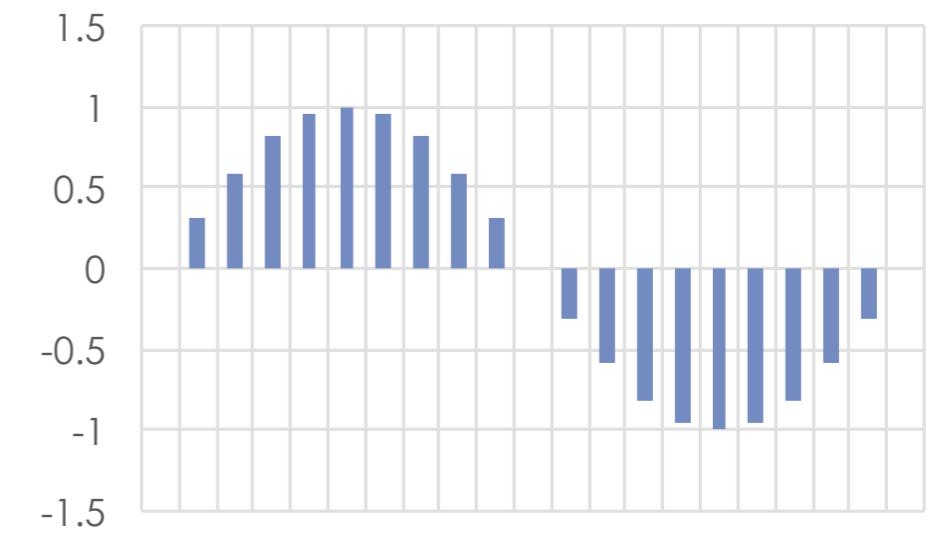
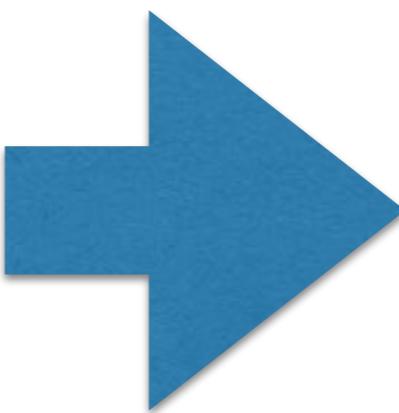
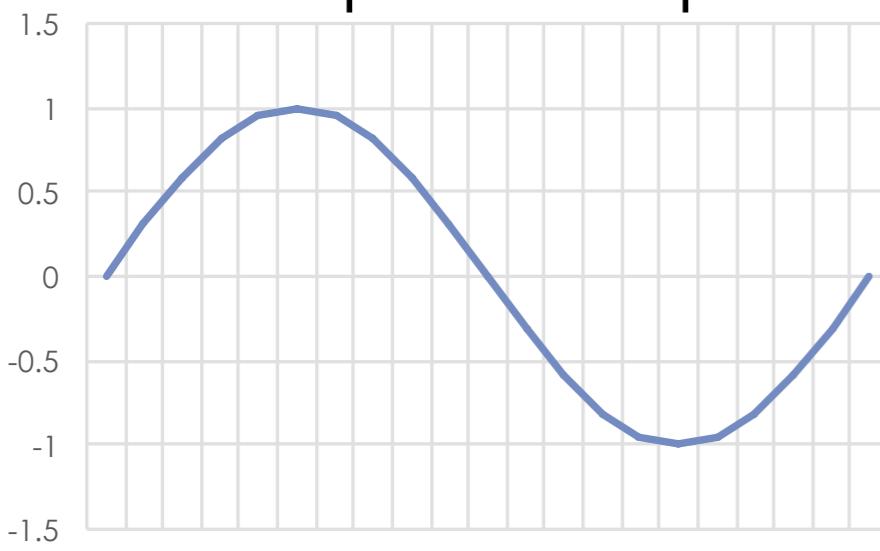
And now its
time for a demo

intro to sampled data

- why is understanding sampled data important?
 - because we'll be dealing with it all semester
 - it's important to understand basic mistakes that can be made
- there are entire courses dedicated to sampled time series
 - actually entire courses on analyzing frequency content
 - we'll touch on a few guidelines to help you design your projects better

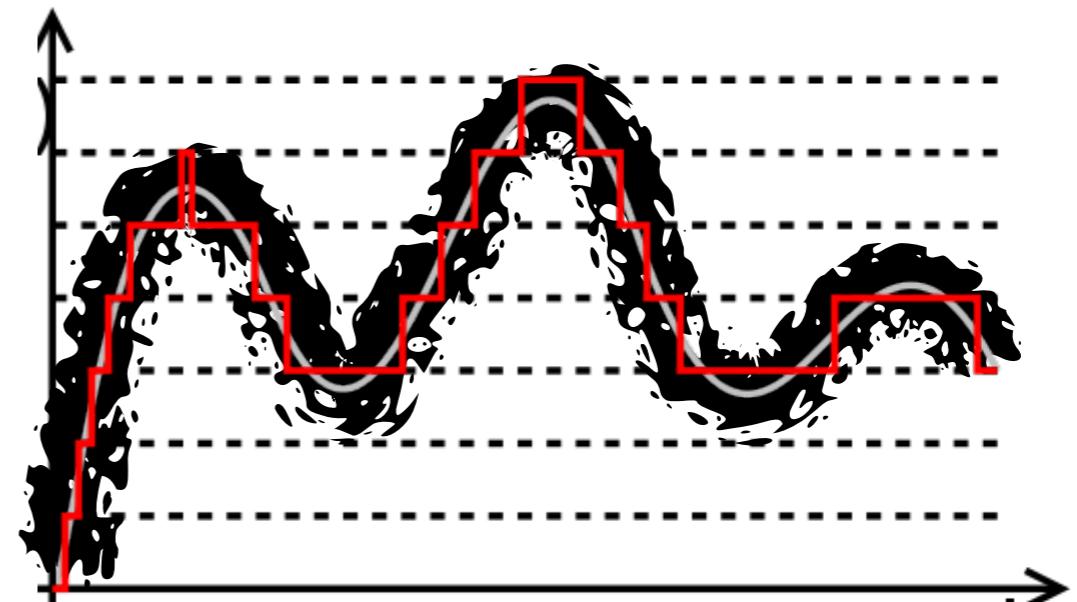
intro to sampled data

- physical processes are continuous
 - to process with computers, we must digitize it
 - **digitization can change how we understand the signal**
- digitization occurs in time and amplitude
 - time: sampling
 - amplitude: quantization



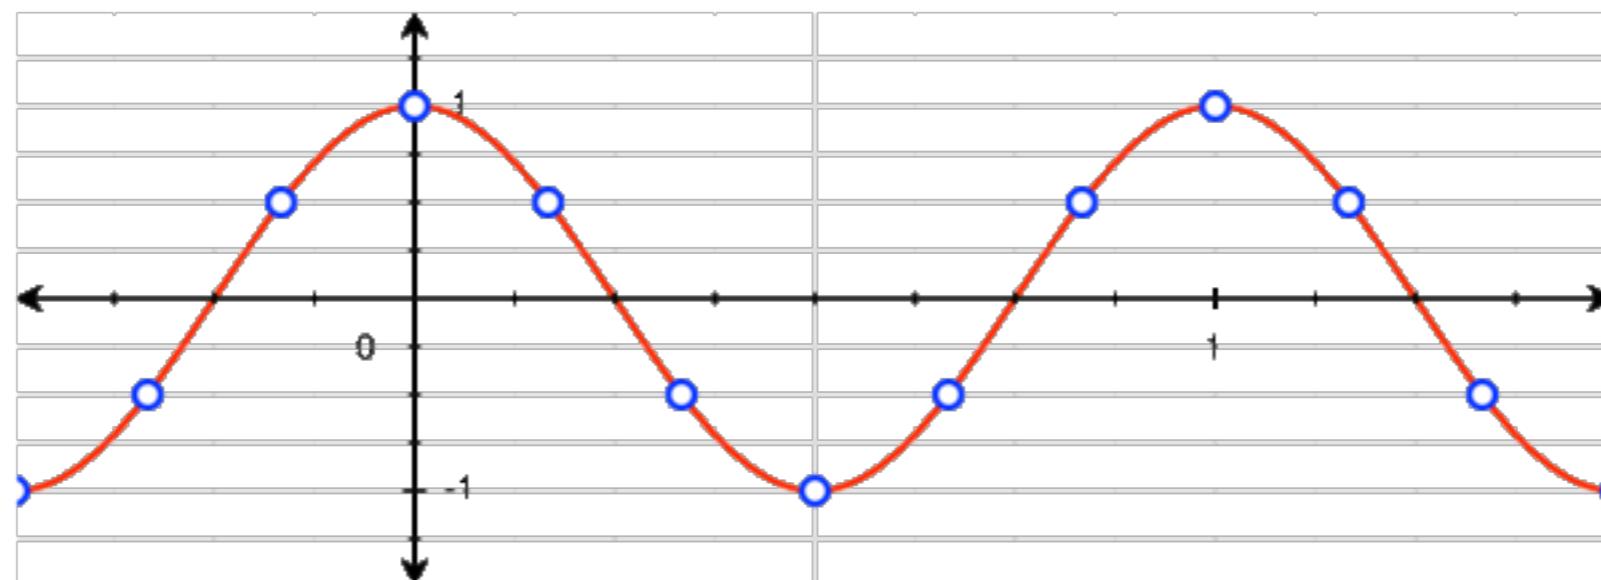
sampled data

- quantization (amplitude)
 - introduces error in estimating amplitude of a signal
 - error can be reduced by adding more “bits per sample”
- most ADCs are 16 bits, considered “good enough”
- sufficient for most uses
 - not for others!



sampling errors

- sampling in time
 - introduces errors through ‘aliasing’
 - limits the range of frequencies able to be accurately captured
 - root of most common mistakes with sampled data



so how do I sample?

- heuristics
 - don't try to sample extremely small increments or values!
 - if capturing an "X"Hz signal, need to sample at least 2"X" Hz
 - changing sample rates is complicated, don't just drop every other sample
- for example, speech
 - majority of necessary energy in speech is located < 8000Hz
 - phones (for speech) typically capture at 16KHz or lower
 - good enough for speech, not music!

sanity check

- I need to detect an 80Hz signal
 - what sampling rate should we use?
- I want to detect a feather dropping next to the microphone
 - can the sound be detected?

making a sine wave

- we want to create a sine wave and play it to the speakers

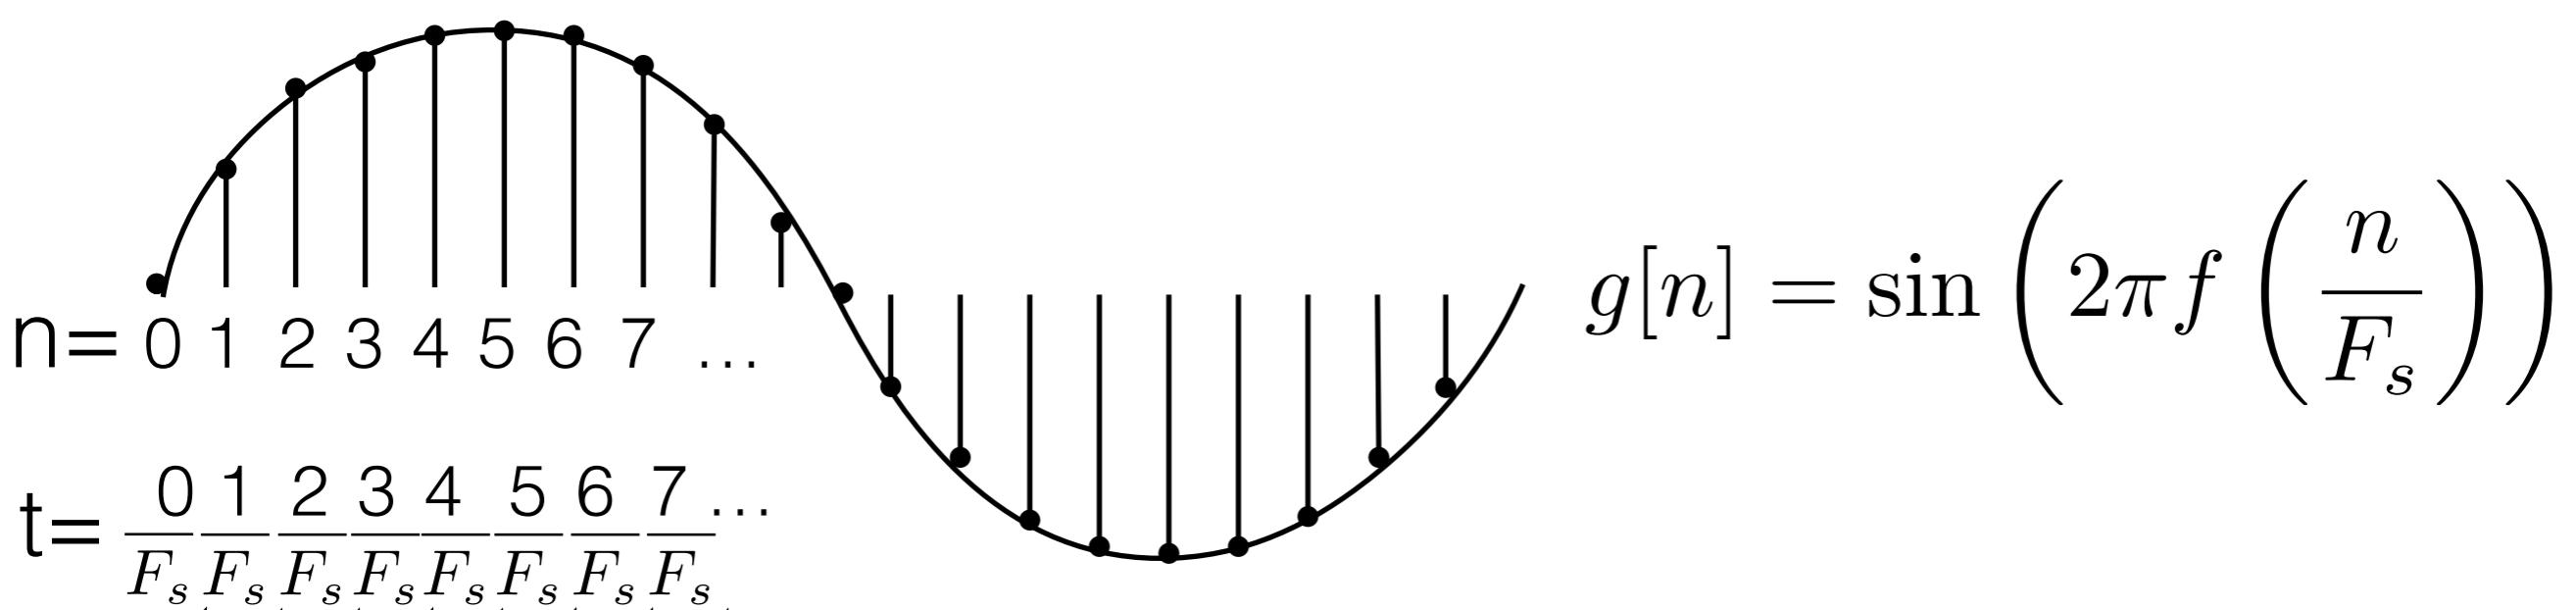
$$g(t) = \sin(2\pi ft)$$

equation for sine wave

frequency in Hz

time in “seconds”

but we are working digitally, so we have an “index” in an array, not time!



making a sine wave

$$g[n] = \sin\left(2\pi f \left(\frac{n}{F_s}\right)\right)$$

how to program this?

```
for (int n=0; n < numFrames; ++n)
{
    data[n] = sin(2*M_PI*frequency*n/samplingRate);
}
```

is this efficient?

```
float phase = 0.0;
double phaseIncrement = 2*M_PI*frequency/samplingRate;
for (int n=0; n < numFrames; ++n)
{
    data[n] = sin(phase);
    phase += phaseIncrement;
}
```

making a sine wave

- bringing it all together

$$g[n] = \sin\left(2\pi f \left(\frac{n}{F_s}\right)\right)$$

```
var frequency = 18000.0; //starting frequency
var phase = 0.0;
var samplingRate = audioManager.samplingRate;

outputBlockFunction(data:(...),numFrames:(UInt32),numChannels:(UInt32))
{
    var phaseIncrement = 2*Double.pi*frequency/samplingRate
    var i=0;
    var sineWaveRepeatMax = 2*Double.pi;
    while (i < numFrames)
    {
        data![i] = sin(phase);
        i += 1
        phase += phaseIncrement;
        if (phase >= sineWaveRepeatMax) phase -= sineWaveRepeatMax;
    }
}
```

data: UnsafeMutablePointer<Float>?

play samples to speakers

- demo, play sine wave



And now its
time for a demo



the accelerate framework

- very powerful digital signal processing (DSP) library
 - look at vDSP Programming Guide on developer.apple.com for the complete API
- provides mathematics for performing fast DSP

```
vDSP_vsmul(data, 1, &mult, data, 1, numFrames*numChannels);
```

```
void vDSP_vsmul (
    const float __vDSP_input1[],
    vDSP_Stride __vDSP_stride1,
    const float *__vDSP_input2,
    float __vDSP_result[],
    vDSP_Stride __vDSP_strideResult,
    vDSP_Length __vDSP_size
);
```

examples

what do each of these implement?

```
inputBlockFunction(data:(...),numFrames:(UInt32),numChannels:(UInt32)) {  
    var volume = userSetMultiplyFromSlider;  
    vDSP_vsmul(data, 1, &volume, data, 1, numFrames*numChannels)  
    ringBuffer.addNewInterleavedFloatData(data,withNumFrames:numFrames)  
}
```

```
inputBlockFunction(data:(...),numFrames:(UInt32),numChannels:(UInt32)) {  
    // get the max  
    var maxVal = 0.0;  
    vDSP_maxv(data, 1, &maxVal, numFrames*numChannels);  
  
    print("Max Audio Value: %f\n", maxVal);  
}
```

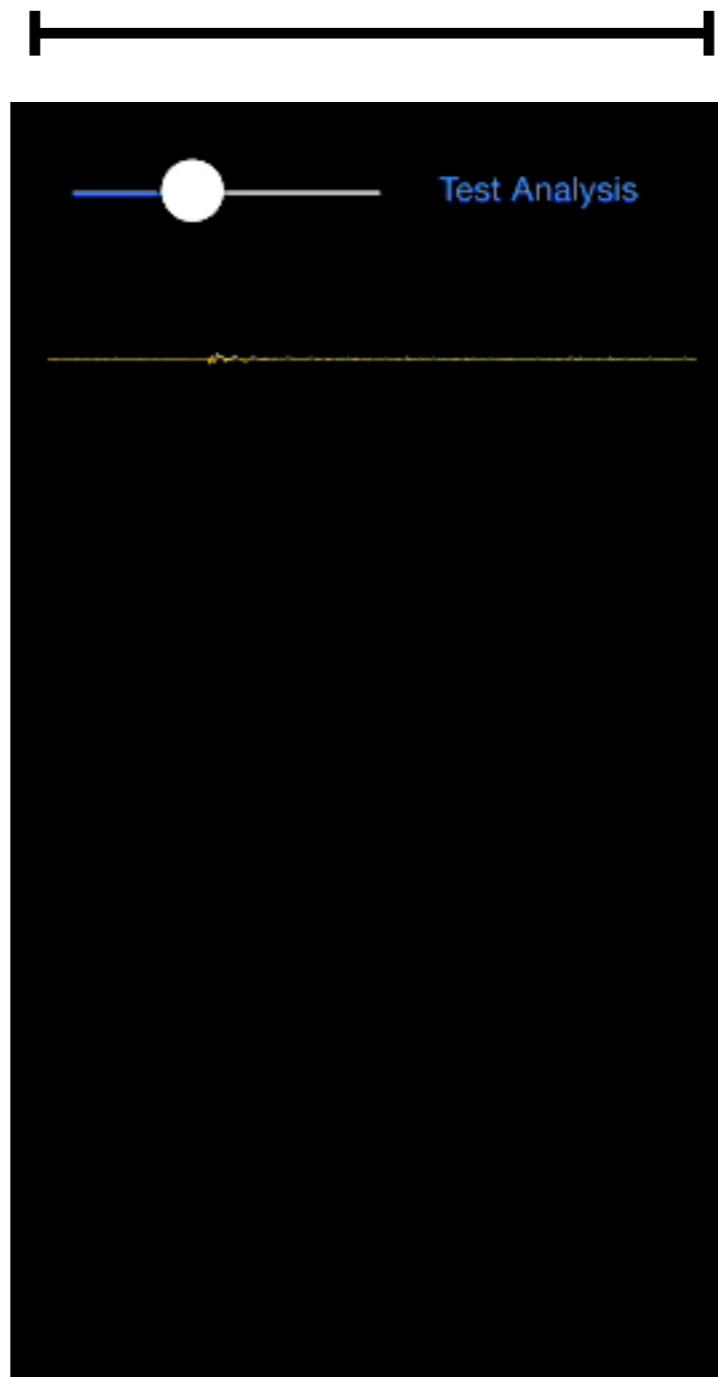
```
inputBlockFunction(data:(...),numFrames:(UInt32),numChannels:(UInt32)) {  
    vDSP_vsq(data, 1, data, 1, numFrames*numChannels);  
    var meanVal = 0.0;  
    vDSP_meanv(data, 1, &meanVal, numFrames*numChannels);  
}
```

audio graphing

- we want to see the incoming samples
 - good for debugging
 - equalizers
 - oscilloscope type applications

how much data to show?

- sampling at 44.1kHz == 44100 samples per second



0.5 seconds is
22050 samples

display is >640
pixels wide

what if we want
lots of graphs?

solution

- use the GPU
- set vectors of data on a 2D plane
- let the renderer perform scaling, anti-aliasing, and bit blitting to screen
- ...this is not a graphics course
- but we need to use the Metal API

Metal	
	Apple used the mobile multiplayer online battle arena game Vainglory to demonstrate Metal's graphics capabilities at the iPhone 6's September 2014 announcement event ^[1]
Developer(s)	Apple Inc.
Initial release	June 2014; 6 years ago
Stable release	3 / June 2019; 1 year ago
Written in	Shading Language: C++14 , Runtime/API: Objective-C
Operating system	iOS , iPadOS , macOS , tvOS
Type	3D graphics and compute API
License	proprietary
Website	developer.apple.com/metal/ ↗

the metalGraph class

ViewController.swift
MetalGraph.swift
Shaders.metal

drag class/shaders
into project, if needed

```
lazy var graph: MetalGraph? = {  
    return MetalGraph(mainView: self.view)  
}()
```

declare and init property

```
// add in a graph for displaying the audio  
graph?.addGraph(withName: "time",  
    shouldNormalize: false,  
    numPointsInGraph: AUDIO_BUFFER_SIZE)
```

add graph names to controller
and how many expected points in array

```
// periodically, display the audio data  
graph?.updateGraph(  
    data: timeData,  
    forKey: "time")
```

refresh data for each
named graph key

Properties: automatic screensize (pixel) downsampling
automatic coloring based in iOS scheme,
efficient memory management through vertex buffers
adding functionality is a pain if you are new to graphics

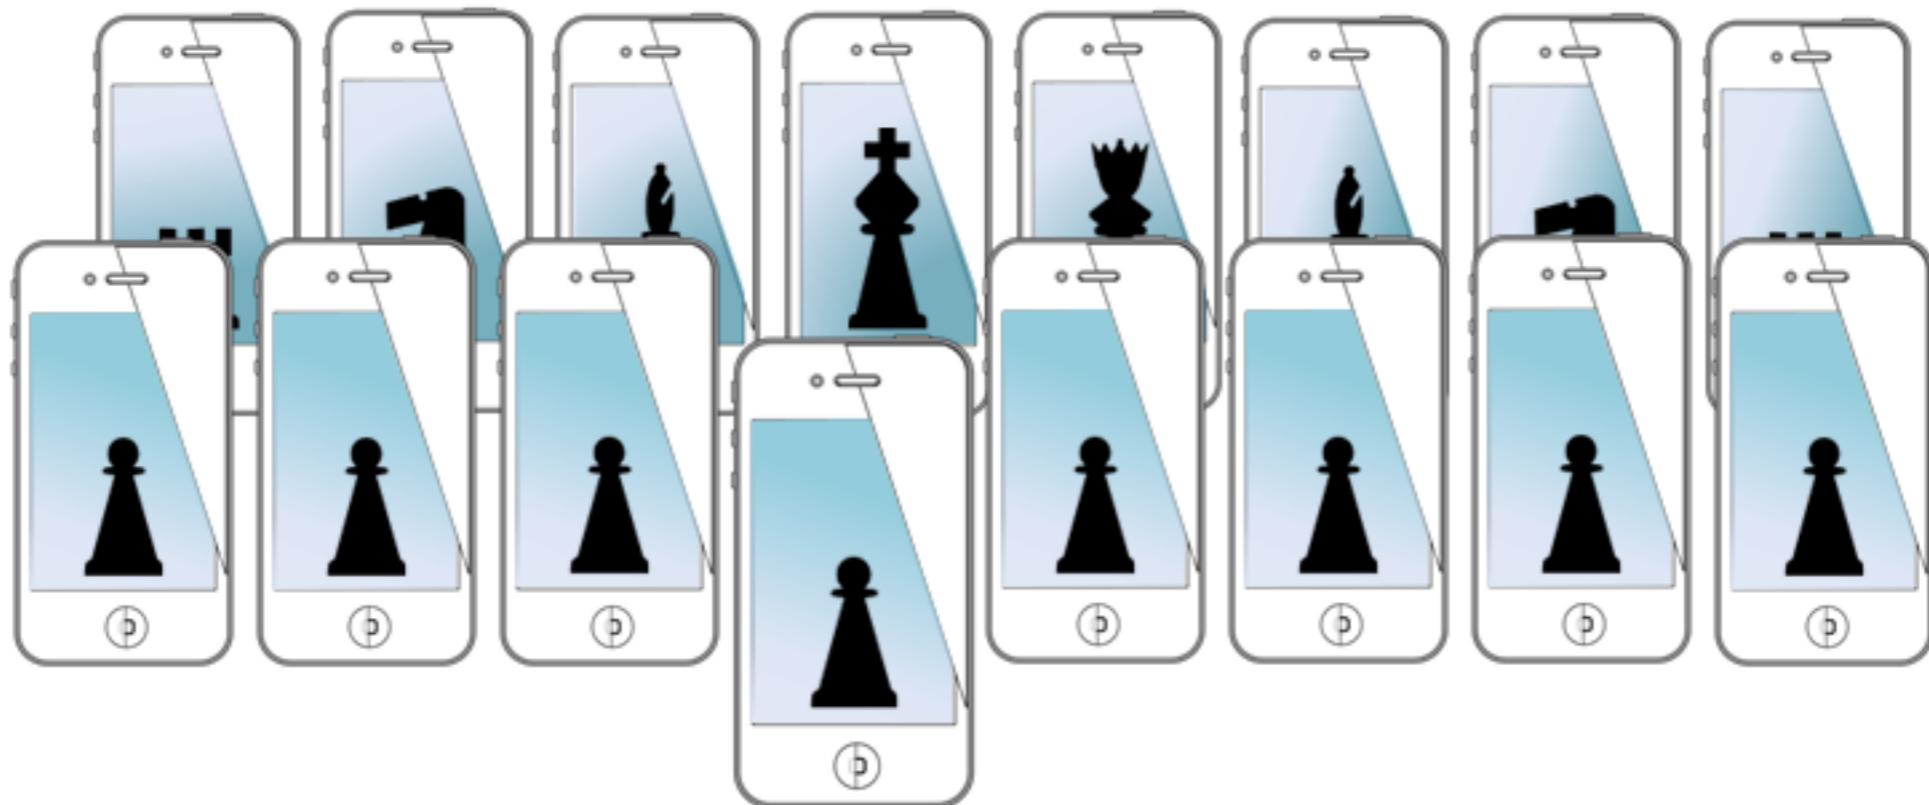
audio graphing demo!



This class needs
better memes!

And now its
time for a demo

MOBILE SENSING LEARNING



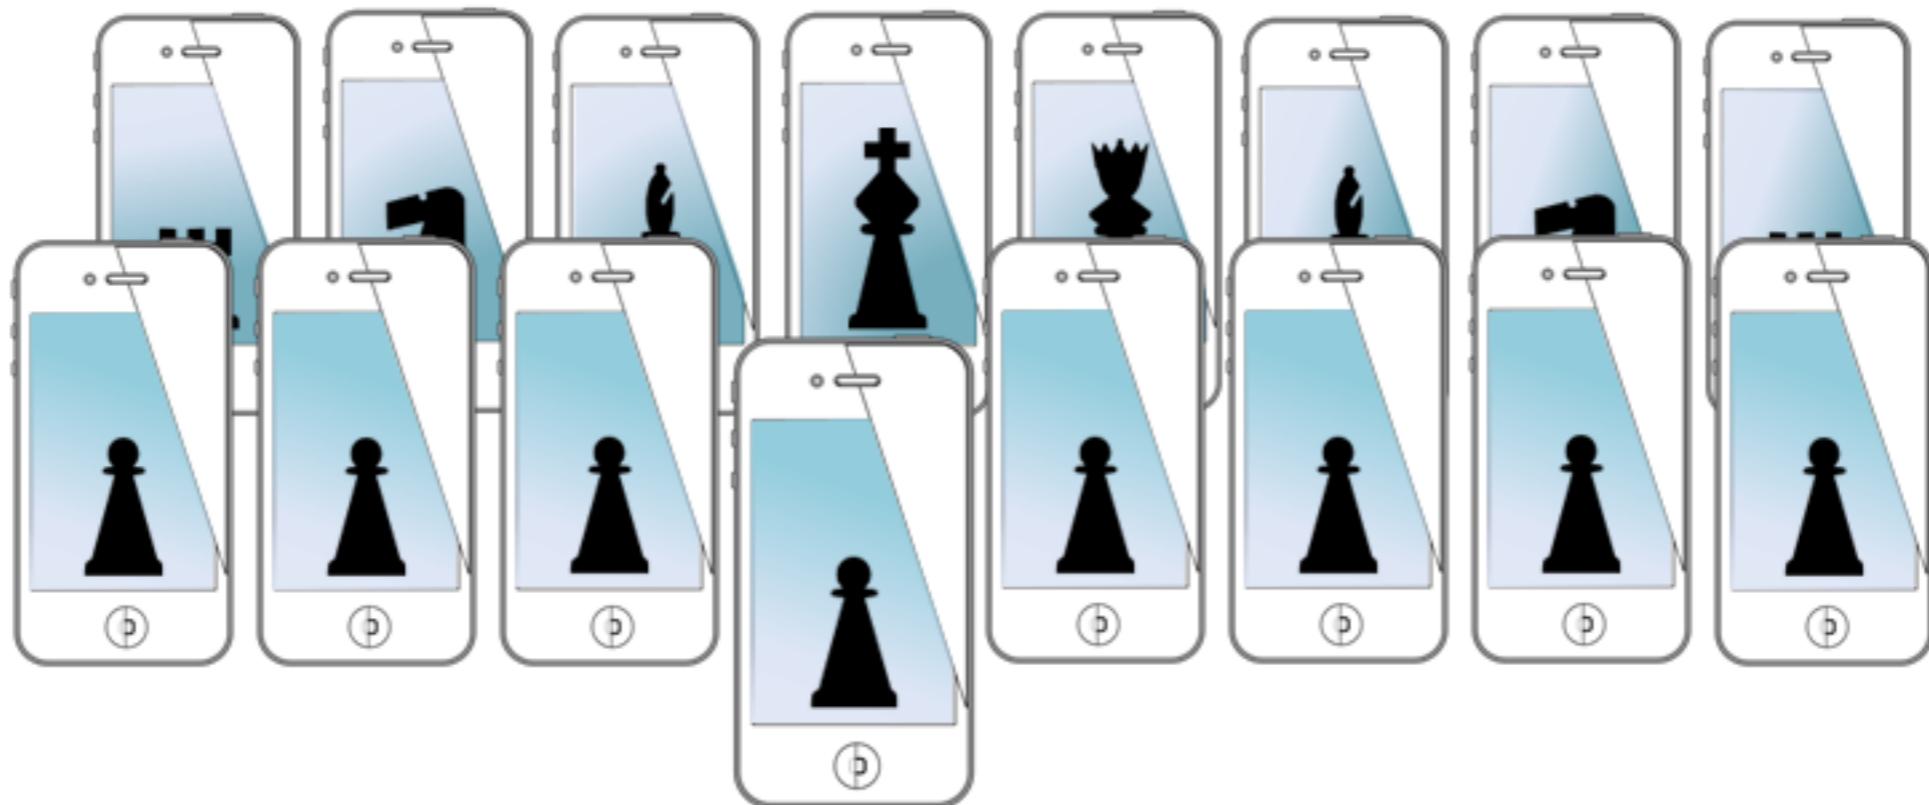
CS5323 & 7323

Mobile Sensing and Learning

audio graphing, sampled data, & accelerate

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

MOBILE SENSING LEARNING



CS5323 & 7323
Mobile Sensing and Learning

video lecture: accelerate & FFT

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

agenda (Video)

- high level: what is audio processing
- audio analysis: the FFT
- how to you use the FFT
- how to program the FFT
- an example: graphing whistles

processing audio

- lots of space to explore, processing signals is big!
 - **great reference:** “DSP First” by McClellan, Schafer, and Yoder
 - <http://www.rose-hulman.edu/DSPFirst/visible3/contents/index.htm>
 - **a reference for CSE’s:** “Signal Computing” by Stiber, Stiber, and Larson
 - <http://faculty.washington.edu/stiber/pubs/Signal-Computing/>
- filtering
 - only let certain frequencies through
- analysis
 - analyze characteristics of signal (like pitch or amplitude)
- synthesis
 - play around with different ideas, and see what sounds good!
 - not just pure synthesis, but also manipulation (like guitar effects)

processing audio

- the accelerate framework is great for analyzing audio
- for instance, what is the max value of an audio chunk?

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels) {  
    // get the max  
    float maxVal = 0.0;  
    vDSP_maxv(data, 1, &maxVal, numFrames*numChannels);  
    printf("Max Audio Value: %f\n", maxVal);  
}
```

input data from microphone

output: max

array length

max of audio might be used to detect:

taps next to **phone**

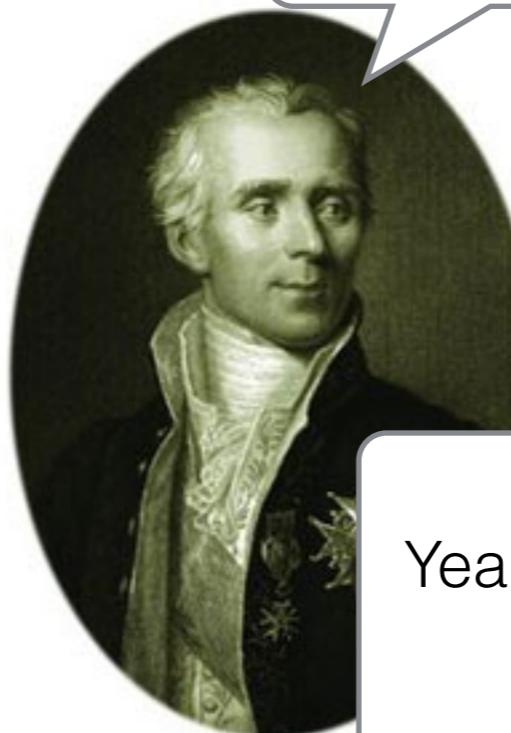
- for now, we're going to stick with analysis
 - specifically, the **Fourier Transform**

Fourier



A mathematician
Worked on
Submitted
in 1822

Did you read that
Fourier paper?



Laplace



Lagrange

Yeah! I was like...
hated it

Fourier



A mat
Worker
Subm
in 18

He said all signals could
be defined as a sum of
simpler functions



Laplace



Lagrange

Simpler functions!? Like
sine waves? Come on!

Fourier



I created the FFT to commit espionage against the Russians!

He developed a method that allowed one to transform a signal into a **sum of sine waves**.

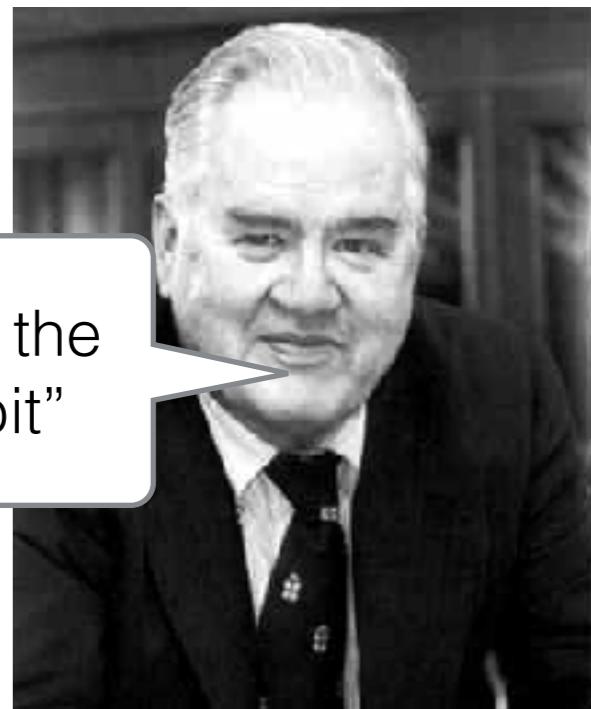
The transform was eventually called The Fourier Transform

~150 years later Cooley and Tukey would discover a way to calculate it really fast on a computer



Cooley

I coined the term “bit”



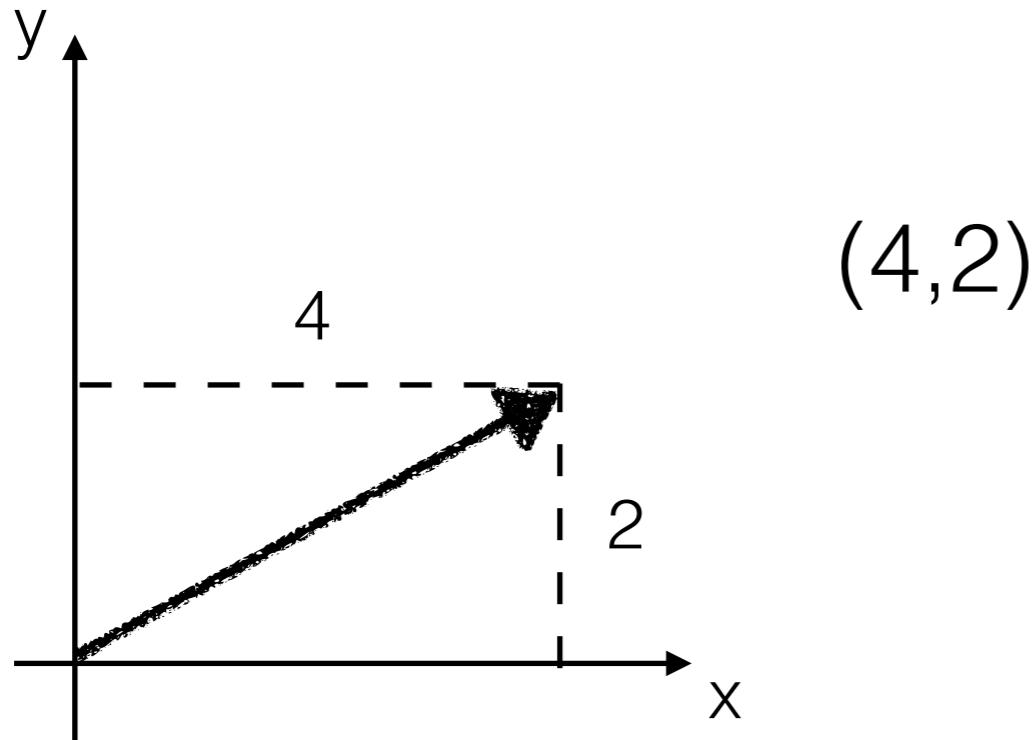
Tukey

the Fourier transform

- extremely useful, not just for signal junkies but also:
 - computer scientists, engineers, physicists, mathematicians, astronomers, oceanographers, health care professionals, etc.
- the Fourier Transform (FT) converts a time series into a frequency spectrum
 - the spectrum is an array of complex numbers which we will represent in polar form (i.e., with magnitude and phase)
 - each complex number represents a sinusoidal wave at a specific frequency
- we will use the FFT in the accelerate framework
 - complexity is $O(N \log_2(N))$ (for radix 2 FFT)

FT by vector intuition

think of it as a vector projection (intuitively)



where did these numbers come from?

$$\vec{x} \quad (1,0) \bullet (4,2) = 4$$
$$\vec{y} \quad (0,1) \bullet (4,2) = 2$$

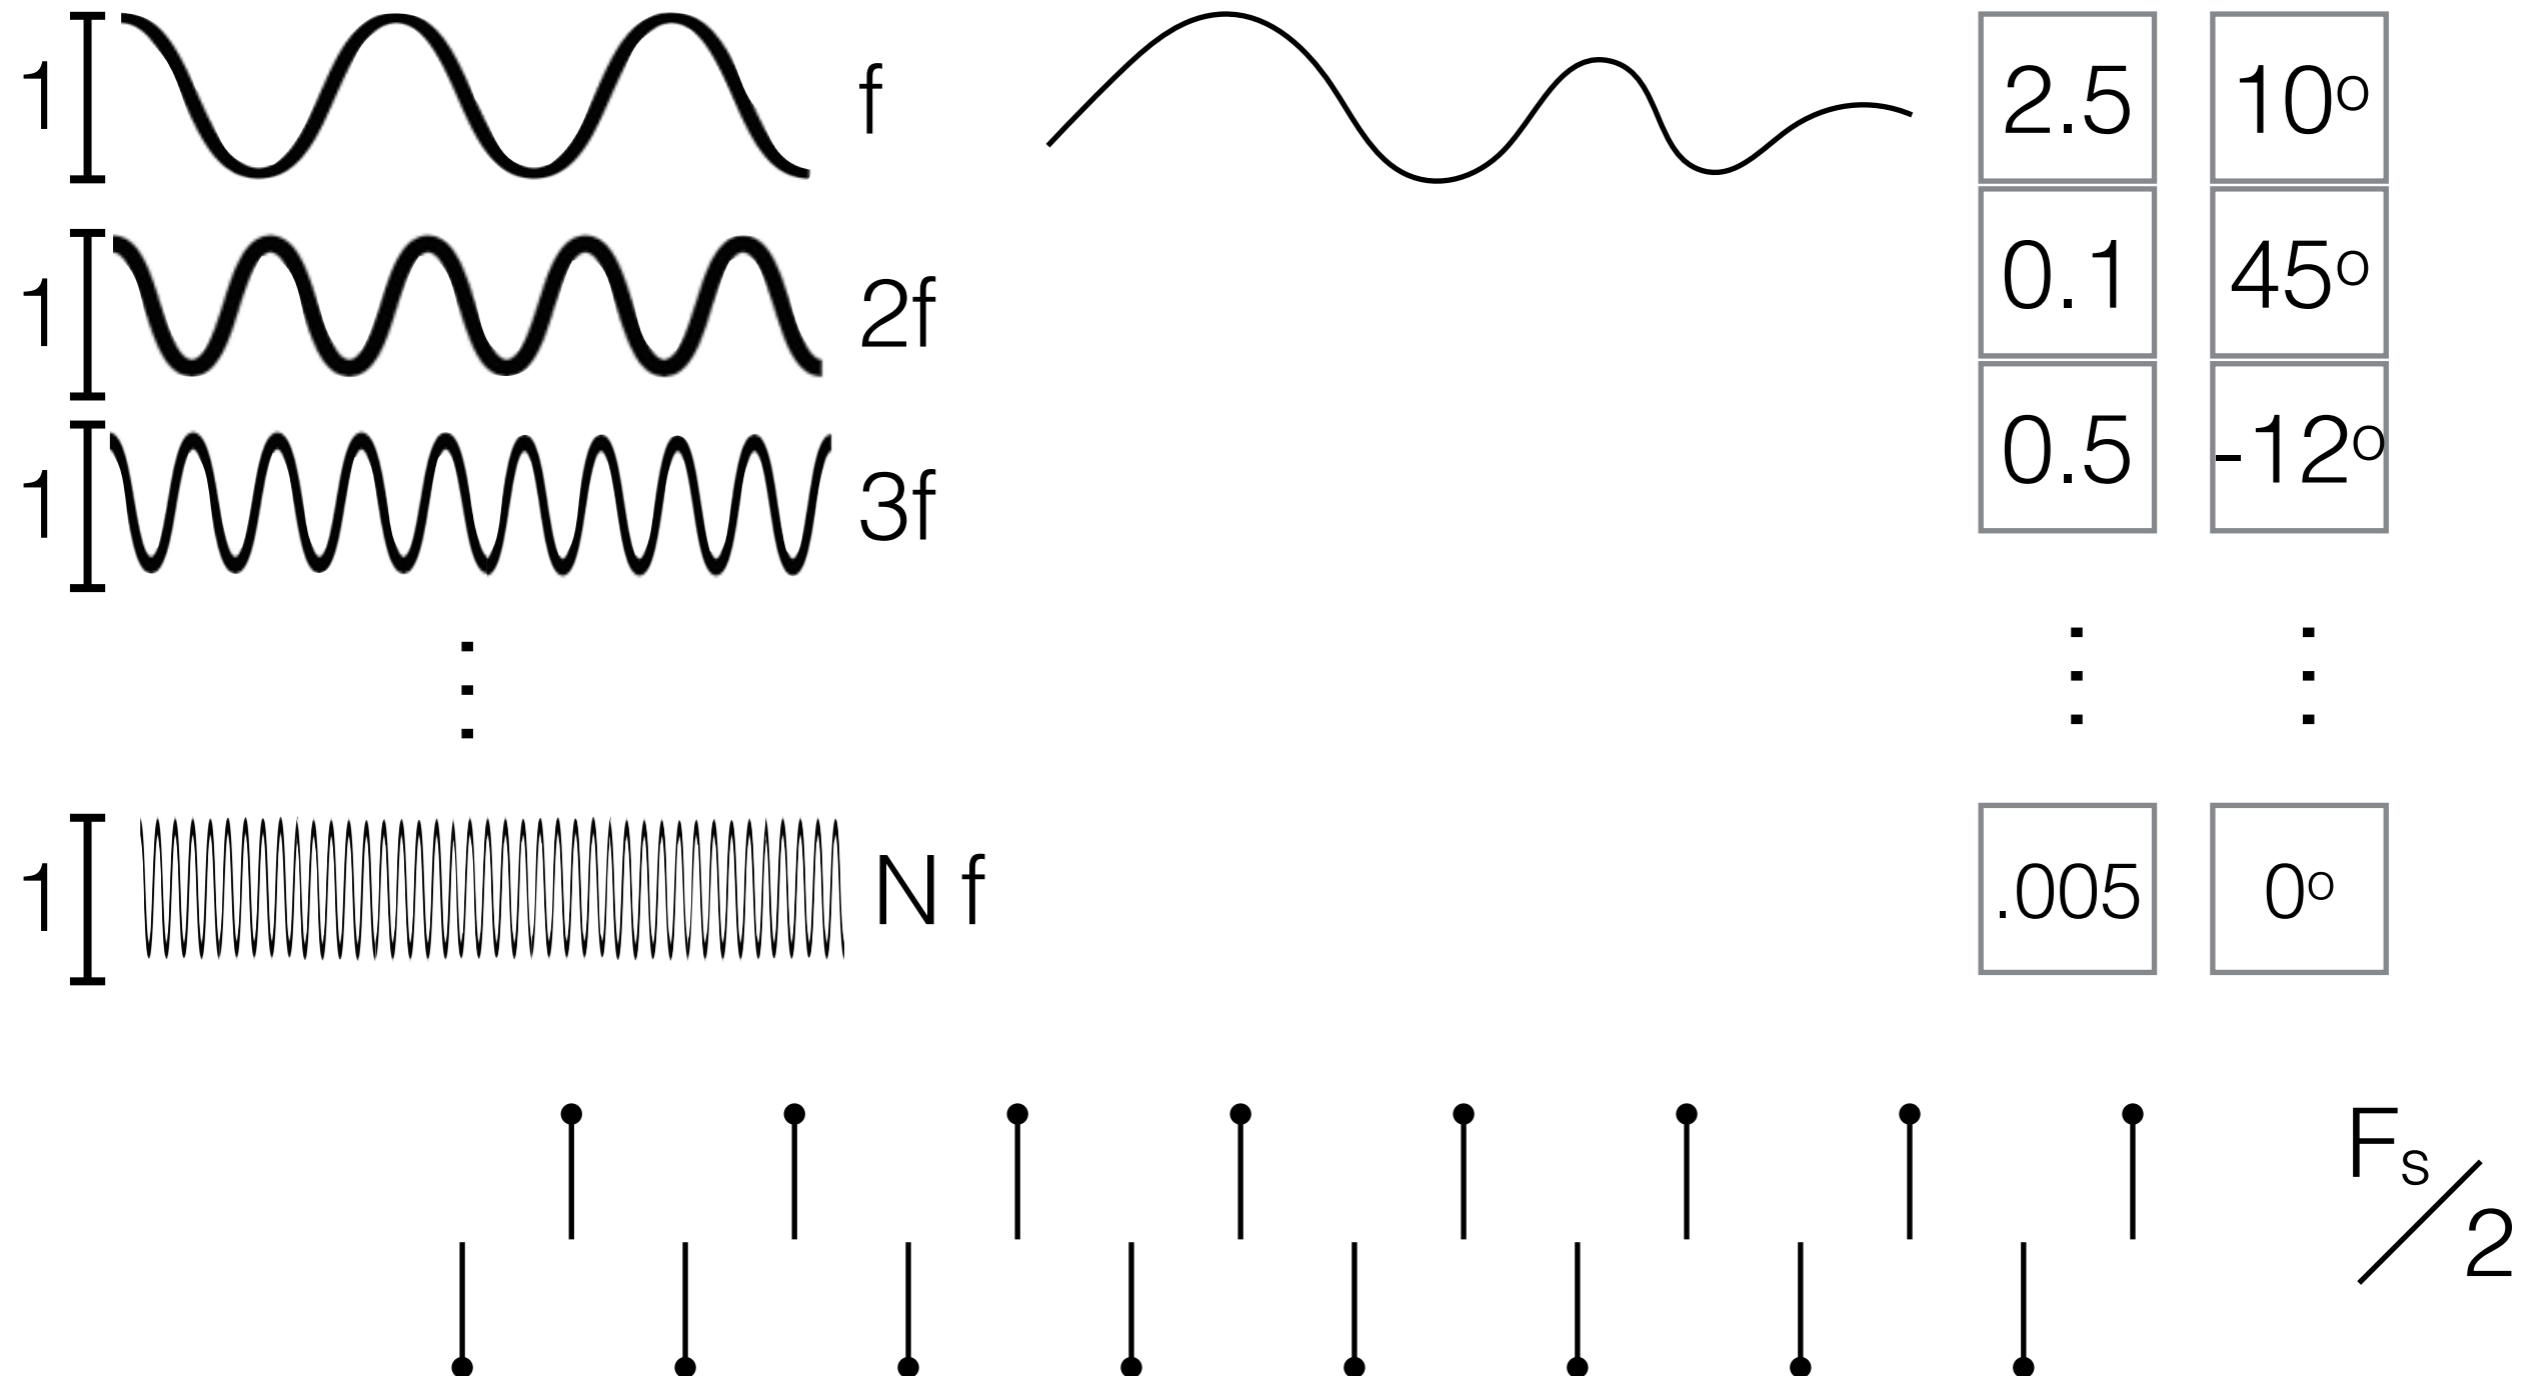
point by point multiplication and add it up!

tells us “how much” of the vector is the result of these vectors

$$\vec{x} \quad \vec{x} \bullet \vec{y} = 0$$
$$\vec{y} \quad |\vec{x}| = |\vec{y}| = 1$$

FT by sine waves

what if the orthogonal vectors were functions?



FT by the numbers

frequency content

0f	[0.7]	[0°]	=	0.7
1f	[2.5]	[10°]	+ [2.5]	-1f ~ 2.5 cos(2pi (f) t+10°)
2f	[0.1]	[45°]	+ [0.1]	-2f ~ 0.1 cos(2pi (2f) t+45°)
3f	[0.5]	[-12°]	+ [0.5]	-3f ~ 0.5 cos(2pi (3f) t-12°)
:	:		:	
N f	[.005]	[0°]	+ [.005]	-N f ~ 0.005 cos(2pi (Nf) t)

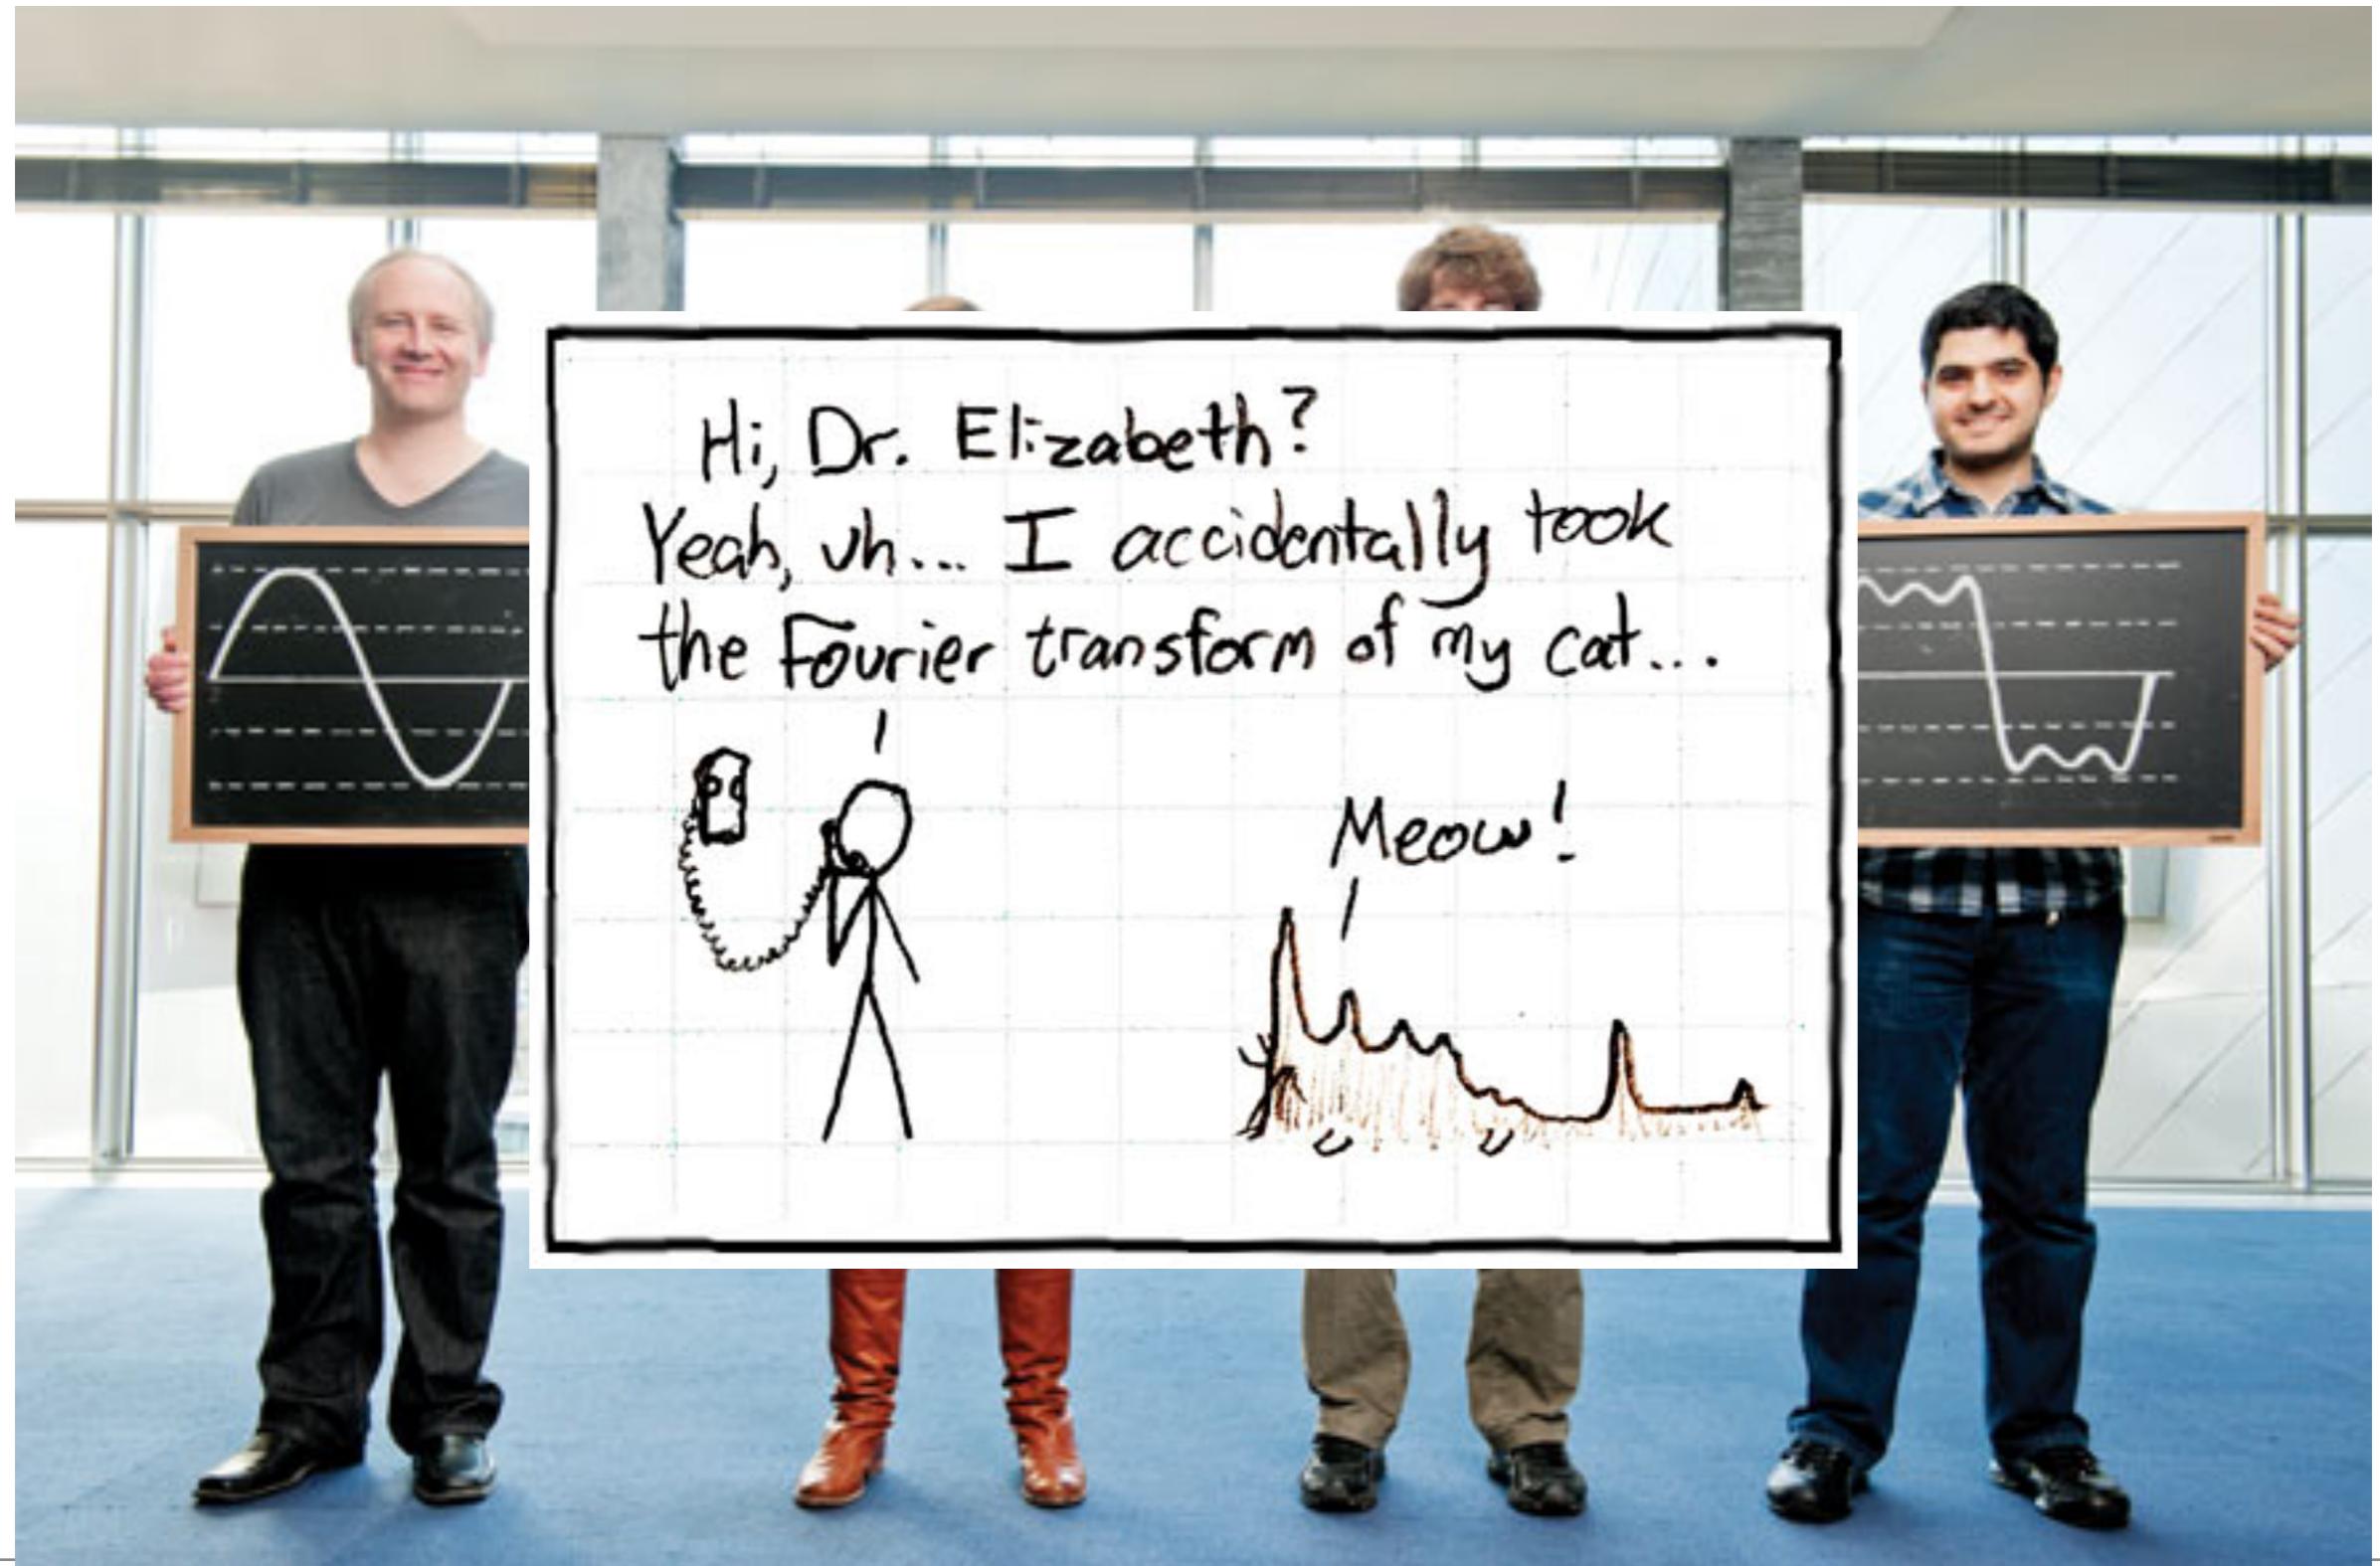
FT by video

<http://en.wikipedia.org/wiki/User:LucasVB/Gallery>



FT by hilarity

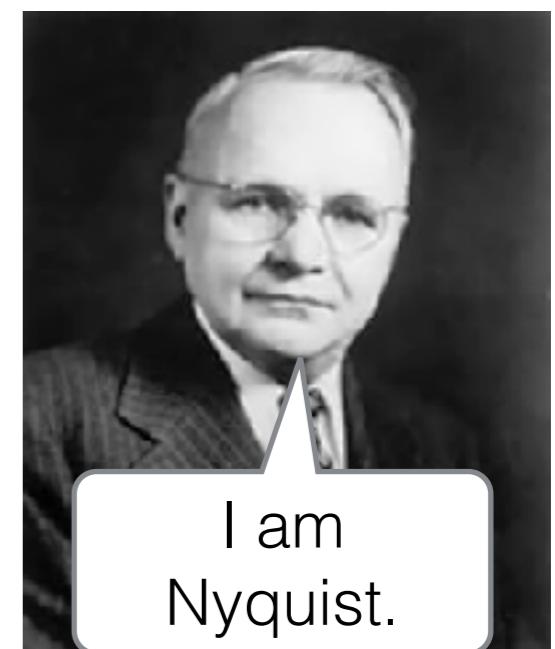
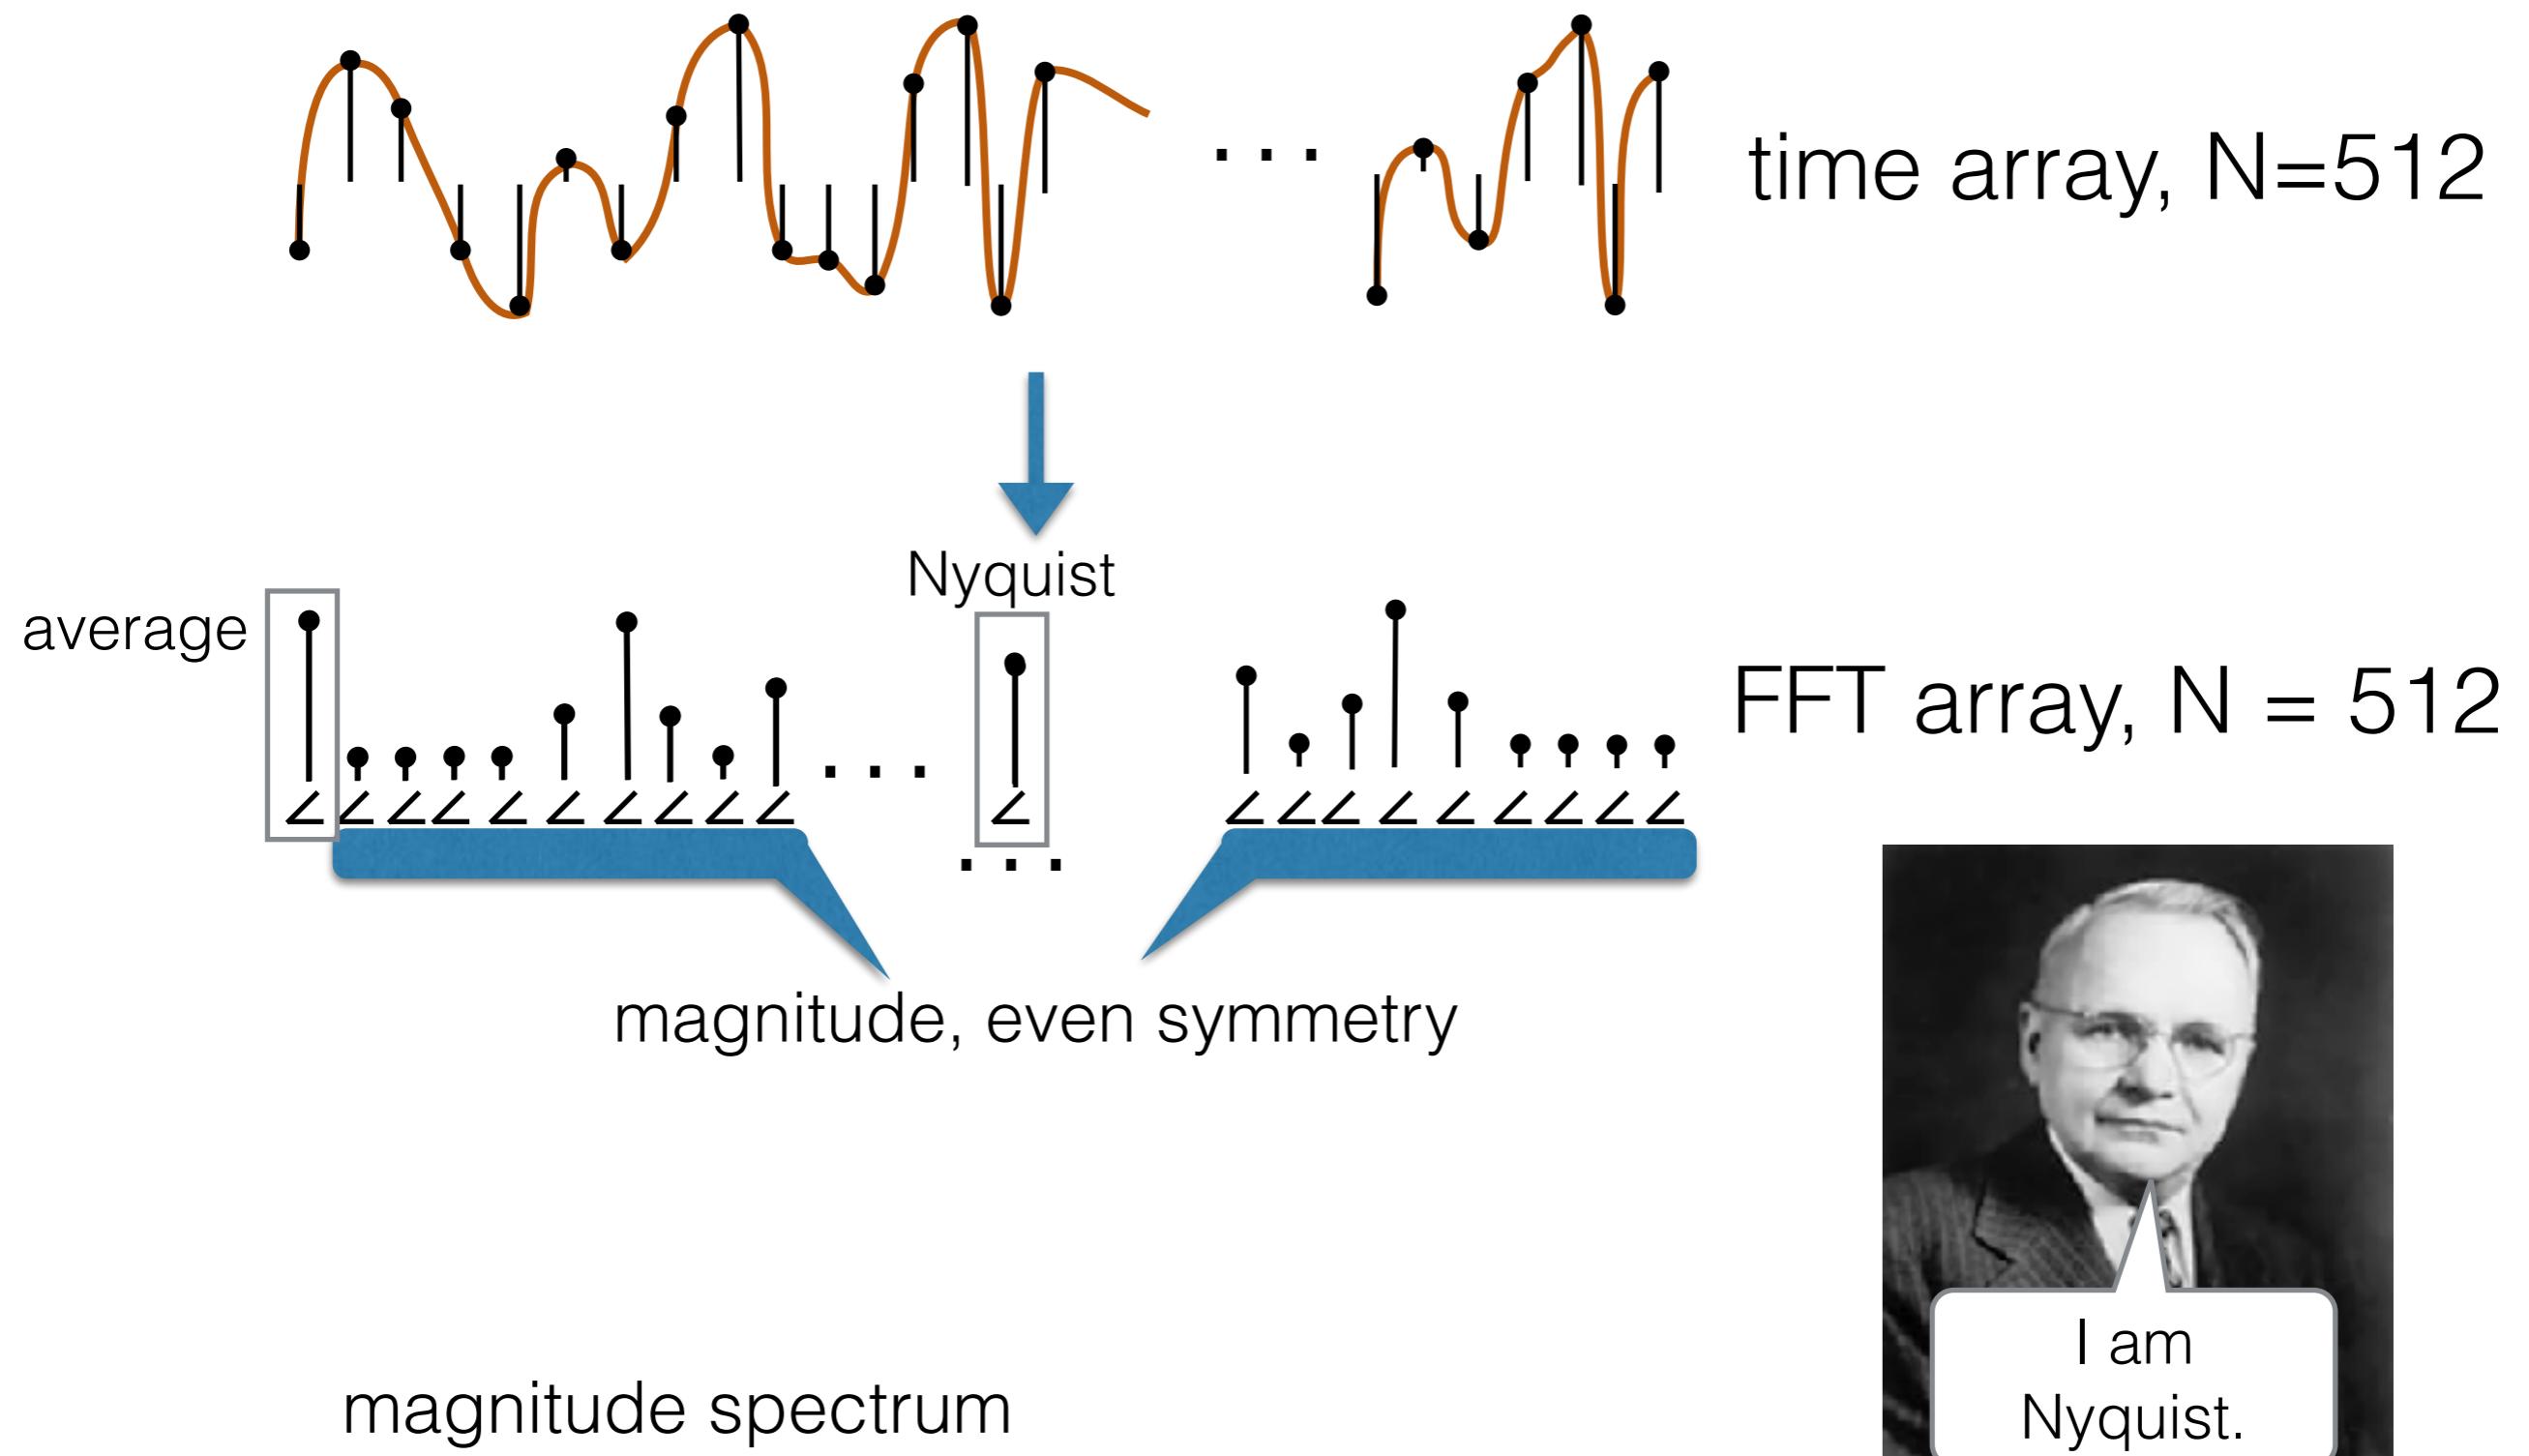
<http://www.preposterousuniverse.com/blog/2014/11/27/thanksgiving-9/>



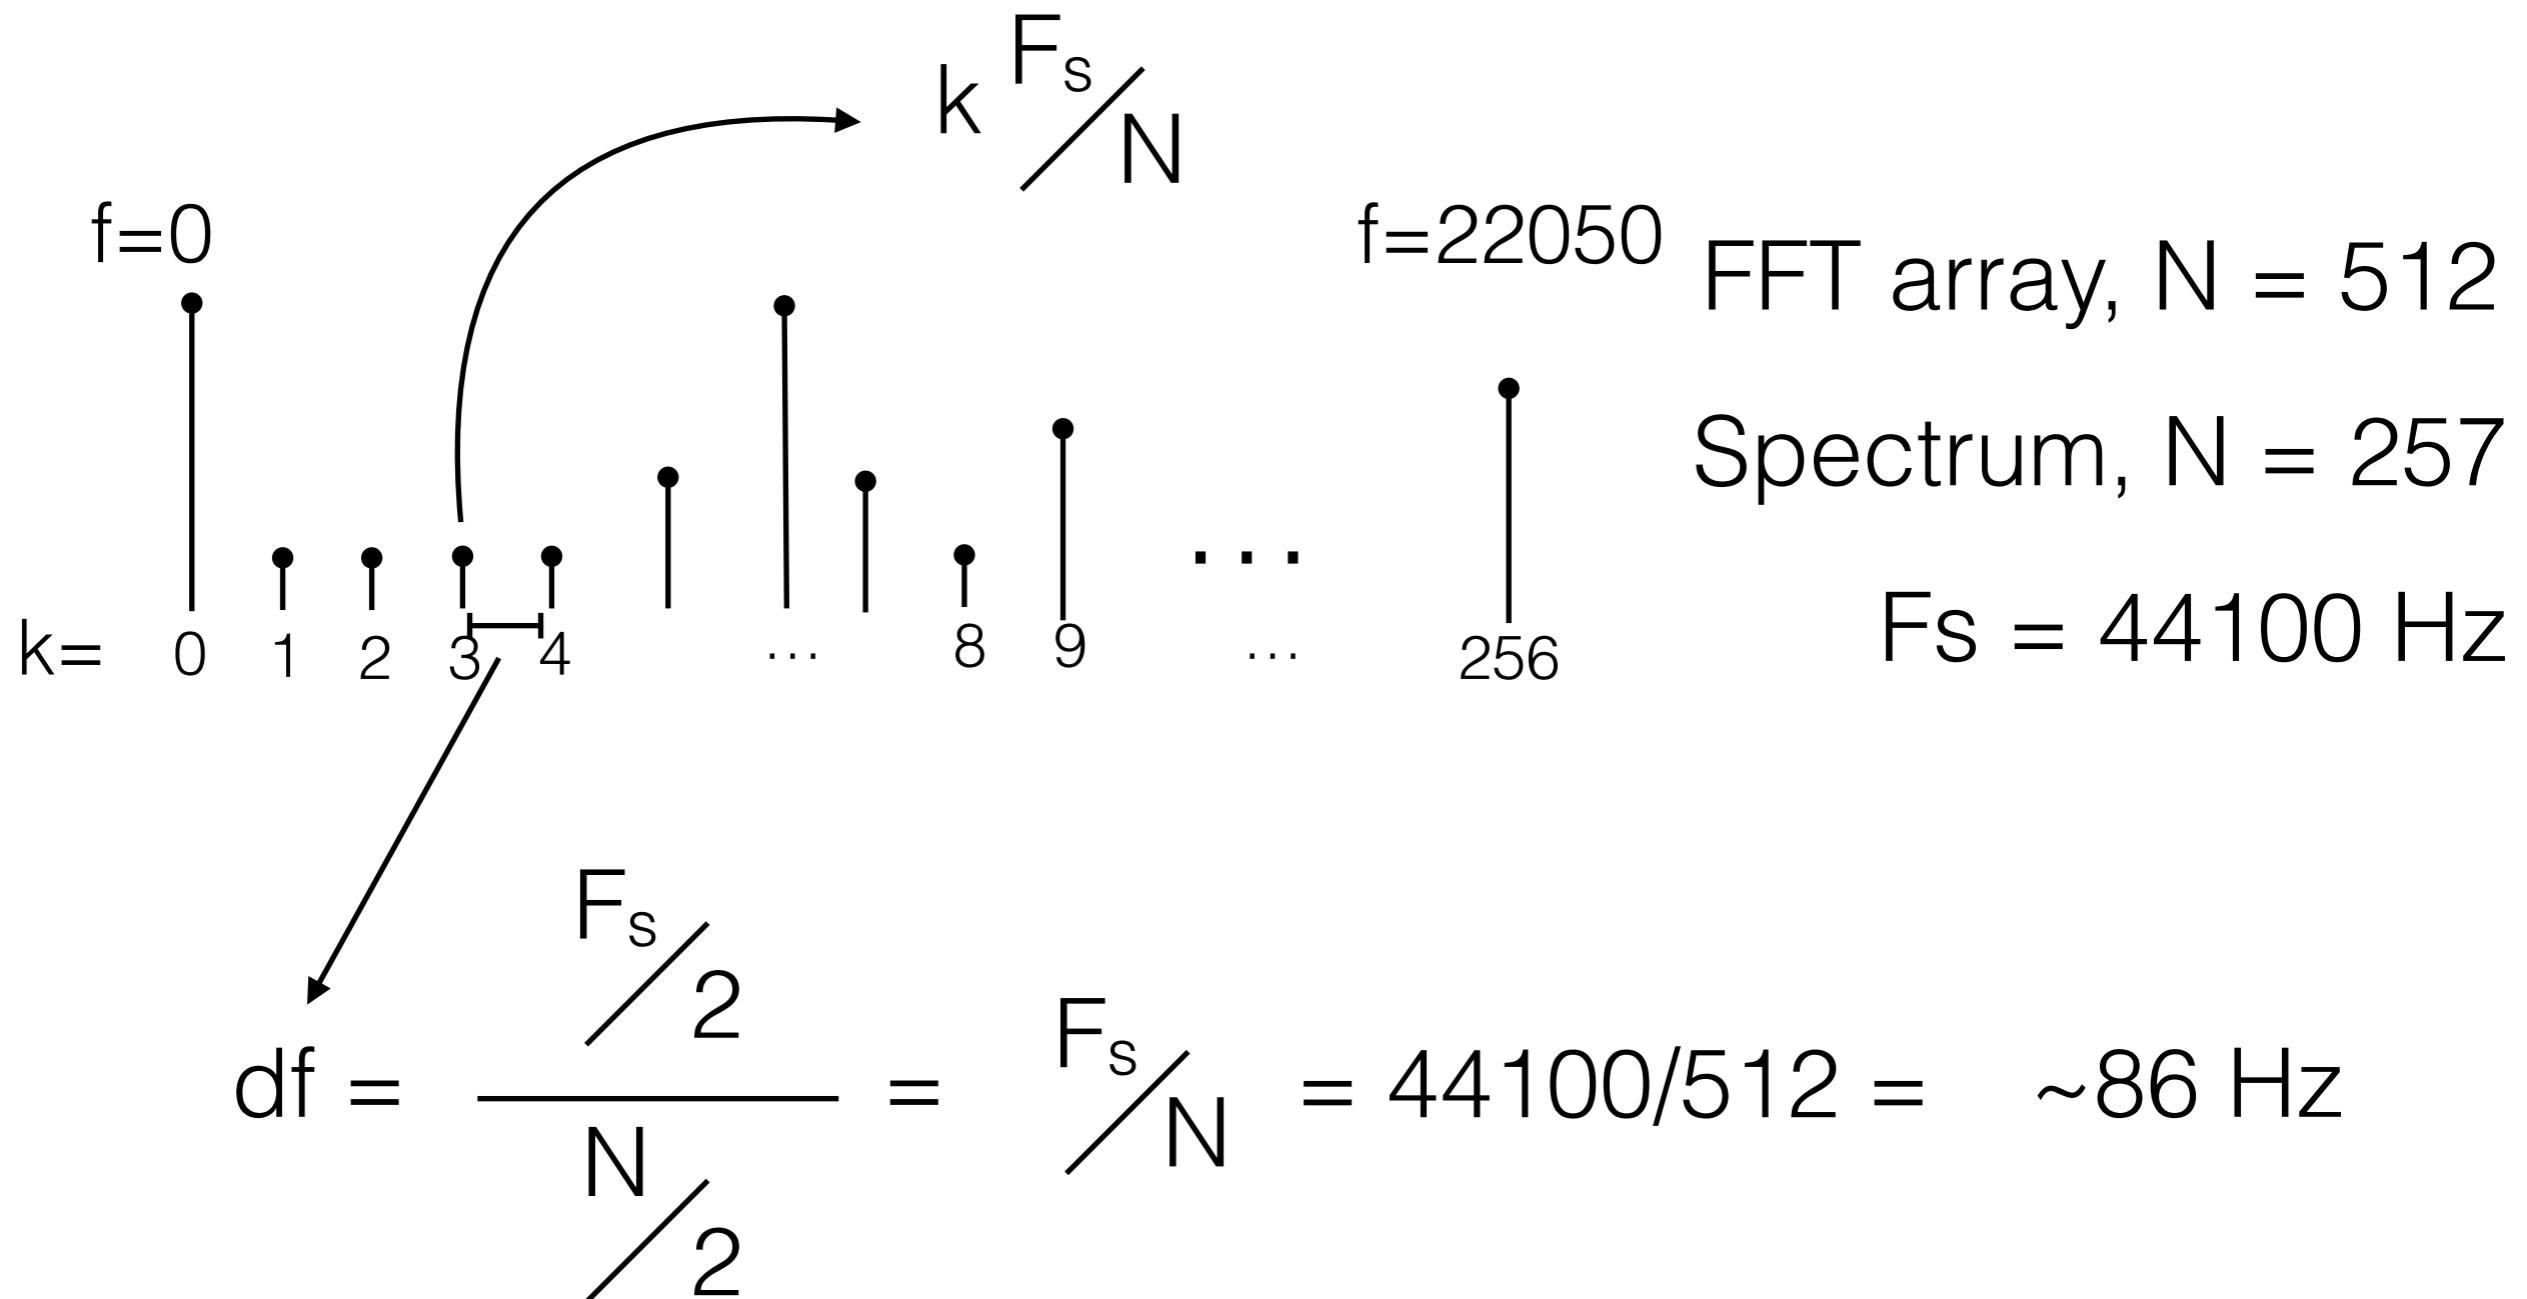
the FFT

- we will use the FFT in the accelerate framework
 - complexity is $O(N \log_2(N))$
- the FFT takes an array of numbers and gives you back an array of complex numbers
 - both input and output arrays are the same length
 - the indices of the input array denote increasing time
 - the indices of the output array denote increasing frequency
- don't quite get the math? so what? understand the output array...

time and frequency

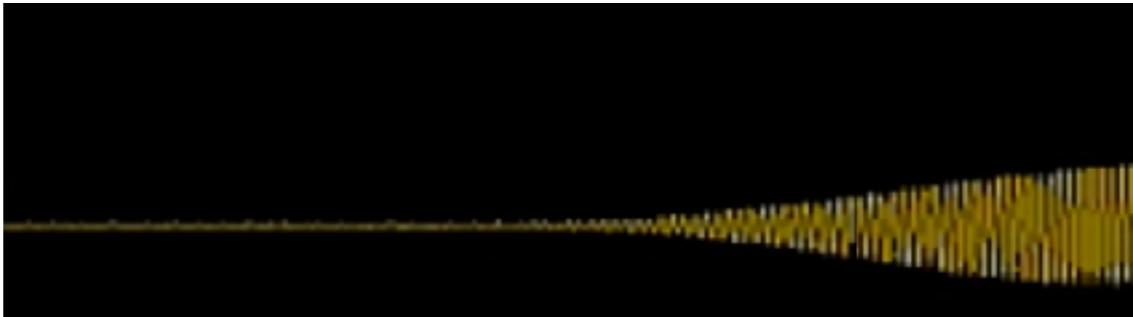


time and frequency



using the FFT

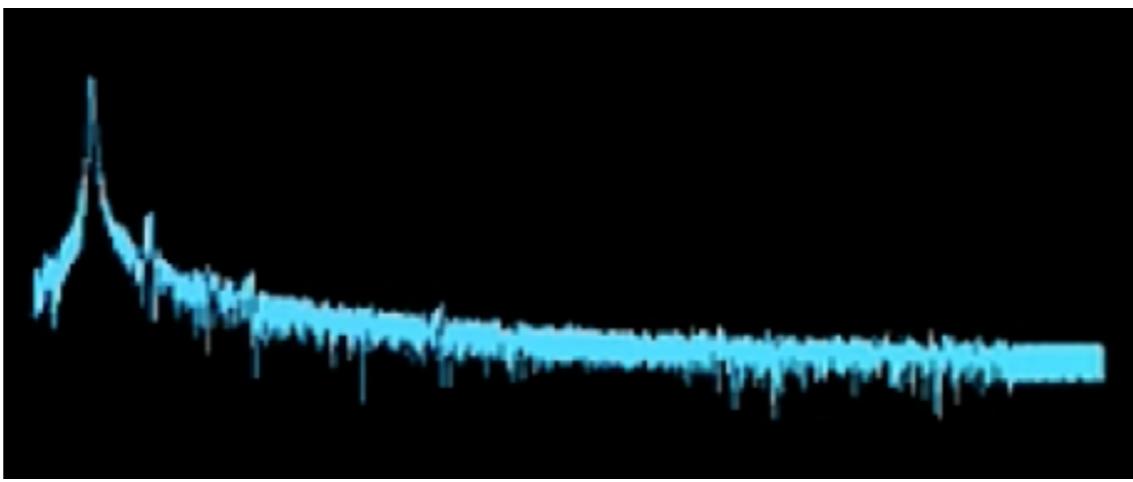
raw audio



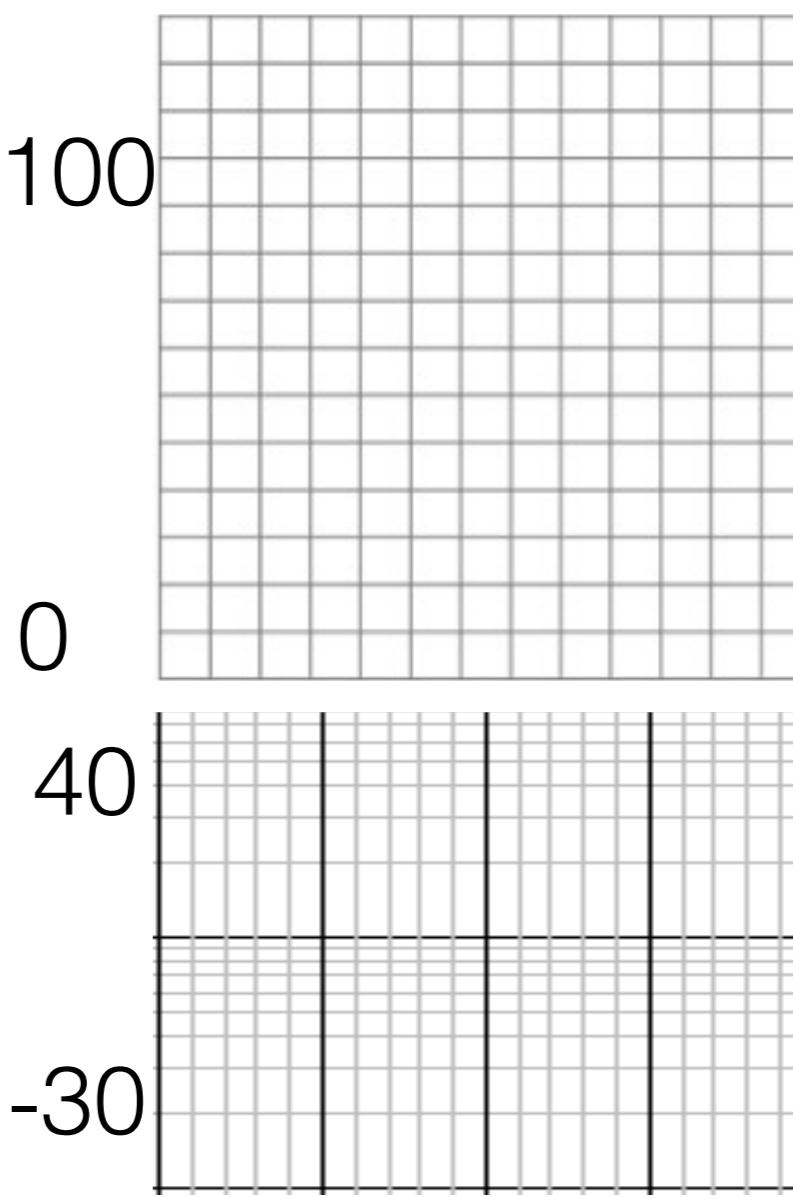
magnitude FFT



magnitude FFT
in dB



$$20 \log_{10}(|\text{FFT}|)$$



some fft examples

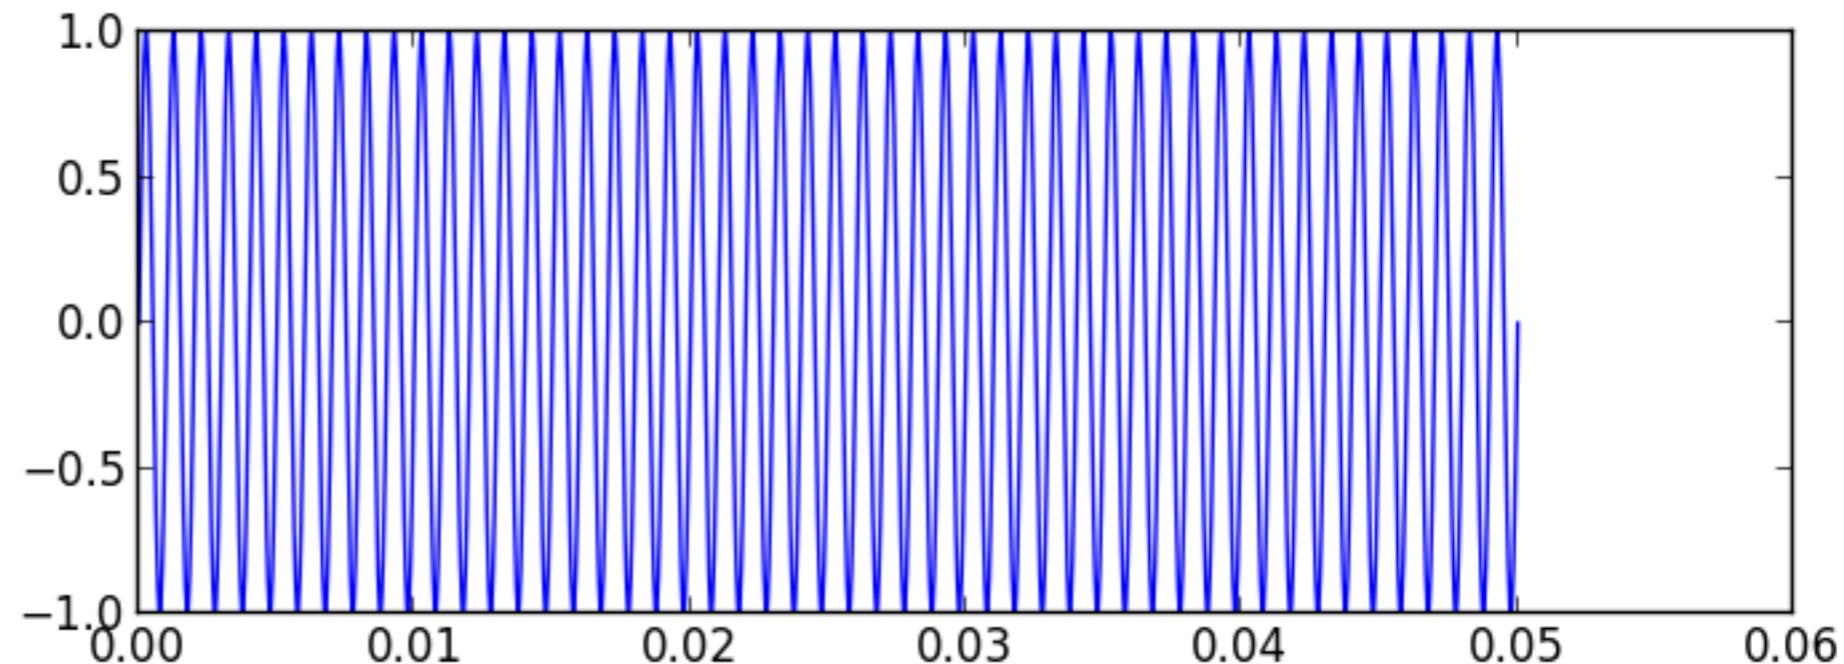
this is just a plot of
an array of values
sampled

**over equal time
intervals**

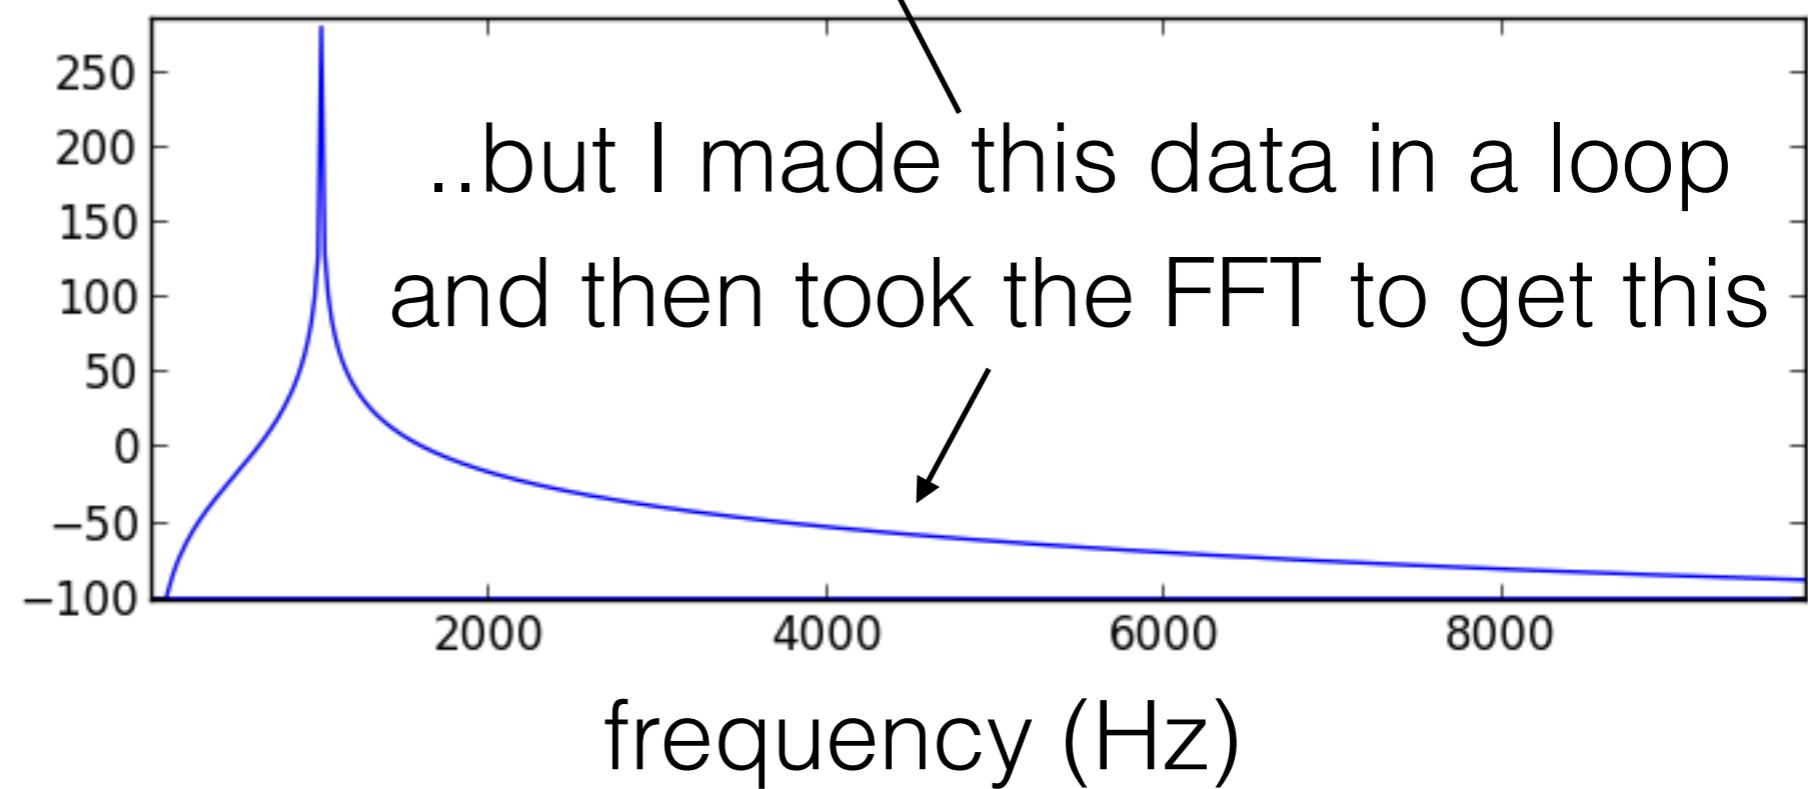
this is just a plot of
an array of FFT
values

**its indices are
equally sampled
over frequency**

1kHz sine wave

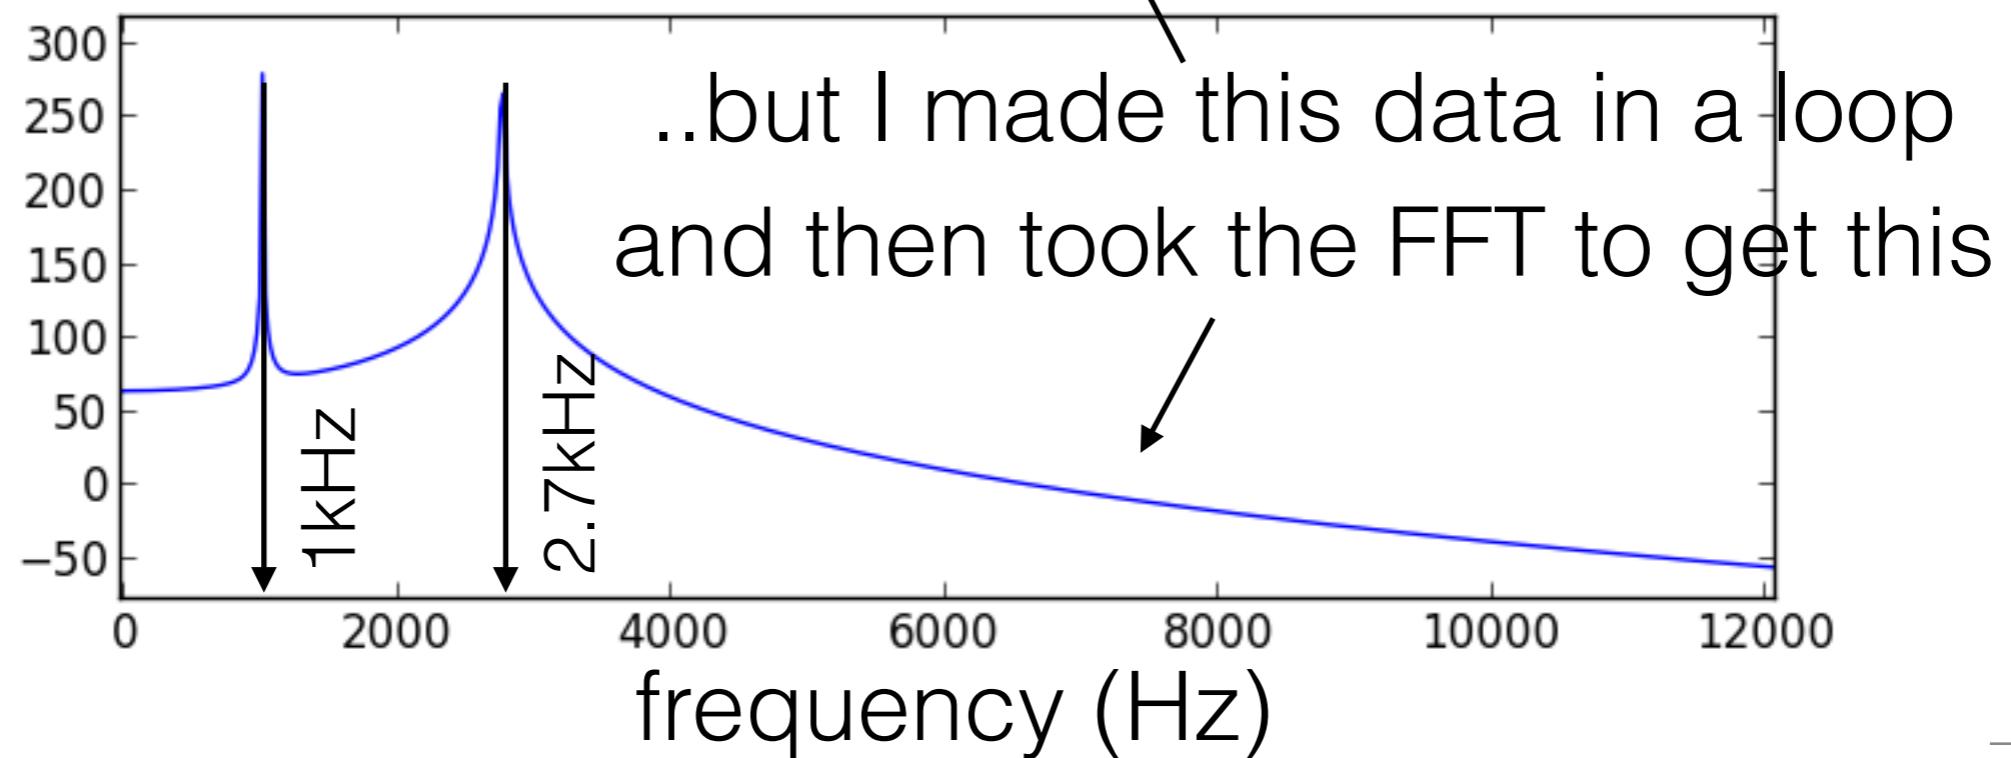
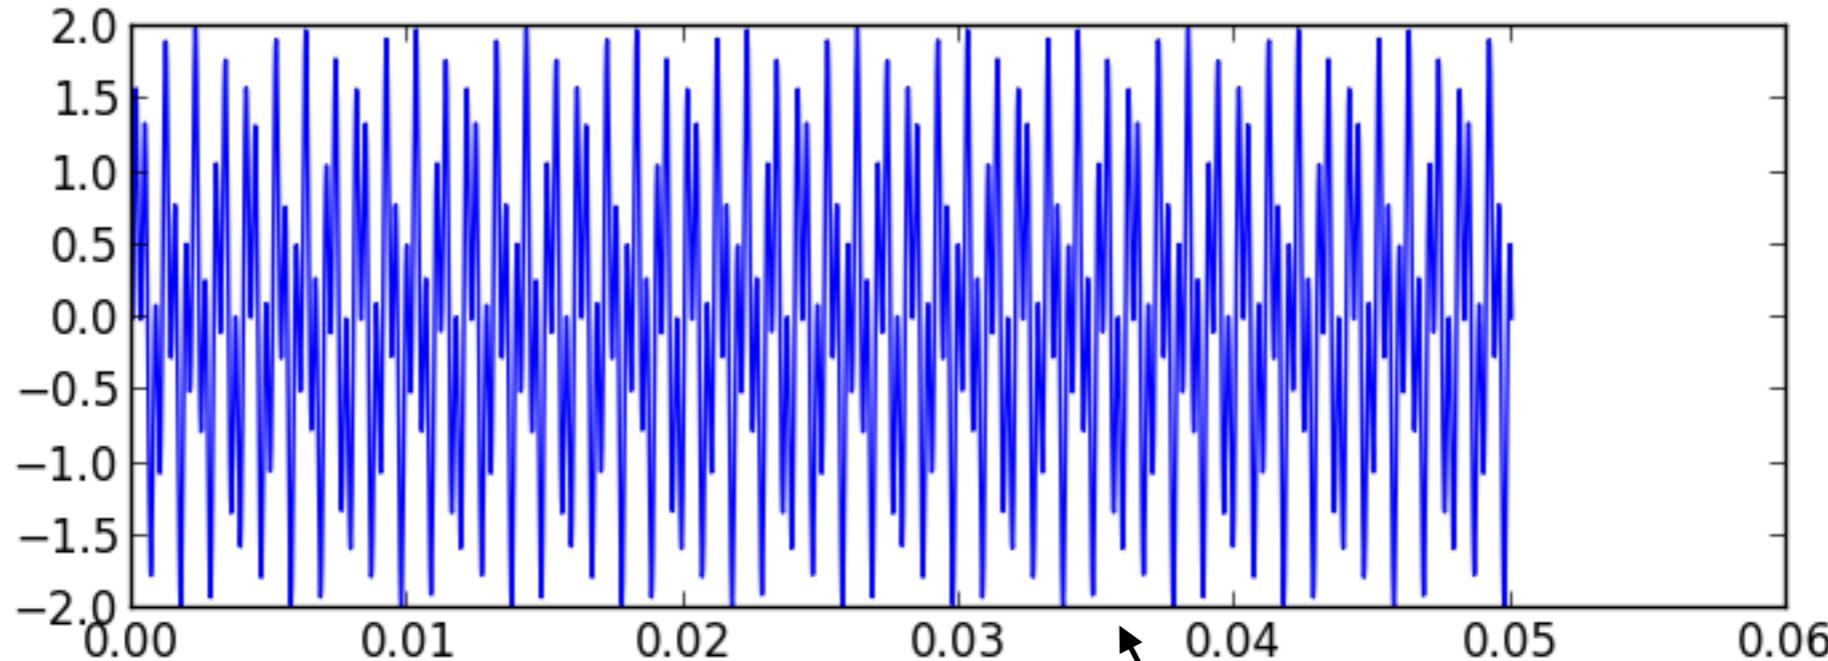


..but I made this data in a loop
and then took the FFT to get this



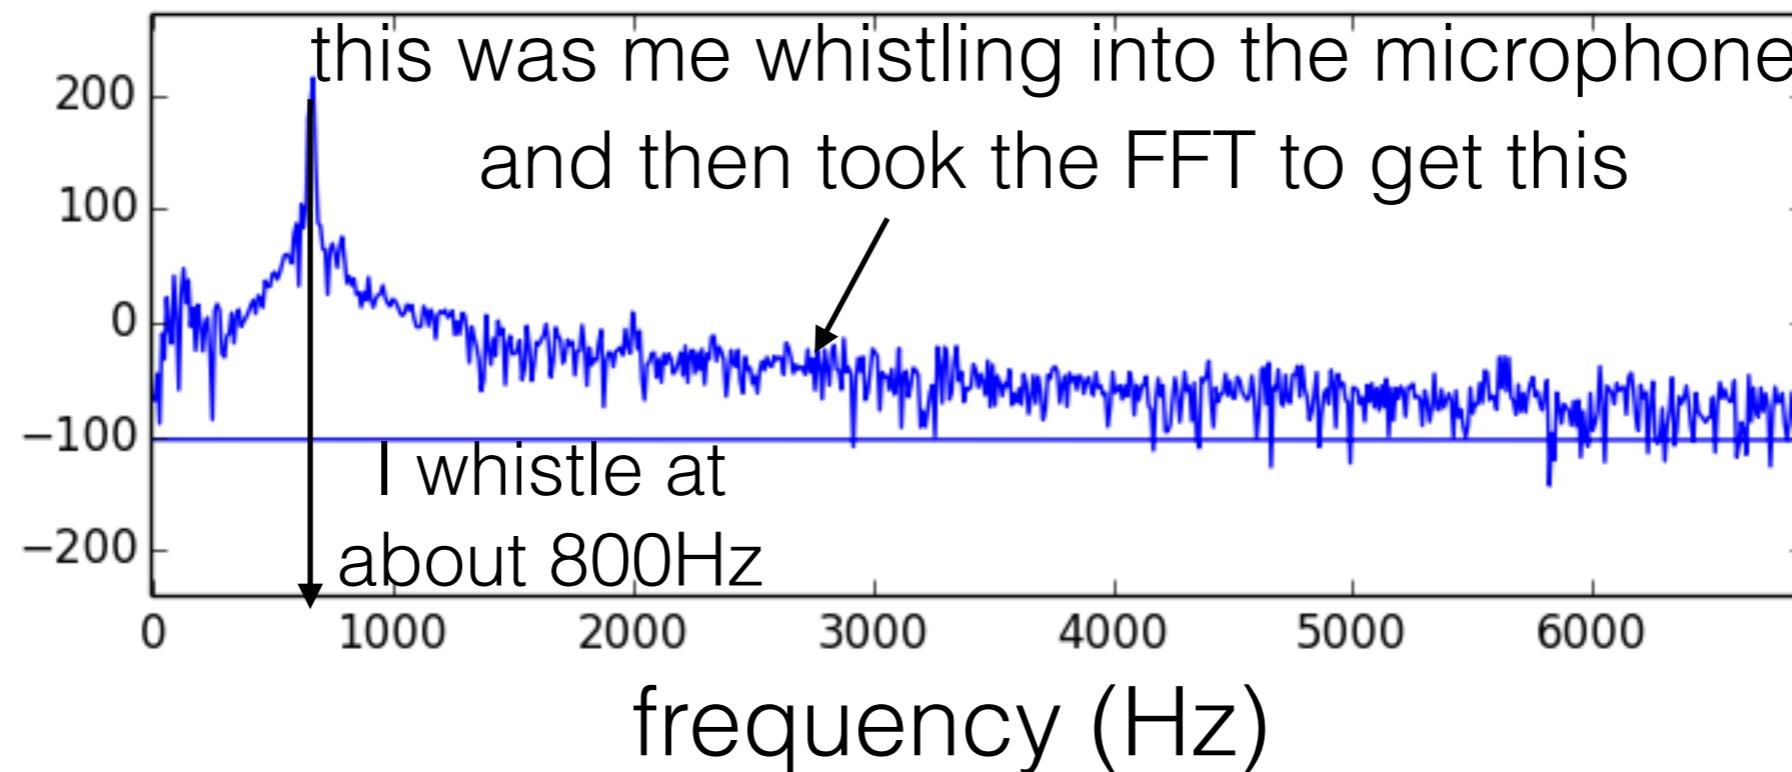
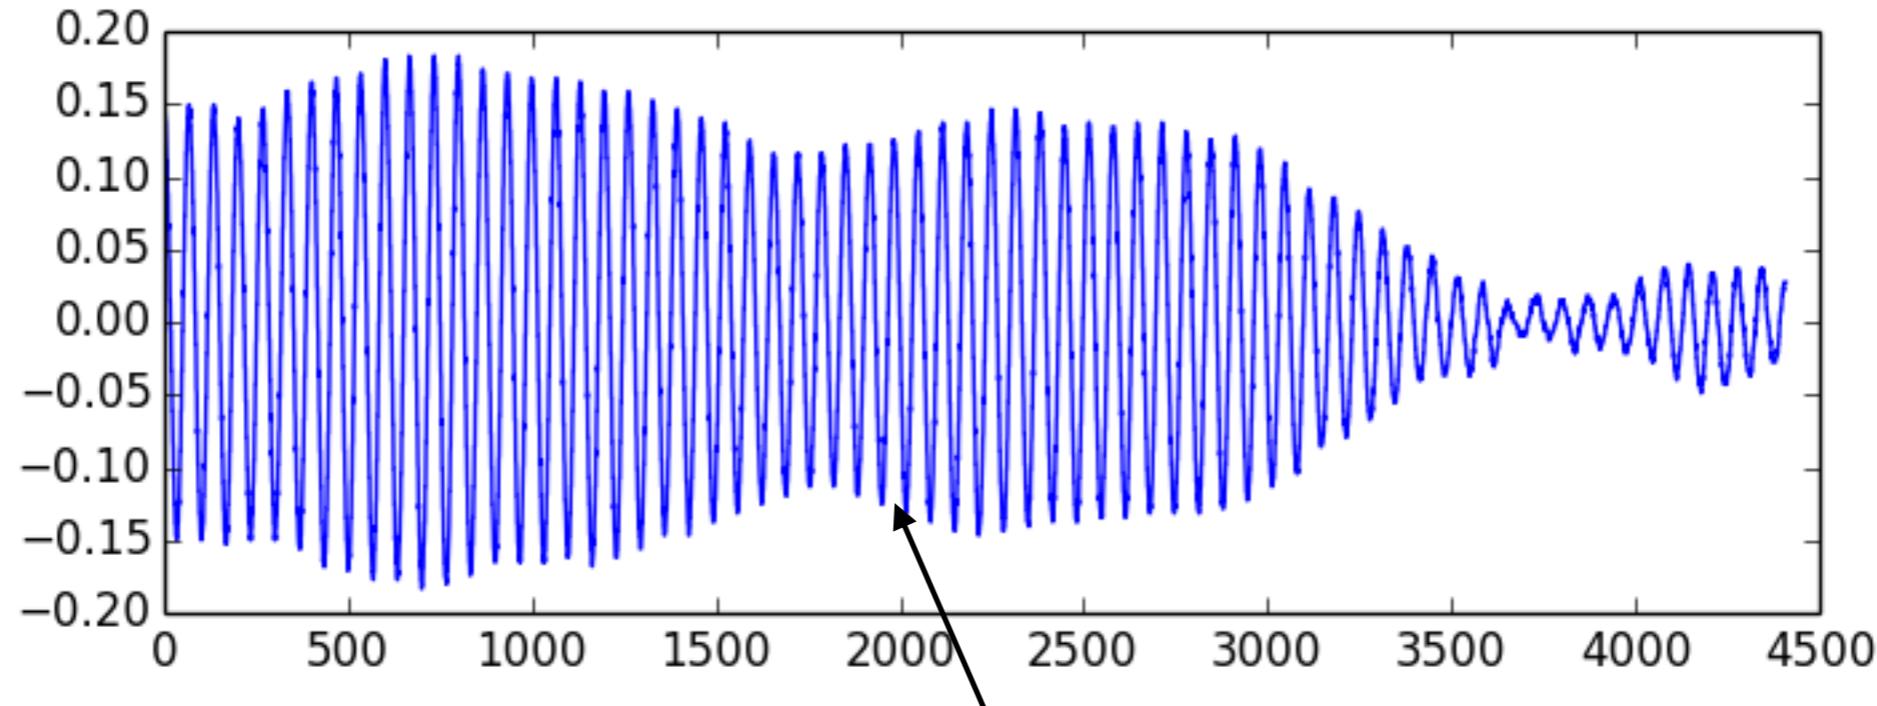
some fft examples

1kHz sine wave + 2.7kHz sine wave



some fft examples

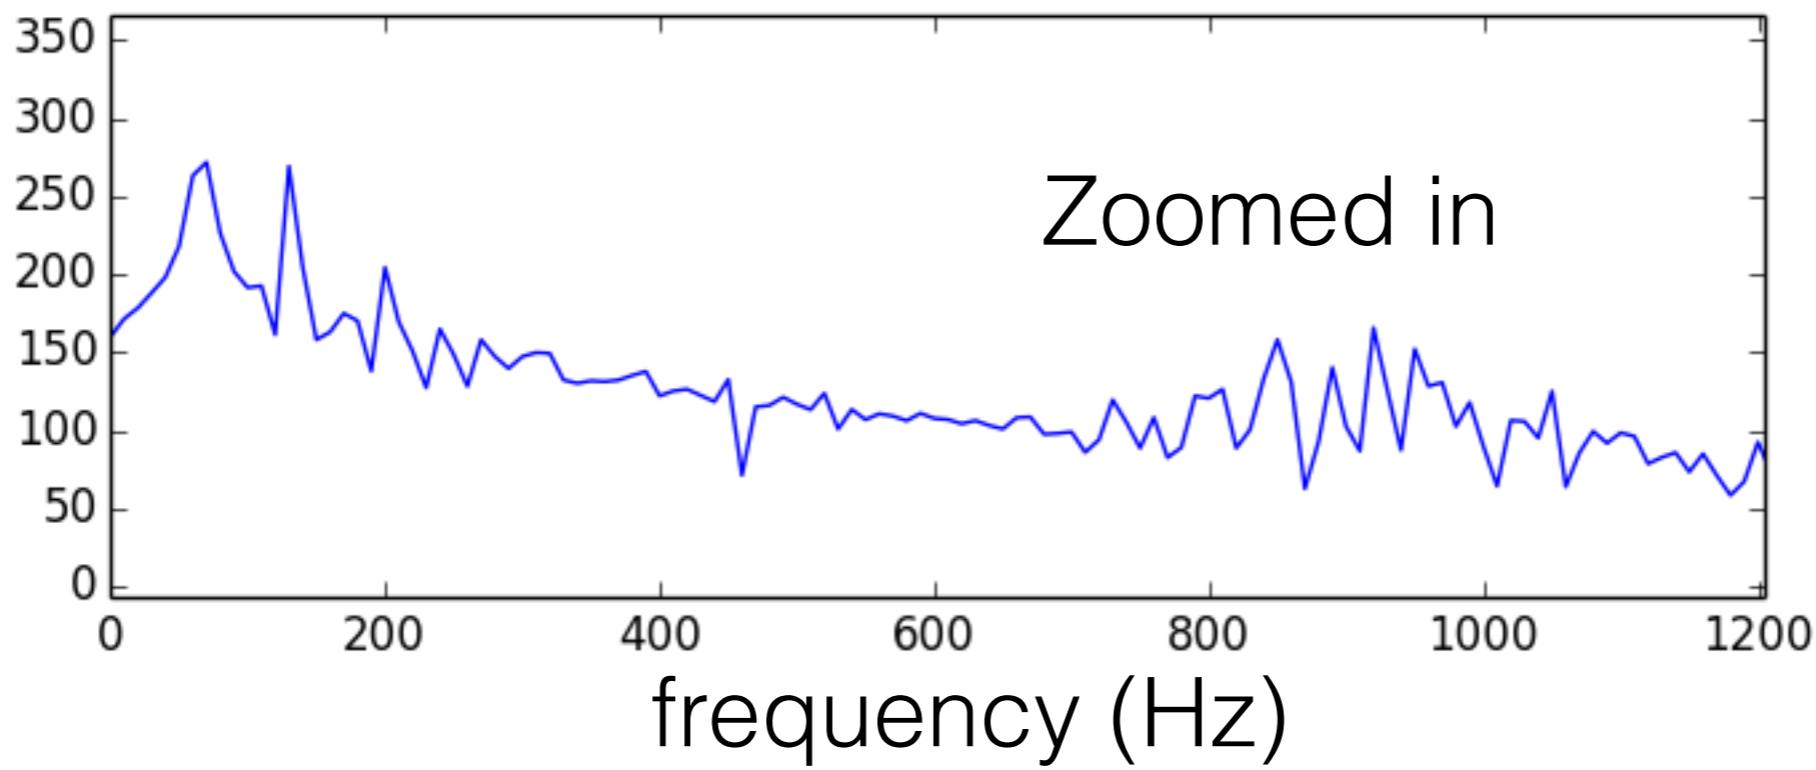
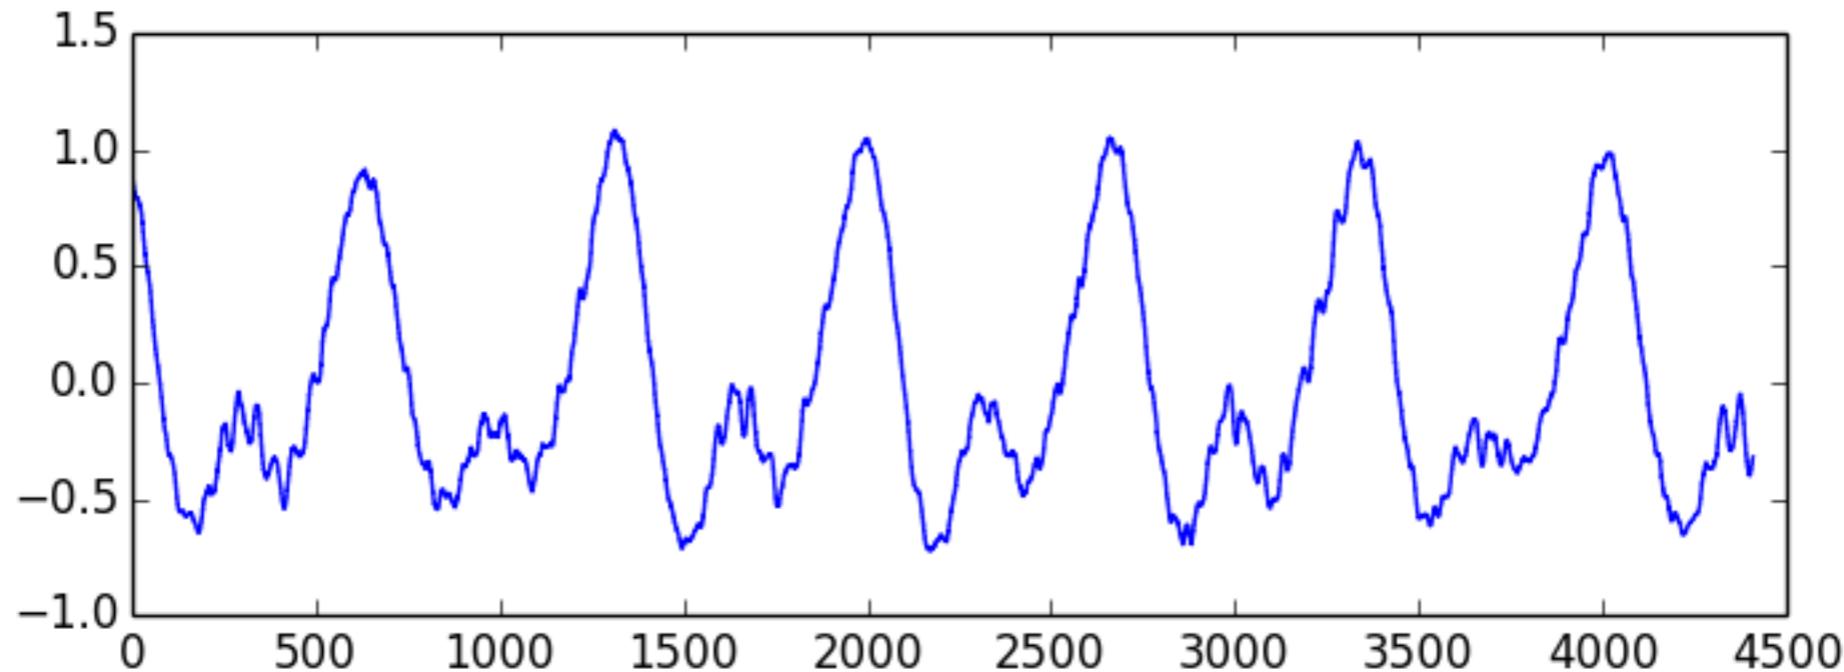
a person whistling



some fft examples

humming

oooooooooooooo



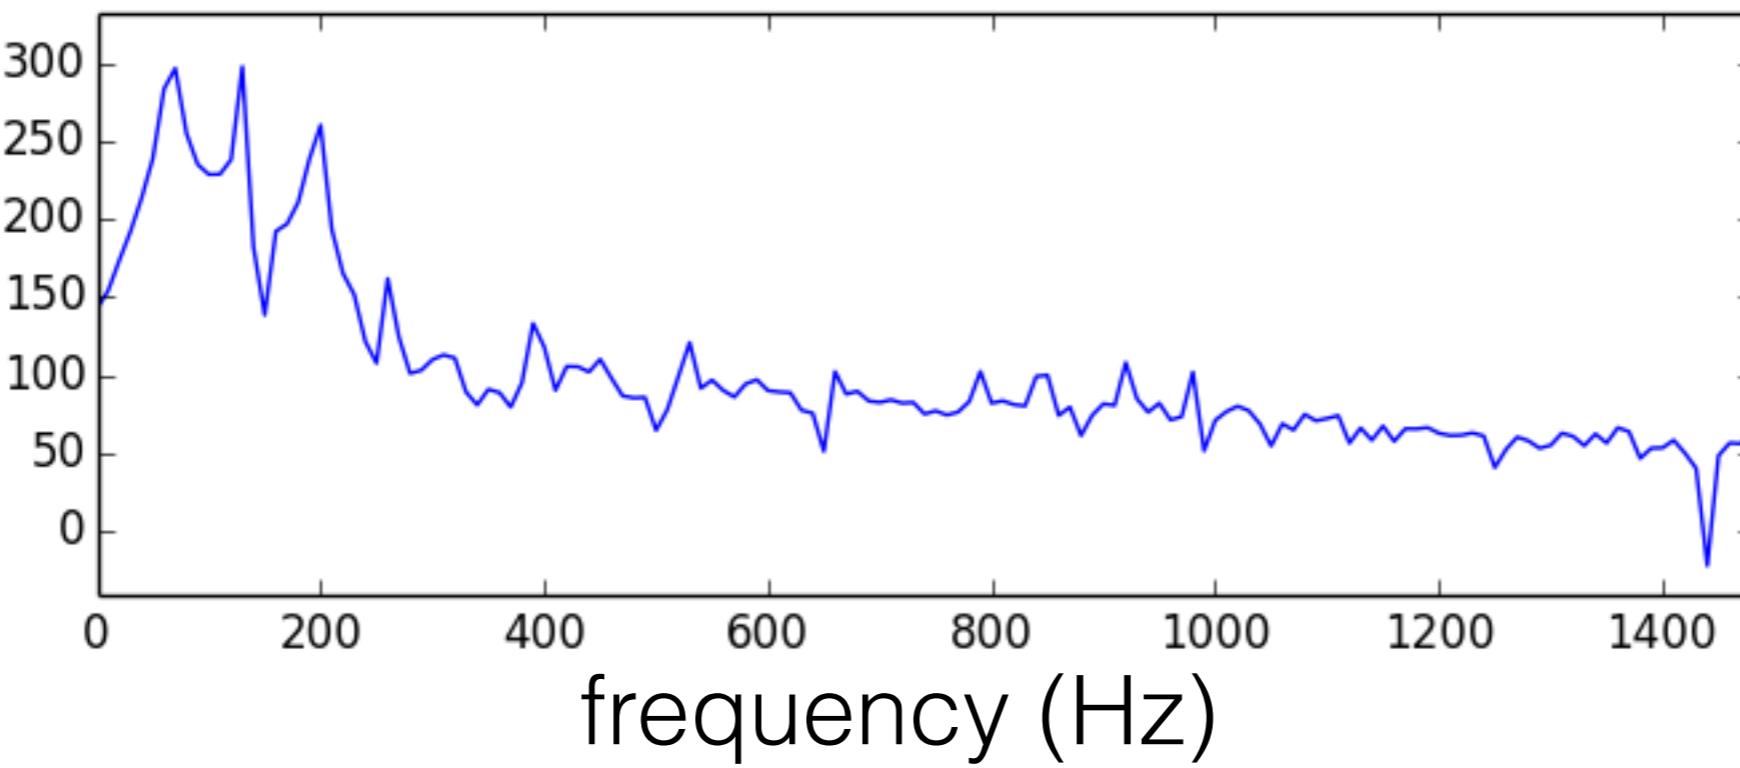
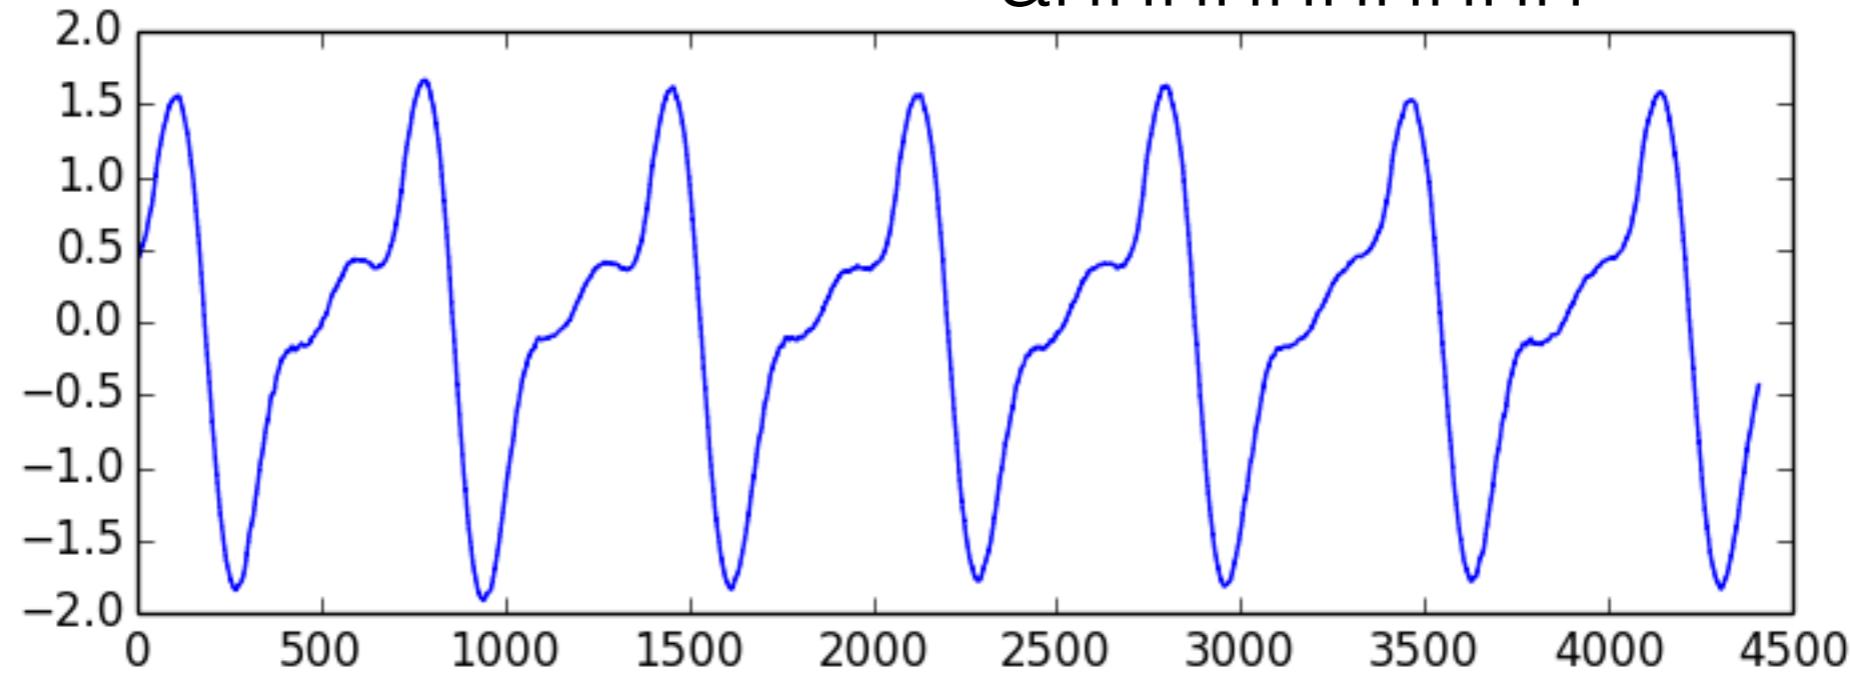
Zoomed in

frequency (Hz)

some fft examples

humming again

ahhhhhhhhhh

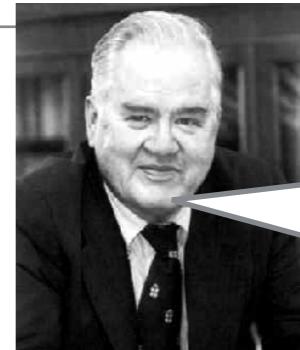


frequency (Hz)



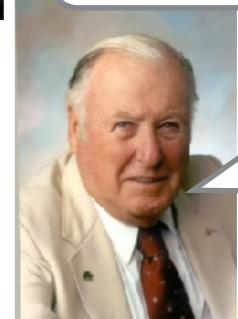
programming the FFT

```
//setup the fft
private lazy var fftHelper:FFTHelper? = {
    return FFTHelper.init(fftSize: Int32(BUFFER_SIZE))
}
```



This is how to use our stuff

fft size == buffer size



Yeah!

```
var fftData:[Float]
fftData = Array.init(repeating: 0.0,
                     count: BUFFER_SIZE/2)
```

Magnitude buffer is half the size!

```
// copy time data to swift array
self.inputBuffer!.fetchFreshData(&timeData,
                                withNumSamples: Int64(BUFFER_SIZE))
// now take FFT
fftHelper!.performForwardFFT(withData: &timeData,
                             andCopydBMagnitudeToBuffer: &fftData)
```

input array

magnitude out
half window size

highly optimized!! even using tricks we have not discussed

programming the FFT

```
#import "SMUFFTHelper.h"  
  
@property (strong, nonatomic) FFTHelper *fftHelper;
```

fft size == window size

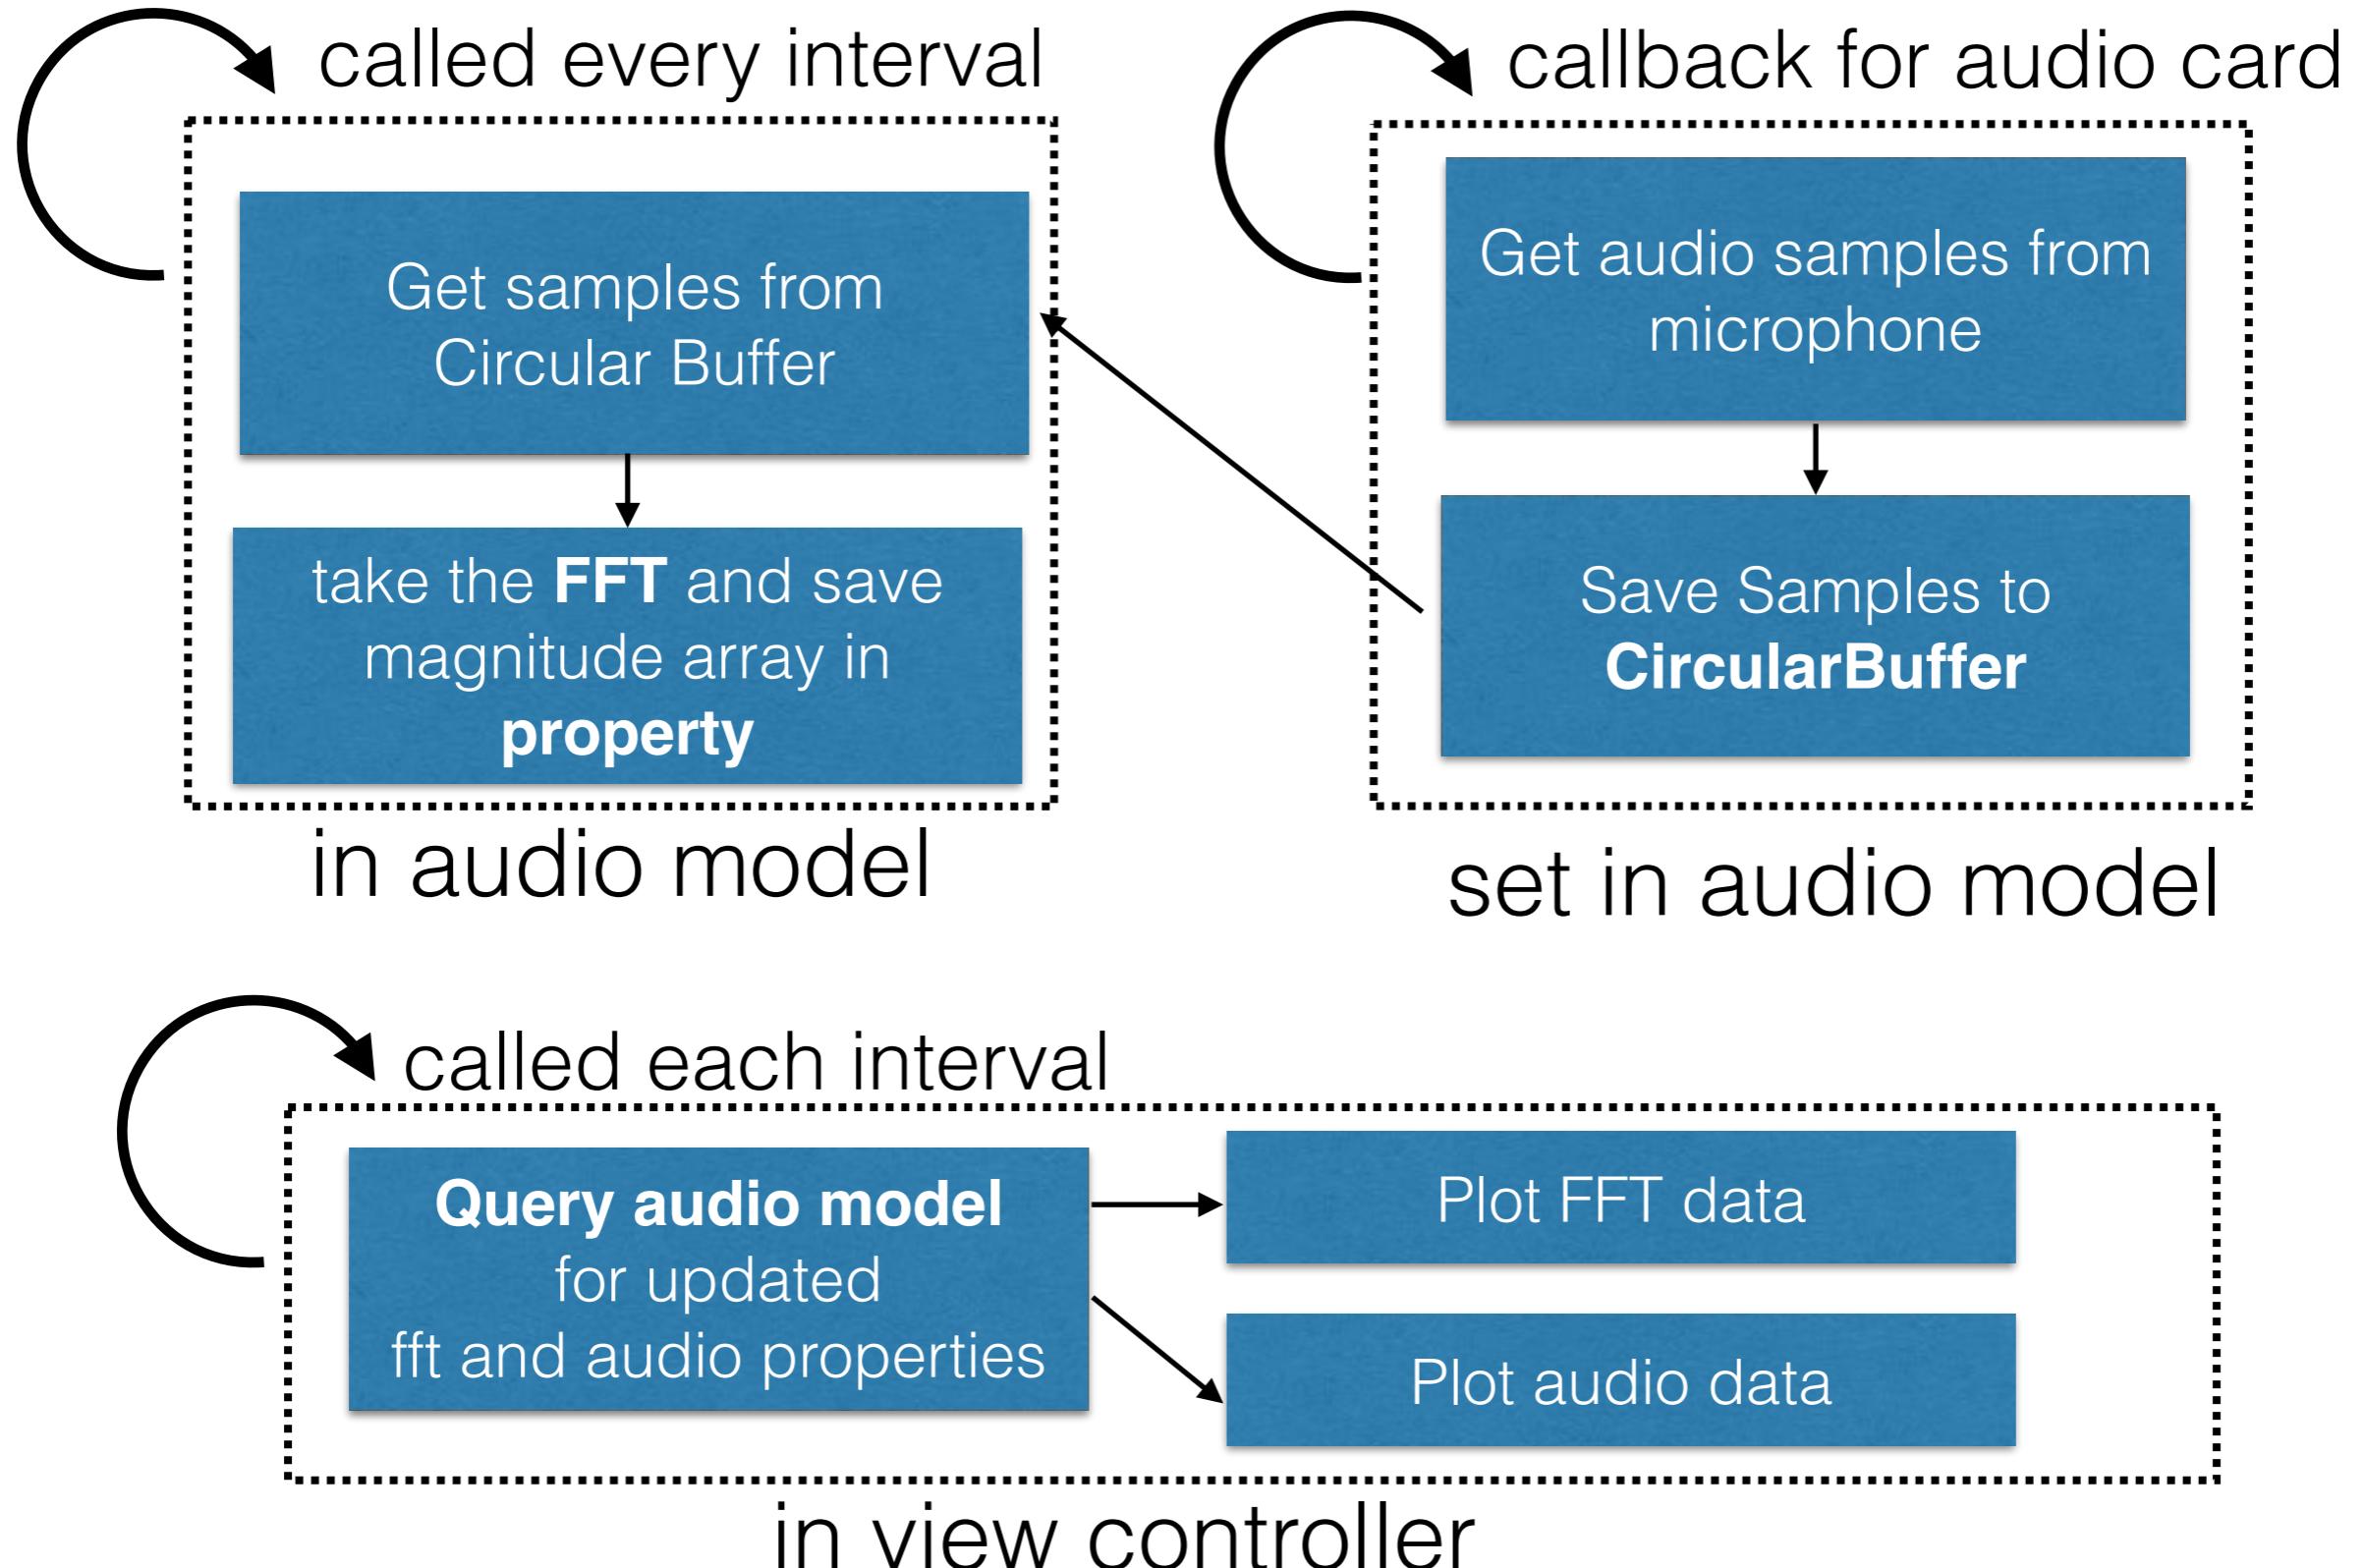
```
//setup the fft  
_fftHelper = [[FFTHelper alloc] initWithFFTSize:BUFFER_SIZE];
```

```
float* fftMagnitude = malloc(sizeof(float)*BUFFER_SIZE/2);  
  
// take forward FFT  
[self.fftHelper performForwardFFTWithData:arrayData  
    andCopydBMagnitudeToBuffer:fftMagnitude];  
  
free(fftMagnitude);
```

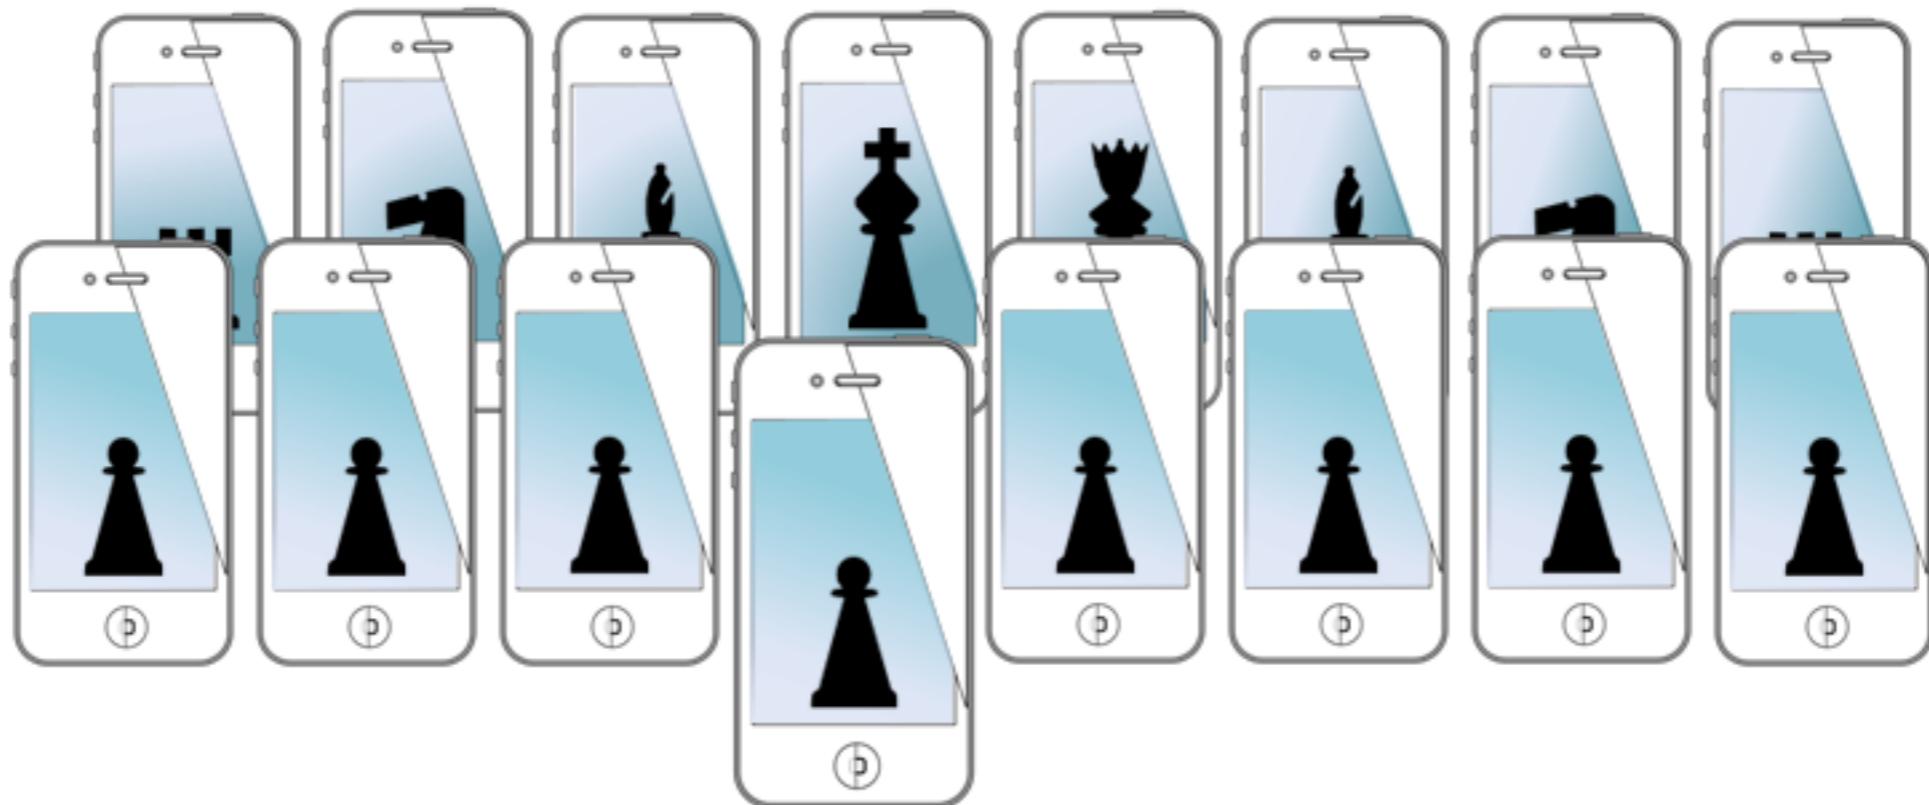
input array
N=window size

magnitude out
N=half window size

The program



MOBILE SENSING LEARNING

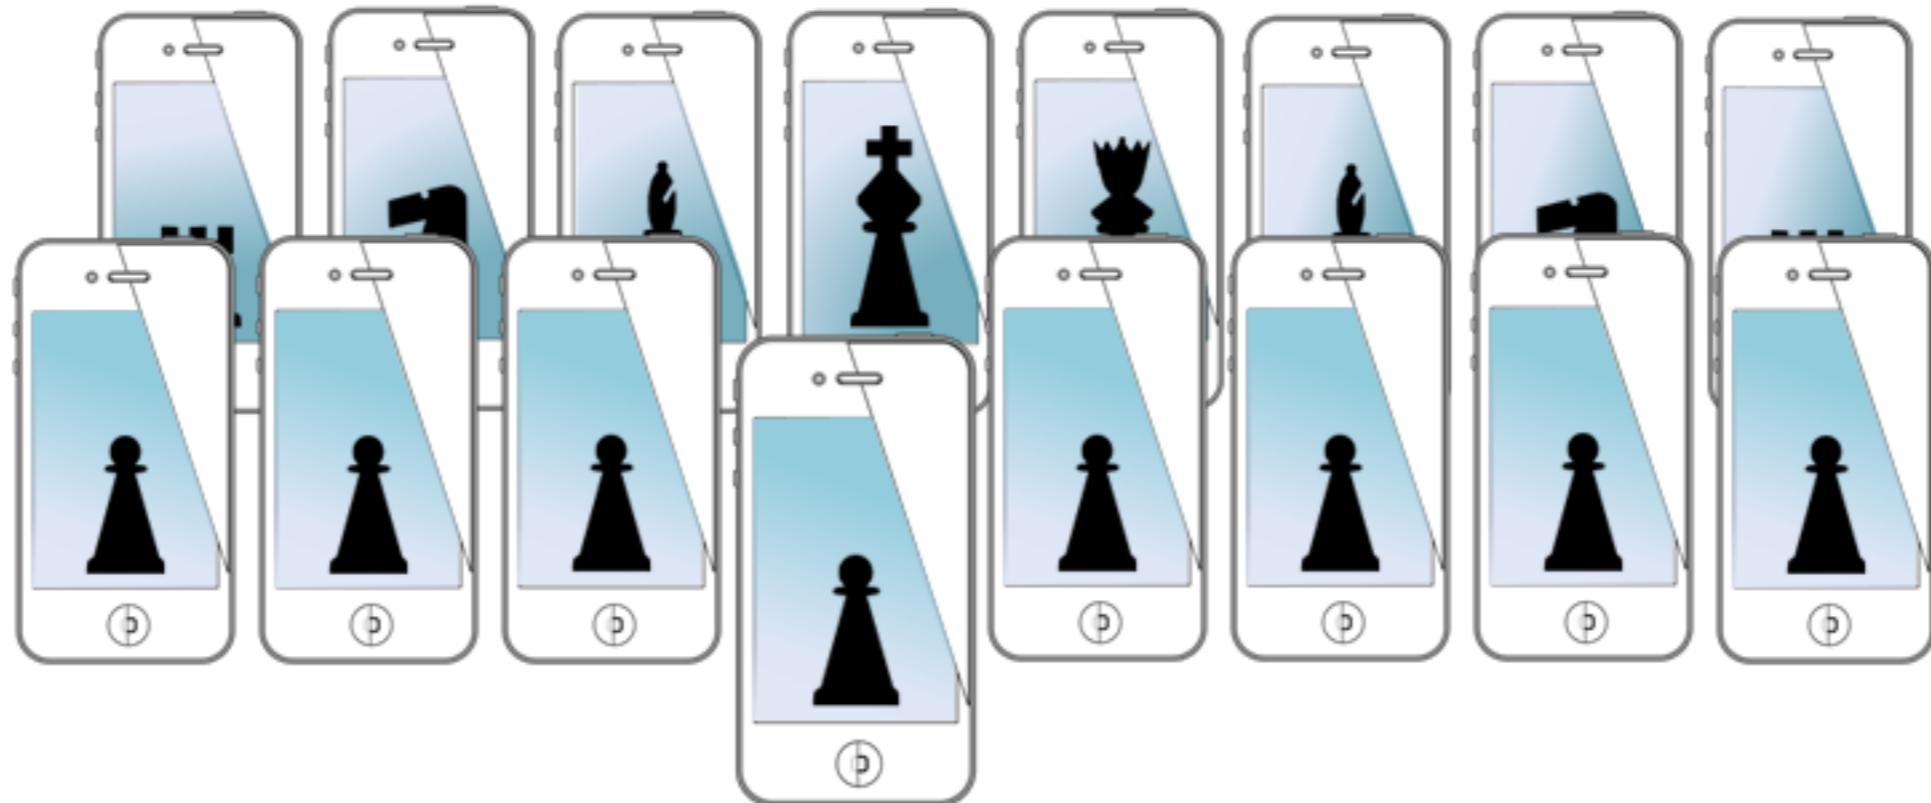


CS5323 & 7323
Mobile Sensing and Learning

video lecture: accelerate & FFT

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

MOBILE SENSING LEARNING



CS5323 & 7323
Mobile Sensing and Learning

Supplemental Slides: filtering and windowing

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

Supplemental Slides

- these slides were removed from the course because of their complexity
- you need a good background in signal representation and time series analysis to really understand these slides

Optional Concepts not Covered in Course Anymore

filters!

- we will cover what we can...

signals and systems

- signals are collections of sampled data (arrays)
 - such as audio, accelerometer, etc.
 - can also be 2D, like images
- systems are objects which manipulate signals
 - characterized by their “input/output” relationships
 - we say “ $x[n]$ is passed through H , resulting in $y[n]$ ”

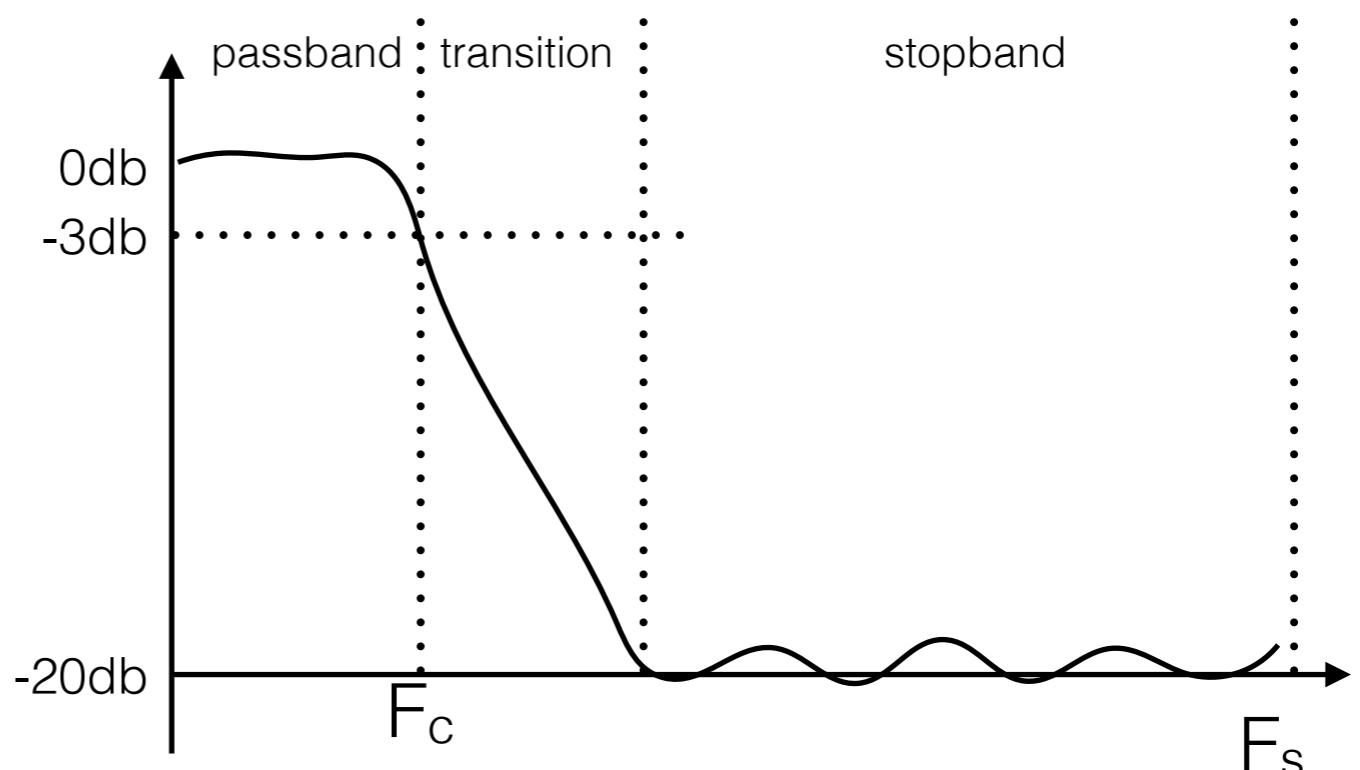


filters

- filters are systems which manipulate frequencies
 - certain frequencies to pass through, but not others
 - “lowpass” filter allows low frequencies to pass through
 - “highpass” filter likewise allows high frequencies through
- keep in mind: no filter is perfect!!
 - no filter will pass everything you want while stopping everything you don’t
 - everything is a balance between different parameters you can control
- we won’t study how to design filters
 - we will study properties of filters and how to use them
 - so we need to know what filters can and cannot do

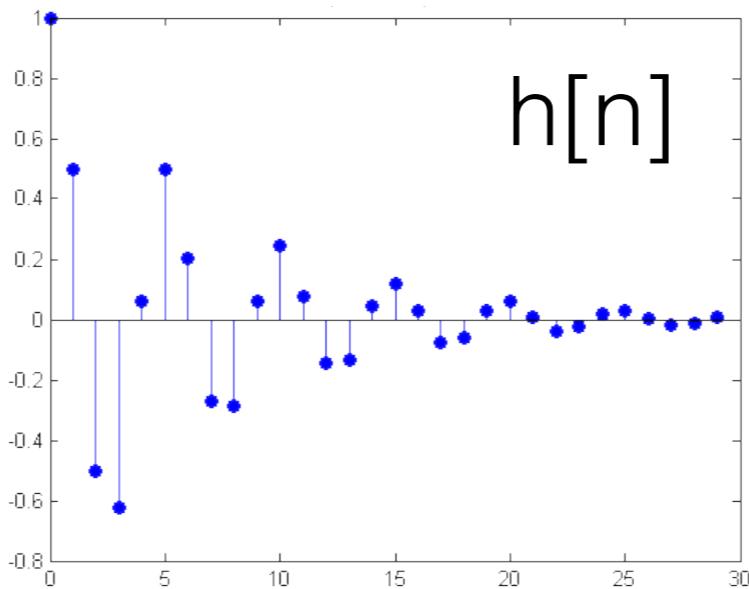
filters in frequency

- filters can be characterized a few different ways
 - let's start by looking at their properties in the frequency domain
- filters have the following frequency-domain attributes:
 - passband gain
 - passband bandwidth
 - stopband attenuation
 - transition bandwidth



filters in time

- filters are also signals (time series)
 - the series is called the “impulse response” of the filter
 - the frequency-domain plots are just Fourier transforms of the impulse response (magnitude)
- the time-domain property we care about is length
 - everything else is best left to a filter design course



so how to design a filter?

- scipy.signal in python (try to use remez)
- decent tutorial:
 - <http://mpastell.com/2010/01/18/fir-with-scipy/>
- matlab
 - fdatool
- lots of other places!

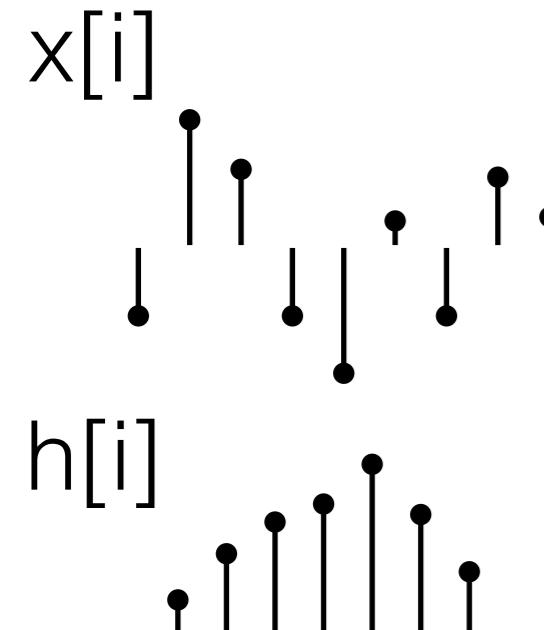
```
from pylab import *
import scipy.signal as signal
n = 61
a = signal.firwin(n, cutoff = 0.3,
                  window = "hamming")
```

filtering by convolution

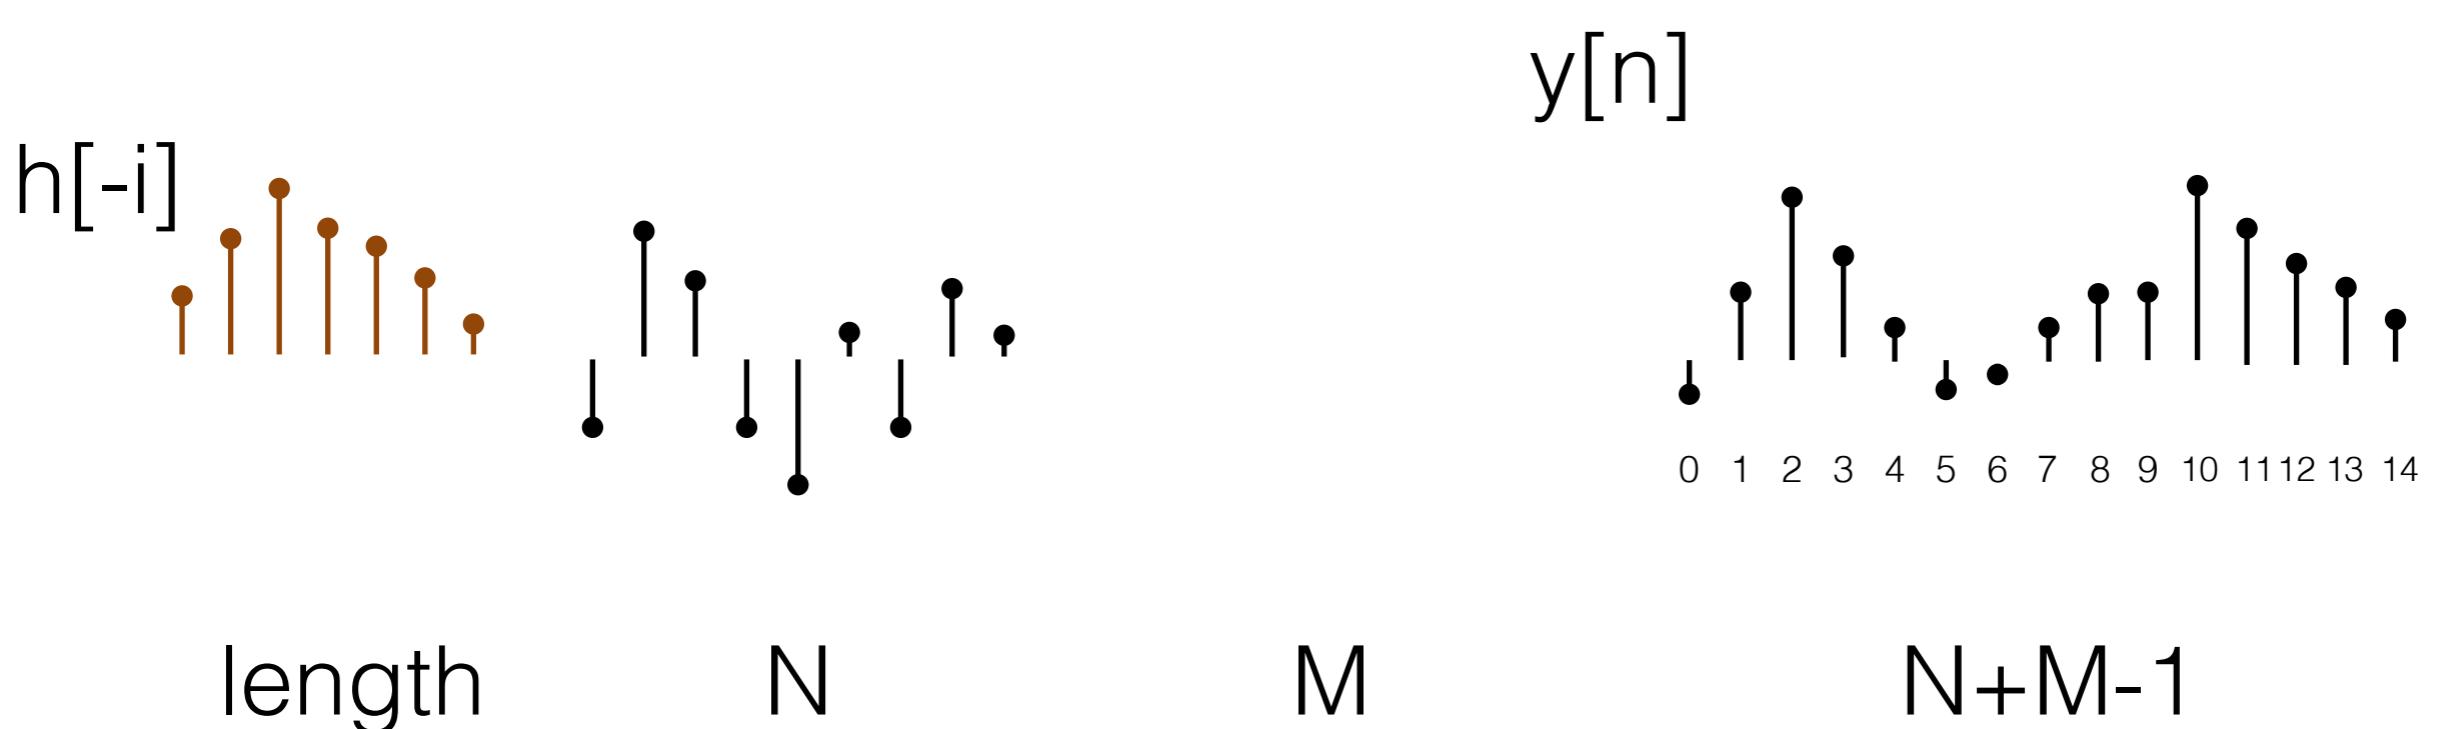
- we apply a filter using **convolution**
 - convolution allows us to combine frequency properties of two signals without taking an FFT
- basic principle:
 - convolution in time is multiplication in frequency
 - so the filter's frequency response will be multiplied by the frequency response of the signal

$$y[n] = \sum_{i=0}^{N-1} h[n-i]x[i]$$

convolution

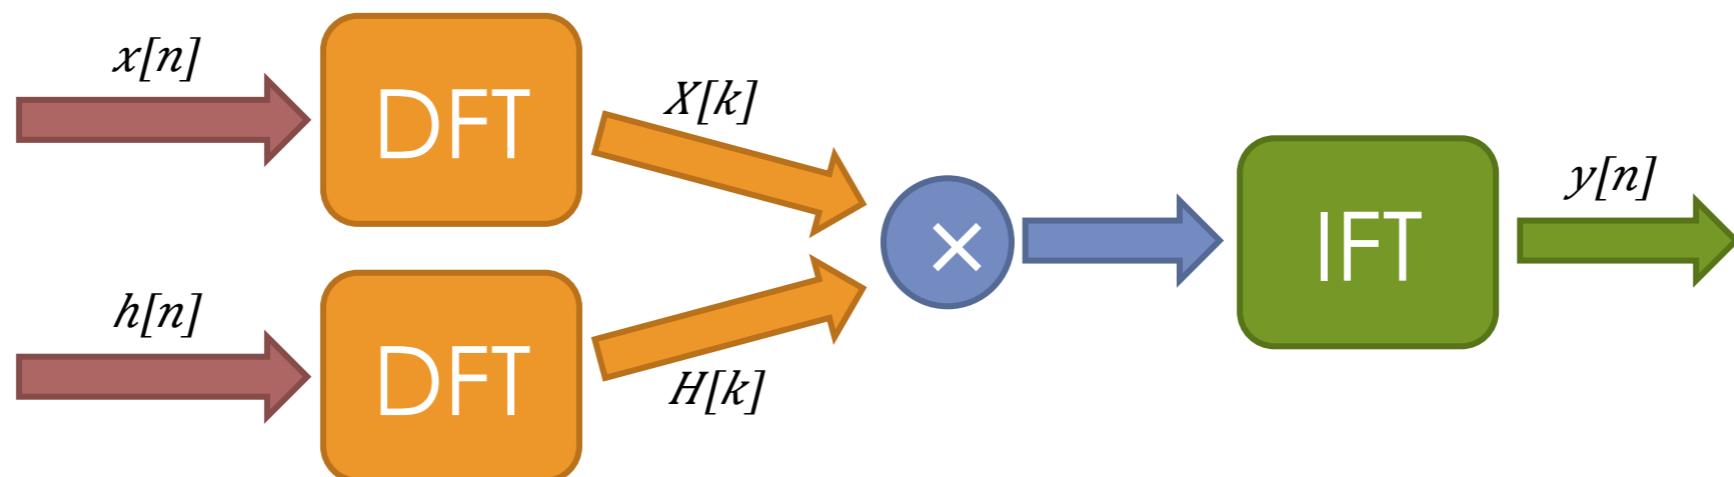


$$y[n] = \sum_{i=0}^{N-1} h[n-i]x[i]$$



convolution efficiency

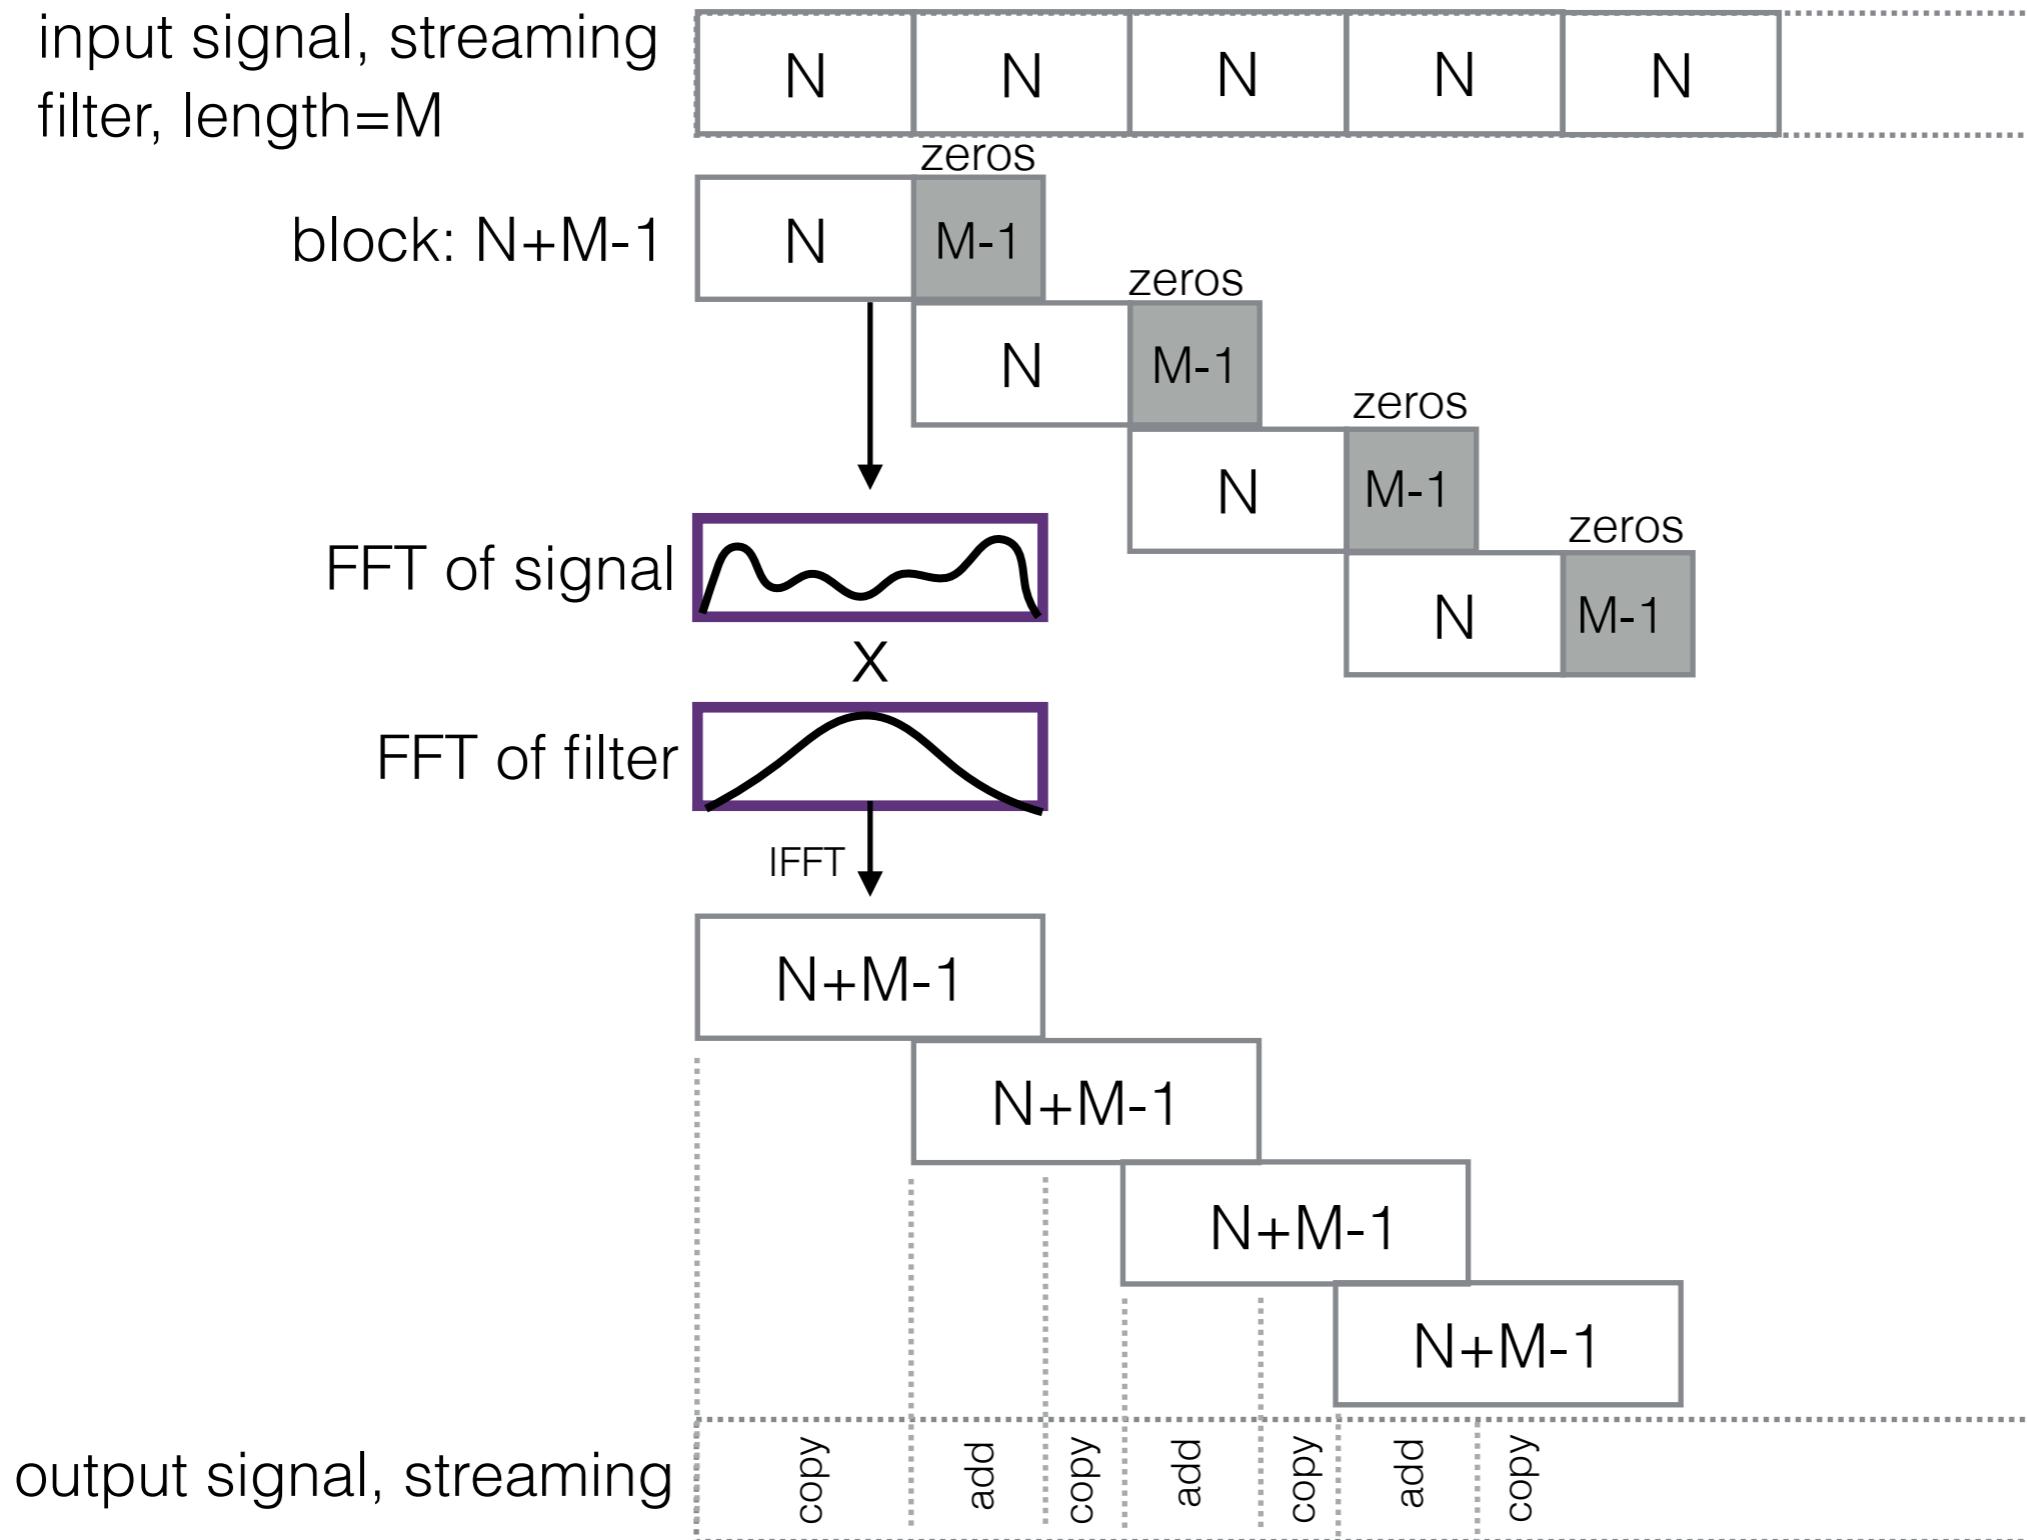
- algorithmic complexity
 - convolution is not particularly efficient, $O(N \times M)$
- to convolve faster, use that Fourier property:
 - “convolution in time is multiplication in frequency”
- why not just multiply to begin with!



its circular

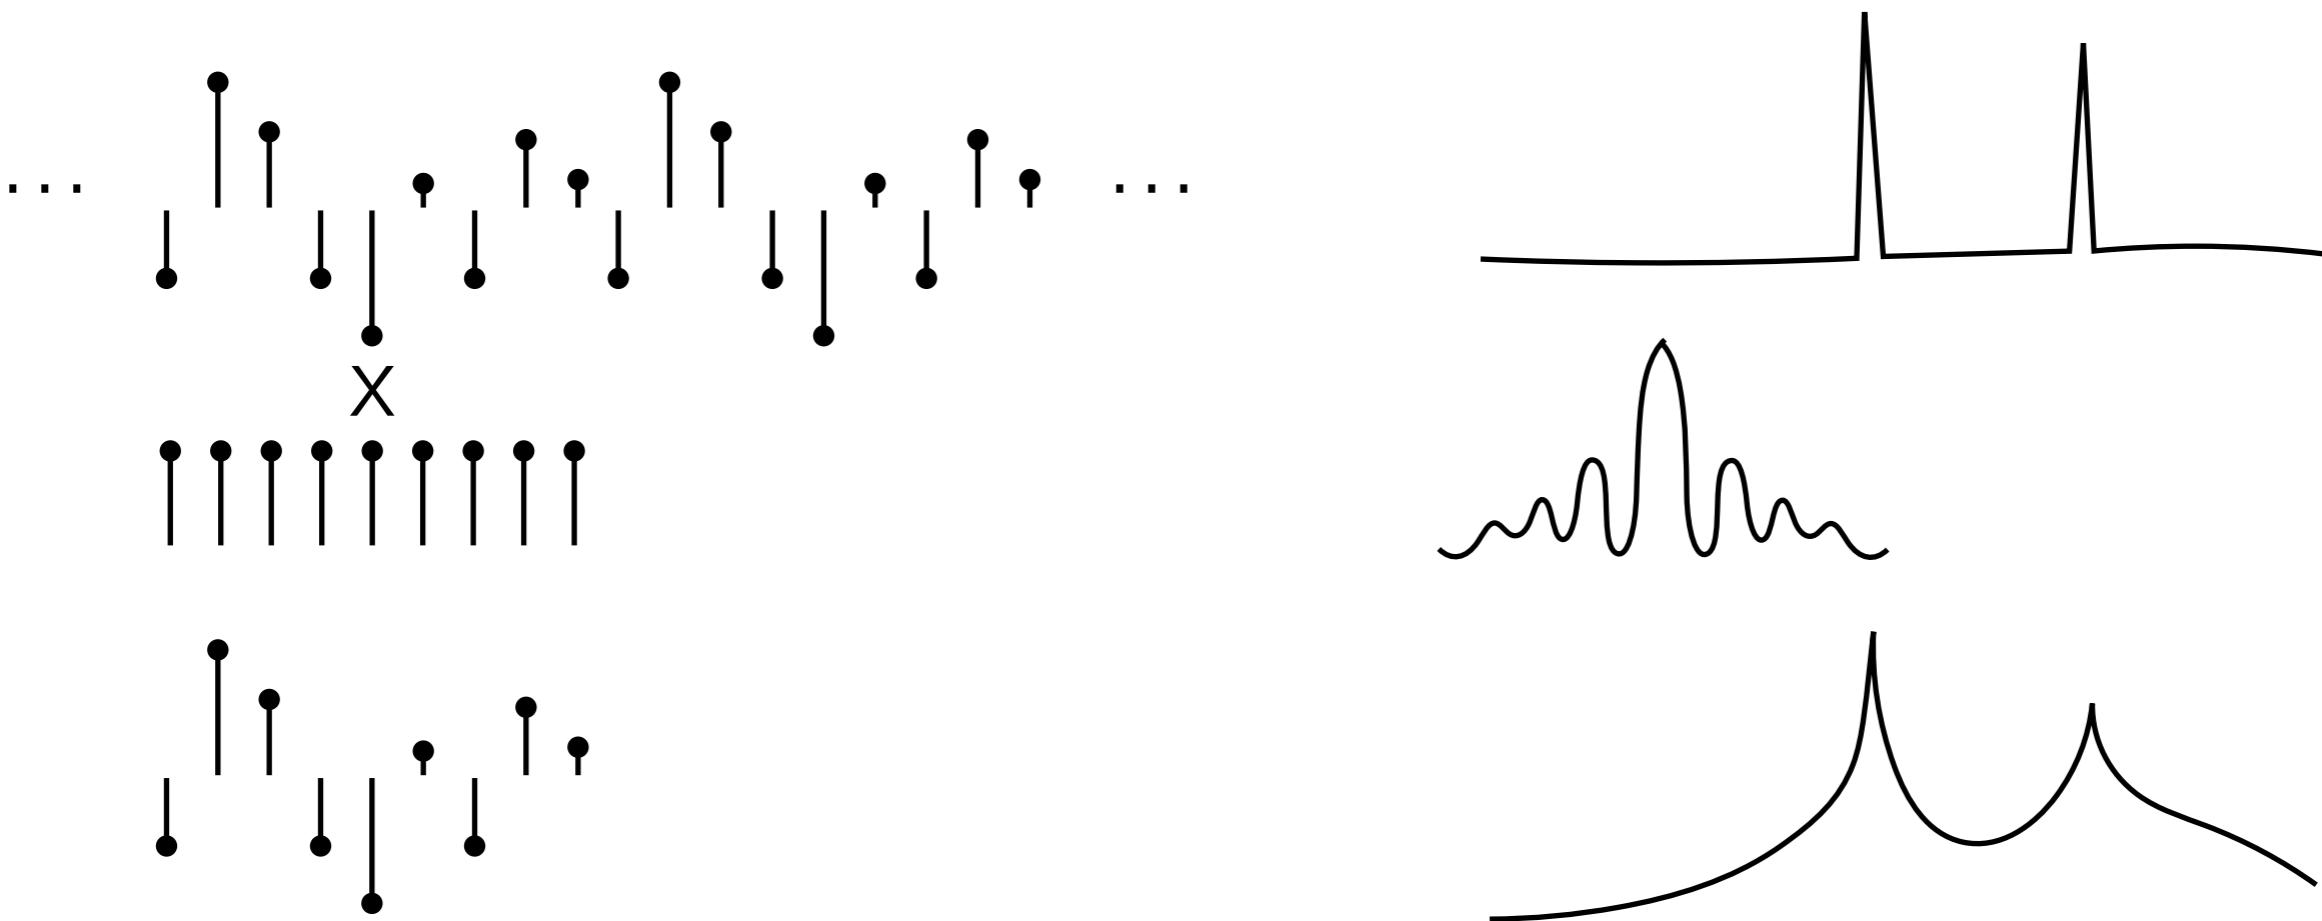
- just using N point FFT performs “Circular Convolution”
 - which is not linear convolution
 - causes the tail end of the convolution to “wrap around” to the beginning
 - FFT assumes the function is periodic (we did not talk about this)
- be aware of circularity when filtering your signal with the FFT
 - zero-padding can solve this for you!
 - zero-pad both signals to a length that will contain the entire convolution, $N+M-1$
 - for streaming, you must use overlap-and-add!
 - [http://en.wikipedia.org/wiki/Overlap–add method](http://en.wikipedia.org/wiki/Overlap–add_method)

overlap and add

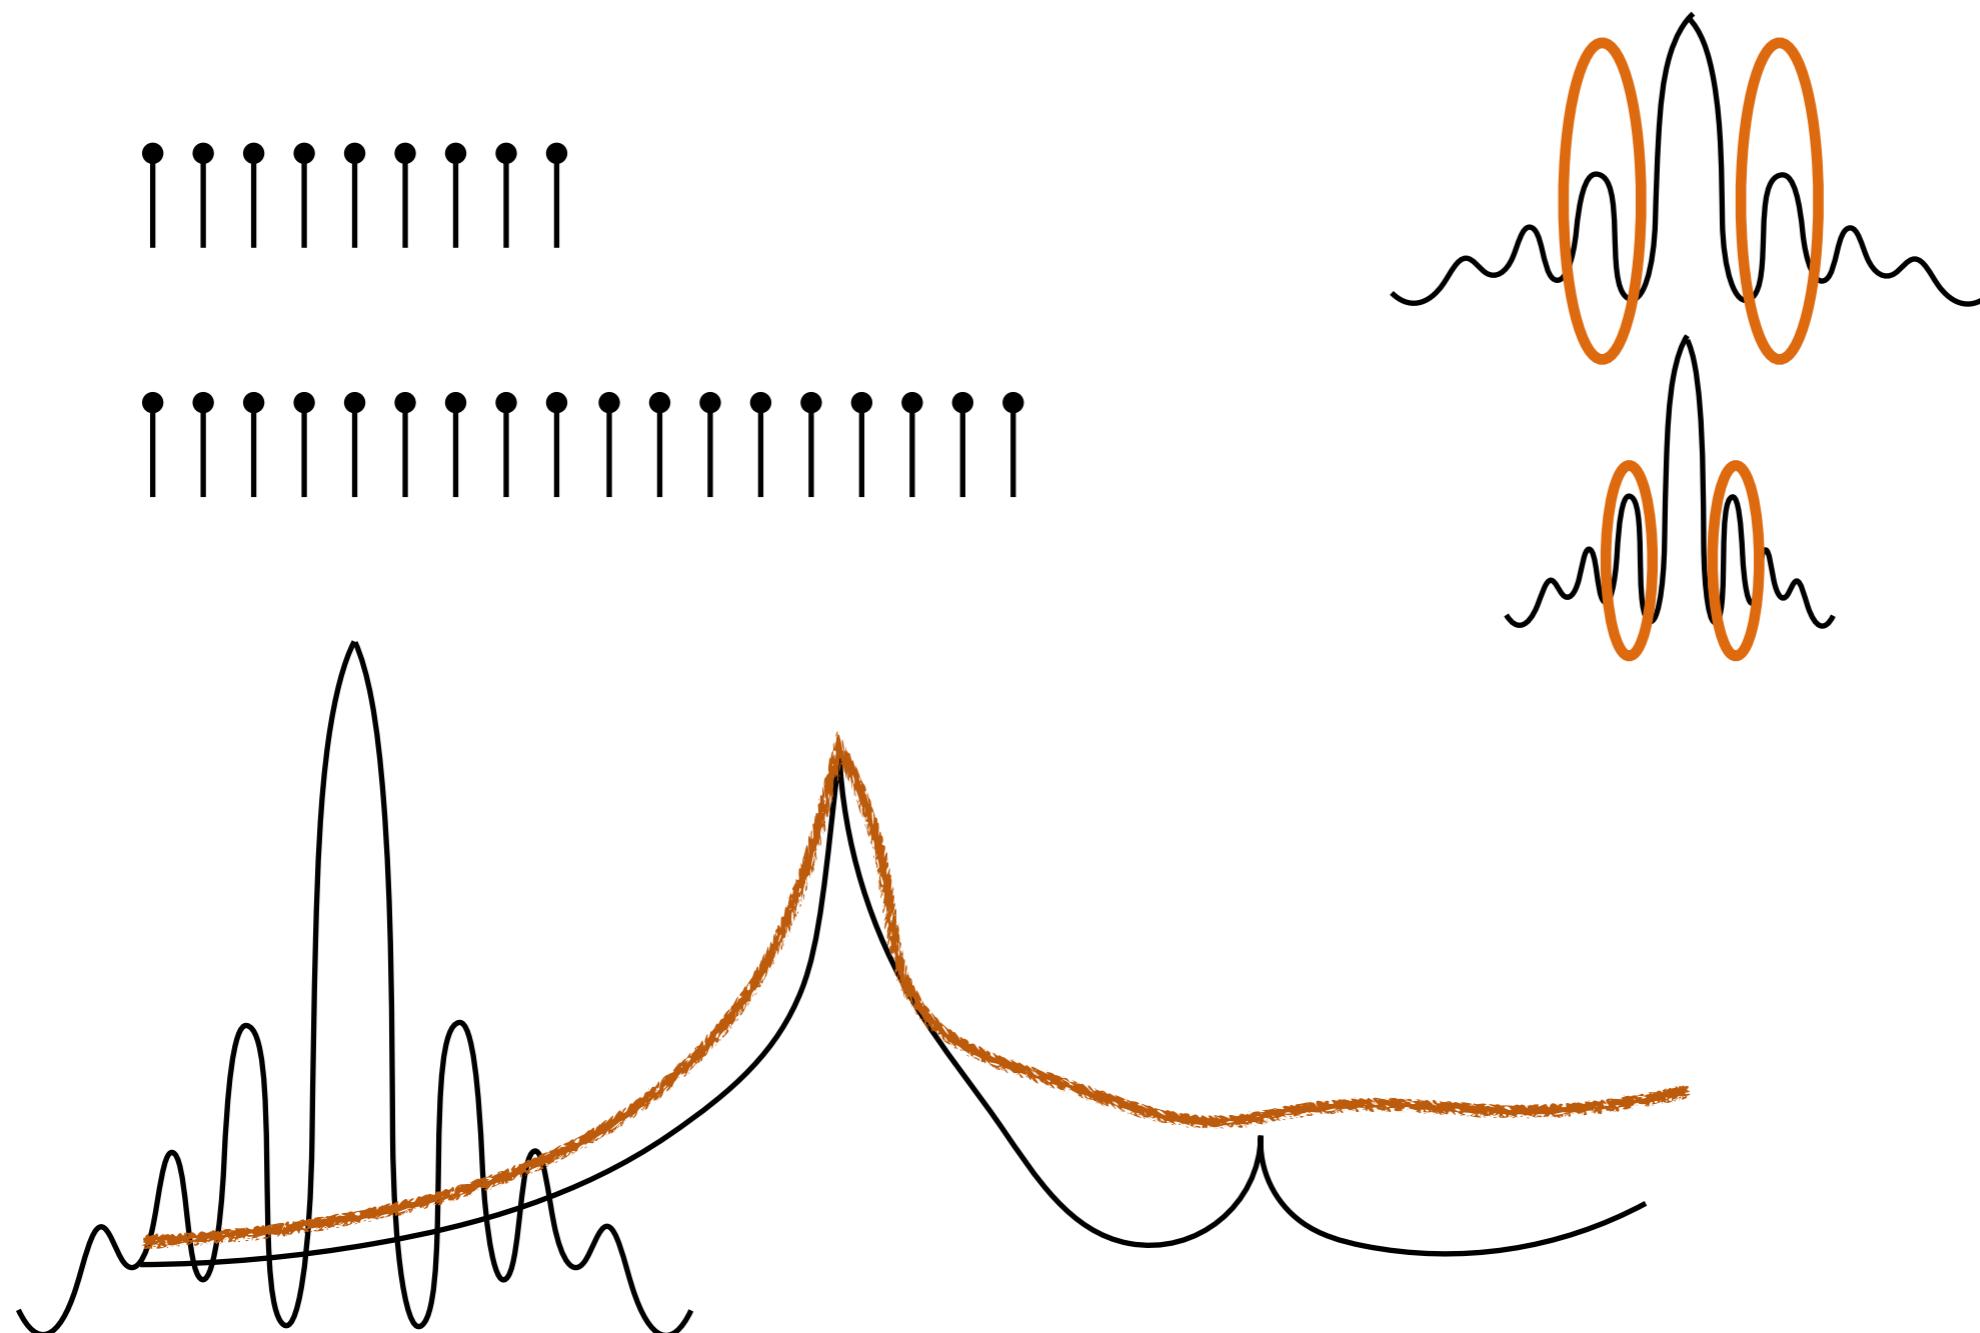


windowing: spectral BW widening

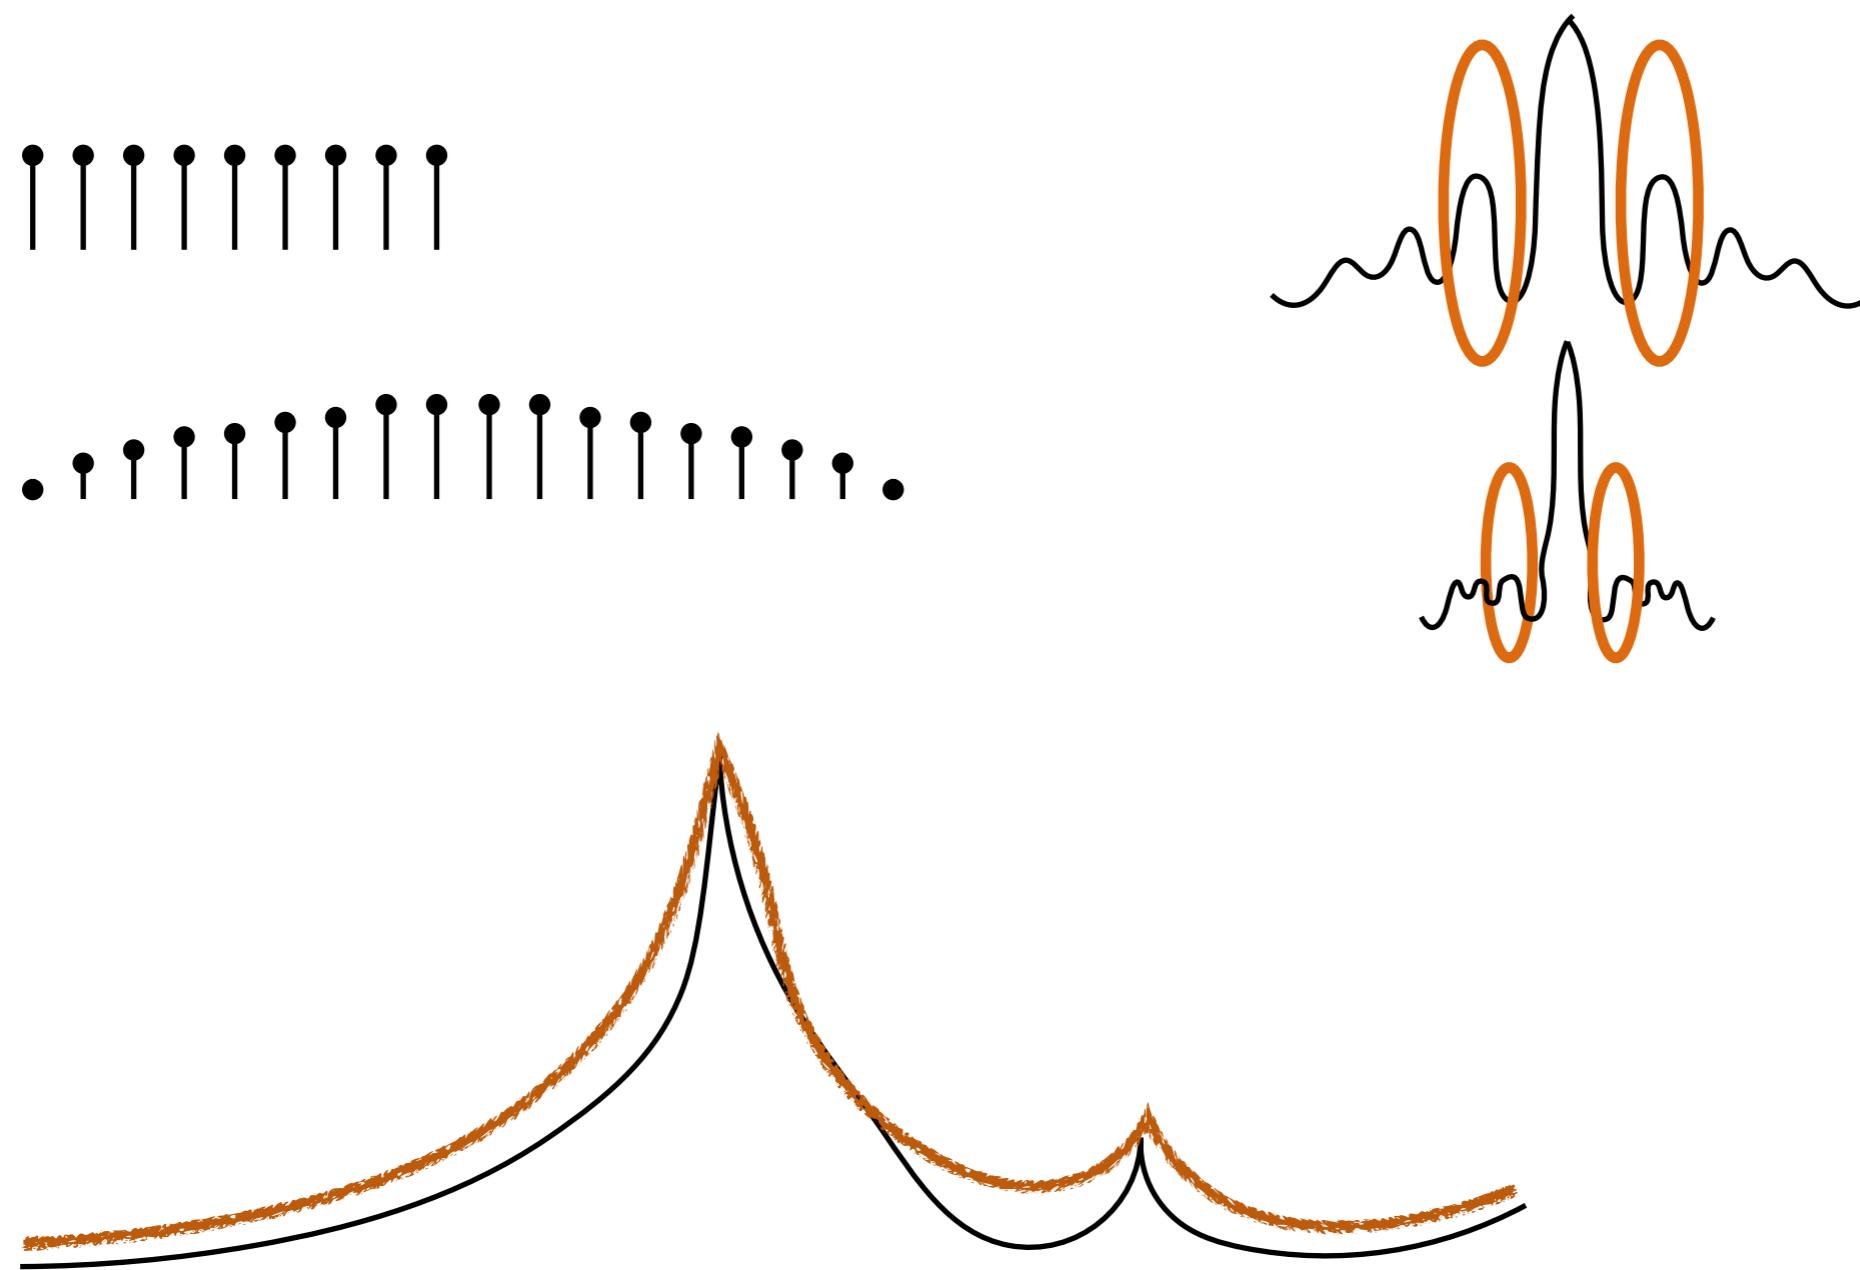
- multiplication in time is convolution in frequency
- a window is something we multiply in time with our signal
- windowing is unavoidable
 - why? we cannot take an infinite FFT...



windowing: spectral leakage



windowing: spectral leakage



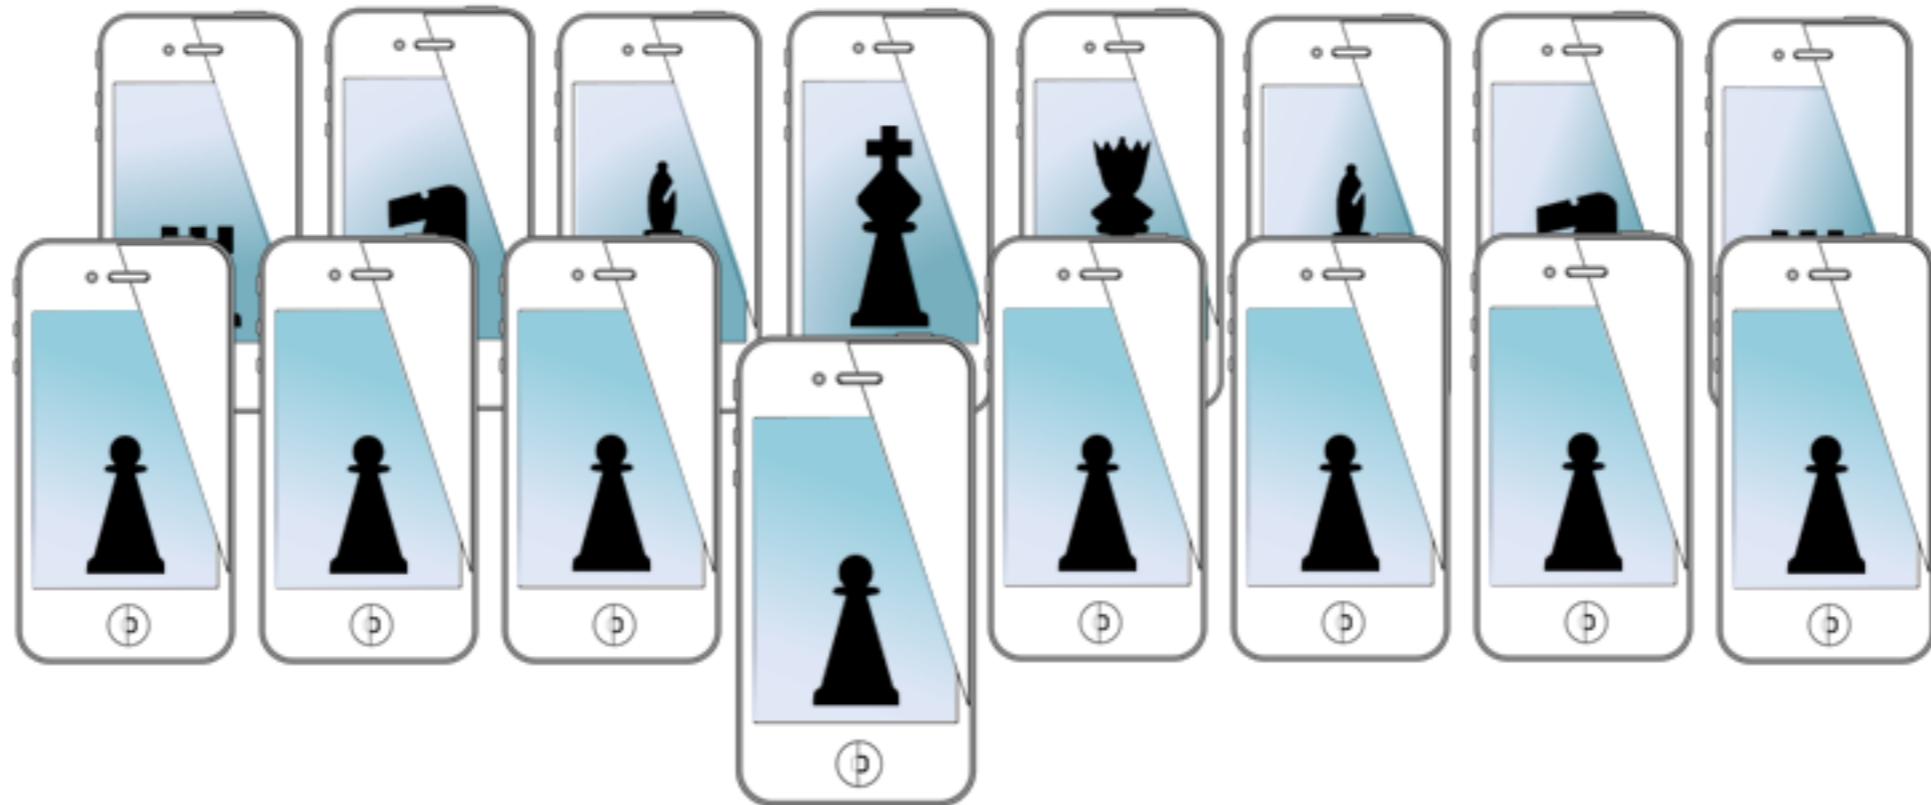
which window to use?

- depends
 - narrowest main lobe: rect
 - good tradeoff: hamming (or Von Hann)
 - optimal tradeoff for a given bandwidth:
 - discrete prolate spheroidal sequence (dpss, Slepian taper)

FFT review

- sampling rate
 - dictates the time between each sample, ($1 / \text{sampling rate}$)
 - max frequency we can measure is half of sampling rate
- resolution in frequency
 - tradeoff between length of FFT and sampling rate
 - each frequency “bin” is an index in the FFT array
 - each bin represents (F_s / N) Hz
 - what does that mean for 12 Hz accuracy?
- windowing is a result of “convolution” in frequency
 - some windows prevent “leakage” at the cost of frequency resolution

MOBILE SENSING LEARNING



CS5323 & 7323
Mobile Sensing and Learning

Supplemental Slides: filtering and windowing

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University