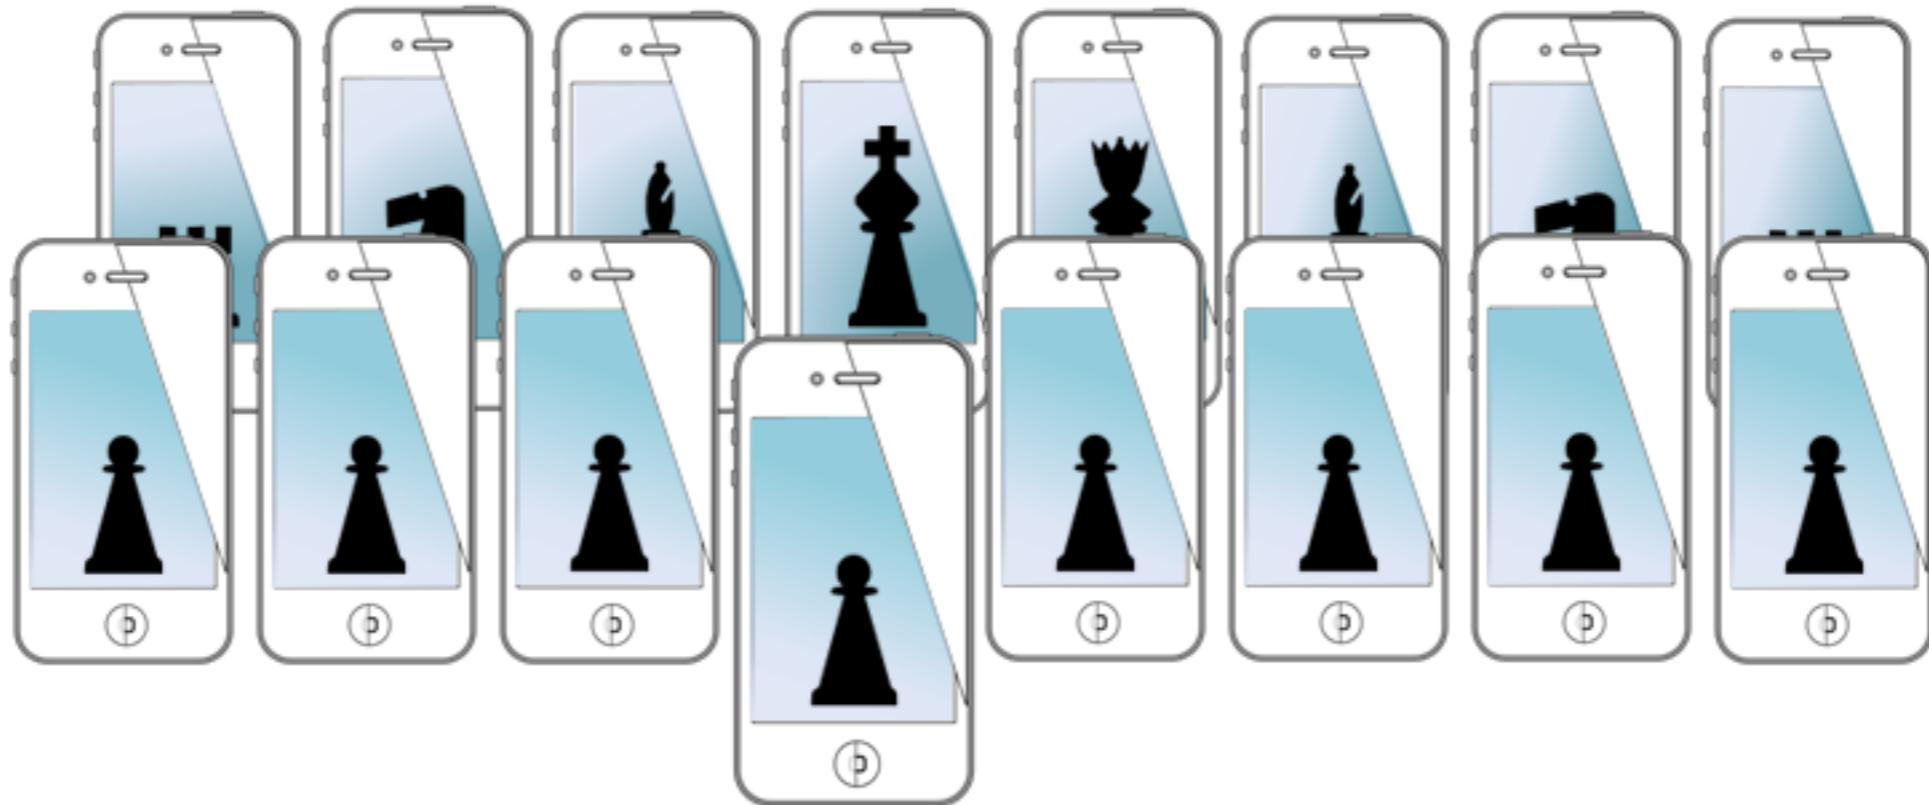


MOBILE SENSING LEARNING



CSE5323 & 7323
Mobile Sensing and Learning

week 7: computer vision with OpenCV

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University

video agenda

- OpenCV in iOS
 - we will look at **using** the tool
 - focus on
 - **outputs** of each algorithm
 - **how to use** each method
 - **ignore** most of what is **under the hood**

OpenCV in iOS

- open computer vision library
- released by intel
- many common functions are implemented
- **written in c++**, but has many wrappers
 - EMGU for .NET (c#,VC++,etc.), pycv2 for python, Java API for Android, and many, many more (**but not swift**)
- some hardware accelerations on iOS
 - but not the GPU, **only accelerate** functions
 - expect **slower** processing, but still **pretty fast!**

OpenCV installation

- download the opencv framework for iOS
 - might need to build from scratch... (or not)
- drag into project
- manually add a bunch of dependencies
- step by step instructions:
 - http://docs.opencv.org/doc/tutorials/ios/video_processing/video_processing.html
- remember to rename your model to .mm
- **alternatively:** do a git checkout of **ImageLab**

OpenCV video

- online tutorials will show you how to setup video capture
 - you can use the delegate protocol and a cvVideoCamera
 - or keep using **VideoAnalgesic!**

```
let bridge = OpenCVBridge()

self.bridge.setImage(retImage, withBounds: retImage.extent, andContext:
                     self.videoManager.getCIContext())
self.bridge.processImage()
retImage = self.bridge.getImage() ← not deprecated anymore!

let f = getFaces(inputImage)
self.bridge.setImage(retImage,
                     withBounds: f[0].bounds, // the first face
                     andContext: self.videoManager.getCIContext())

self.bridge.processImage()
retImage = self.bridge.getImageComposite() // get back (overlaid on original)
```

OpenCV video

- now head over to OpenCVBridge.mm

```
-(void)processImage{  
  
    cvtColor( _image, frame_gray, CV_BGR2GRAY );  
    bitwise_not(frame_gray, _image);  
  
}
```

```
-(void)processImage{  
  
    cvtColor(_image, image_copy, COLOR_BGRA2GRAY);  
    cv::threshold(image_copy, image_copy, 0, 255, CV_THRESH_BINARY | CV_THRESH_OTSU);  
    cvtColor(image_copy, _image, CV_GRAY2BGR); //add back for display  
  
}
```

OpenCV video

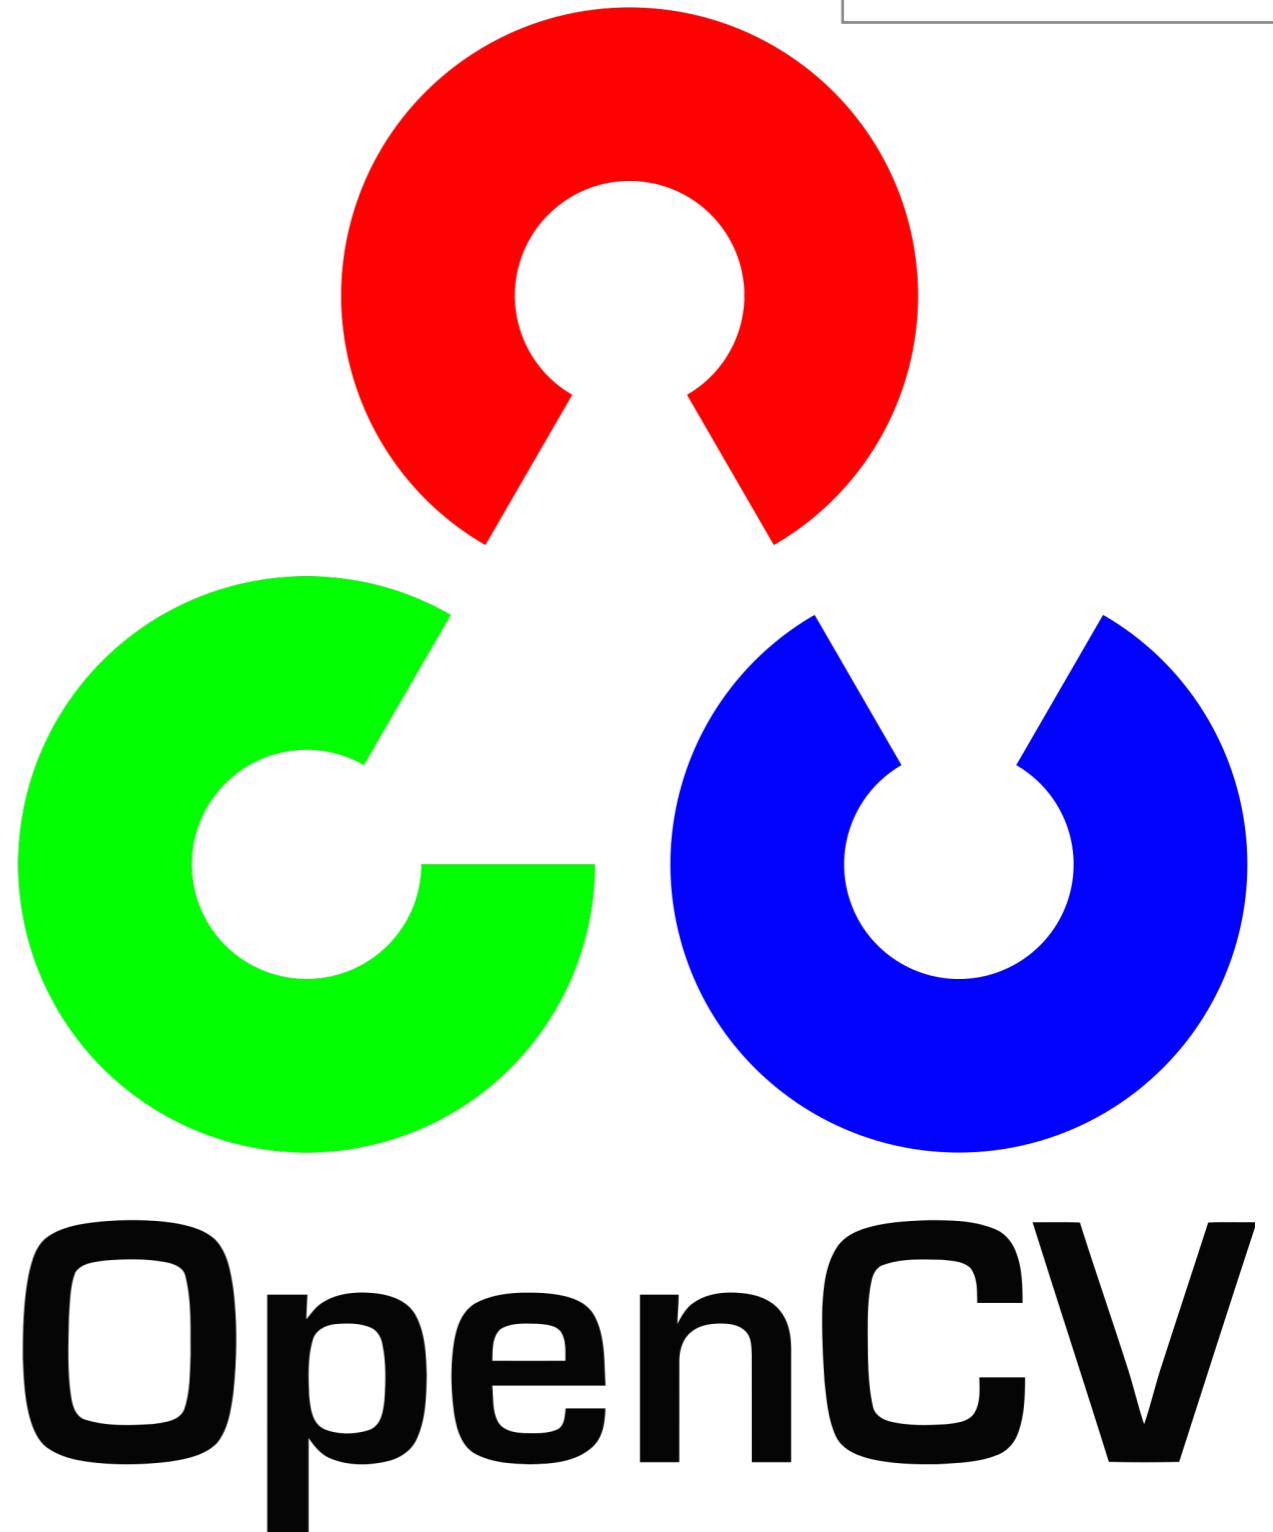
- you are not on the main queue here
- I am using an older version of OpenCV 2
 - but you can build OpenCV 3.0 and use it if you want!
- OpenCV is **mostly updated** for iOS
 - some functions are deprecated, but currently work, some transforms are not quite right

```
-(void)processImage{  
  
    cvtColor( _image, frame_gray, CV_BGR2GRAY );  
    bitwise_not(frame_gray, _image);  
  
}
```

OpenCV Setup Demo



And now its
time for a demo



access torch

before we get too far, some convenience methods

```
@IBAction func toggleFlash(sender: AnyObject) {  
    self.videoManager.toggleFlash()  
}
```

```
@IBAction func setFlashLevel(sender: UISlider) {  
    if(sender.value>0.0){  
        self.videoManager.turnOnFlashwithLevel(sender.value)  
    }  
    else if(sender.value==0.0){  
        self.videoManager.turnOffFlash()  
    }  
}
```

OpenCV operations

- your input is a matrix
- data is interleaved BGR
 - if setting color conversion to CV_BGRA2BGR
 - must get a pointer in the array for the row and column

```
image.ptr(row, column) starts from upper left  
  
for(int i=0;i<50;i++){  
    image.ptr(i, i)[0] = 255;  
    image.ptr(i, i)[1] = 0;  
    image.ptr(i, i)[2] = 0;  
}  
  
for(int i=0;i<50;i++){  
    uchar *pt = image.ptr(i, i);  
    pt[0] = 255;  
    pt[1] = 0;  
    pt[2] = 0;  
}  
  
for(int i=0;i<50;i++){  
    uchar *pt = image.ptr(i, i);  
    pt[0] = 255;  
    pt[1] = 0;  
    pt[2] = 0;  
    pt[3] = 255;  
    pt[4] = 0;  
    pt[5] = 0;  
}
```

blue

green

red

pixel at i,i

next pixel in row

OpenCV operations

- filtering

```
Mat gauss = cv::getGaussianKernel(25, 3);
cv::filter2D(image_copy, image_copy, -1, gauss);
GaussianBlur(image_copy, image_copy, cv::Size(3, 3), 2, 2 );
```

- inversion

```
bitwise_not(image_copy, image_copy);
```

- statistics

```
Scalar avgPixelIntensity = cv::mean( image_copy );
avgPixelIntensity.val[0]
avgPixelIntensity.val[1]
avgPixelIntensity.val[2]
```

- color conversion

```
cvtColor(image, image_copy, CV_BGRA2BGR);
cvtColor(image, image_copy, CV_GRAY2BGR);
cvtColor(image, image_copy, CV_RGB2BGR);
cvtColor(image, image_copy, CV_BGR2HSV);
cvtColor(image, image_copy, CV_BGR2Lab);
cvtColor(image, image_copy, CV_BGR2Lab);
cvtColor(image, image_copy, CV_BGR2YCrCb);
cvtColor(image, image_copy, CV_BGR2YUV);
```

OpenCV Demo

- basic operations



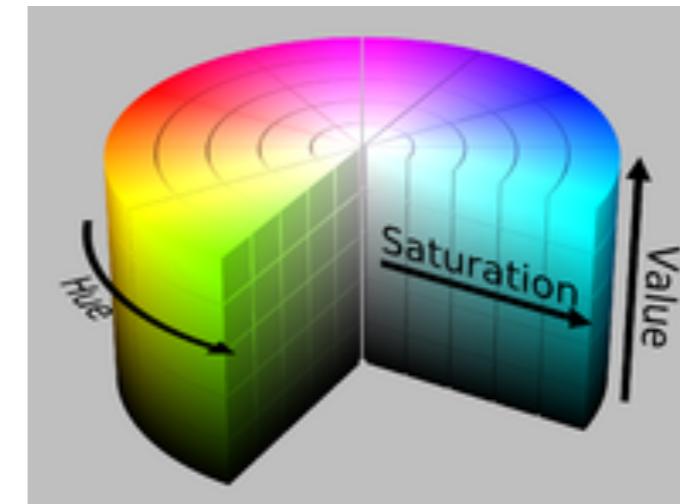
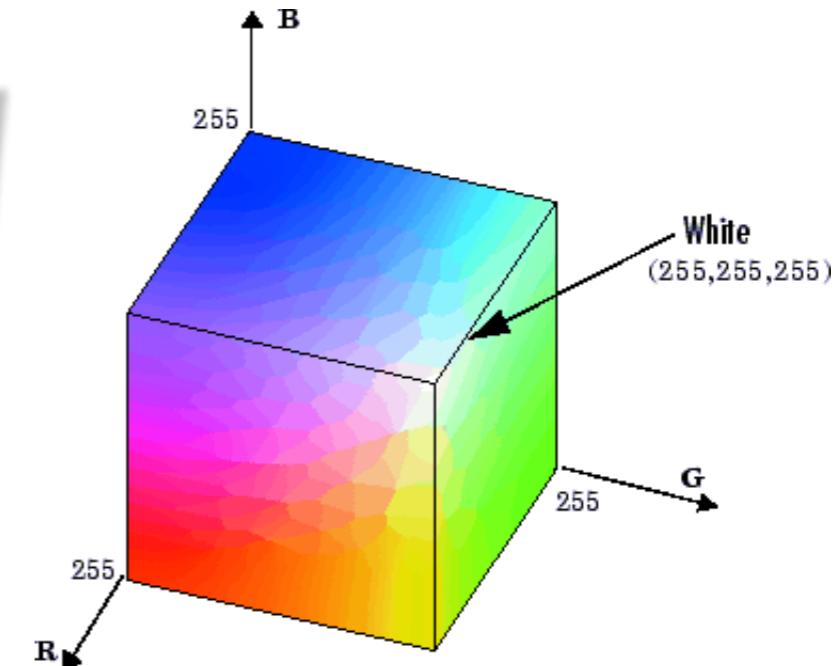
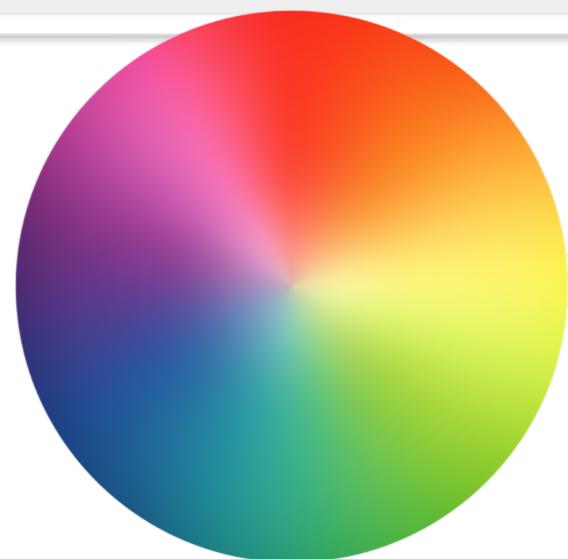
And now its
time for a demo



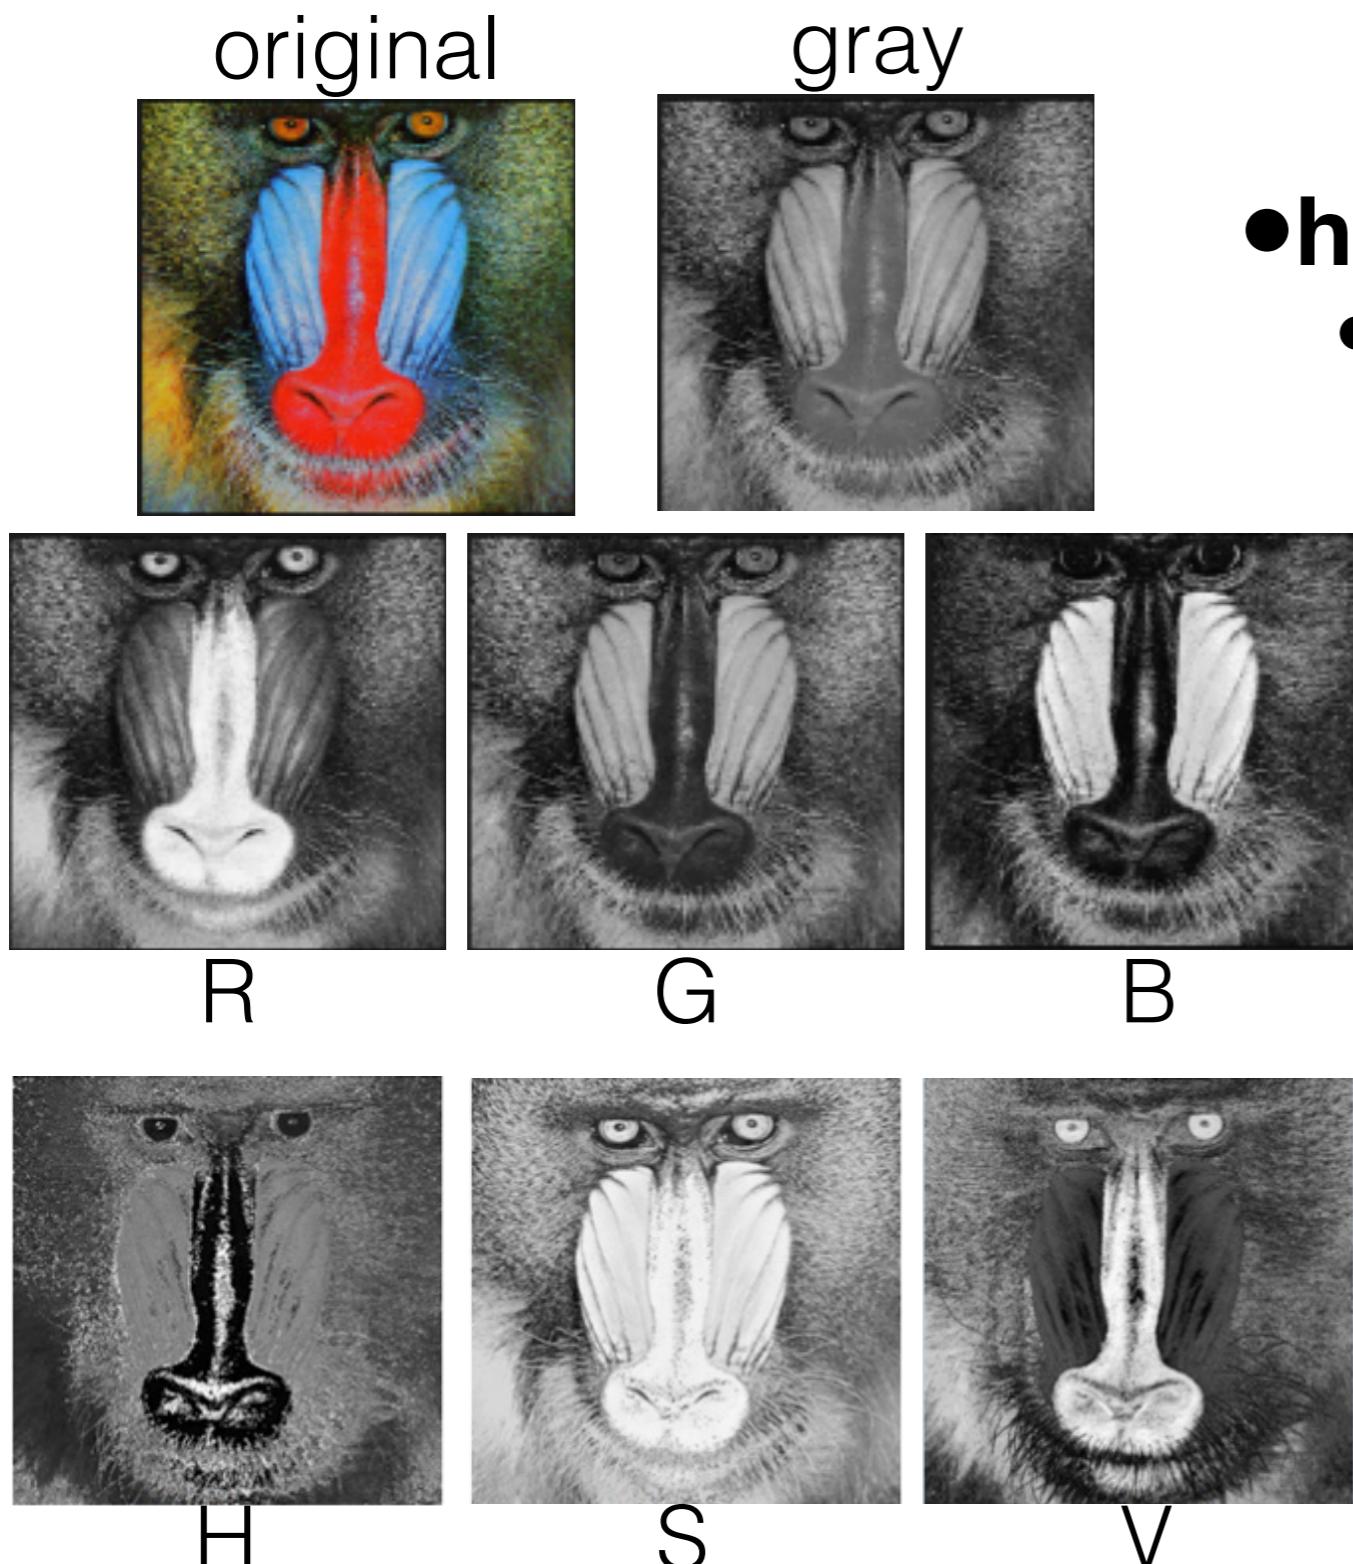
color conversion

- to display properly, use BGRA

```
cvtColor(image, image_copy, CV_BGRA2BGR);  
  
cvtColor(image_copy, image, CV_BGR2BGRA);  
  
cvtColor(image, image_copy, CV_BGRA2GRAY);  
cvtColor(image_copy, image, CV_GRAY2BGRA);  
  
cvtColor(image, image_copy, CV_BGRA2HSV);  
cvtColor(image_copy, image, CV_HSV2BGRA);
```

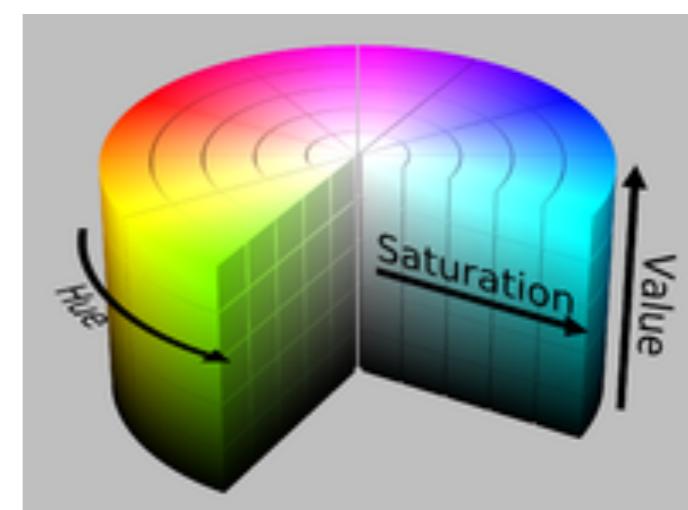


color conversion



•hsv

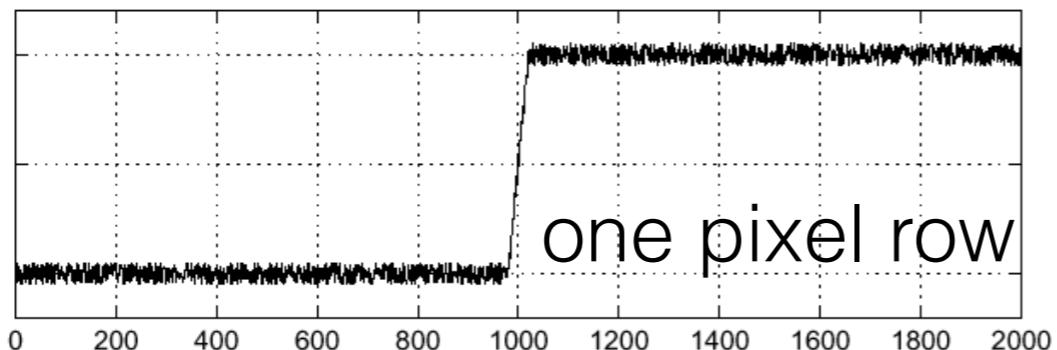
- what we perceive as color, rather than “sense” as color (sort of)
 - hue: the color value
 - saturation: the richness of the color relative to brightness
 - value: the intensity



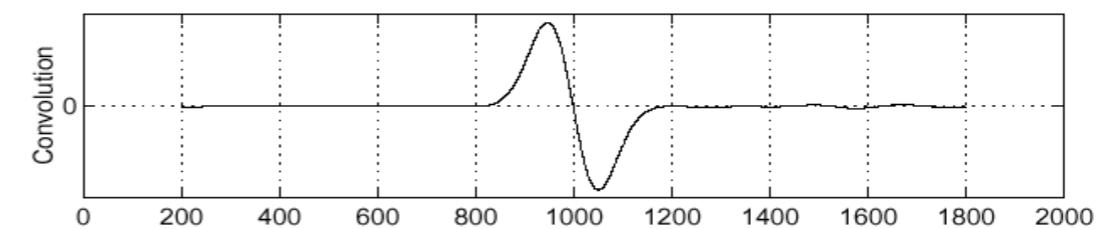
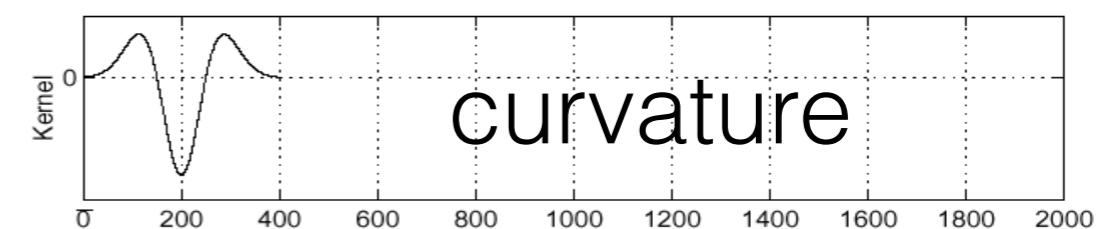
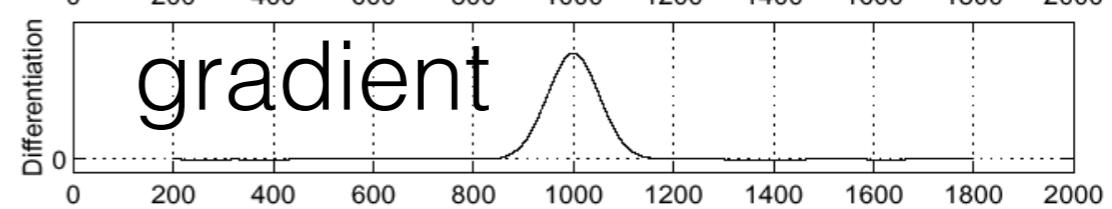
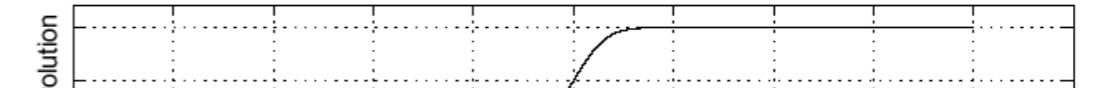
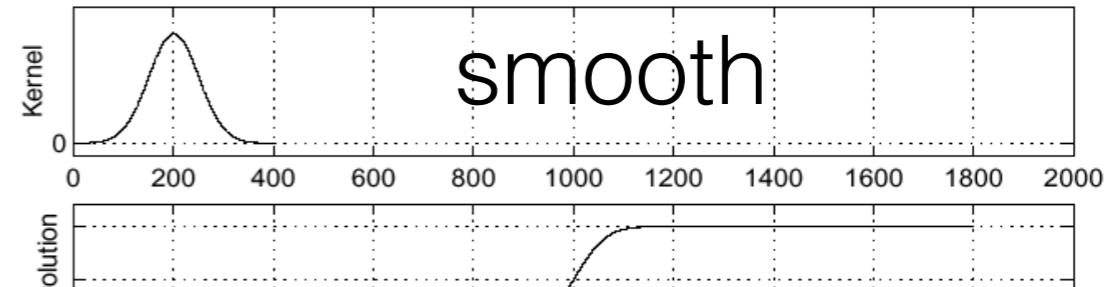
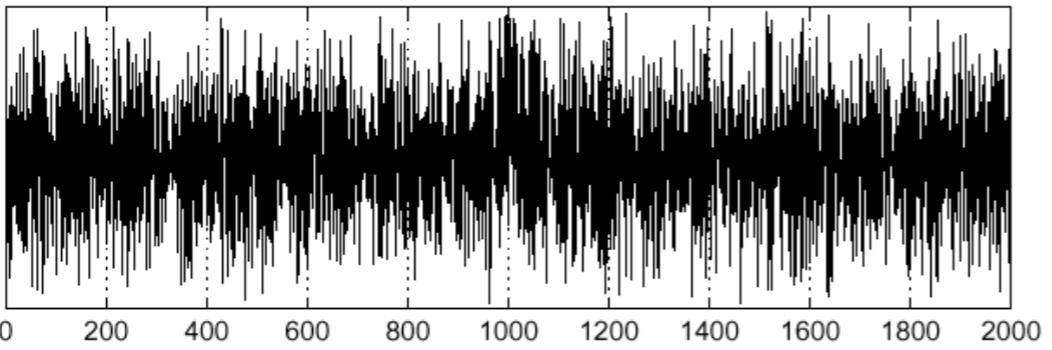
edge detection

- can use linear filters to get gradient

$$\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline 1 & 2 & 1 \\ \hline -2 & 0 & 2 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$



$$\frac{d}{dx}f(x)$$



images courtesy of S. Narasimhan

gradient example



can we do better?

check local maxima

canny edge detection

- get local maxima of gradient

- see if above first threshold

- if yes, then add to a list parallel to



$G_{x,y}$



From magnitude
and direction of
 $G_{x,y}$

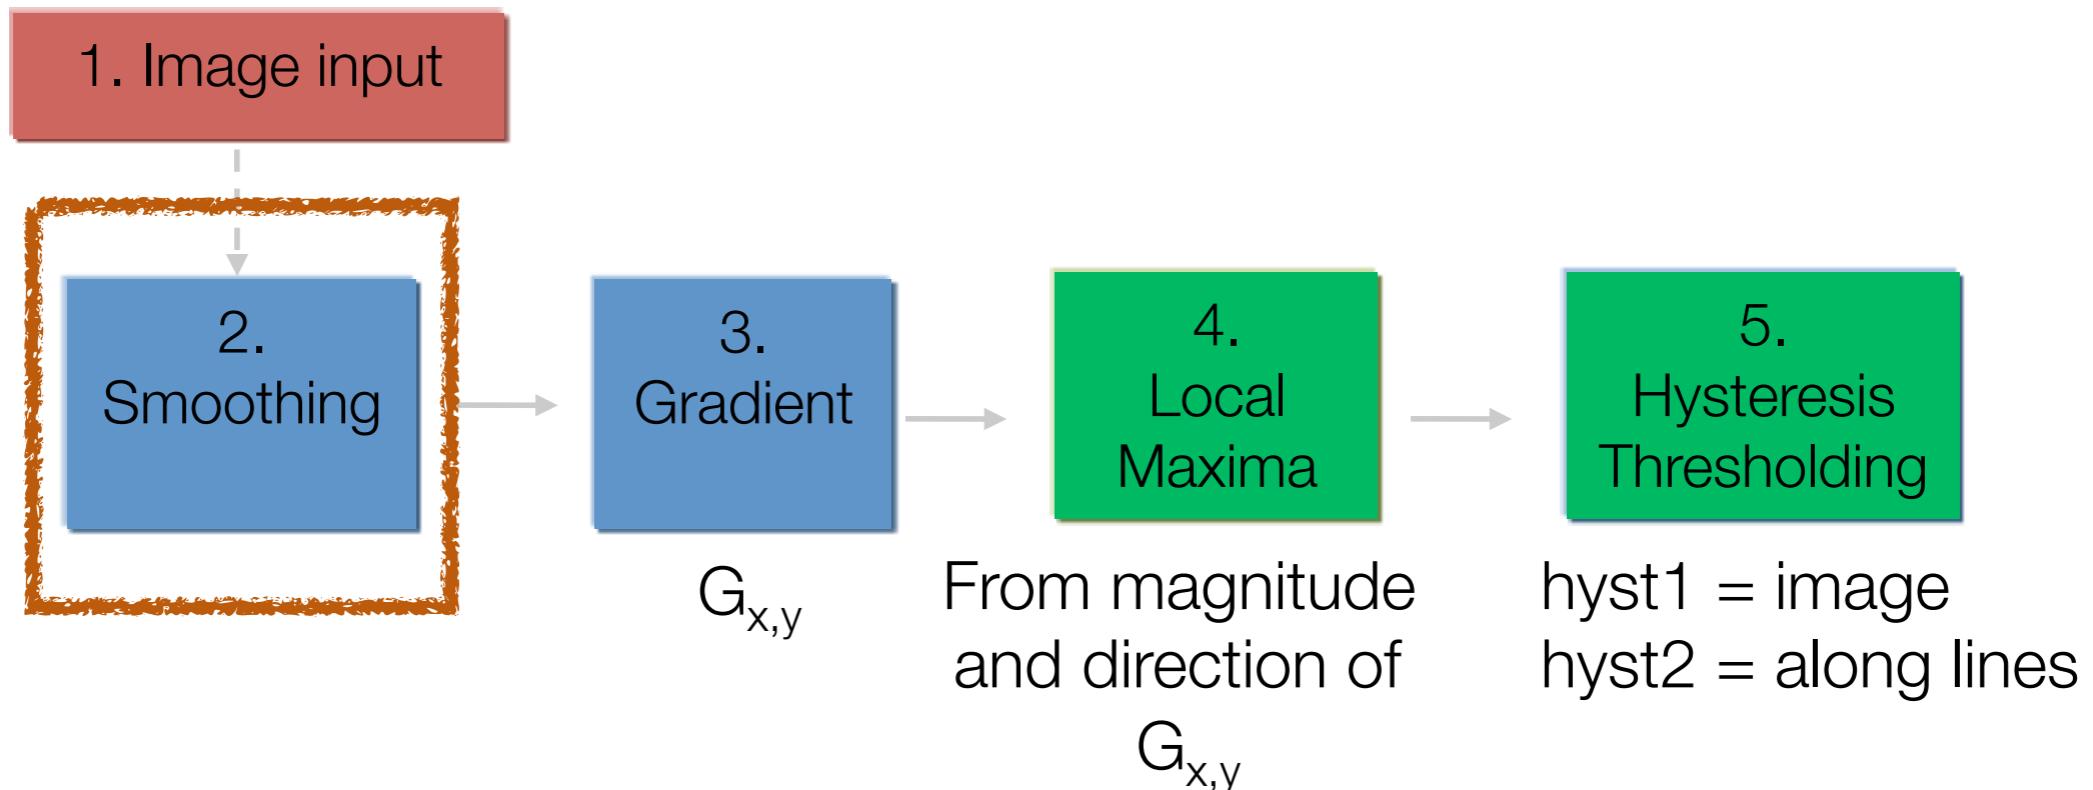
hyst1 = image
hyst2 = along lines

canny edge detection

```
const int kCannyLowThreshold = 300;  
cvtColor(image, grayFrame, COLOR_RGB2GRAY);  
  
// Perform Canny edge detection using s1  
Canny(grayFrame, output,  
      kCannyLowThreshold,  
      kCannyLowThreshold*7,  
      kFilterKernelSize);  
  
// copy back for further processing  
cvtColor(output, image, CV_GRAY2BGRA); //add back for display
```

Annotations:

- hyst2: small threshold
- hyst1: large threshold
- size of gradient filters



edges to contours

- connected components search
- contour detection from “outside” of component

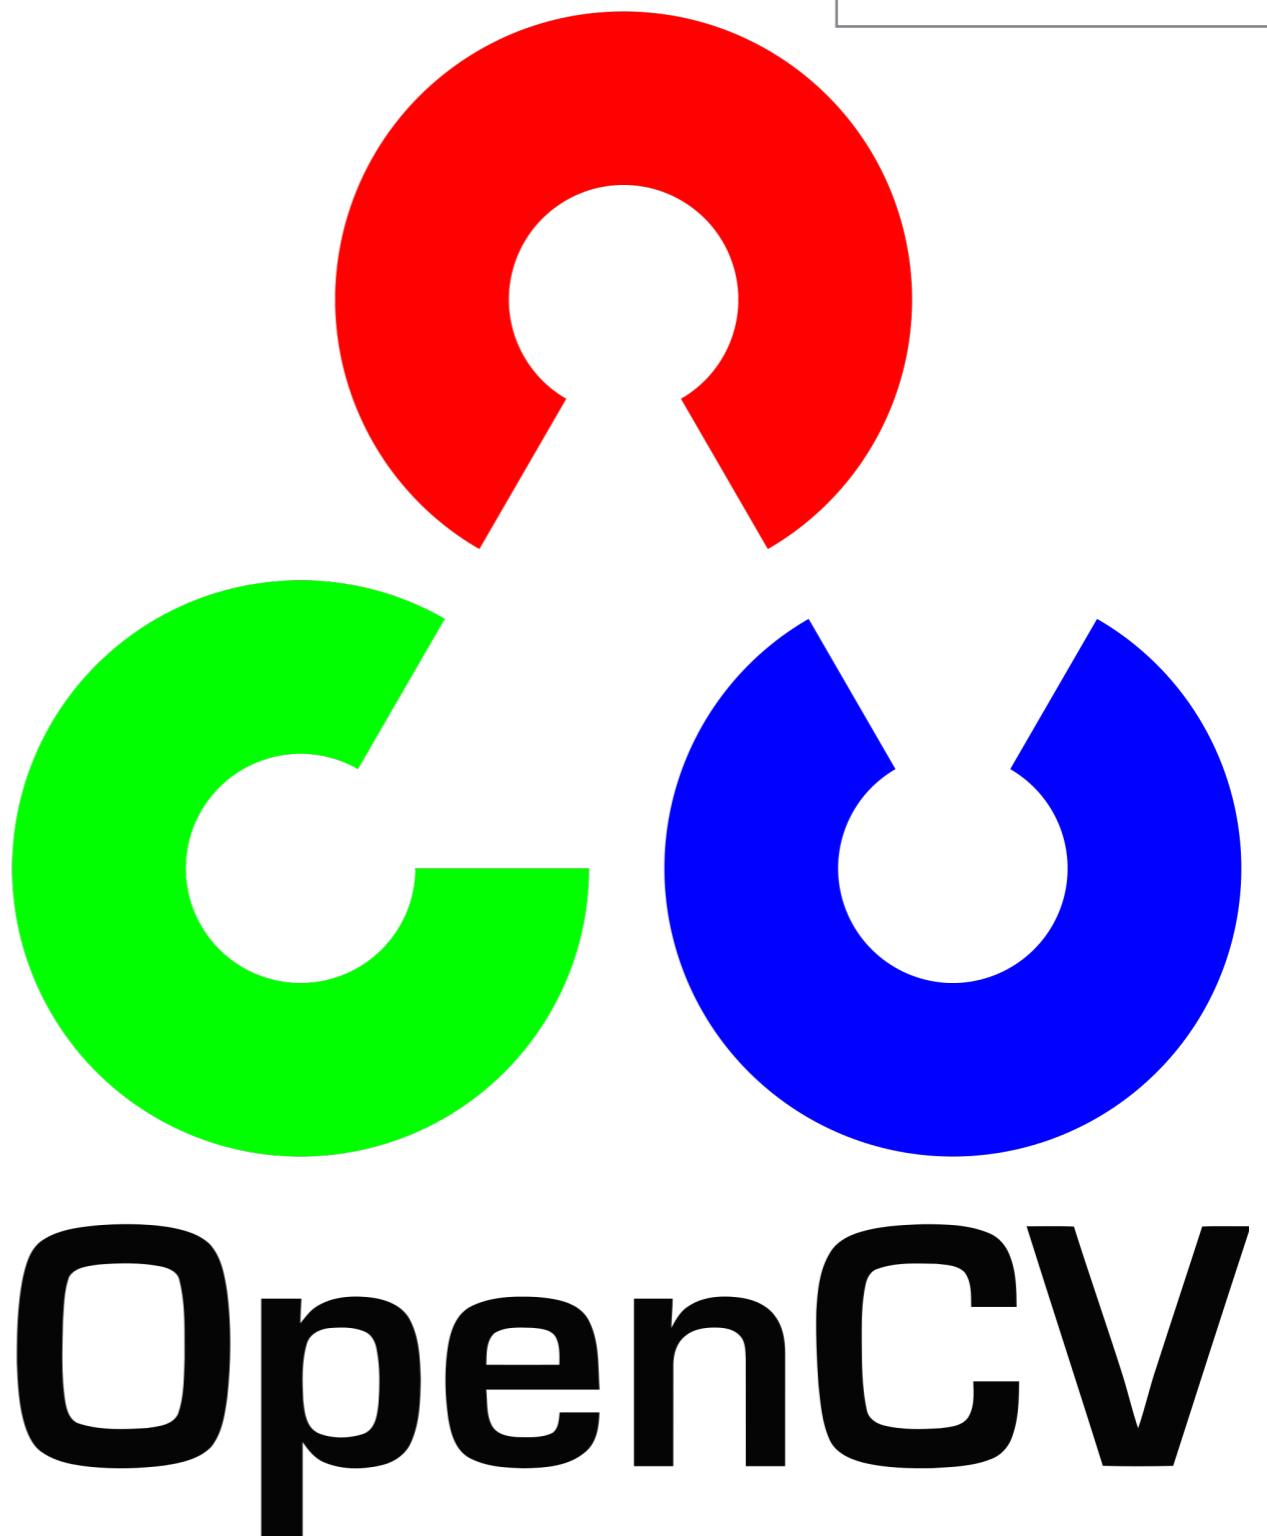
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	1	1	1	1	0	1	1	1	1	1	1	0	0
0	1	1	0	0	1	0	0	0	0	1	1	1	1	1	1	0	0
0	1	1	0	0	1	1	1	0	0	1	1	1	1	1	1	1	1
0	1	1	0	0	1	0	0	0	0	1	1	1	1	1	1	1	1
0	1	0	0	0	1	1	1	1	0	1	1	1	1	1	1	1	1
0	1	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

OpenCV edge demo

- edges
- contours

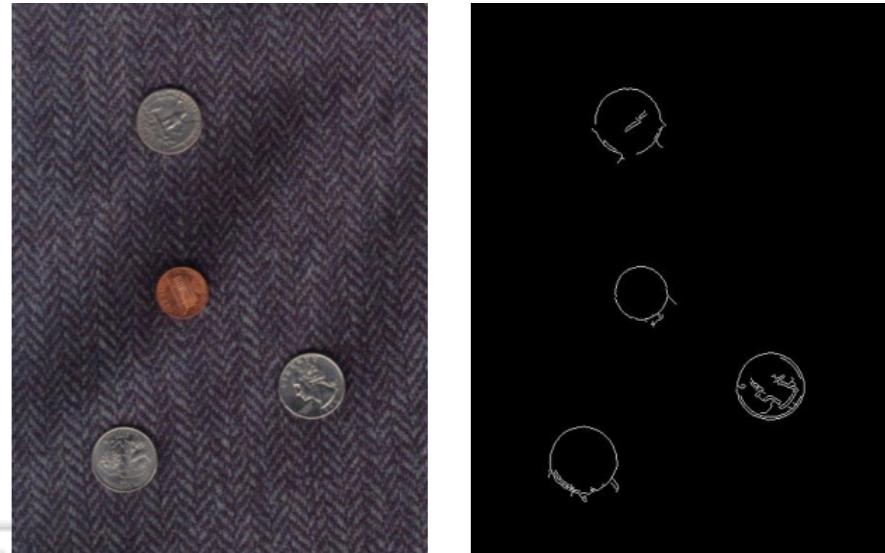


And now its
time for a demo



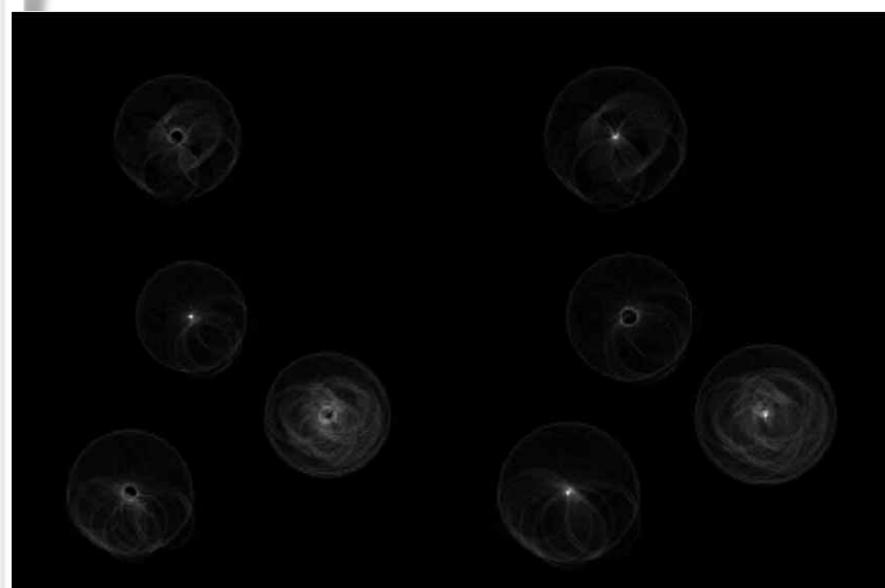
hough transform

- In general, hough transform consists of
 - edge detection
 - for each detected point,
 - draw shape with different parameter
 - accumulation in parameter space
 - look for local maxima
 - for a circle, look for maxima in (x,y,R)



```
vector<Vec3f> circles;

// Apply the Hough Transform to find the circles
HoughCircles( image_copy, circles,
               CV_HOUGH_GRADIENT,
               1, // downsample factor
               image_copy.rows/30, // distance between centers
               kCannyLowThreshold/2, // canny upper thresh
               40, // magnitude thresh for hough param space
               0, 0 ); // min/max centers
```



generic Haar cascade

- remember Haar cascade filtering?
- you can use any trained cascade of classifiers
 - just need a trained file to load
- get some trained xml files, here is a start:
 - <http://alereimondo.no-ip.org/OpenCV/34/>
 - or train your own (this is not trivial)
 - http://docs.opencv.org/doc/user_guide/ug_traincascade.html
 - database of positive example images
 - database of negative example images

Haar syntax for iOS

```
cv::CascadeClassifier classifier;

// load in custom trained Haar Cascade filter
// This one is a famous trained face detector from Rainer Lienhart
// http://www.lienhart.de/Prof._Dr._Rainer_Lienhart/Welcome.html
NSString *fileName = [[NSBundle mainBundle]
                      pathForResource:@"haarcascade_frontalface_alt2" ofType:@"xml"];

classifier = cv::CascadeClassifier([fileName UTF8String]);

cvtColor(image, grayFrame, CV_BGRA2GRAY);
vector<cv::Rect> objects;

// run classifier
classifier.detectMultiScale(grayFrame, objects);

// display bounding rectangles around the detected objects
for( vector<cv::Rect>::const_iterator r = objects.begin(); r != objects.end(); r++)
{
    cv::rectangle( image,
                  cvPoint( r->x, r->y ),
                  cvPoint( r->x + r->width, r->y + r->height ),
                  Scalar(0,0,255,255));
}
```

have fun with it!

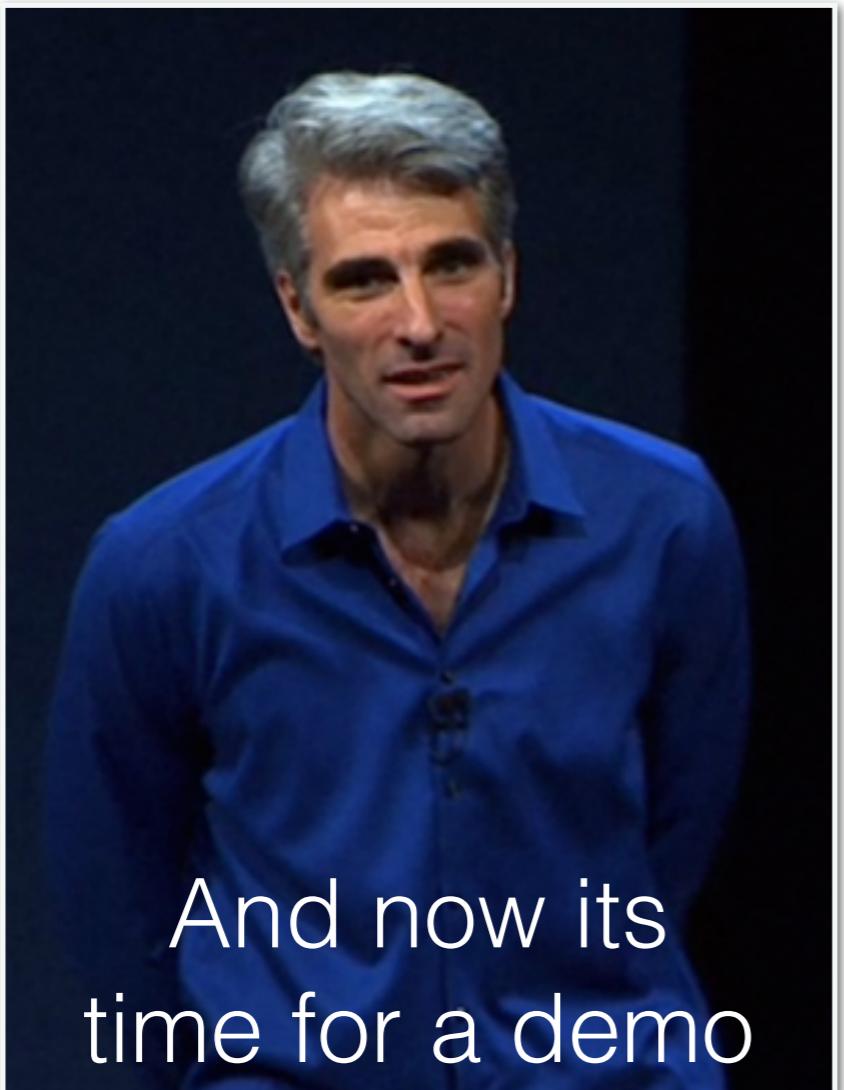
- OpenCV is a powerful framework
- many contributors
 - each algorithm has (possibly) different semantics
 - highly comprehensive and updated
 - lots of examples
 - OpenCV is absolutely an industry standard

for class

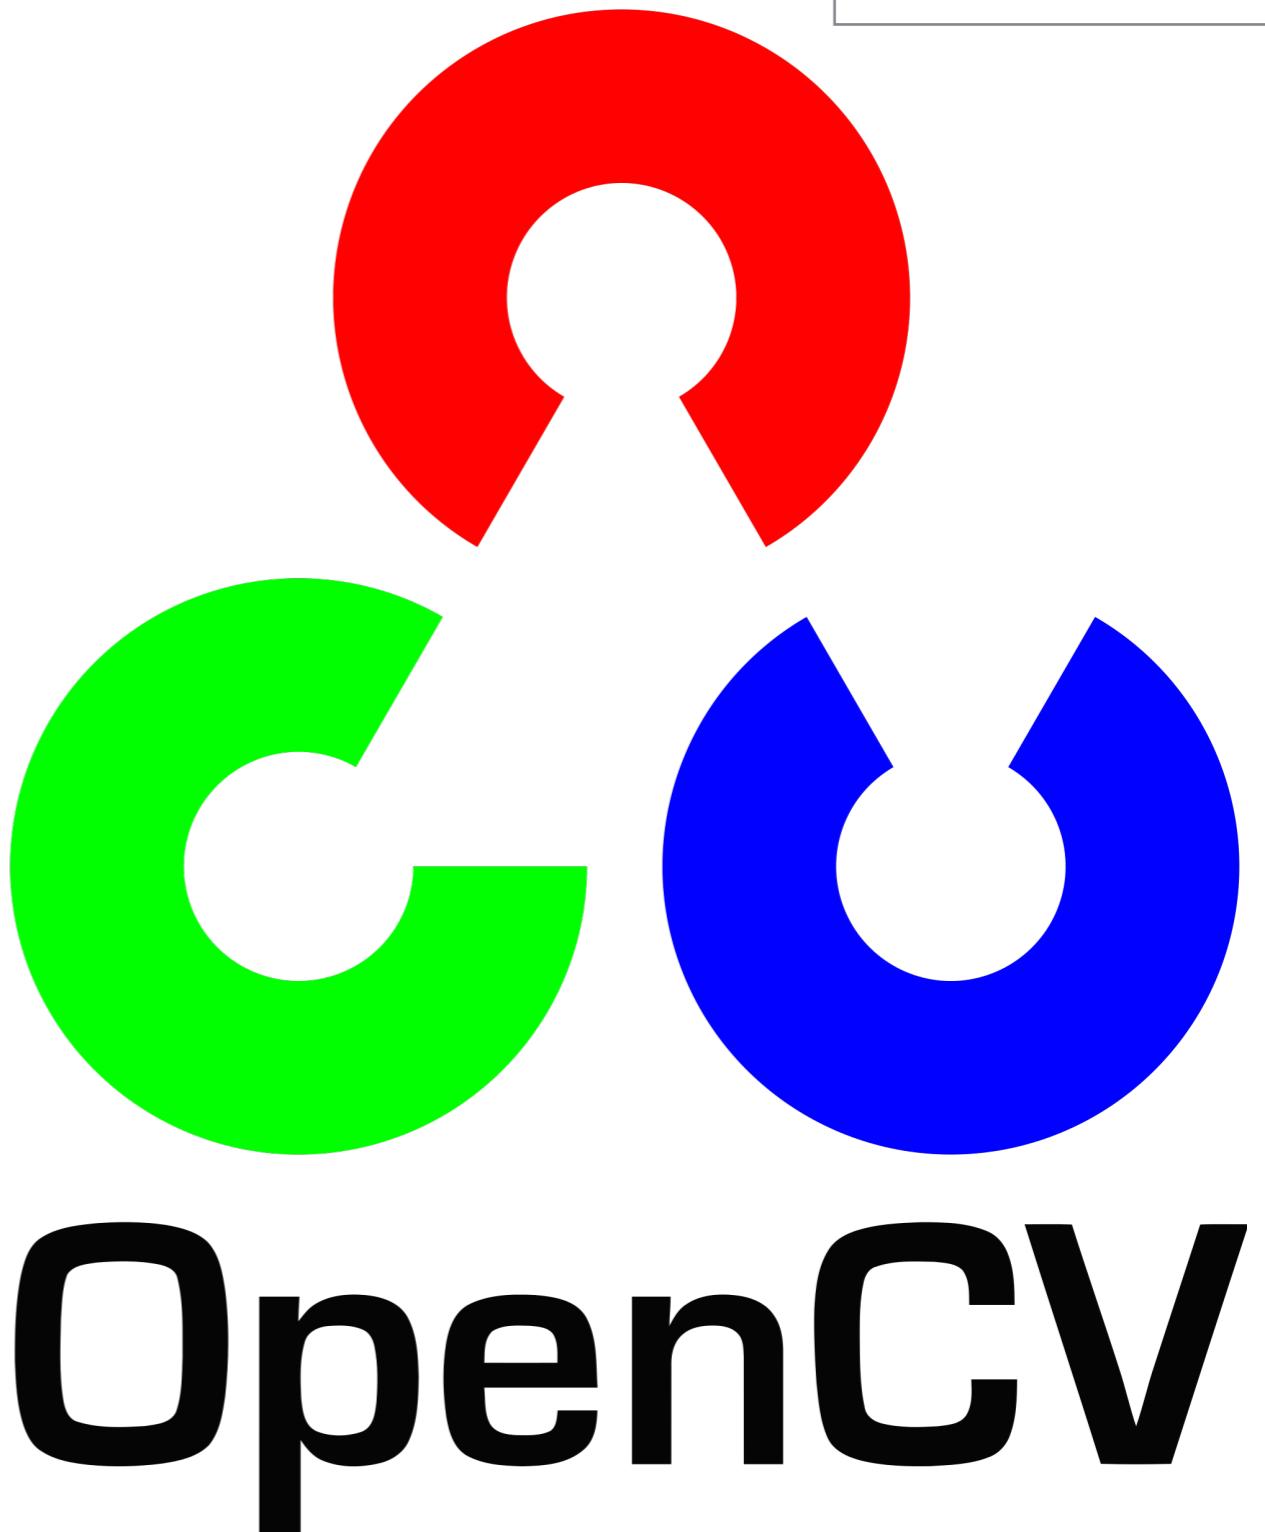
- download this example project from GitHub
- come ready to perform some computer vision

OpenCV Demo

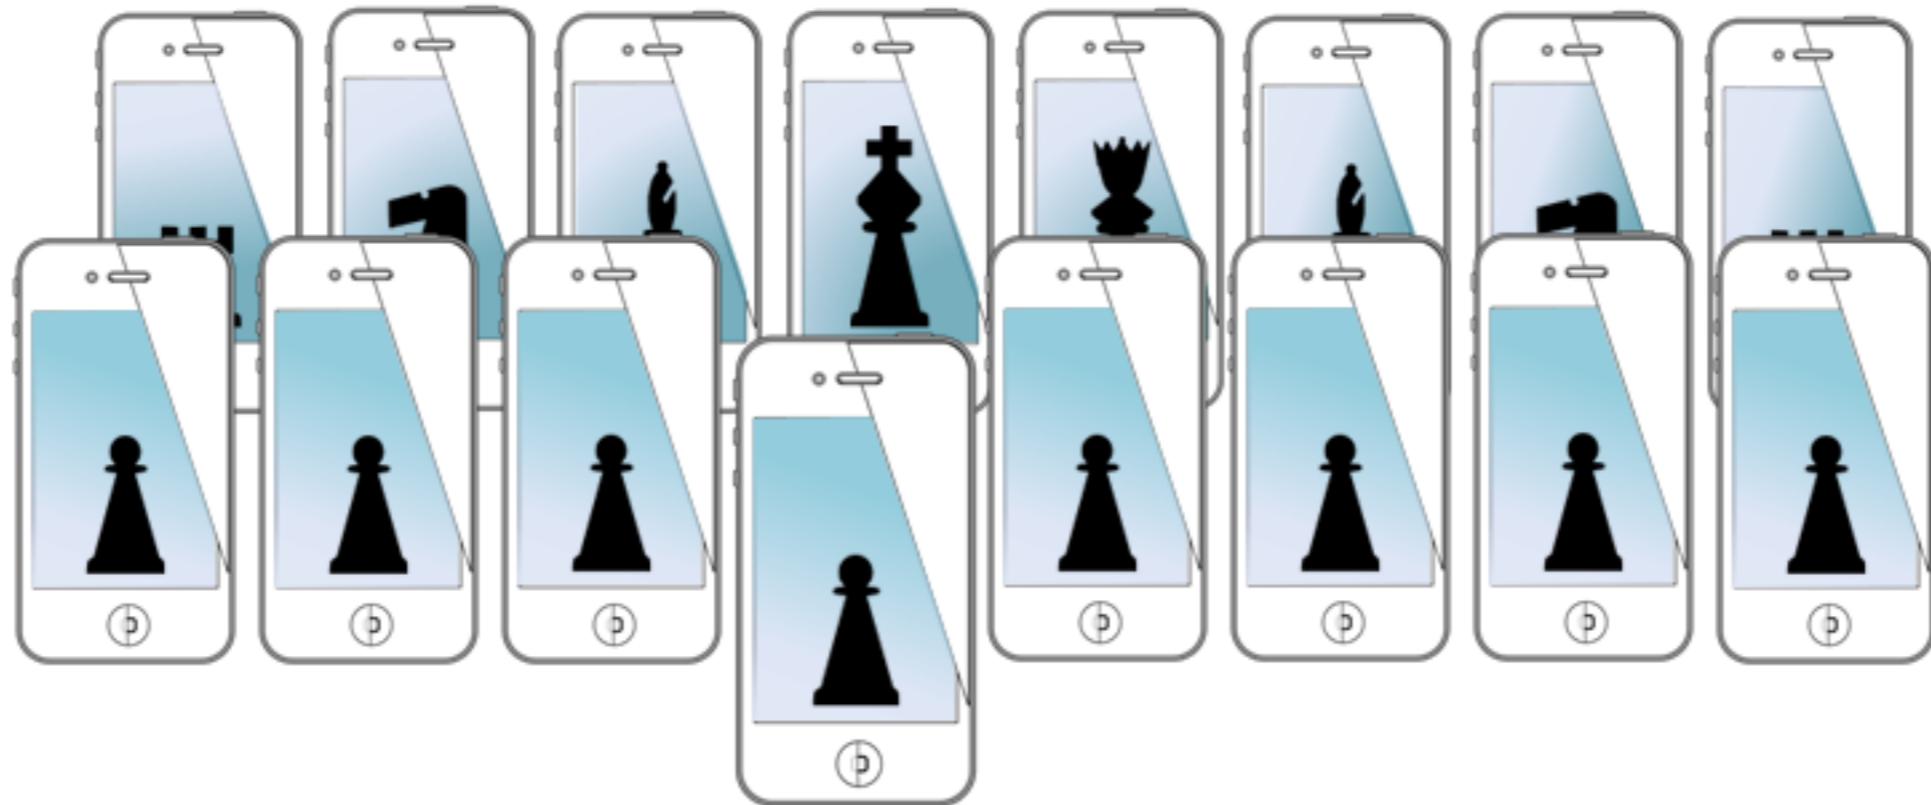
- pre-trained Haar cascades
- nose.xml



And now its
time for a demo



MOBILE SENSING LEARNING



CSE5323 & 7323
Mobile Sensing and Learning

week 7: computer vision with OpenCV

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University